

Timing Models

Suresh Purini, IIIT-H

When we study the design of logic circuits and understand their behavior by looking at their timing diagrams, we take an idealized view of things. We treat each circuit component as a mathematical object with well-defined behavioral model as given by its truth table. We also have a well-defined model for interconnections between various circuit components. For example refer the Timing Diagram for the positive edge triggered D-FlipFlop in the Figure 1.

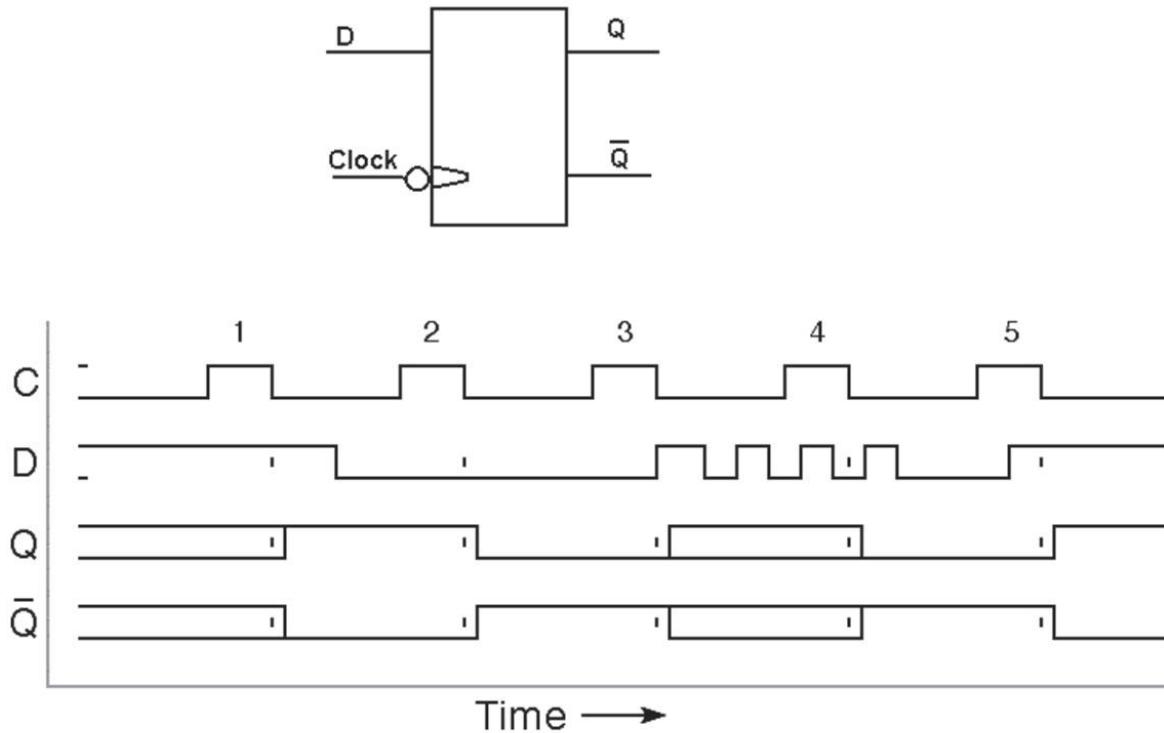


Figure 1: Timing Diagram for an Ideal Negative Edge Triggered D-Flip Flop

At the falling clock edge of the Clock Pulse 1, the D-Flip Flop input is equal to 1 and the D-Flip Flop reads the value and the output Q changes from 0 to 1 instantaneously. But that is not how it happens in any physical realization of the D-Flip Flop as against its mathematical model. In any physical realization of the D-flop the output Q will reflect the new value after certain delay. The following are three important characteristic parameters of the circuit components which are technology dependent (refer Figure 2).

1. **Set Up Time:** Setup time (denoted as t_{su}) is the minimum amount of time the data signal

should be held steady before the clock event so that the data are reliably sampled by the clock.

2. **Hold Time:** Hold time (denoted as t_h) is the minimum amount of time the data signal should be held steady after the clock event so that the data are reliably sampled.
3. **Clock-to-Q Time:** Clock-to-Q Time (denoted as t_{cq}) is the time it takes for the flip-flop to change its output after the clock edge.

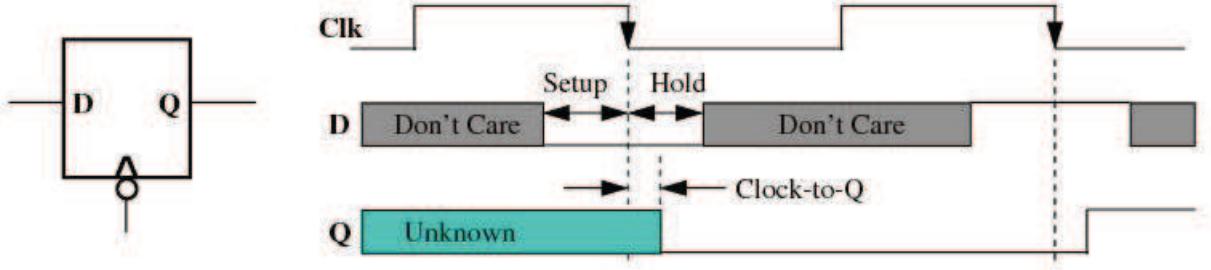


Figure 2: Diagram illustration Setup, Hold and Clock-to-Q Times.

You might wonder whether there is any relationship between the three parameters t_{su} , t_h and t_{cq} . Especially between the parameters t_h and t_{cq} , intuitively we feel there could be a relationship¹ For the sake of our discussion we can safely assume that all the three parameters are independent and their values are dictated by the underlying device physics. We shall now see how we have to factor in these parameters while designing logic circuits. Consider the Figure 3 where the output of the register R_1 is connected to the input of the register R_2 through some combinational logic circuit.

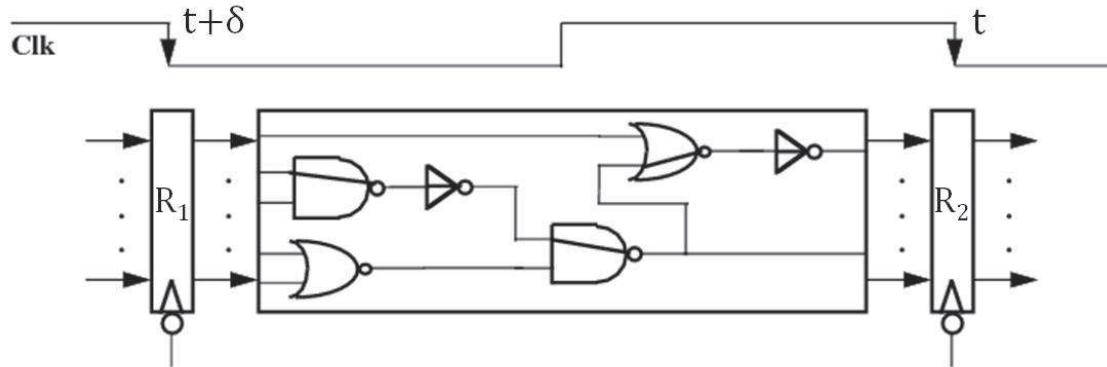


Figure 3: Two registers connected through intermediate combinational logic

At time instant t a falling clock edge arrives at both the registers R_1 and R_2 simultaneously (assuming no clock skew). The registers R_1 and R_2 sample their respective input signals at that falling clock edge and update their states accordingly and the updated state is available on the

¹I am not aware of any such relation. However I have to double check before I give a definite answer to you all.

output after the Clock-2-Q time which would be at the time instant $t + t_{cq}$. Remember the input signals to the registers R_1 and R_2 should be stable during the interval $[t - t_{su}, t + t_h]$ for the proper update of the registers. Now let us focus on how the signals flow between the registers R_1 and R_2 alone. The updated output signals of the register R_1 which are available at the time instant $t + t_{cq}$ flow through the combination logic and reach the input of the register R_2 . The combinational logic would induce some delay before the input signals to the register R_2 gets stabilized. This delay value is equal to the sum of the logic gate delays along the longest path through the combinational logic. This longest path is also called as the critical path. Let t_{cl} be the delay along this critical path. So the new input signals to the register R_2 will be available by the time instant $t + t_{cq} + t_{cl}$. It is important to note that during the time interval $[t + t_{cq}, t + t_{cq} + t_{cl}]$, the input signals to the register R_2 may oscillate (why?).

We expect that the newly available input signals to be latched onto the register R_2 at the next falling clock edge which arrives at the time instant $t + \delta$ where δ is the clock cycle length. If we factor in the setup time requirements for the proper update of the register R_2 , we will get the following necessary condition on the clock cycle time.

$$\begin{aligned} t + t_{cq} + t_{cl} + t_{su} &\leq t + \delta \\ \delta &\geq t_{cq} + t_{cl} + t_{su} \end{aligned}$$

If the above condition is not satisfied, then we say there is a **setup time violation**. In general if t_{cl}^{ij} is the critical path length in the combinational logic connecting two registers R_i and R_j , then the following condition has to be satisfied.

$$\delta \geq t_{cq} + t_{cl}^{ij} + t_{su}$$

Assuming that the Clock-to-Q and setup times are same for all the registers, the clock cycle length is governed by the largest critical path length. For example if the largest critical path corresponds to the combinational logic between two registers R_m and R_n then we can choose the clock cycle length as $\delta = t_{cq} + t_{cl}^{mn} + t_{su}$.

There is another interesting problem that we can face during logic circuit design. Let us consider a path in the combinational logic circuit connecting the registers R_1 and R_2 . Let t_p be the delay along this path. So after the falling clock edge occurring at time instant t , a signal can reach the input of the register R_2 by time $t + t_{cq} + t_p$. If it turns out that $t + t_{cq} + t_p \in [t - t_{su}, t + t_h]$ then the new signal value will over write the old input signals to register R_2 causing an instability of the input. We call this as **hold time violation**. To prevent hold time violation we need to make sure that for every path in the combinational logic circuit with delay t_p the following condition is true.

$$t + t_{cq} + t_p > t + t_h$$

If there is a path in the combinational circuit which does not satisfy the aforementioned condition, then the delay along the path has to be increased by adding buffer elements.

In the previous discussion we assumed that a following clock edge arrives at both the registers exactly at the same time instant. However in practice due to wire delays arising out of wire resistance and other physical characteristics of the wires, the same falling clock edge could reach different registers at different time instants (refer Figure 4). Let t_{cs} be the clock skew. Due to the clock skew the falling clock edge which is arriving at register R_2 at time instant t would now arrive at time instant $t + t_{cs}$. However it can so happen that during this time interval $[t, t + t_{cs}]$, along some path in the combination logic the new signal would flow and overwrite the old input signal to the register R_2 . This can happen if

$$t + t_{cq} + t_p \leq t + t_{cs} + t_h$$

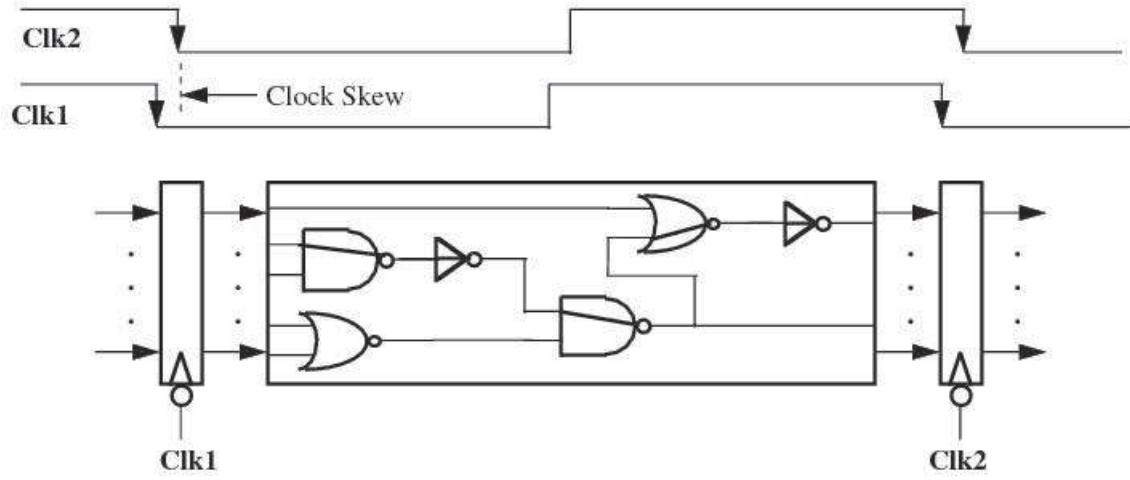


Figure 4: Clock skew due to clock signal propagation delays

where t_p is delay along the combination logic path under consideration. Refer to the Figures 5 and 6 which illustrate this problem.

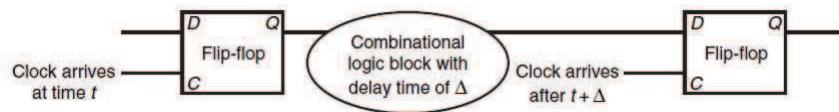


Figure 5: Potential problems due to clock skew

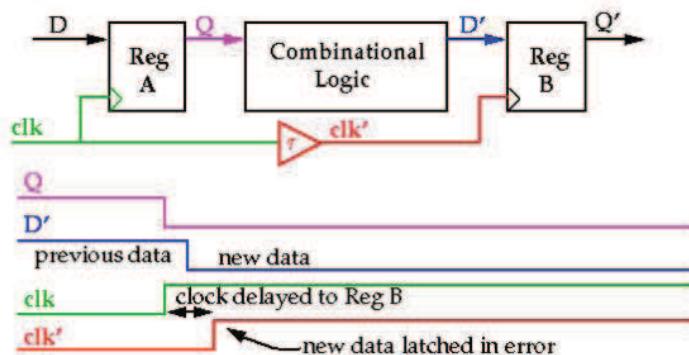


Figure 6: Potential problems due to clock skew

Representation of Integers and their Arithmetic

Suresh Purini, IIIT-H

What does the 8-bit string 11100000 represent? It could mean 224, -96, -31 and -32 when treated as an unsigned integer, sign-magnitude integer, one's complement integer and two's complement integer respectively. Or it could be mean the ASCII character α . So what a bit string means depends on the semantics or the definition we associate with it. In this write-up we shall study binary representation of unsigned and signed integers which is a primitive data structure supported by all modern processors.

1 Unsigned Integers

Consider the bijective function $B2U_w : \{0, 1\}^w \rightarrow \{0, \dots, 2^w - 1\}$ which maps w -bit binary strings to unsigned integers as follows.

$$B2U_w(\vec{b}) = \sum_{i=0}^{w-1} b_i 2^i$$

For example $B2U_4(0101) = 5$ and $B2U_4(1101) = 13$. You can observe that the function $B2U_w$ and its inverse are efficiently computable. In other words, we can easily compute the binary representation of an unsigned integer in the range of the function making it a viable representation.

In C-language all variables of type unsigned integers are allocated a fixed number of bytes (or equivalent number of bits) for storage which is typically 4 bytes or 32 bits. You can check this by running the following C-program on your machine.

```
#include <stdio.h>
main()
{
    printf("Size of Unsigned Integer: %d\n", sizeof(unsigned int));
}
```

Having represented unsigned integers in binary, we would like to figure out how to perform addition and multiplication operations. Let us just focus on addition operation in our discussion and the relevant ideas can be applied to multiplication operations too with suitable modifications. We presume that you know the algorithm to add two binary numbers as illustrated in the Figure 1¹. We also know the analogous algorithm for addition in the unsigned integer domain. Now the beauty of the mapping function $B2U_w$ is that it shows the isomorphic structure between the unsigned integers and their binary representation with respect to the addition operation (and also multiplication operation). To elaborate more on this idea let us define the w -bit addition of two numbers as the regular binary addition except that we ignore the carry-out bit from the MSB if at all there is one. With this definition, when we add two w -bit numbers, the result is always a w -bit number. The key claim here is whether we do addition of two unsigned integers in decimal notation

¹Recall the w -bit ripple carry adder circuit.

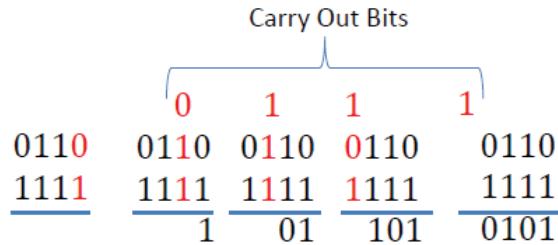


Figure 1: Addition of binary numbers

Row No	3-bit Binary	Unsigned Integer
R_0	000	0
R_1	001	1
R_2	010	2
R_3	011	3
R_4	100	4
R_5	101	5
R_6	110	6
R_7	111	7

Table 1: Isomorphic structure of 3-bit binary numbers and their unsigned interpretation

which we are familiar with or we do addition of their respective w -bit binary representations, the net result is just the same except for the difference in their notational representation. This claim is true so far as the result of the addition operation does cause an overflow, in other words the result would fit into w -bits. Consider the following Table 1 with three columns. If we want to add 1 and 4, whether we carry out addition in column 2 or in column 3, the respective results would fall in Row 5. However if we want to add 4 and 5, then the result wouldn't fall in the range in the column 2 and the result in the column 3 would fall in Row-1 (recall how w -bit addition is defined). It can be observed though that there is isomorphism between $(\text{mod } 2^w)$ addition of decimal numbers and w -bit addition of binary numbers without worrying about overflow at all since it would never happen in modular arithmetic. It has to be noted here that we can use any other function (preferably bijective) from the w -bit strings to unsigned numbers and create a isomorphism between the decimal domain and the binary domain by appropriately defining the addition operations in the binary domain. We leave it to you to ponder whether such an alternate function has any utility. It is easy to note here that the addition of two w -bit unsigned numbers would cause an overflow if and only if the carry-out bit is 1.

2 Signed Integers

The following are three different ways of representing signed integers.

1. Sign-Magnitude Representation. The mapping function here is:

$$B2S_w(b_{w-1} \dots b_0) = (-1)^{b_{w-1}} * (2^{w-2} * b_{w-2} + \dots + 2^0 * b_0)$$

Row No	3-bit String	Sign-Magnitude	1's Complement	2's Complement
R_0	000	0	0	0
R_1	001	1	1	1
R_2	010	2	2	2
R_3	011	3	3	3
R_4	100	-0	-3	-4
R_5	101	-1	-2	-3
R_6	110	-2	-1	-2
R_7	111	-3	0	-1

Table 2: Isomorphic structure of 3-bit binary numbers and 2's complement signed integers

2. 1's Complement Representation. The mapping function here is:

$$B2O_w(b_{w-1} \dots b_0) = -b_{w-1} * (2^{w-1} - 1) + b_{w-2} * 2^{w-2} + \dots + b_0 * 2^0$$

3. 2's Complement Representation. The mapping function here is:

$$B2T_w(b_{w-1} \dots b_0) = -b_{w-1} * 2^{w-1} + b_{w-2} * 2^{w-2} + \dots + b_0 * 2^0$$

Pretty much all systems use 2's complement representation for signed integers. We shall see the rationale behind such a choice in the following discussion. First you can verify that among the 3 mapping functions only the $B2T_w$ function corresponding to 2's complement representation is bijective. Let us stick to our definition of w -bit addition of binary numbers and we shall see that there is an isomorphic structure between signed integers and their 2's complement representation with respect to addition. It has to be noted that this isomorphism holds if and only if the results of addition does not cause overflow or underflow. Sign-magnitude and 1's complement representation of signed integers doesn't carry this isomorphic structure with respect to the canonical binary addition rules. It is worth noting that we can create an isomorphic structure even with these representations by suitable modifying the rules of binary addition. To understand these ideas consider the Table 2. For example if we add Row3 with Row4, the resulting binary number is 111 which lies in Row 7, whereas if we perform the addition on Sign-Magnitude numbers in Column 2, we get a value in Row 3 indicating the lack of isomorphic structure with respect to addition between the binary and sign-magnitude representation of numbers. It can be verified that there is no isomorphic structure between binary and one's complement representation of numbers by adding elements in Row 5 and Row 6. In binary addition we get an element in Row 3, whereas in the one's complement representation we get an element in Row 4 in Column 3. However it can be verified that as long as there is no overflow there is a perfect isomorphism with respect to addition between binary and two's complement representation of numbers.

3 Unsigned versus 2's Complement Addition

From the previous discussion it could have been noted that the rules of binary addition for both Unsigned and 2's Complement Addition is exactly the same. It means that we could use the same k -bit ripple carry to add any 2 unsigned or 2's complement numbers and we need not tell the k -bit ripple adder whether we are doing signed arithmetic or unsigned arithmetic. To illustrate this point further let us that I have a k -bit adder circuit with me, some of the students in the class want to

do 2's complement addition and some of you may want to perform unsigned addition over k -bit numbers using my k -bit adder circuit. But you don't want to reveal me whether you are performing signed or unsigned arithmetic for whatever reasons you have. It is no big deal for my k -bit adder circuit as the rules of addition remains the same for both signed and unsigned numbers. However there is a catch here. The catch is that overflow conditions for signed and unsigned arithmetic are different.

Predication in ARM ISA

Suresh Purini, IIIT-H

January 29, 2012

In the ARM ISA, for almost all instructions we can add a predicate as a suffix which tells the processor to execute the corresponding instruction only if the respective conditions associated with the predicate are satisfied (refer Table 1). For example when we add the predicate MI to an ADD instruction to get an instruction of the form ADDMI, then the following things happen:

- The assembler packs the bits 0100 in the most significant nibble while constructing the 32-bit opcode of the ADDMI instruction. This is an *offline* task that happens at the program assembly time.
- When the processor encounters the ADDMI instruction, it executes it if and only if the N flag is equal to 1 (or set in other words). This is an *online* task that happens at the program execution time.

Usually we use this predicate table in conjunction with the CMP instruction ¹. A brief description of the CMP instruction follows. When the processor encounters CMP r1, r2 instruction, it will subtract the register r_2 from r_1 and uses the result to affect the N, Z, C, V bits of the CPSR register as follows. Let $\text{alu_out} = r_1 - r_2$.

- $N = \text{alu_out}[31]$
- $Z = \text{if } \text{alu_out} == 0 \text{ then } 1 \text{ else } 0$
- $C = \text{if } r_1 \geq r_2 \text{ then } 1 \text{ else } 0$ (assuming r_1 and r_2 as unsigned integers)
- $V = 1$ if the result of subtraction does not fit into 32 bits else 0 (assuming r_1 and r_2 as signed integers)

Finally alu_out will be discarded without storing it in any register.

After the CMP instruction, we have to use the right predicate suffix to indicate the condition we would like to test. For example if we treat r_1 and r_2 as unsigned integers and want to check if $r_1 \leq r_2$, then we have to use the suffix LS (for example BLS). However if treat r_1 and r_2 as signed integer to test the same condition we use the suffix LE (for example BLE). As assembly language programmers this much knowledge is sufficient for us to write our programs correctly. However as processor designers and computer scientists, we are interested in checking whether the conditions mentioned in the Column 3 and Column 4 of the Table 1 are equivalent or not. For example is the following statement true after the execution of the CMP r_1, r_2 instruction.

$$r_1 \geq r_2 \text{ (signed)} \iff N = V$$

Question: Check the equivalence of the conditions in Column 3 and Column 4 of the Table 1 are equivalent (after the execution of the CMP instruction).

¹We can use these predicates in conjunction with other instructions also. For example SUBS instruction followed by ADDEQ.

Opcode [31:28]	Extension	Interpretation	Status flag state for execution
0000	EQ	Equal>equals zero	Z set
0001	NE	Not equal	Z clear
0010	CS/HS	Carry set/unsigned higher or same	C set
0011	CC/LO	Carry clear/unsigned lower	C clear
0100	MI	Minus/negative	N set
0101	PL	Plus/positive or zero	N clear
0110	VS	Overflow	V set
0111	VC	No overflow	V clear
1000	HI	Unsigned higher	C set and Z clear
1001	LS	Unsigned lower or same	C clear or Z set
1010	GE	Signed greater than or equal	N equals V
1011	LT	Signed less than	N is not equal to V
1100	GT	Signed greater than	Z clear and N equals V
1101	LE	Signed less than or equal	Z set or N is not equal to V
1110	AL	Always	any
1111	NV	Never (do not use!)	none

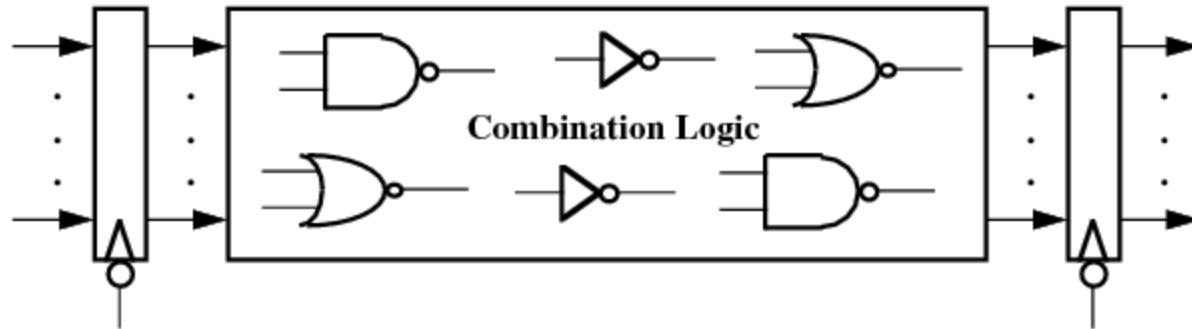
Table 1: ARM Predicate Table

Acknowledgment: Almost all of these slides are based on
Dave Patterson's CS152 Lecture Slides at UC, Berkeley.

COMPUTER SYSTEMS ORGANIZATION

Timing Model and Register File Design -- Spring 2012 --
IIIT-H -- Suresh Purini

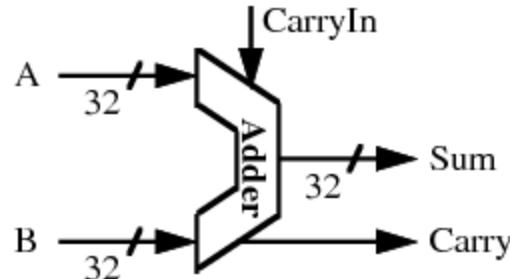
Sequential and Combinational Circuits



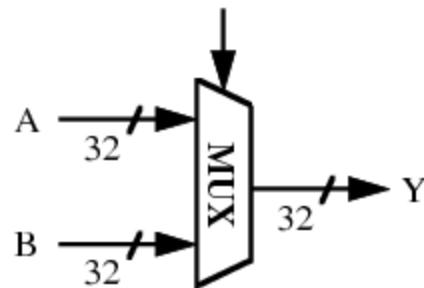
- ❑ What's the difference between sequential and combinational circuits?

Combinational Logic Elements

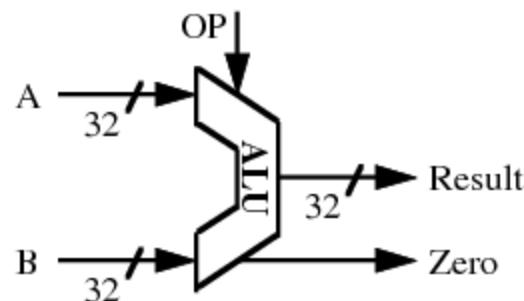
- ° **Adder**



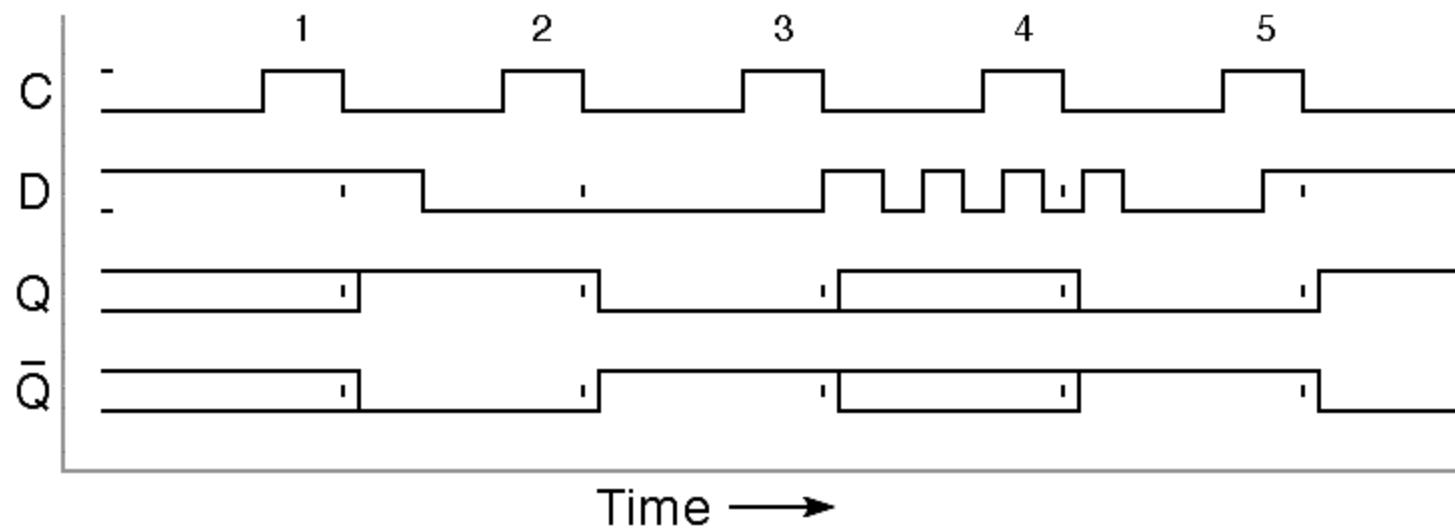
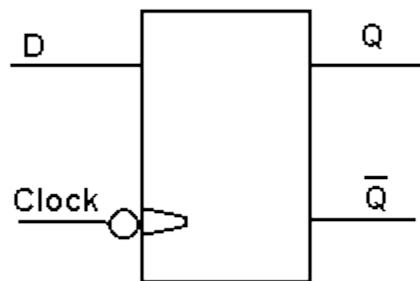
- ° **MUX**



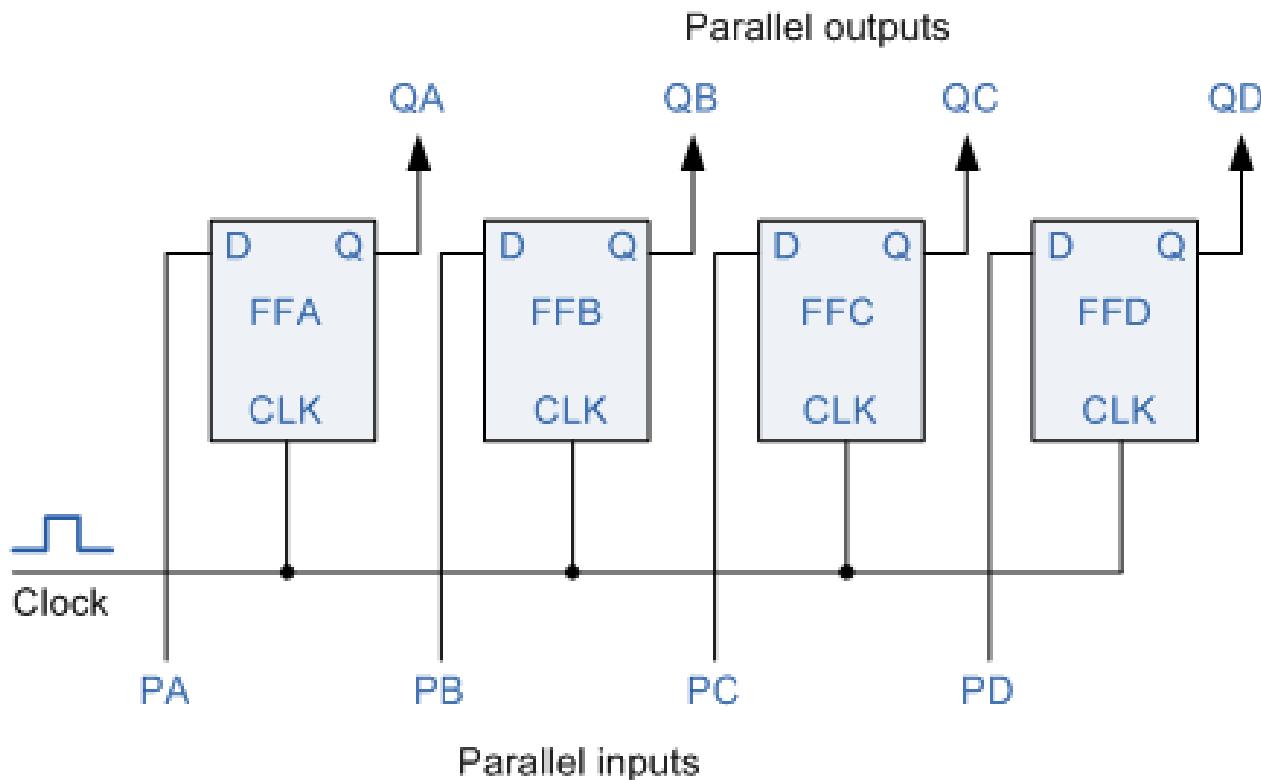
- ° **ALU**



Sequential Element: Negative Edge Triggered D-Flip Flop

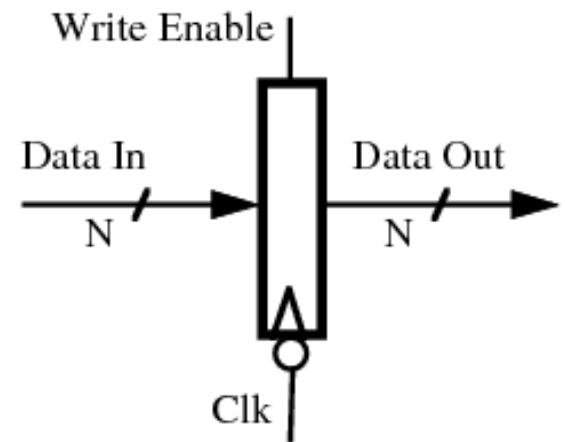


4-bit Register

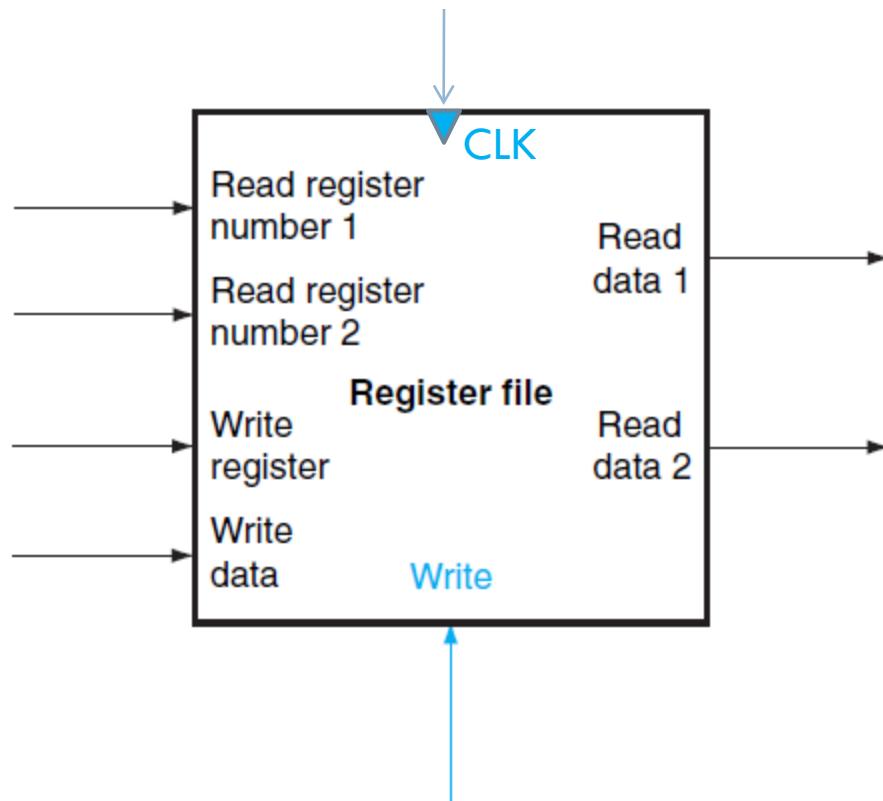


Sequential Element: Register

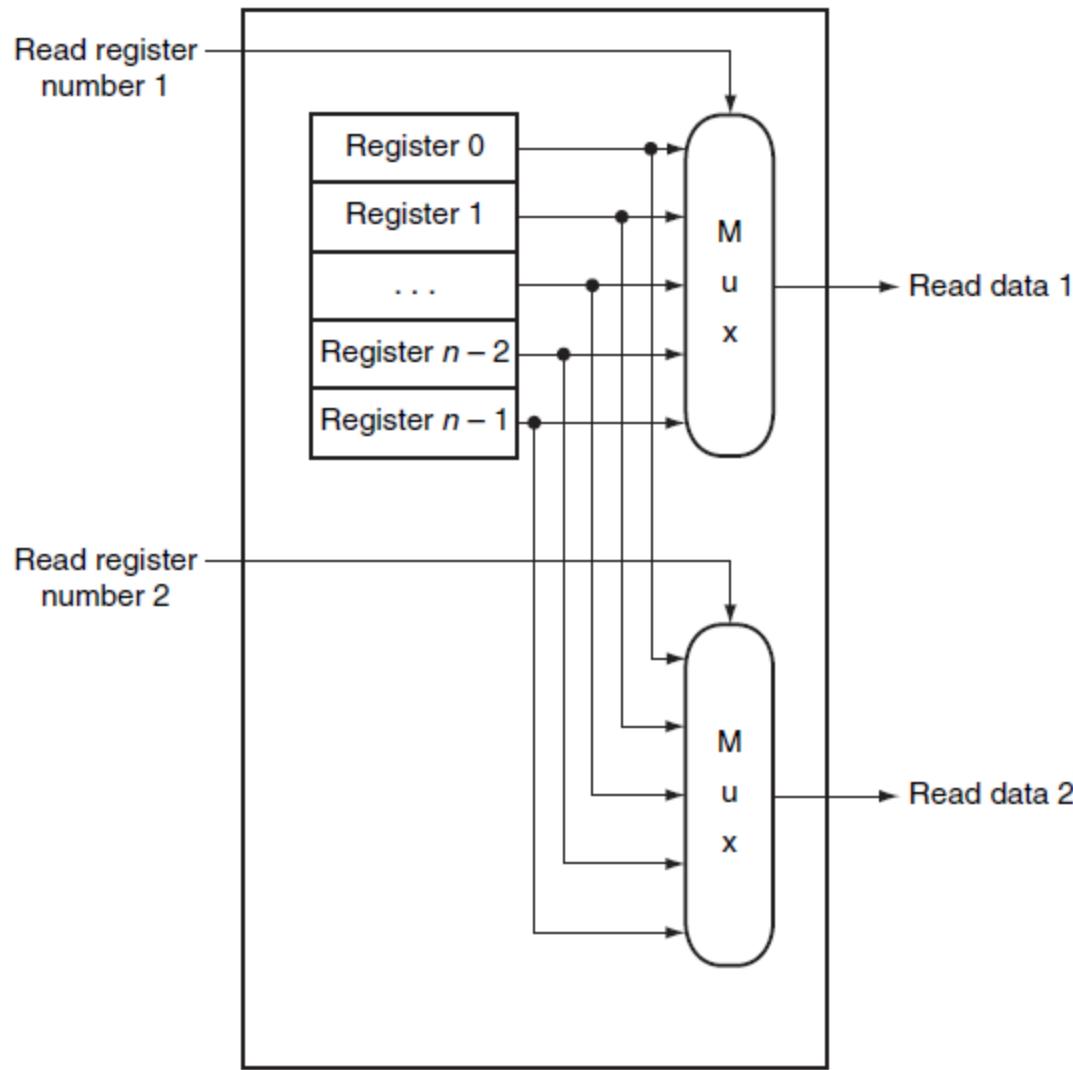
- Register
 - Similar to D Flip Flop except
 - N bit input and output
 - Write Enable input
 - Write Enable
 - 0: Data out will not change
 - 1: Data out will become Data In



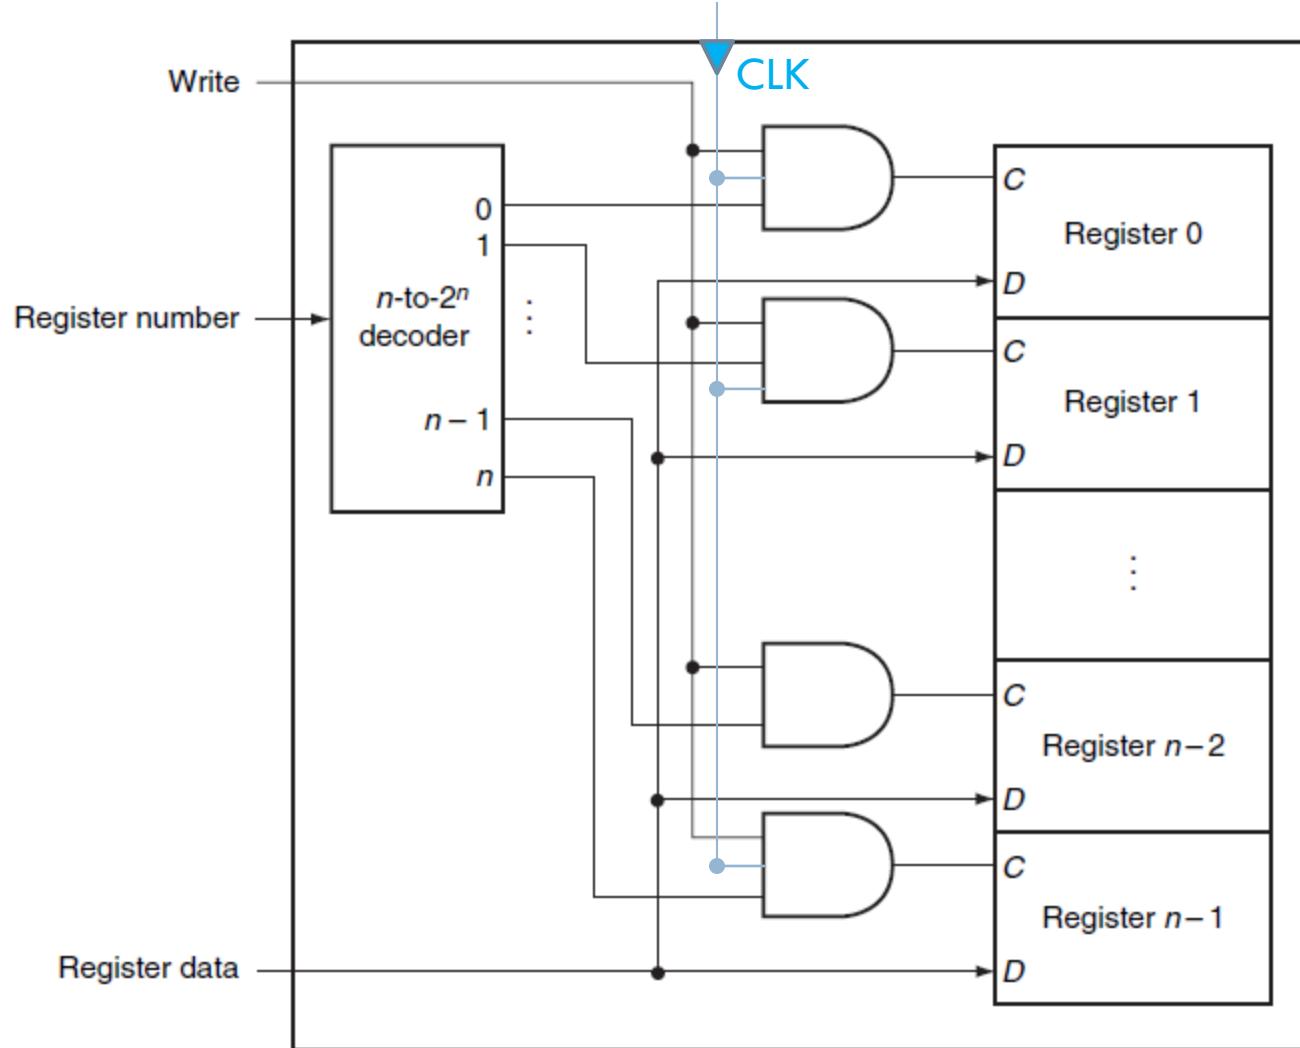
Register File



Register File

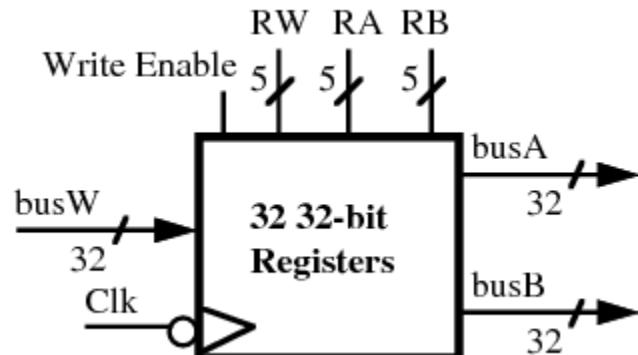


Register File



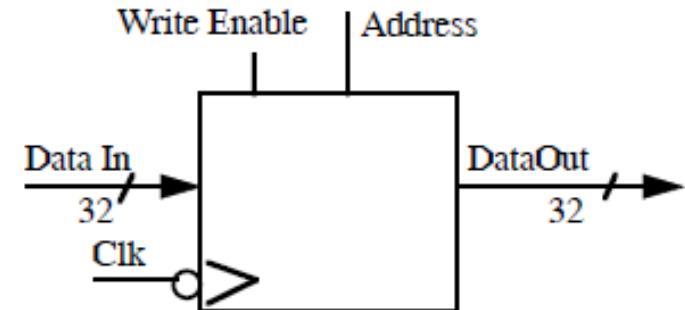
Storage Element: Register File

- Register File consists of 32 registers
 - Two 32-bit output busses: busA and busB
 - one 32-bit input bus: busW
- Register is selected by
 - RA selects the register to put on busA
 - RB selects the register to put on busB
 - RW selects the register to be written via busW when Write Enable is 1
- Clock input (CLK)
 - The CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block
 - RA or RB valid => busA or busB valid after access time

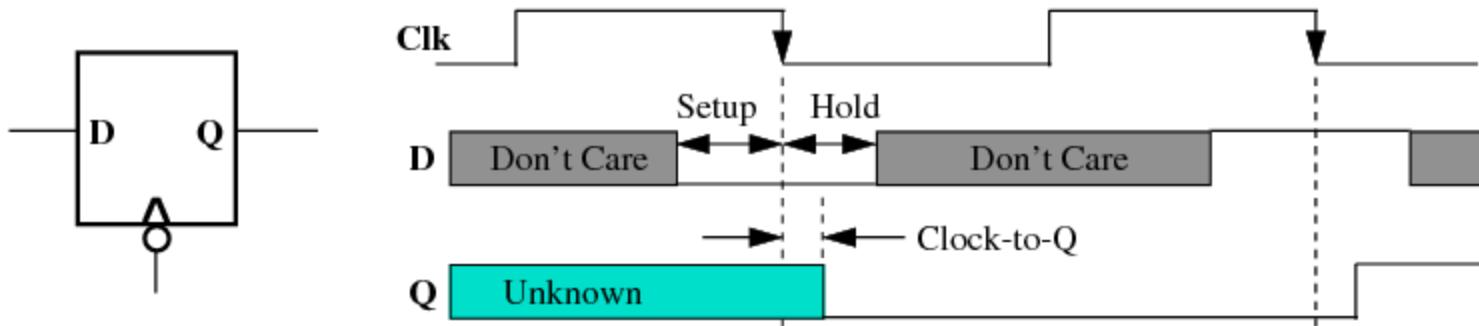


Storage Element: Memory

- **Memory**
 - One input bus: Data In
 - One output bus: Data Out
- **Memory word is selected by:**
 - Address selects the word to put on Data Out
 - Write Enable = 1: address selects the memory word to be written via the Data In bus
- **Clock input (CLK)**
 - The CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block:
 - Address valid => Data Out valid after “access time.”

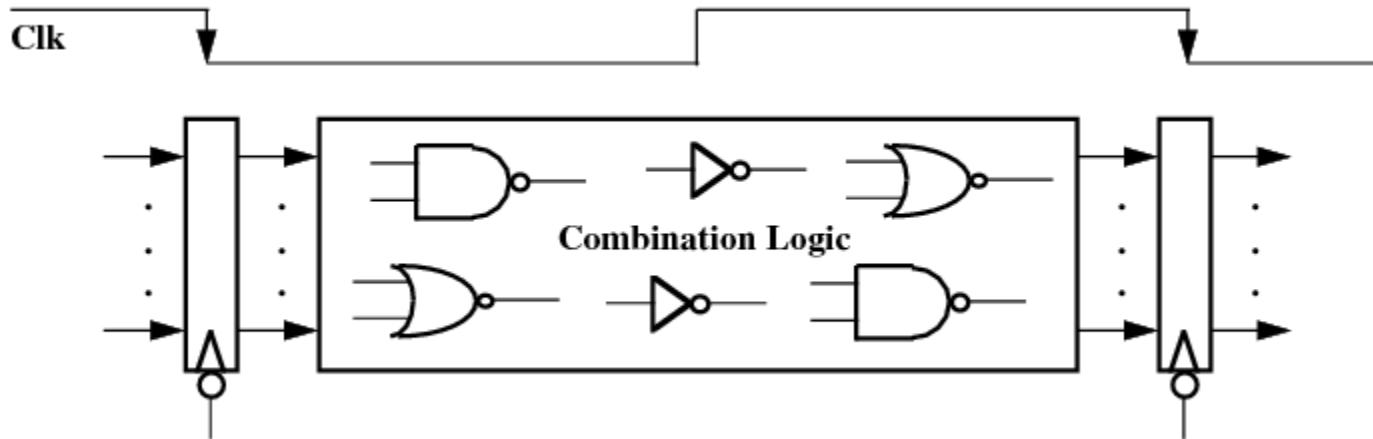


Storage Element Timing Model – Negative Edge Triggered D-Flip Flop



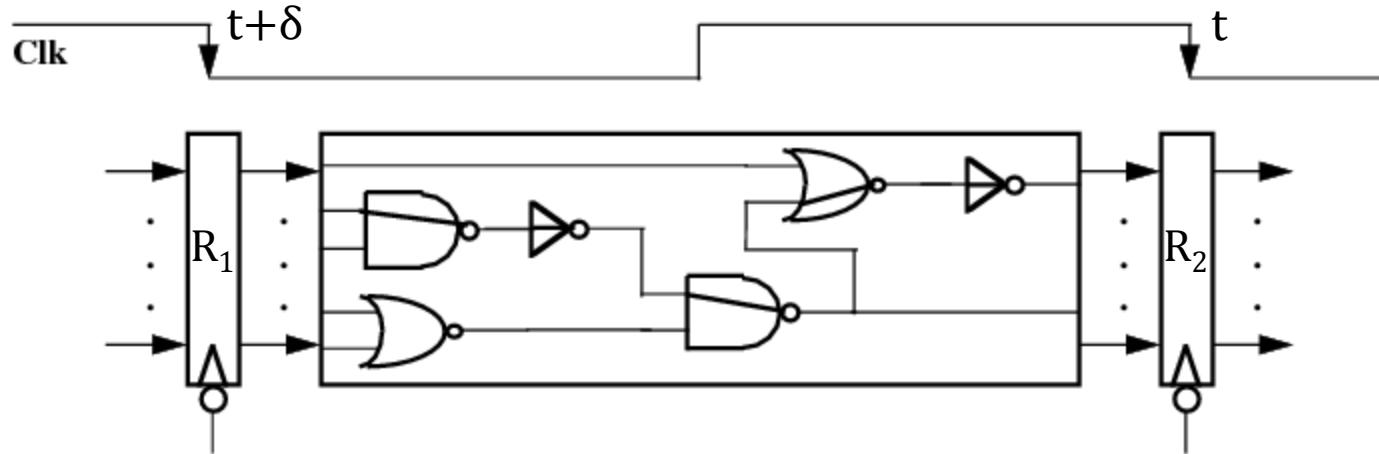
- **Setup Time:** Input must be stable BEFORE the trigger clock edge.
- **Hold Time:** Input must be stable AFTER the trigger clock edge.
- **Clock-to-Q time:** Output cannot change instantaneously at the trigger clock edge.
 - Similar to delay in logic gates.

Clocking Methodology



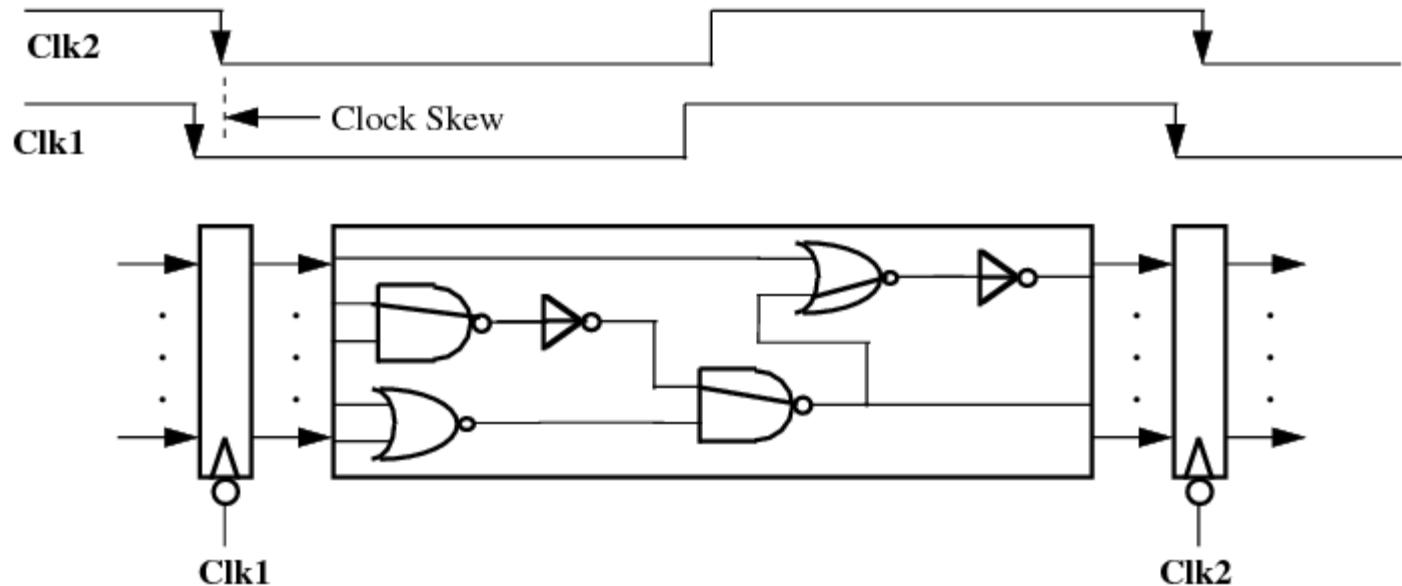
- All storage elements are clocked by the same clock edge
- The combination logic block's:
 - Inputs are updated at each clock tick
 - All outputs MUST be stable before the next clock tick

Critical Path and Cycle Time

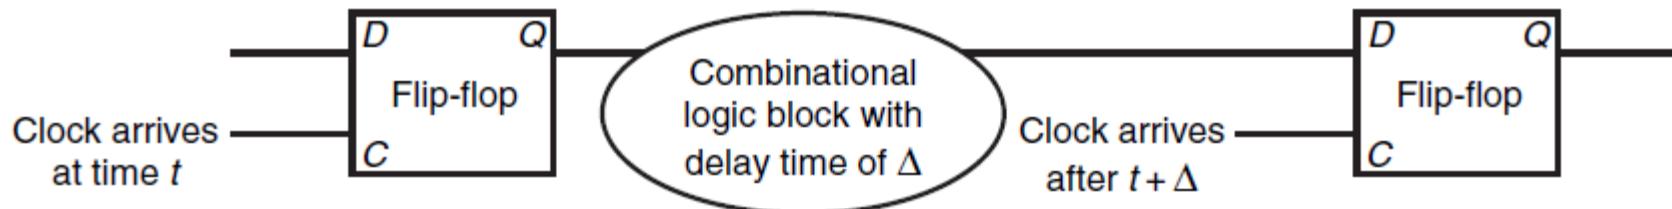


- **Critical Path:** Slowest path between any two storage devices
- Cycle time is a function of critical path
- More specifically, the cycle time must be greater than:
 - Clock-to-Q + Longest Path through the Combinational Logic + Setup Time

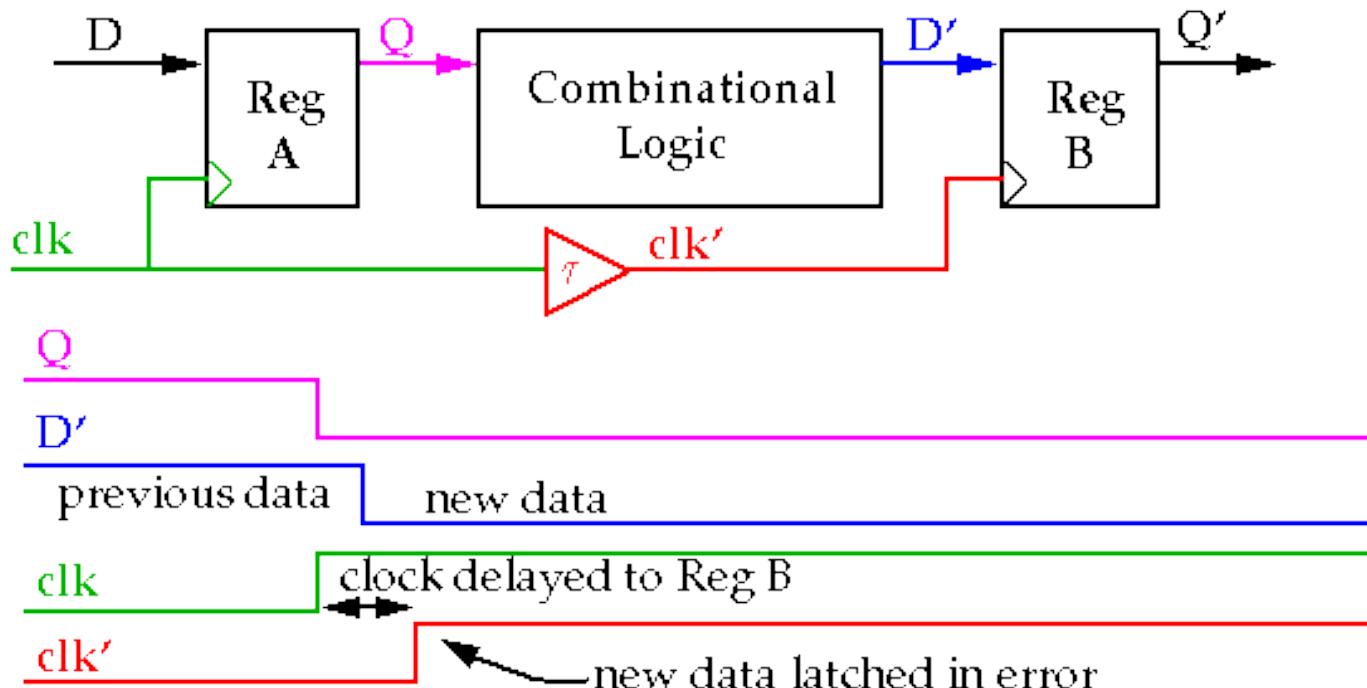
Clock Skew's Effect on Cycle Time



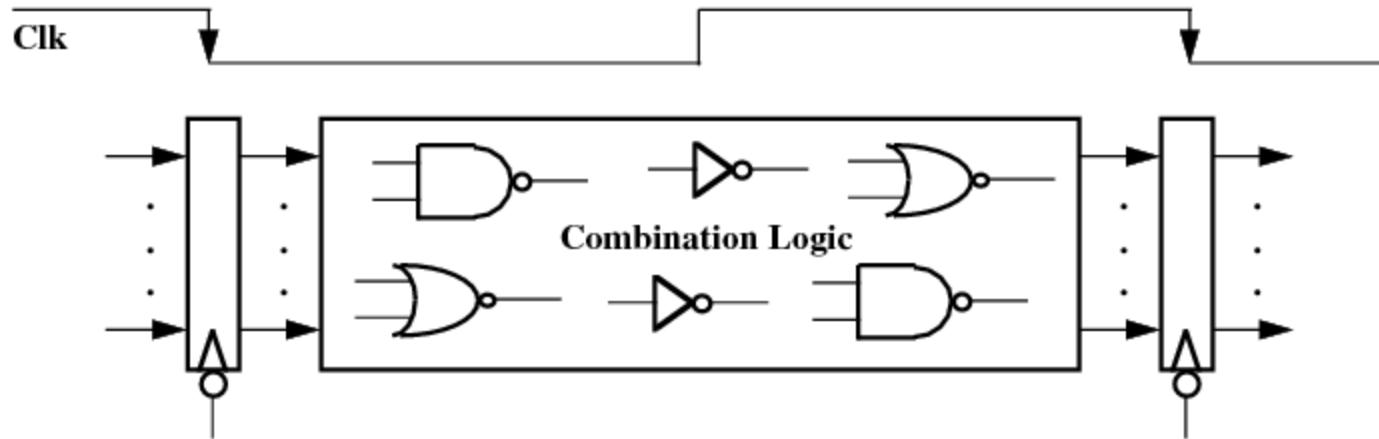
- ❑ How to take care of Clock Skew?
- ❑ We shall assume there is not Clock Skew.



Clock Skew



How to Avoid Hold Time Violation?



- Hold time requirement:
 - Input to register must NOT change immediately after the clock tick.
- CLK-to-Q + Shortest Delay Path must be greater than Hold Time

MIT 6.004 ISA Reference Card: Instructions

Instruction	Syntax	Description	Execution
LUI	lui rd, immU	Load Upper Immediate	$\text{reg}[rd] \leq \text{immU} \ll 12$
JAL	jal rd, immJ	Jump and Link	$\text{reg}[rd] \leq \text{pc} + 4$ $\text{pc} \leq \text{pc} + \text{immJ}$
JALR	jalr rd, rs1, immI	Jump and Link Register	$\text{reg}[rd] \leq \text{pc} + 4$ $\text{pc} \leq \{\text{reg}[rs1] + \text{immI}\}[31:1], 1'b0\}$
BEQ	beq rs1, rs2, immB	Branch if =	$\text{pc} \leq (\text{reg}[rs1] == \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
BNE	bne rs1, rs2, immB	Branch if \neq	$\text{pc} \leq (\text{reg}[rs1] != \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
BLT	blt rs1, rs2, immB	Branch if < (Signed)	$\text{pc} \leq (\text{reg}[rs1] <_s \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
BGE	bge rs1, rs2, immB	Branch if \geq (Signed)	$\text{pc} \leq (\text{reg}[rs1] >= \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
BLTU	bltu rs1, rs2, immB	Branch if < (Unsigned)	$\text{pc} \leq (\text{reg}[rs1] <_u \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
BGEU	bgeu rs1, rs2, immB	Branch if \geq (Unsigned)	$\text{pc} \leq (\text{reg}[rs1] >= \text{reg}[rs2]) ? \text{pc} + \text{immB} : \text{pc} + 4$
LW	lw rd, immI(rs1)	Load Word	$\text{reg}[rd] \leq \text{mem}[\text{reg}[rs1] + \text{immI}]$
SW	sw rs2, immS(rs1)	Store Word	$\text{mem}[\text{reg}[rs1] + \text{immS}] \leq \text{reg}[rs2]$
ADDI	addi rd, rs1, immI	Add Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] + \text{immI}$
SLTI	slti rd, rs1, immI	Compare < Immediate (Signed)	$\text{reg}[rd] \leq (\text{reg}[rs1] <_s \text{immI}) ? 1 : 0$
SLTIU	sltiu rd, rs1, immI	Compare < Immediate (Unsigned)	$\text{reg}[rd] \leq (\text{reg}[rs1] <_u \text{immI}) ? 1 : 0$
XORI	xori rd, rs1, immI	Xor Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] ^ \text{immI}$
ORI	ori rd, rs1, immI	Or Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] \text{immI}$
ANDI	andi rd, rs1, immI	And Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] \& \text{immI}$
SLLI	slli rd, rs1, immI	Shift Left Logical Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] \ll \text{immI}$
SRLI	srli rd, rs1, immI	Shift Right Logical Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] \gg_u \text{immI}$
SRAI	srai rd, rs1, immI	Shift Right Arithmetic Immediate	$\text{reg}[rd] \leq \text{reg}[rs1] \gg_s \text{immI}$
ADD	add rd, rs1, rs2	Add	$\text{reg}[rd] \leq \text{reg}[rs1] + \text{reg}[rs2]$
SUB	sub rd, rs1, rs2	Subtract	$\text{reg}[rd] \leq \text{reg}[rs1] - \text{reg}[rs2]$
SLL	sll rd, rs1, rs2	Shift Left Logical	$\text{reg}[rd] \leq \text{reg}[rs1] \ll \text{reg}[rs2]$
SLT	slt rd, rs1, rs2	Compare < (Signed)	$\text{reg}[rd] \leq (\text{reg}[rs1] <_s \text{reg}[rs2]) ? 1 : 0$
SLTU	sltu rd, rs1, rs2	Compare < (Unsigned)	$\text{reg}[rd] \leq (\text{reg}[rs1] <_u \text{reg}[rs2]) ? 1 : 0$
XOR	xor rd, rs1, rs2	Xor	$\text{reg}[rd] \leq \text{reg}[rs1] ^ \text{reg}[rs2]$
SRL	srl rd, rs1, rs2	Shift Right Logical	$\text{reg}[rd] \leq \text{reg}[rs1] \gg_u \text{reg}[rs2]$
SRA	sra rd, rs1, rs2	Shift Right Arithmetic	$\text{reg}[rd] \leq \text{reg}[rs1] \gg_s \text{reg}[rs2]$
OR	or rd, rs1, rs2	Or	$\text{reg}[rd] \leq \text{reg}[rs1] \text{reg}[rs2]$
AND	and rd, rs1, rs2	And	$\text{reg}[rd] \leq \text{reg}[rs1] \& \text{reg}[rs2]$

NOTE: All immediate values (immU , immJ , immI , immB , and immS) are sign-extended to 32-bits.

MIT 6.004 ISA Reference Card: Pseudoinstructions

Pseudoinstruction	Description	Execution
li rd, constant	Load Immediate	$\text{reg}[rd] \leq \text{constant}$
mv rd, rs1	Move	$\text{reg}[rd] \leq \text{reg}[rs1] + 0$
not rd, rs1	Logical Not	$\text{reg}[rd] \leq \text{reg}[rs1] ^ -1$
neg rd, rs1	Arithmetic Negation	$\text{reg}[rd] \leq 0 - \text{reg}[rs1]$
j label	Jump	$\text{pc} \leq \text{label}$
jal label	Jump and Link (with ra)	$\text{reg}[ra] \leq \text{pc} + 4$ $\text{pc} \leq \text{label}$
call label		
jr rs	Jump Register	$\text{pc} \leq \text{reg}[rs1] \& \sim 1$
jalr rs	Jump and Link Register (with ra)	$\text{reg}[ra] \leq \text{pc} + 4$ $\text{pc} \leq \text{reg}[rs1] \& \sim 1$
ret	Return from Subroutine	$\text{pc} \leq \text{reg}[ra]$
bgt rs1, rs2, label	Branch $>$ (Signed)	$\text{pc} \leq (\text{reg}[rs1] >_s \text{reg}[rs2]) ? \text{label} : \text{pc} + 4$
ble rs1, rs2, label	Branch \leq (Signed)	$\text{pc} \leq (\text{reg}[rs1] \leq_s \text{reg}[rs2]) ? \text{label} : \text{pc} + 4$
bgtu rs1, rs2, label	Branch $>$ (Unsigned)	$\text{pc} \leq (\text{reg}[rs1] >_u \text{reg}[rs2]) ? \text{label} : \text{pc} + 4$
bleu rs1, rs2, label	Branch \leq (Unsigned)	$\text{pc} \leq (\text{reg}[rs1] \leq_u \text{reg}[rs2]) ? \text{label} : \text{pc} + 4$
beqz rs1, label	Branch = 0	$\text{pc} \leq (\text{reg}[rs1] == 0) ? \text{label} : \text{pc} + 4$
bnez rs1, label	Branch $\neq 0$	$\text{pc} \leq (\text{reg}[rs1] != 0) ? \text{label} : \text{pc} + 4$
bltz rs1, label	Branch < 0 (Signed)	$\text{pc} \leq (\text{reg}[rs1] <_s 0) ? \text{label} : \text{pc} + 4$
bgez rs1, label	Branch ≥ 0 (Signed)	$\text{pc} \leq (\text{reg}[rs1] \geq_s 0) ? \text{label} : \text{pc} + 4$
bgtz rs1, label	Branch > 0 (Signed)	$\text{pc} \leq (\text{reg}[rs1] >_s 0) ? \text{label} : \text{pc} + 4$
blez rs1, label	Branch ≤ 0 (Signed)	$\text{pc} \leq (\text{reg}[rs1] \leq_s 0) ? \text{label} : \text{pc} + 4$

MIT 6.004 ISA Reference Card: Calling Convention

Registers	Symbolic names	Description	Saver
x0	zero	Hardwired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5-x7	t0-t2	Temporary registers	Caller
x8-x9	s0-s1	Saved registers	Callee
x10-x11	a0-a1	Function arguments and return values	Caller
x12-x17	a2-a7	Function arguments	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporary registers	Caller

MIT 6.004 ISA Reference Card: Instruction Encodings

31	25	24	20	19	15	14	12	11	7	6	0	
funct7		rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]				rs1		funct3		rd		opcode		I-type
imm[11:5]		rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]		rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
		imm[31:12]						rd		opcode		U-type
		imm[20 10:1 11 19:12]						rd		opcode		J-type

RV32I Base Instruction Set (MIT 6.004 subset)

imm[31:12]				rd	0110111	LUI
imm[20 10:1 11 19:12]				rd	1101111	JAL
imm[11:0]		rs1	000	rd	1100111	JALR
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	BGEU
imm[11:0]		rs1	010	rd	0000011	LW
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW
imm[11:0]		rs1	000	rd	0010011	ADDI
imm[11:0]		rs1	010	rd	0010011	SLTI
imm[11:0]		rs1	011	rd	0010011	SLTIU
imm[11:0]		rs1	100	rd	0010011	XORI
imm[11:0]		rs1	110	rd	0010011	ORI
imm[11:0]		rs1	111	rd	0010011	ANDI
0000000	shamt	rs1	001	rd	0010011	SLLI
0000000	shamt	rs1	101	rd	0010011	SRLI
0100000	shamt	rs1	101	rd	0010011	SRAI
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

RISC-V REFERENCE

RISC-V Instruction Set

Core Instruction Formats

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
														opcode
	funct7		rs2		rs1		funct3			rd				R-type
				imm[11:0]		rs1		funct3		rd				I-type
					imm[11:5]	rs2	rs1		funct3	imm[4:0]				S-type
				imm[12:10:5]	rs2	rs1		funct3	imm[4:1][11]	imm[4:1][11]				B-type
							imm[31:12]			rd				U-type
										rd				J-type
											rd			

RV32I Base Integer Instructions

Inst	Name	FMT	Opcode	funct3	funct7	Description (C)	Note
add	ADD	R	0110011	0x0	0x00	rd = rs1 + rs2	
sub	SUB	R	0110011	0x0	0x20	rd = rs1 - rs2	
xor	XOR	R	0110011	0x4	0x00	rd = rs1 ^ rs2	
or	OR	R	0110011	0x6	0x00	rd = rs1 rs2	
and	AND	R	0110011	0x7	0x00	rd = rs1 & rs2	
sll	Shift Left Logical	R	0110011	0x1	0x00	rd = rs1 << rs2	
srl	Shift Right Logical	R	0110011	0x5	0x00	rd = rs1 >> rs2	
sra	Shift Right Arith*	R	0110011	0x5	0x20	rd = rs1 >> rs2	
slt	Set Less Than	R	0110011	0x2	0x00	rd = (rs1 < rs2)?1:0	
sltu	Set Less Than (U)	R	0110011	0x3	0x00	rd = (rs1 < rs2)?1:0	zero-extends
addi	ADD Immediate	I	0010011	0x0		rd = rs1 + imm	
xori	XOR Immediate	I	0010011	0x4		rd = rs1 ^ imm	
ori	OR Immediate	I	0010011	0x6		rd = rs1 imm	
andi	AND Immediate	I	0010011	0x7		rd = rs1 & imm	
slli	Shift Left Logical Imm	I	0010011	0x1	imm[5:11]=0x00	rd = rs1 << imm[0:4]	
srli	Shift Right Logical Imm	I	0010011	0x5	imm[5:11]=0x00	rd = rs1 >> imm[0:4]	
srai	Shift Right Arith Imm	I	0010011	0x5	imm[5:11]=0x20	rd = rs1 >> imm[0:4]	msb-extends
slti	Set Less Than Imm	I	0010011	0x2		rd = (rs1 < imm)?1:0	
sltiu	Set Less Than Imm (U)	I	0010011	0x3		rd = (rs1 < imm)?1:0	zero-extends
lb	Load Byte	I	0000011	0x0		rd = M[rs1+imm][0:7]	
lh	Load Half	I	0000011	0x1		rd = M[rs1+imm][0:15]	
lw	Load Word	I	0000011	0x2		rd = M[rs1+imm][0:31]	
lbu	Load Byte (U)	I	0000011	0x4		rd = M[rs1+imm][0:7]	zero-extends
lhu	Load Half (U)	I	0000011	0x5		rd = M[rs1+imm][0:15]	zero-extends
sb	Store Byte	S	0100011	0x0		M[rs1+imm][0:7] = rs2[0:7]	
sh	Store Half	S	0100011	0x1		M[rs1+imm][0:15] = rs2[0:15]	
sw	Store Word	S	0100011	0x2		M[rs1+imm][0:31] = rs2[0:31]	
beq	Branch ==	B	1100011	0x0		if(rs1 == rs2) PC += imm	
bne	Branch !=	B	1100011	0x1		if(rs1 != rs2) PC += imm	
blt	Branch <	B	1100011	0x4		if(rs1 < rs2) PC += imm	
bge	Branch ≥	B	1100011	0x5		if(rs1 >= rs2) PC += imm	
bltu	Branch < (U)	B	1100011	0x6		if(rs1 < rs2) PC += imm	zero-extends
bgeu	Branch ≥ (U)	B	1100011	0x7		if(rs1 >= rs2) PC += imm	zero-extends
jal	Jump And Link	J	1101111			rd = PC+4; PC += imm	
jalr	Jump And Link Reg	I	1100111	0x0		rd = PC+4; PC = rs1 + imm	
lui	Load Upper Imm	U	0110111			rd = imm << 12	
auipc	Add Upper Imm to PC	U	0010111			rd = PC + (imm << 12)	
ecall	Environment Call	I	1110011	0x0	imm=0x0	Transfer control to OS	
ebreak	Environment Break	I	1110011	0x0	imm=0x1	Transfer control to debugger	

Standard Extensions

RV32M Multiply Extension

Inst	Name	FMT	Opcode	funct3	funct7	Description (C)
mul	MUL	R	0110011	0x0	0x01	$rd = (rs1 * rs2)[31:0]$
mulh	MUL High	R	0110011	0x1	0x01	$rd = (rs1 * rs2)[63:32]$
mulsu	MUL High (S) (U)	R	0110011	0x2	0x01	$rd = (rs1 * rs2)[63:32]$
mulu	MUL High (U)	R	0110011	0x3	0x01	$rd = (rs1 * rs2)[63:32]$
div	DIV	R	0110011	0x4	0x01	$rd = rs1 / rs2$
divu	DIV (U)	R	0110011	0x5	0x01	$rd = rs1 / rs2$
rem	Remainder	R	0110011	0x6	0x01	$rd = rs1 \% rs2$
remu	Remainder (U)	R	0110011	0x7	0x01	$rd = rs1 \% rs2$

RV32A Atomic Extension

31	27	26	25	24	20 19	15 14	12 11	7 6	0
funct5	aq	rl		rs2	rs1	funct3	rd		opcode
5	1	1		5	5	3	5		7

Inst	Name	FMT	Opcode	funct3	funct5	Description (C)
lr.w	Load Reserved	R	0101111	0x2	0x02	$rd = M[rs1]$, reserve $M[rs1]$
sc.w	Store Conditional	R	0101111	0x2	0x03	if (reserved) { $M[rs1] = rs2$; $rd = 0$ } else { $rd = 1$ }
amoswap.w	Atomic Swap	R	0101111	0x2	0x01	$rd = M[rs1]$; swap(rd , $rs2$); $M[rs1] = rd$
amoadd.w	Atomic ADD	R	0101111	0x2	0x00	$rd = M[rs1] + rs2$; $M[rs1] = rd$
amoand.w	Atomic AND	R	0101111	0x2	0x0C	$rd = M[rs1] \& rs2$; $M[rs1] = rd$
amoor.w	Atomic OR	R	0101111	0x2	0x0A	$rd = M[rs1] rs2$; $M[rs1] = rd$
amoxor.w	Atomix XOR	R	0101111	0x2	0x04	$rd = M[rs1] ^ rs2$; $M[rs1] = rd$
amomax.w	Atomic MAX	R	0101111	0x2	0x14	$rd = \max(M[rs1], rs2)$; $M[rs1] = rd$
amomin.w	Atomic MIN	R	0101111	0x2	0x10	$rd = \min(M[rs1], rs2)$; $M[rs1] = rd$

RV32F / D Floating-Point Extensions

Inst	Name	FMT	Opcode	funct3	funct5	Description (C)
flw	Flt Load Word	*				$rd = M[rs1 + imm]$
fsw	Flt Store Word	*				$M[rs1 + imm] = rs2$
fmadd.s	Flt Fused Mul-Add	*				$rd = rs1 * rs2 + rs3$
fmsub.s	Flt Fused Mul-Sub	*				$rd = rs1 * rs2 - rs3$
fnmadd.s	Flt Neg Fused Mul-Add	*				$rd = -rs1 * rs2 + rs3$
fnmsub.s	Flt Neg Fused Mul-Sub	*				$rd = -rs1 * rs2 - rs3$
fadd.s	Flt Add	*				$rd = rs1 + rs2$
fsub.s	Flt Sub	*				$rd = rs1 - rs2$
fmul.s	Flt Mul	*				$rd = rs1 * rs2$
fdiv.s	Flt Div	*				$rd = rs1 / rs2$
fsqrt.s	Flt Square Root	*				$rd = \sqrt{rs1}$
fsgnj.s	Flt Sign Injection	*				$rd = \text{abs}(rs1) * \text{sgn}(rs2)$
fsgnjn.s	Flt Sign Neg Injection	*				$rd = \text{abs}(rs1) * -\text{sgn}(rs2)$
fsgnjx.s	Flt Sign Xor Injection	*				$rd = rs1 * \text{sgn}(rs2)$
fmin.s	Flt Minimum	*				$rd = \min(rs1, rs2)$
fmax.s	Flt Maximum	*				$rd = \max(rs1, rs2)$
fcvt.s.w	Flt Conv from Sign Int	*				$rd = (\text{float}) rs1$
fcvt.s.wu	Flt Conv from Uns Int	*				$rd = (\text{float}) rs1$
fcvt.w.s	Flt Convert to Int	*				$rd = (\text{int32_t}) rs1$
fcvt.w.u.s	Flt Convert to Int	*				$rd = (\text{uint32_t}) rs1$
fmv.x.w	Move Float to Int	*				$rd = *((\text{int}*) \&rs1)$
fmv.w.x	Move Int to Float	*				$rd = *((\text{float}*) \&rs1)$
feq.s	Float Equality	*				$rd = (rs1 == rs2) ? 1 : 0$
flt.s	Float Less Than	*				$rd = (rs1 < rs2) ? 1 : 0$
fle.s	Float Less / Equal	*				$rd = (rs1 \leq rs2) ? 1 : 0$
fclass.s	Float Classify	*				$rd = 0..9$

RV32C Compressed Extension

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0			
				funct4				rd/rs1				rs2		op				
				funct3	imm				rd/rs1				imm		op			
				funct3					imm				rs2		op			
				funct3					imm				rd'		op			
				funct3	imm		rs1'		imm		rd'		op					
				funct3	imm		rd'/rs1'		imm		rs2'		op					
				funct3	imm		rs1'		imm				op					
				funct3	offset								op					

Inst	Name	FMT	OP	Funct	Description
c.lwsp	Load Word from SP	CI	10	010	lw rd, (4*imm)(sp)
c.swsp	Store Word to SP	CSS	10	110	sw rs2, (4*imm)(sp)
c.lw	Load Word	CL	00	010	lw rd', (4*imm)(rs1')
c.sw	Store Word	CS	00	110	sw rs1', (4*imm)(rs2')
c.j	Jump	CJ	01	101	jal x0, 2*offset
c.jal	Jump And Link	CJ	01	001	jal ra, 2*offset
c.jr	Jump Reg	CR	10	1000	jalr x0, rs1, 0
c.jalr	Jump And Link Reg	CR	10	1001	jalr ra, rs1, 0
c.beqz	Branch == 0	CB	01	110	beq rs', x0, 2*imm
c.bnez	Branch != 0	CB	01	111	bne rs', x0, 2*imm
c.li	Load Immediate	CI	01	010	addi rd, x0, imm
c.lui	Load Upper Imm	CI	01	011	lui rd, imm
c.addi	ADD Immediate	CI	01	000	addi rd, rd, imm
c.addi16sp	ADD Imm * 16 to SP	CI	01	011	addi sp, sp, 16*imm
c.addi4spn	ADD Imm * 4 + SP	CIW	00	000	addi rd', sp, 4*imm
c.slli	Shift Left Logical Imm	CI	10	000	slli rd, rd, imm
c.srli	Shift Right Logical Imm	CB	01	100x00	srlti rd', rd', imm
c.srai	Shift Right Arith Imm	CB	01	100x01	srai rd', rd', imm
c.andi	AND Imm	CB	01	100x10	andi rd', rd', imm
c.mv	MoVe	CR	10	1000	add rd, x0, rs2
c.add	ADD	CR	10	1001	add rd, rd, rs2
c.and	AND	CS	01	10001111	and rd', rd', rs2
c.or	OR	CS	01	10001110	or rd', rd', rs2
c.xor	XOR	CS	01	10001101	xor rd', rd', rs2
c.sub	SUB	CS	01	10001100	sub rd', rd', rs2
c.nop	No OPeration	CI	01	000	addi x0, x0, 0
c.ebreak	Environment BREAK	CR	10	1001	ebreak

Pseudo Instructions

Pseudoinstruction	Base Instruction(s)	Meaning
la rd, symbol	auipc rd, symbol[31:12] addi rd, rd, symbol[11:0]	Load address
l{b h w d} rd, symbol	auipc rd, symbol[31:12] l{b h w d} rd, symbol[11:0](rd)	Load global
s{b h w d} rd, symbol, rt	auipc rt, symbol[31:12] s{b h w d} rd, symbol[11:0](rt)	Store global
f{w d} rd, symbol, rt	auipc rt, symbol[31:12] f{w d} rd, symbol[11:0](rt)	Floating-point load global
fs{w d} rd, symbol, rt	auipc rt, symbol[31:12] fs{w d} rd, symbol[11:0](rt)	Floating-point store global
nop	addi x0, x0, 0	No operation
li rd, immediate	<i>Myriad sequences</i>	Load immediate
mv rd, rs	addi rd, rs, 0	Copy register
not rd, rs	xori rd, rs, -1	One's complement
neg rd, rs	sub rd, x0, rs	Two's complement
negw rd, rs	subw rd, x0, rs	Two's complement word
sext.w rd, rs	addiw rd, rs, 0	Sign extend word
seqz rd, rs	sltiu rd, rs, 1	Set if = zero
snez rd, rs	sltu rd, x0, rs	Set if \neq zero
sltz rd, rs	slt rd, rs, x0	Set if < zero
sgtz rd, rs	slt rd, x0, rs	Set if > zero
fmv.s rd, rs	fsgnj.s rd, rs, rs	Copy single-precision register
fabs.s rd, rs	fsgnjx.s rd, rs, rs	Single-precision absolute value
fneg.s rd, rs	fsgnjn.s rd, rs, rs	Single-precision negate
fmv.d rd, rs	fsgnj.d rd, rs, rs	Copy double-precision register
fabs.d rd, rs	fsgnjx.d rd, rs, rs	Double-precision absolute value
fneg.d rd, rs	fsgnjn.d rd, rs, rs	Double-precision negate
beqz rs, offset	beq rs, x0, offset	Branch if = zero
bnez rs, offset	bne rs, x0, offset	Branch if \neq zero
blez rs, offset	bge x0, rs, offset	Branch if \leq zero
bgez rs, offset	bge rs, x0, offset	Branch if \geq zero
bltz rs, offset	blt rs, x0, offset	Branch if < zero
bgtz rs, offset	blt x0, rs, offset	Branch if > zero
bgt rs, rt, offset	blt rt, rs, offset	Branch if >
ble rs, rt, offset	bge rt, rs, offset	Branch if \leq
bgtu rs, rt, offset	bltu rt, rs, offset	Branch if >, unsigned
bleu rs, rt, offset	bgeu rt, rs, offset	Branch if \leq , unsigned
j offset	jal x0, offset	Jump
jal offset	jal x1, offset	Jump and link
jr rs	jalr x0, rs, 0	Jump register
jalr rs	jalr x1, rs, 0	Jump and link register
ret	jalr x0, x1, 0	Return from subroutine
call offset	auipc x1, offset[31:12] jalr x1, x1, offset[11:0]	Call far-away subroutine
tail offset	auipc x6, offset[31:12] jalr x0, x6, offset[11:0]	Tail call far-away subroutine
fence	fence iorw, iorw	Fence on all memory and I/O

Registers

Register	ABI Name	Description	Saver
x0	zero	Zero constant	—
x1	ra	Return address	Callee
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5-x7	t0-t2	Temporaries	Caller
x8	s0 / fp	Saved / frame pointer	Callee
x9	s1	Saved register	Callee
x10-x11	a0-a1	Fn args/return values	Caller
x12-x17	a2-a7	Fn args	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporaries	Caller
f0-7	ft0-7	FP temporaries	Caller
f8-9	fs0-1	FP saved registers	Callee
f10-11	fa0-1	FP args/return values	Caller
f12-17	fa2-7	FP args	Caller
f18-27	fs2-11	FP saved registers	Callee
f28-31	ft8-11	FP temporaries	Caller

Base Integer Instructions: RV32I, RV64I, and RV128I							RV Privileged Instructions									
Category	Name	Fmt	RV32I Base			+RV{64,128}			Category	Name	RV mnemonic					
Loads	Load Byte	I	LB	rd,rs1,imm		L{D Q} rd,rs1,imm			CSR Access	Atomic R/W	CSRRW	rd,csr,rs1				
	Load Halfword	I	LH	rd,rs1,imm						Atomic Read & Set Bit	CSRRS	rd,csr,rs1				
	Load Word	I	LW	rd,rs1,imm						Atomic Read & Clear Bit	CSRRC	rd,csr,rs1				
	Load Byte Unsigned	I	LBU	rd,rs1,imm						Atomic R/W Imm	CSRRWI	rd,csr,imm				
	Load Half Unsigned	I	LHU	rd,rs1,imm						Atomic Read & Set Bit Imm	CSRSSI	rd,csr,imm				
Stores	Store Byte	S	SB	rs1,rs2,imm		S{D Q} rs1,rs2,imm			Change Level	Env. Call	ECALL					
	Store Halfword	S	SH	rs1,rs2,imm						Environment Breakpoint	EBREAK					
	Store Word	S	SW	rs1,rs2,imm						Environment Return	ERET					
Shifts	Shift Left	R	SLL	rd,rs1,rs2		SLL{W D} rd,rs1,rs2			Trap Redirect to Supervisor	MRTS						
	Shift Left Immediate	I	SLLI	rd,rs1,shamt		SLLI{W D} rd,rs1,shamt				Redirect Trap to Hypervisor	MRTH					
	Shift Right	R	SRL	rd,rs1,rs2		SRL{W D} rd,rs1,rs2				Hypervisor Trap to Supervisor	HRTS					
	Shift Right Immediate	I	SRLI	rd,rs1,shamt		SRLI{W D} rd,rs1,shamt				Interrupt	WFI					
	Shift Right Arithmetic	R	SRA	rd,rs1,rs2		SRA{W D} rd,rs1,rs2				MMU	Supervisor FENCE	SFENCE.VM rs1				
	Shift Right Arith Imm	I	SRAI	rd,rs1,shamt		SRAI{W D} rd,rs1,shamt										
Arithmetic	ADD	R	ADD	rd,rs1,rs2		ADD{W D} rd,rs1,rs2			Optional Compressed (16-bit) Instruction Extension: RVC							
	ADD Immediate	I	ADDI	rd,rs1,imm		ADDI{W D} rd,rs1,imm										
	SUBtract	R	SUB	rd,rs1,rs2		SUB{W D} rd,rs1,rs2										
	Load Upper Imm	U	LUI	rd,imm												
Logical	Add Upper Imm to PC	U	AUIPC	rd,imm					Optional Compressed (16-bit) Instruction Extension: RVC							
	XOR	R	XOR	rd,rs1,rs2												
Logical	XOR Immediate	I	XORI	rd,rs1,imm												
	OR	R	OR	rd,rs1,rs2												
	OR Immediate	I	ORI	rd,rs1,imm												
	AND	R	AND	rd,rs1,rs2												
Logical	AND Immediate	I	ANDI	rd,rs1,imm												
	Set <	R	SLT	rd,rs1,rs2												
Compare	Set < Immediate	I	SLTI	rd,rs1,imm												
	Set < Unsigned	R	SLTU	rd,rs1,rs2												
Compare	Set < Imm Unsigned	I	SLTIU	rd,rs1,imm												
Branches	Branch =	SB	BEQ	rs1,rs2,imm												
	Branch ≠	SB	BNE	rs1,rs2,imm												
	Branch <	SB	BLT	rs1,rs2,imm												
	Branch ≥	SB	BGE	rs1,rs2,imm												
	Branch < Unsigned	SB	BLTU	rs1,rs2,imm												
	Branch ≥ Unsigned	SB	BGEU	rs1,rs2,imm												
Jump & Link	J&L	UJ	JAL	rd,imm					Optional Compressed (16-bit) Instruction Extension: RVC							
	Jump & Link Register	UJ	JALR	rd,rs1,imm												
Synch	Synch thread	I	FENCE													
	Synch Instr & Data	I	FENCE.I													
System	System CALL	I	SCALL													
	System BREAK	I	SBREAK													
Counters	ReaD CYCLE	I	RDCYCLE	rd												
	ReaD CYCLE upper Half	I	RDCYCLES	rd												
	ReaD TIME	I	RDTIME	rd												
	ReaD TIME upper Half	I	RDTIMES	rd												
	ReaD INSTR RETired	I	RDINSTRET	rd												
	ReaD INSTR upper Half	I	RDINSTRETH	rd												
32-bit Instruction Formats							16-bit (RVC) Instruction Formats									
R							CR	15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0								
funct7							CI	funct4	rd/rs1	rs2	op					
imm[11:0]							CSS	funct3	imm	rd/rs1	imm	op				
imm[11:5]							CIW	funct3	imm	rs2	op					
imm[12] imm[10:5]							CL	funct3	imm	rd'	op					
imm[31:12]							CS	funct3	imm	rs1'	imm	rd'	op			
imm[20] imm[10:1] imm[11] imm[19:12]							CB	funct3	offset	rs1'	offset	op				
							CJ	funct3	jump target		op					

RISC-V Integer Base (RV32I/64I/128I), privileged, and optional compressed extension (RVC). Registers x1-x31 and the pc are 32 bits wide in RV32I, 64 in RV64I, and 128 in RV128I (x0=0). RV64I/128I add 10 instructions for the wider formats. The RVI base of <50 classic integer RISC instructions is required. Every 16-bit RVC instruction matches an existing 32-bit RVI instruction. See risc.org.

Optional Multiply-Divide Instruction Extension: RVM					
Category	Name	Fmt	RV32M (Multiply-Divide)	+RV{64,128}	
Multiply	MULTiply	R	MUL rd,rs1,rs2	MUL{W D} rd,rs1,rs2	
	MULTiply upper Half	R	MULH rd,rs1,rs2		
	MULTiply Half Sign/Uns	R	MULHSU rd,rs1,rs2		
	MULTiply upper Half Uns	R	MULHU rd,rs1,rs2		
Divide	DIVide	R	DIV rd,rs1,rs2	DIV{W D} rd,rs1,rs2	
	DIVide Unsigned	R	DIVU rd,rs1,rs2		
Remainder	REMAinder	R	REM rd,rs1,rs2	REM{W D} rd,rs1,rs2	
	REMAinder Unsigned	R	REMU rd,rs1,rs2	REMU{W D} rd,rs1,rs2	

Optional Atomic Instruction Extension: RVA					
Category	Name	Fmt	RV32A (Atomic)	+RV{64,128}	
Load	Load Reserved	R	LR.W rd,rs1	LR.{D Q} rd,rs1	
Store	Store Conditional	R	SC.W rd,rs1,rs2	SC.{D Q} rd,rs1,rs2	
Swap	SWAP	R	AMOSWAP.W rd,rs1,rs2	AMOSWAP.{D Q} rd,rs1,rs2	
Add	ADD	R	AMOADD.W rd,rs1,rs2	AMOADD.{D Q} rd,rs1,rs2	
Logical	XOR	R	AMOXOR.W rd,rs1,rs2	AMOXOR.{D Q} rd,rs1,rs2	
	AND	R	AMOAND.W rd,rs1,rs2	AMOAND.{D Q} rd,rs1,rs2	
	OR	R	AMOOR.W rd,rs1,rs2	AMOOR.{D Q} rd,rs1,rs2	
Min/Max	MINimum	R	AMOMIN.W rd,rs1,rs2	AMOMIN.{D Q} rd,rs1,rs2	
	MAXimum	R	AMOMAX.W rd,rs1,rs2	AMOMAX.{D Q} rd,rs1,rs2	
	MINimum Unsigned	R	AMOMINU.W rd,rs1,rs2	AMOMINU.{D Q} rd,rs1,rs2	
	MAXimum Unsigned	R	AMOMAXU.W rd,rs1,rs2	AMOMAXU.{D Q} rd,rs1,rs2	

Three Optional Floating-Point Instruction Extensions: RVF, RVD, & RVQ					
Category	Name	Fmt	RV32{F D Q} (HP/SP,DP,QP Fl Pt)	+RV{64,128}	
Move	Move from Integer	R	FMV.{H S}.X rd,rs1	FMV.{D Q}.X rd,rs1	
	Move to Integer	R	FMV.X.{H S} rd,rs1	FMV.X.{D Q} rd,rs1	
Convert	Convert from Int	R	FCVT.{H S D Q}.W rd,rs1	FCVT.{H S D Q}.{L T} rd,rs1	
	Convert from Int Unsigned	R	FCVT.{H S D Q}.WU rd,rs1	FCVT.{H S D Q}.{L T}U rd,rs1	
	Convert to Int	R	FCVT.W.{H S D Q} rd,rs1	FCVT.{L T}.{H S D Q} rd,rs1	
	Convert to Int Unsigned	R	FCVT.WU.{H S D Q} rd,rs1	FCVT.{L T}U.{H S D Q} rd,rs1	

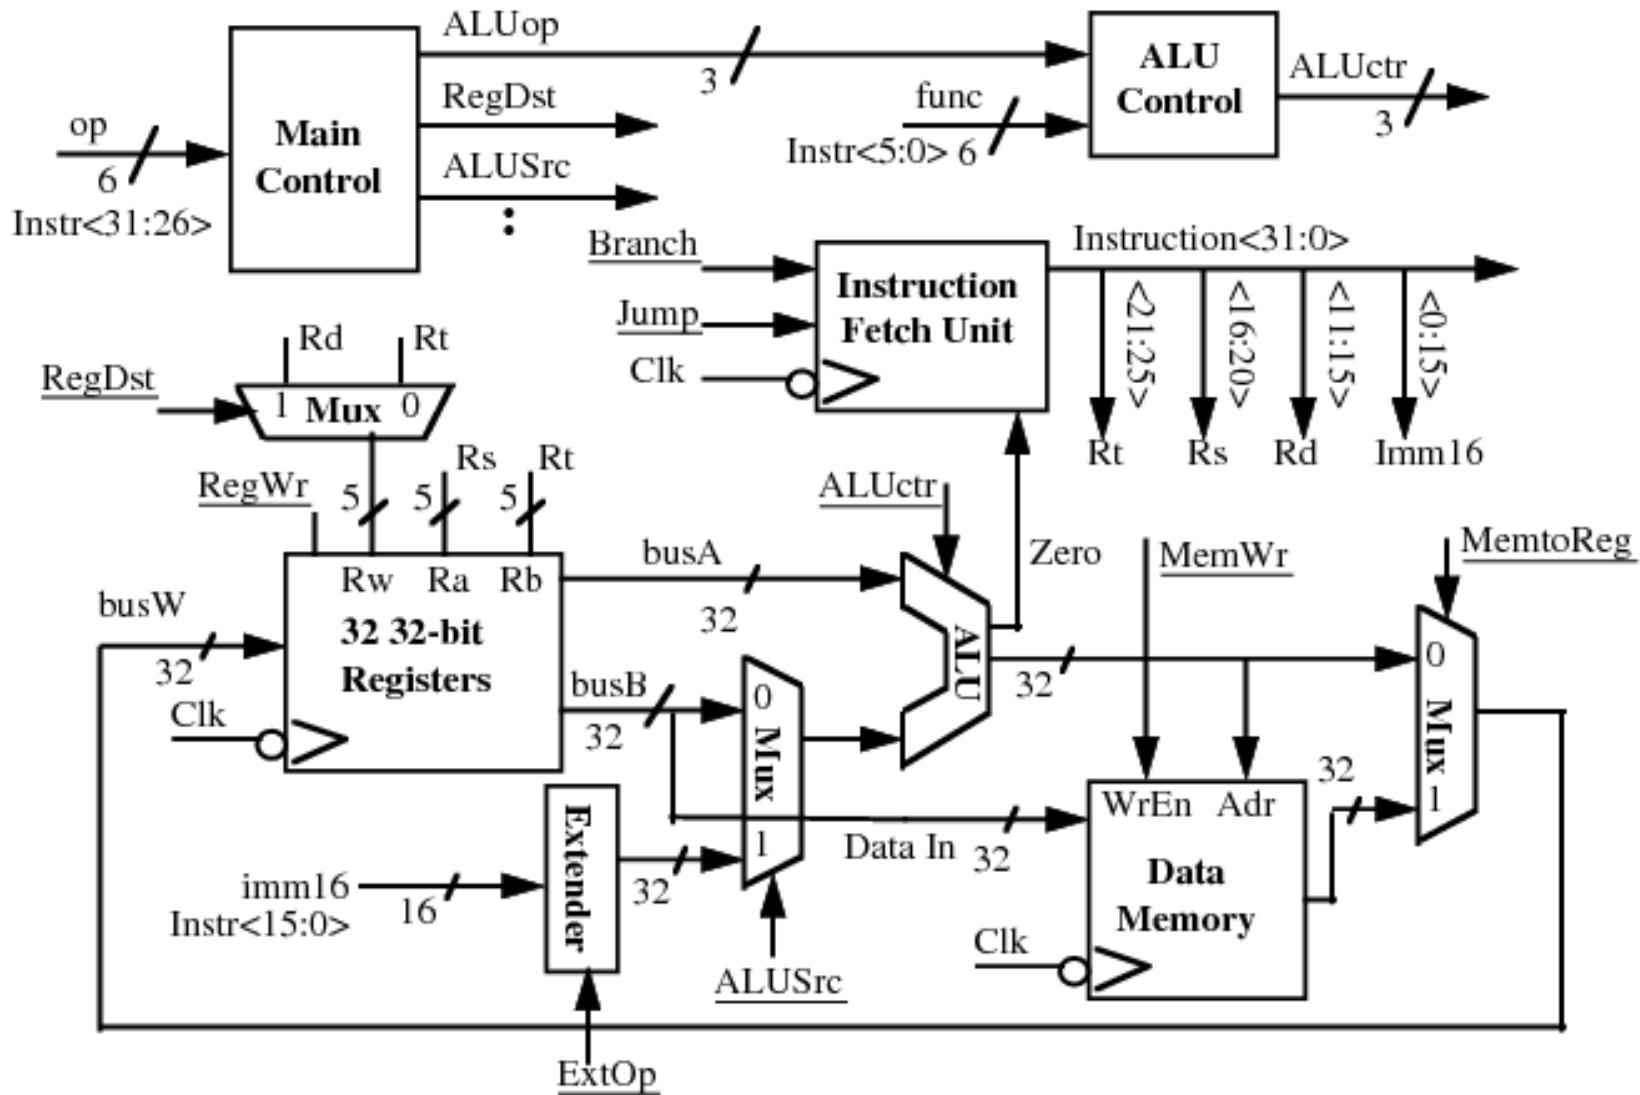
RISC-V Calling Convention					
Register	ABI Name	Saver	Description		
x0	zero	---	Hard-wired zero		
x1	ra	Caller	Return address		
x2	sp	Callee	Stack pointer		
x3	gp	---	Global pointer		
x4	tp	---	Thread pointer		
x5-x7	t0-2	Caller	Temporaries		
x8	s0/fp	Callee	Saved register/frame pointer		
x9	s1	Callee	Saved register		
x10-x11	a0-1	Caller	Function arguments/return values		
x12-x17	a2-7	Caller	Function arguments		
x18-27	s2-11	Callee	Saved registers		
x28-31	t3-t6	Caller	Temporaries		
f0-7	ft0-7	Caller	FP temporaries		
f8-9	fs0-1	Callee	FP saved registers		
f10-11	fa0-1	Caller	FP arguments/return values		
f12-17	fa2-7	Caller	FP arguments		
f18-27	fs2-11	Callee	FP saved registers		
f28-31	ft8-11	Caller	FP temporaries		

RISC-V calling convention and five optional extensions: 10 multiply-divide instructions (RVM); 11 optional atomic instructions (RVA); and 25 floating-point instructions each for single-, double-, and quadruple-precision (RVF, RVD, RVQ). The latter add registers f0-f31, whose width matches the widest precision, and a floating-point control and status register fcsr. Each larger address adds some instructions: 4 for RVM, 11 for RVA, and 6 each for RVF/D/Q. Using regex notation, {} means set, so L{D|Q} is both LD and LQ. See riscv.org. (8/21/15 revision)

Acknowledgment: Almost all of these slides are based on Dave Patterson's CS152 Lecture Slides at UC, Berkeley.

COMPUTER SYSTEMS ORGANIZATION

A Single Cycle Processor



Drawbacks of Single Cycle Processor

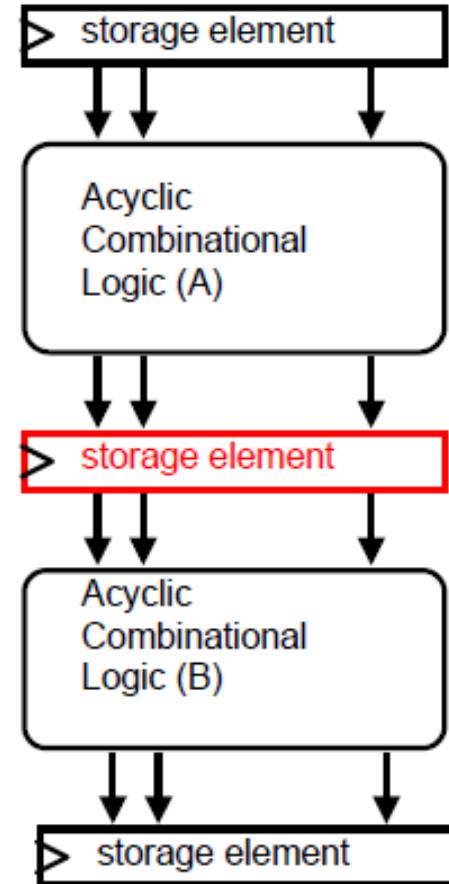
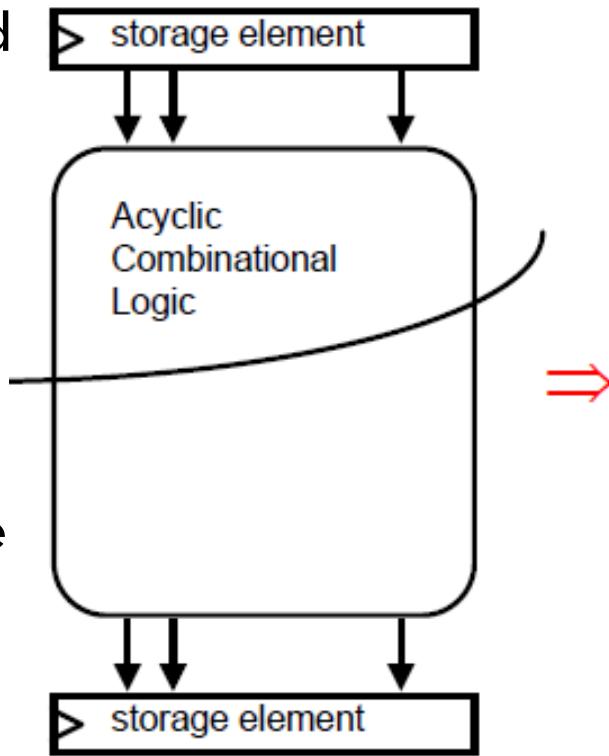
- ❑ Long cycle time: Cycle time must be long enough for the load instruction:
 - ❑ PC's Clock -to-Q +
 - ❑ Instruction Memory Access Time +
 - ❑ Register File Access Time +
 - ❑ ALU Delay (address calculation) +
 - ❑ Data Memory Access Time +
 - ❑ Register File Setup Time
- ❑ Cycle time is much longer than needed for all other instructions.
 - ❑ R-type instructions do not require data memory access
 - ❑ Jump does not require ALU operation nor data memory access

Overview of a Multiple Cycle Implementation

- The root of the single cycle processor's problems:
 - The cycle time has to be long enough for the slowest instruction
- Solution:
 - Break the instruction into smaller steps
 - Execute each step (instead of the entire instruction) in one cycle
 - Cycle time: time it takes to execute the longest step
 - Keep all the steps to have similar length
 - This is the essence of the multiple cycle processor

Reducing Cycle Time

- ❑ Cut combinational dependency graph and insert register / latch
- ❑ Do same work in two fast cycles, rather than one slow one
- ❑ May be able to short-circuit path and remove some components for some instructions!



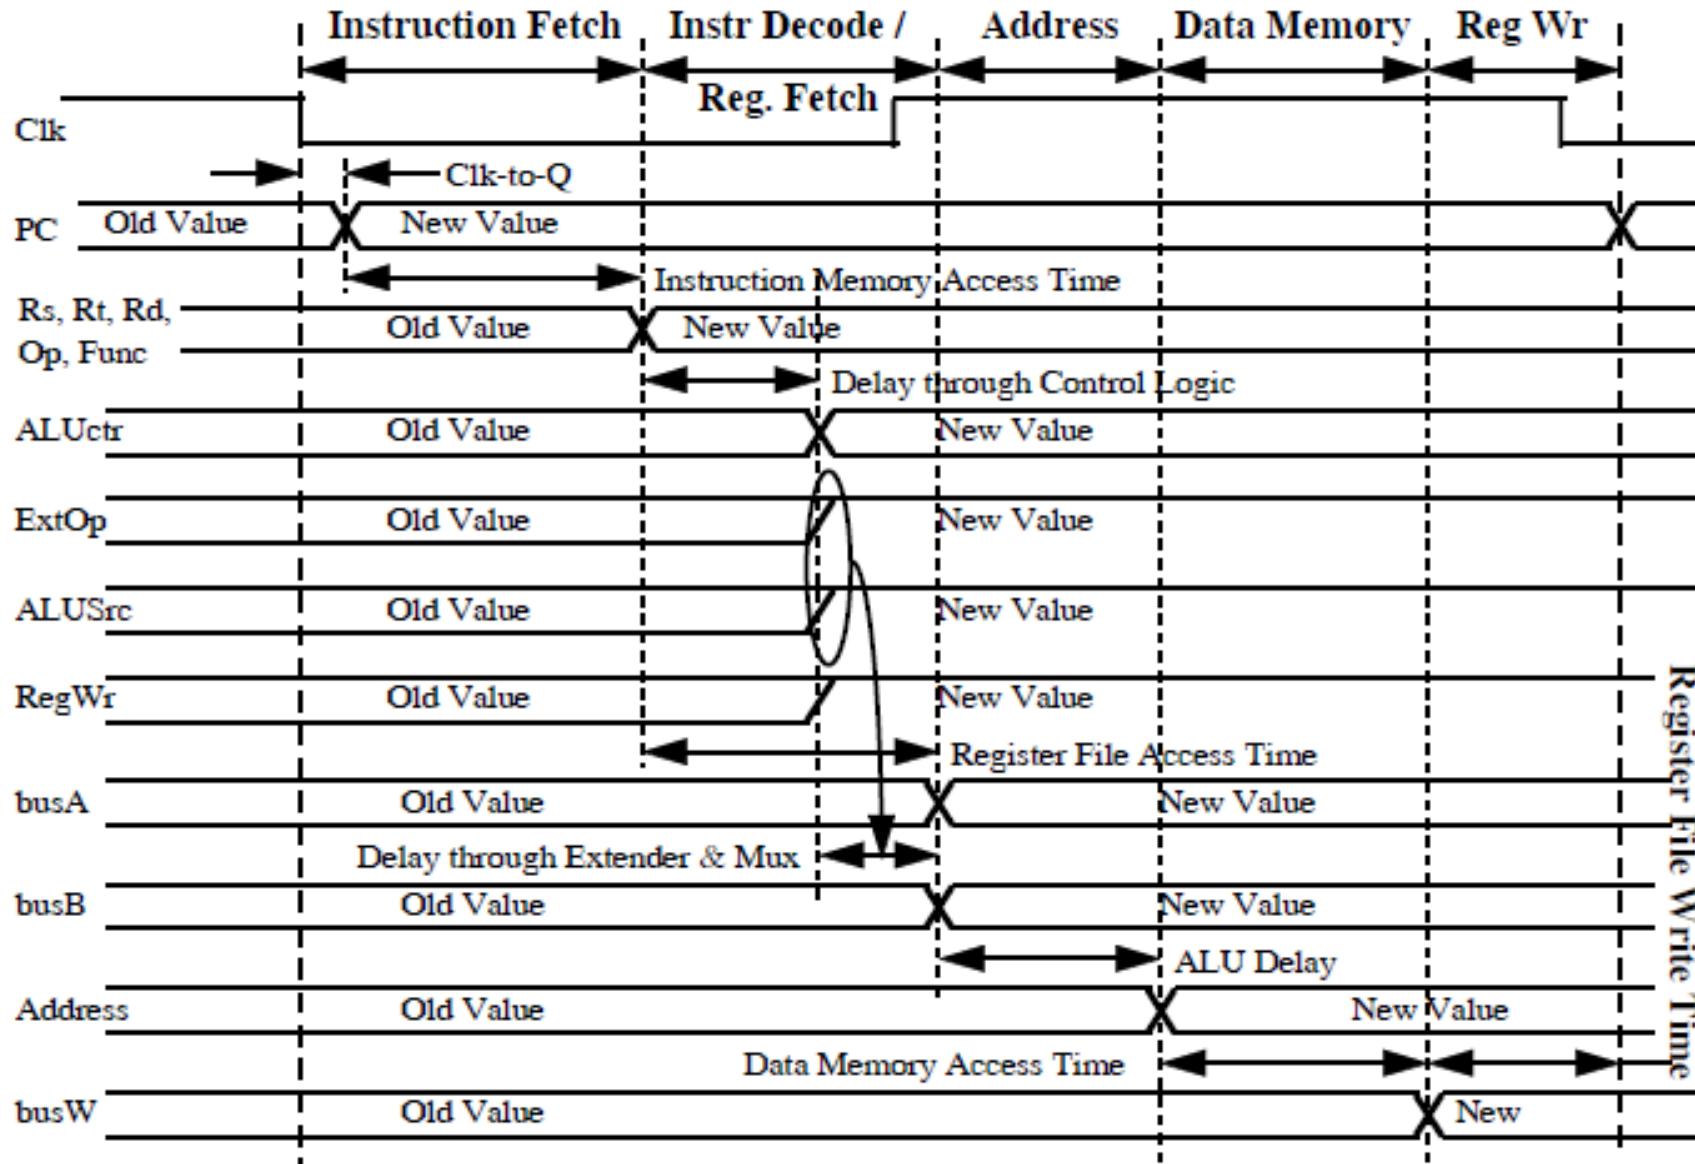
Advantages of Multiple Cycle Processor

- ❑ Cycle time is much shorter
- ❑ Different instructions take different number of cycles to complete
 - ❑ Load takes five cycles
 - ❑ Jump only takes three cycles
- ❑ Allows a functional unit to be used more than once per instruction

The Big Picture: Performance Perspective

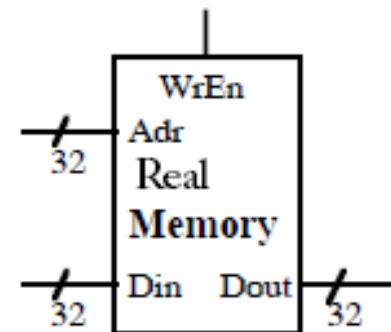
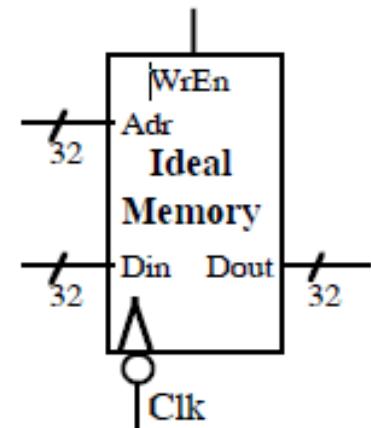
- ❑ Performance of a machine was determined by
 - ❑ Instruction Count
 - ❑ Clock cycle Time
 - ❑ Clock cycles per instruction
- ❑ Processor Design (data path and control) will determine
 - ❑ Clock cycle time
 - ❑ Clock cycles per instruction

The Five Steps of a Load Instruction



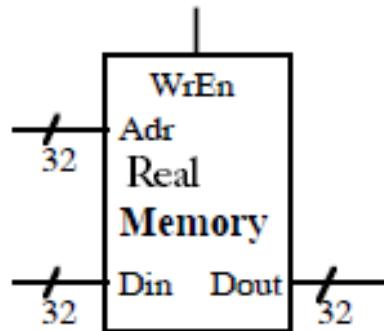
Memory Write Timing: Ideal Vs Reality

- ❑ In the Single Cycle Processor memory module is simplified
 - ❑ Write happens at the clock tick
 - ❑ Address, data, and write enable must be stable one “set-up” time before the clock tick
- ❑ In real life memory module has no clock input
 - ❑ The write path is a combinational logic delay path:
 - ❑ Write enable goes to 1 and Din settles down
 - ❑ Memory write access delay
 - ❑ Din is written into mem[address]
- ❑ **Important:** Address and Data must be stable BEFORE Write Enable goes to 1



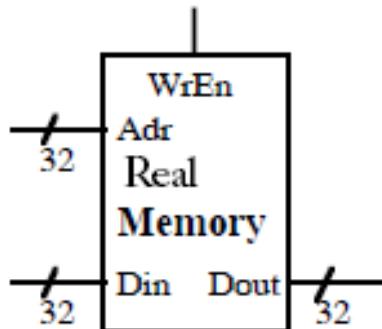
Race Condition Between Address and Write Enable

- ❑ The “real” (no clock input) memory may not work reliably in the single cycle processor because:
 - ❑ We cannot guarantee Address will be stable BEFORE WrEn = 1
 - ❑ There is a race between Adr and WrEn



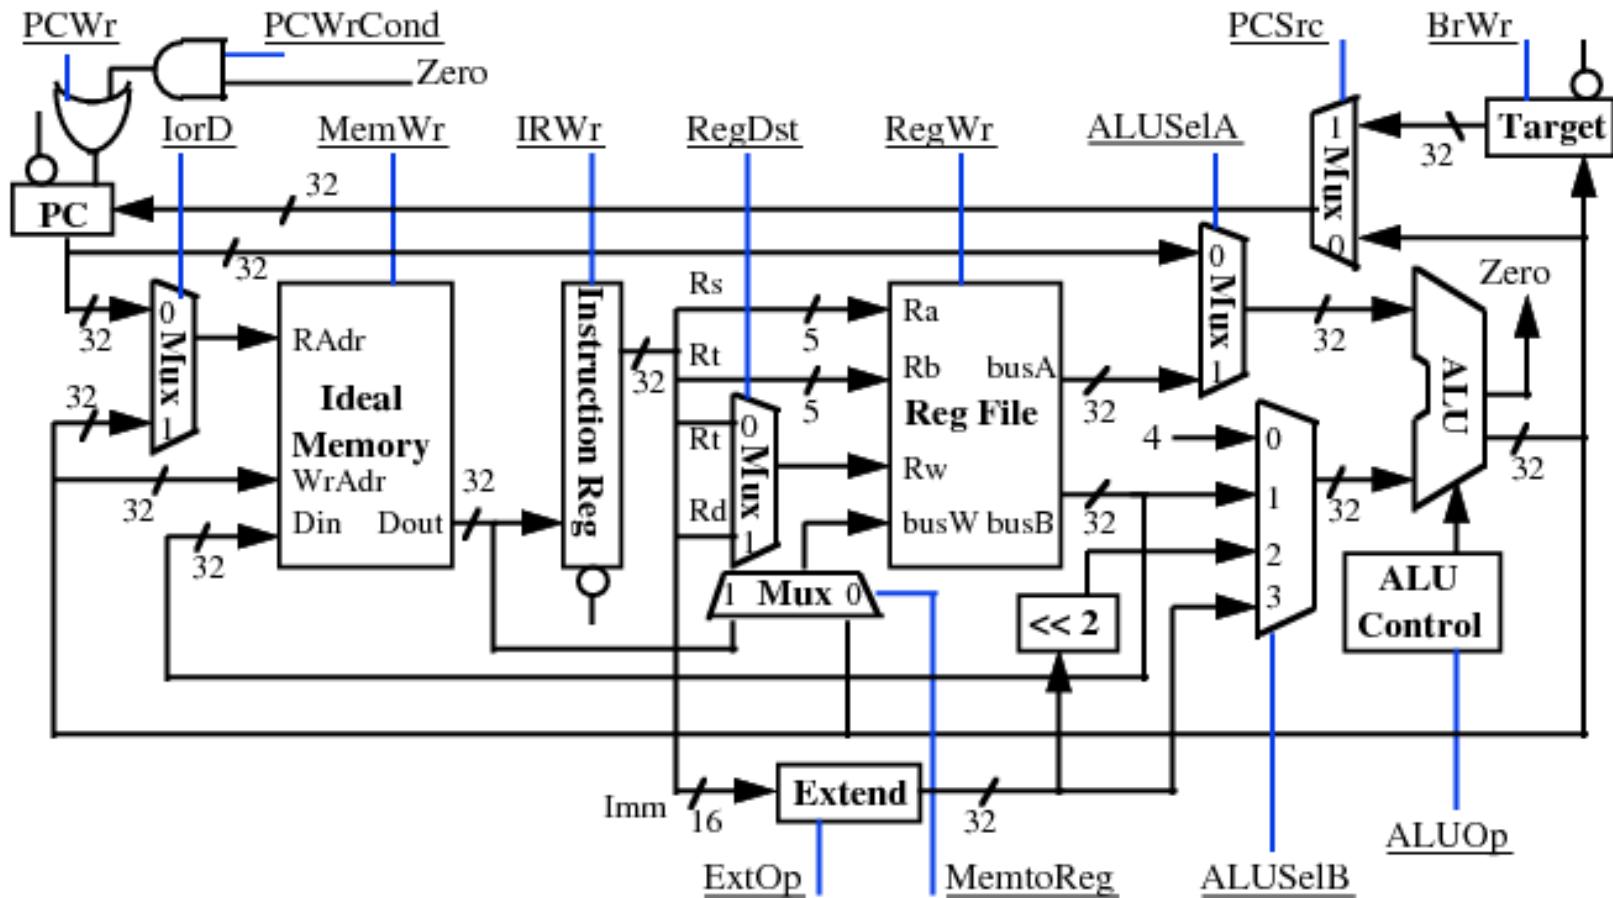
How to Avoid this Race Condition?

- ❑ Solution for the multiple cycle implementation:
 - ❑ Make sure Address is stable by the end of Cycle N
 - ❑ Assert Write Enable signal ONE cycle later at Cycle (N + 1)
 - ❑ Address cannot change until Write Enable is deasserted.

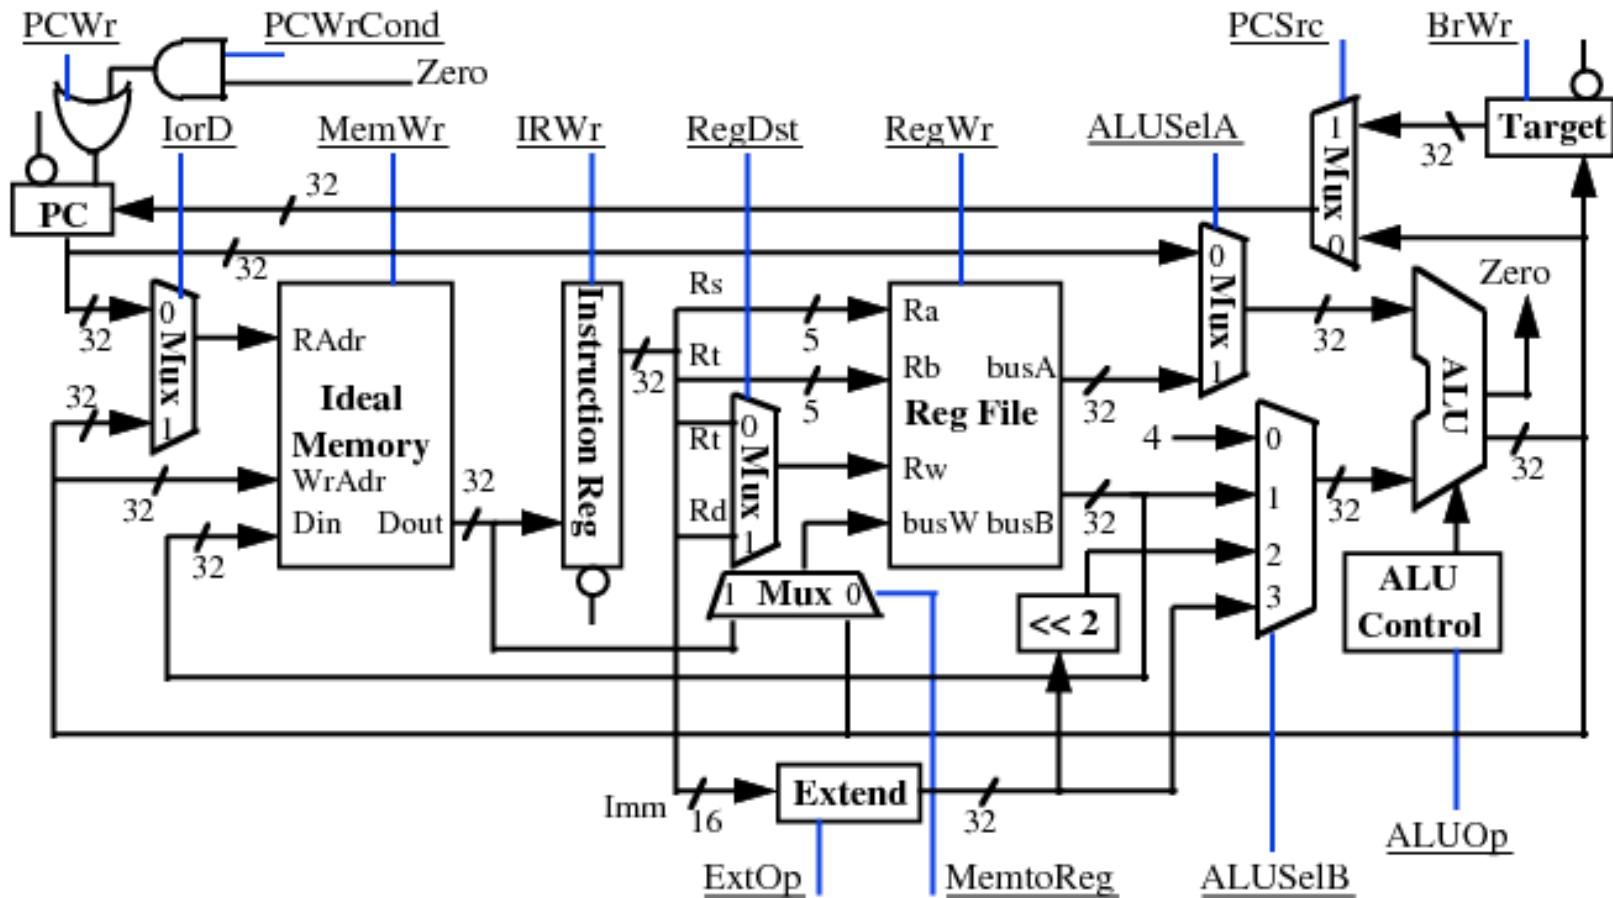


We can make use the same logic if we are using unclocked Register File to build our multiple cycle processor.

Multiple Cycle Data Path



Multiple Cycle Data Path



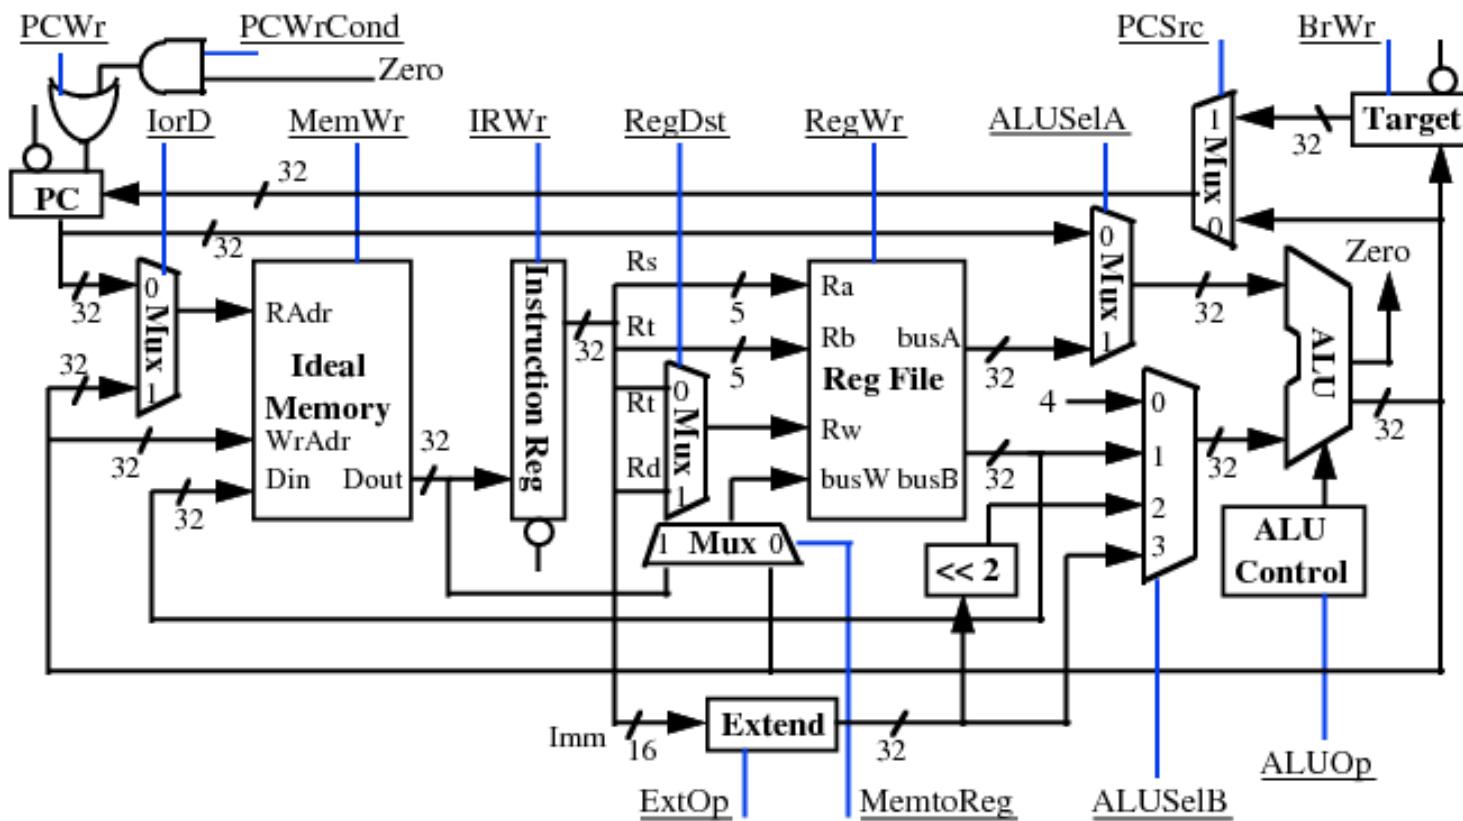
Instruction Fetch Phase

- IR = mem[PC]
- PC = PC + 4

PCWr	IorD	MemWr	IrWr	RegDst	RegWr	ALUSelA
1	0	0	1	x	0	0
PCWrCond	PCSrc	BrWr	ExtOp	MemtoReg	ALUOp	ALUSelB
x	0	0	x	x	ADD	0

Ifetch

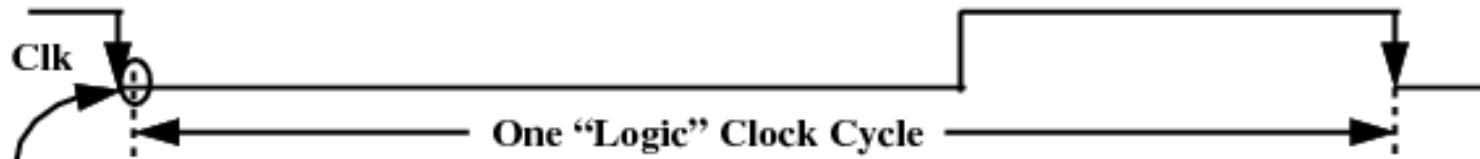
ALUOp=Add
 1: PCWr, IRWr
 x: PCWrCond
 RegDst, Mem2R
 Others: 0s



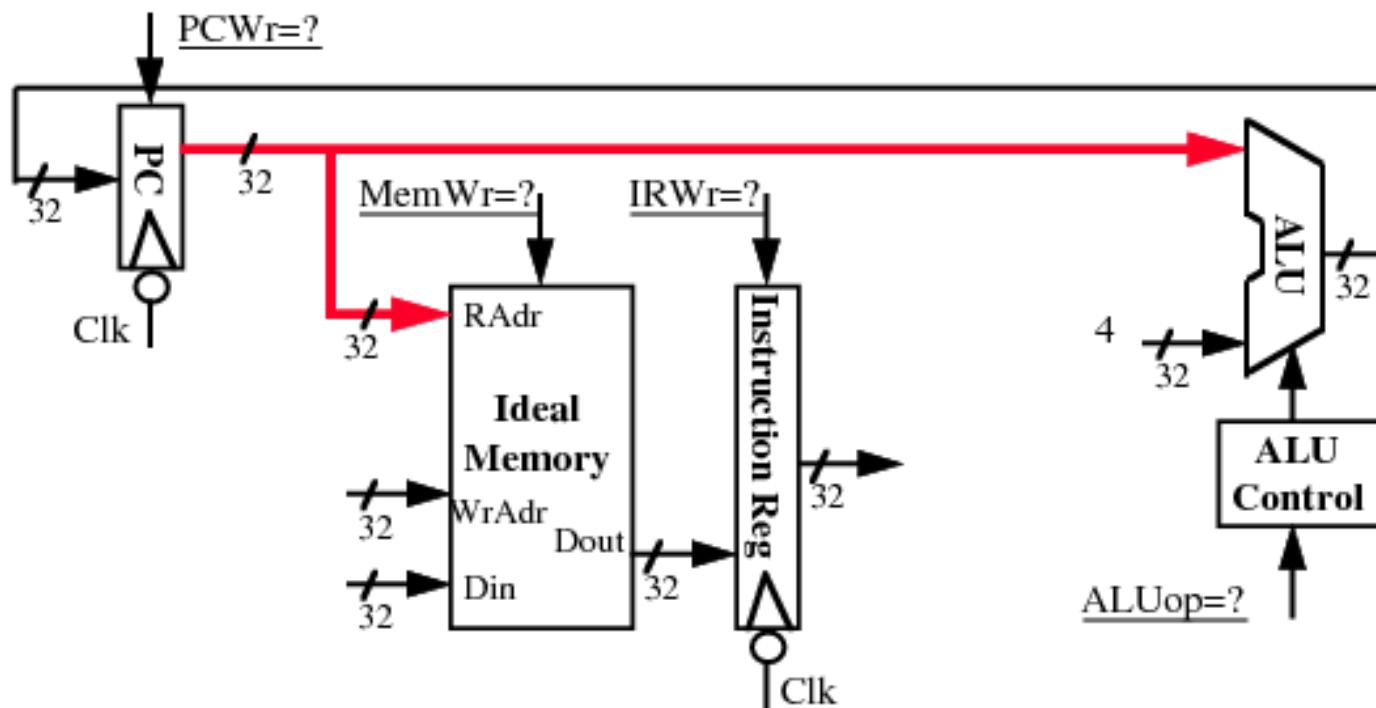
Instruction Fetch Cycle: In the Beginning

- Every cycle begins right AFTER the clock tick:

- $\text{mem}[\text{PC}] \quad \text{PC}\langle 31:0 \rangle + 4$

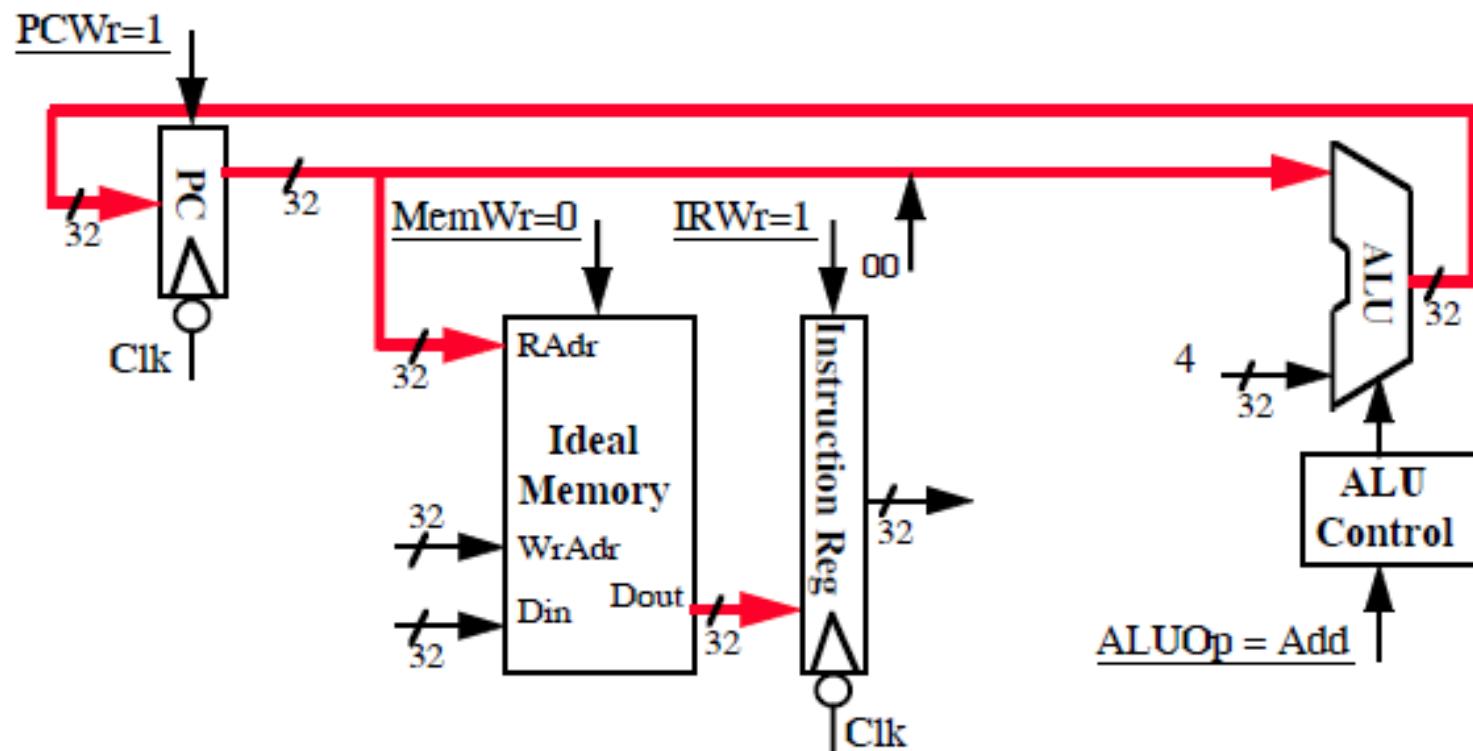
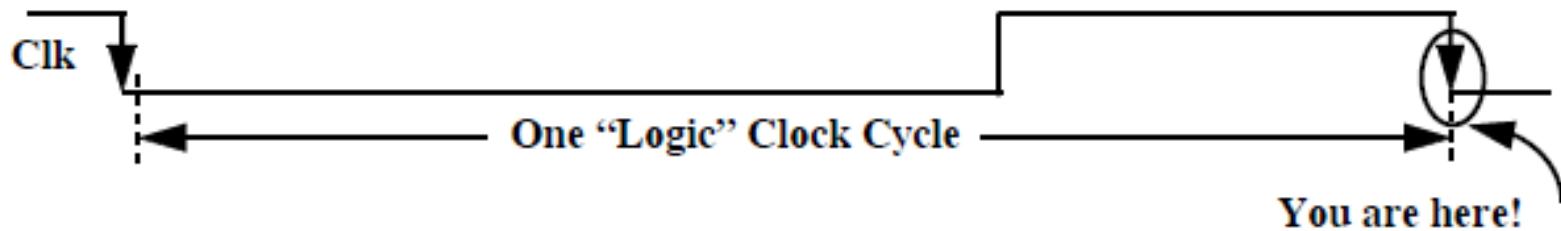


You are here!

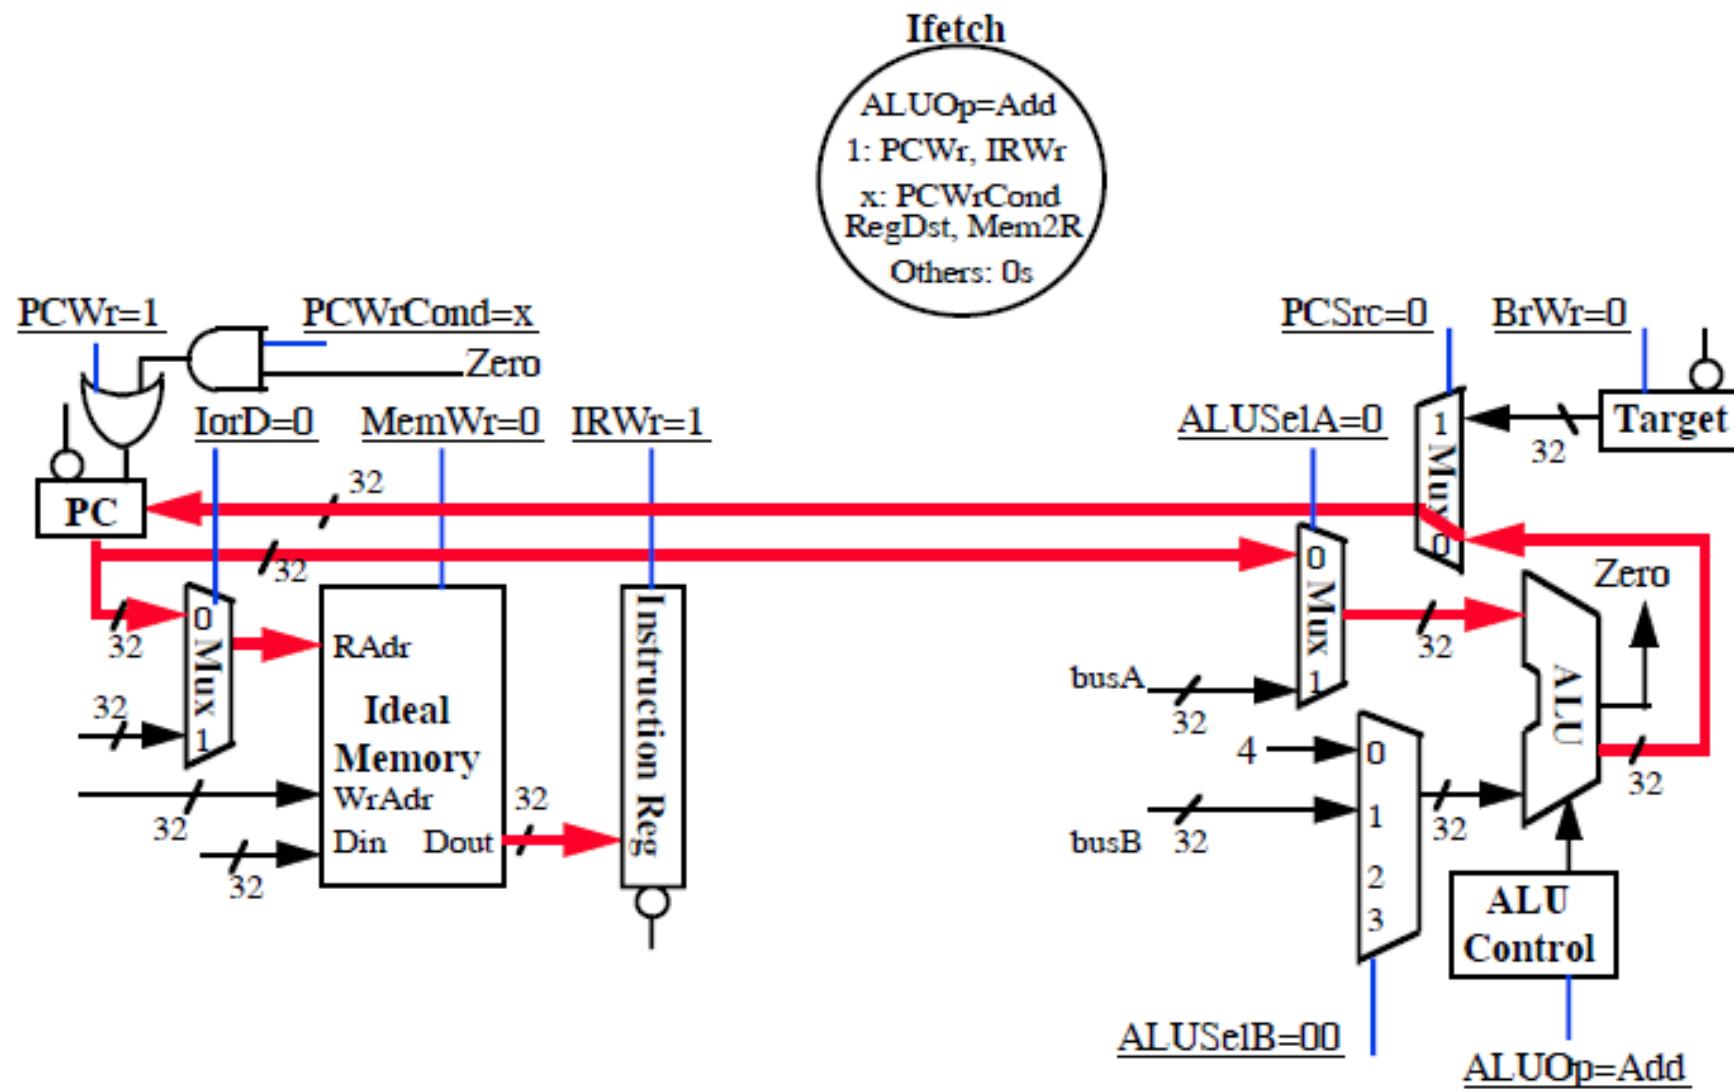


Instruction Fetch Cycle: The End

- Every cycle ends AT the next clock tick (storage element updates):
 - $\text{IR} \leftarrow \text{mem}[\text{PC}]$ $\text{PC}_{31:0} \leftarrow \text{PC}_{31:0} + 4$

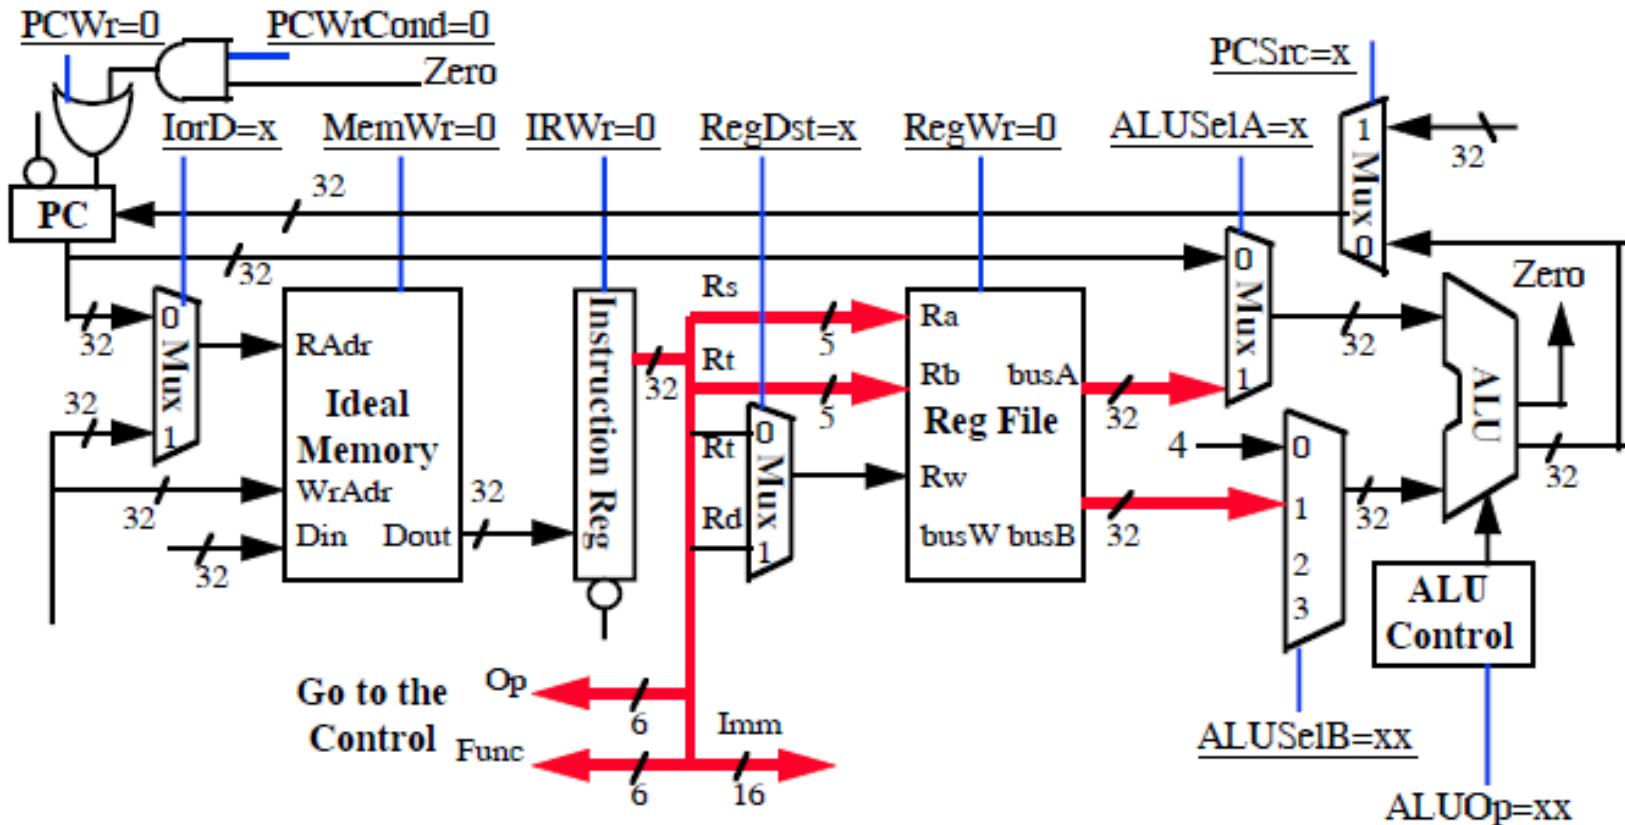


Instruction Fetch Cycle: Overall Picture



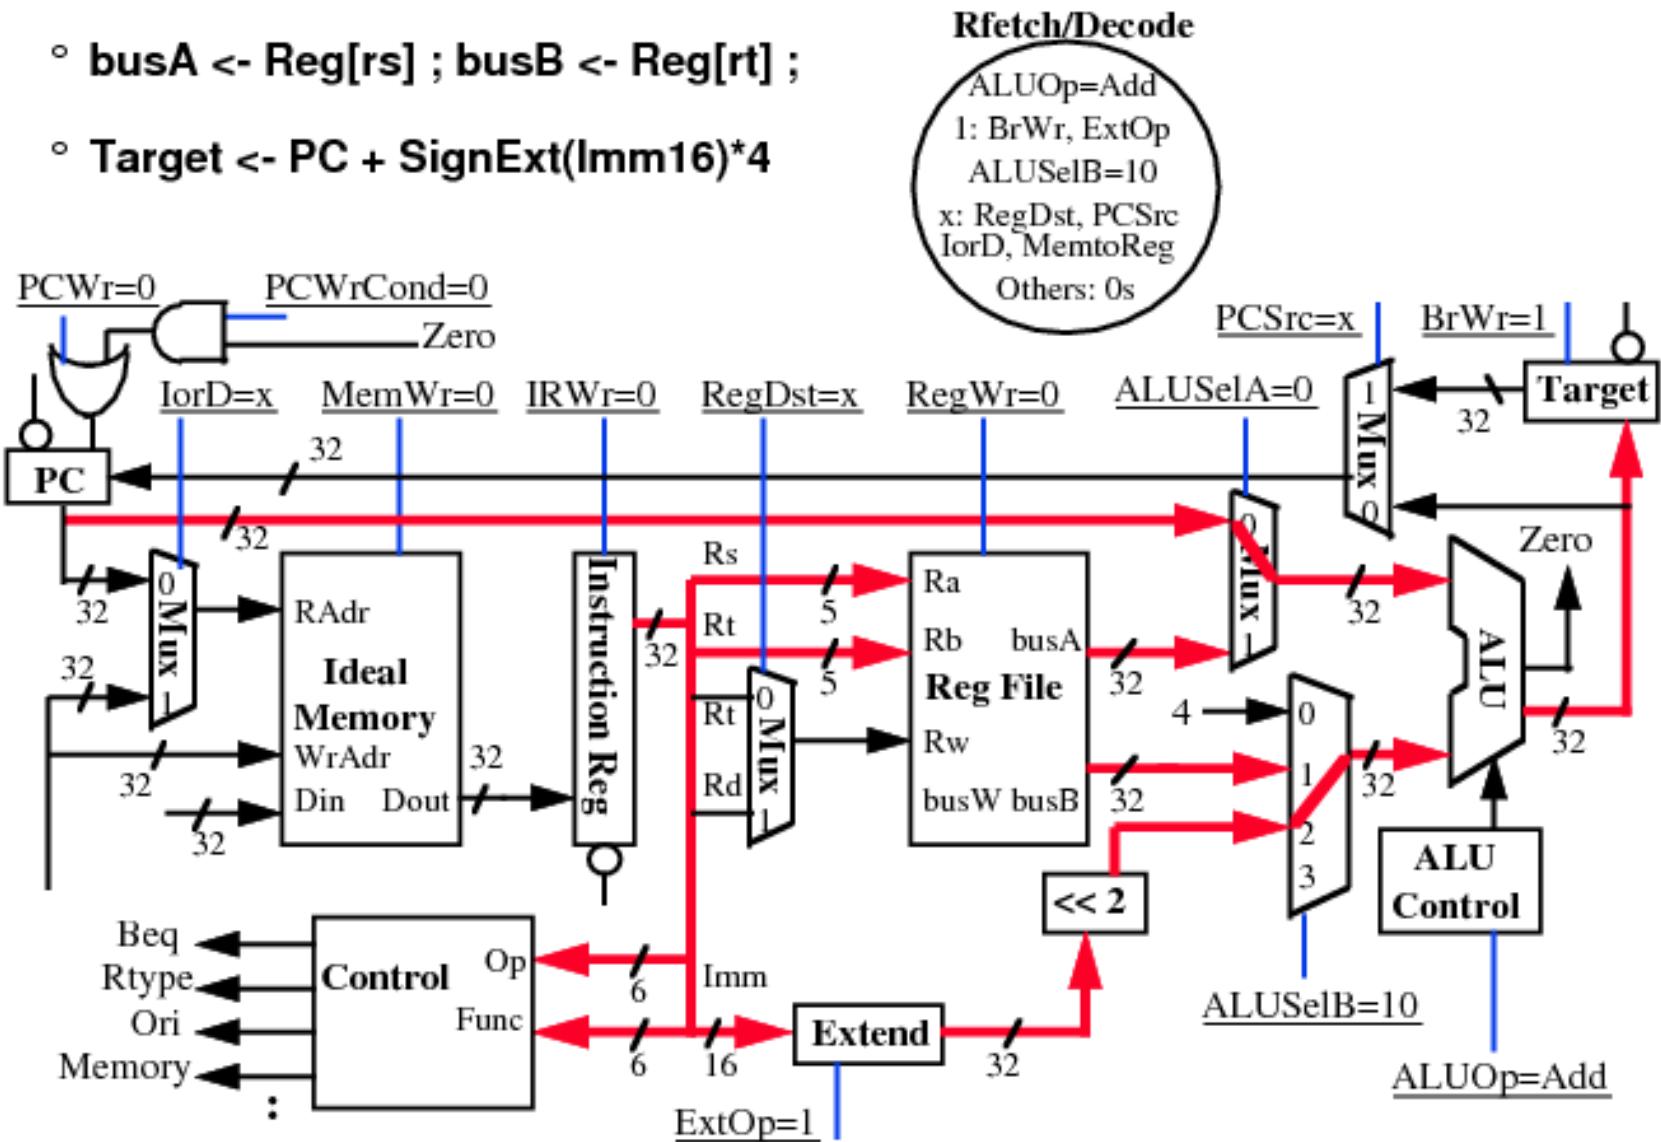
Register Fetch / Instruction Decode

- ° $\text{busA} \leftarrow \text{RegFile}[rs]$; $\text{busB} \leftarrow \text{RegFile}[rt]$;
- ° ALU is not being used: $\text{ALUctr} = \text{xx}$



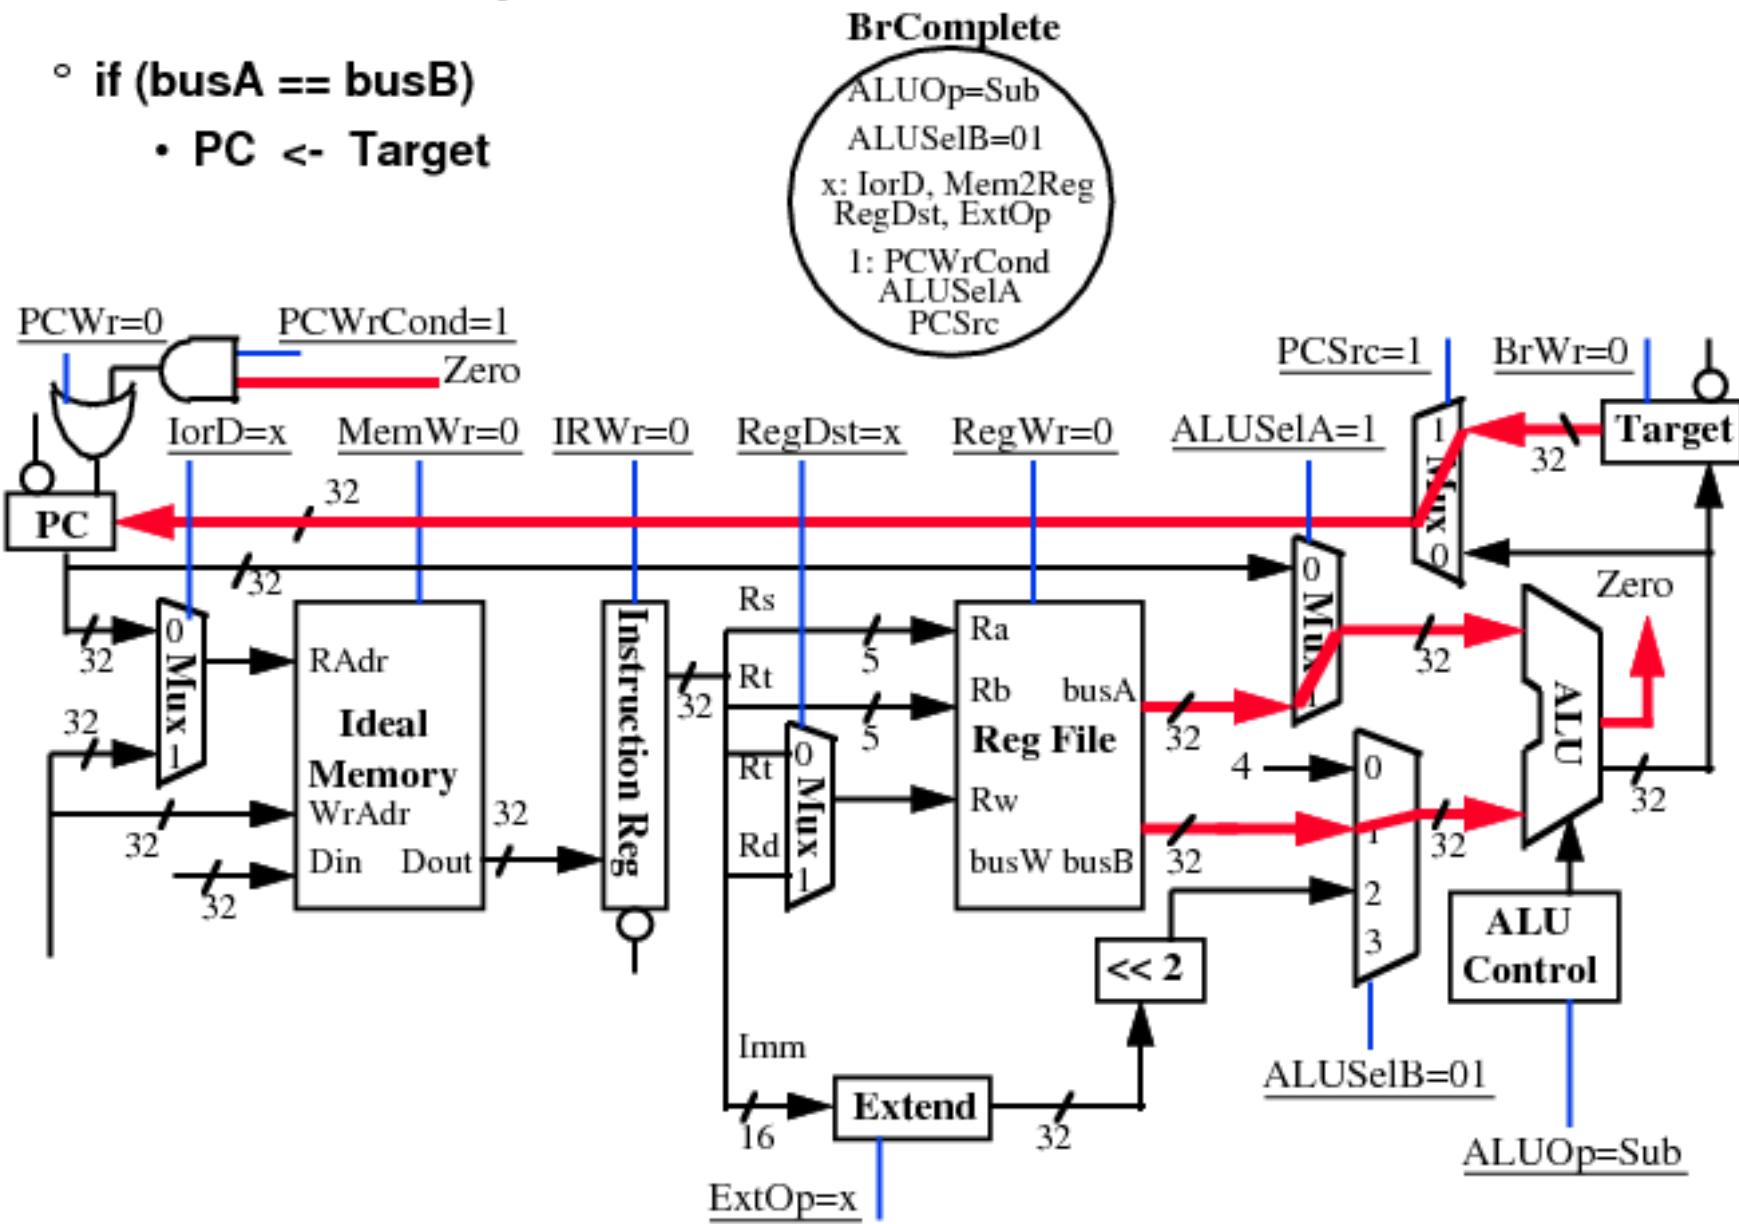
Register Fetch / Instruction Decode (Continue)

- $\text{busA} \leftarrow \text{Reg}[rs]$; $\text{busB} \leftarrow \text{Reg}[rt]$;
- Target $\leftarrow \text{PC} + \text{SignExt}(\text{Imm16})^*4$



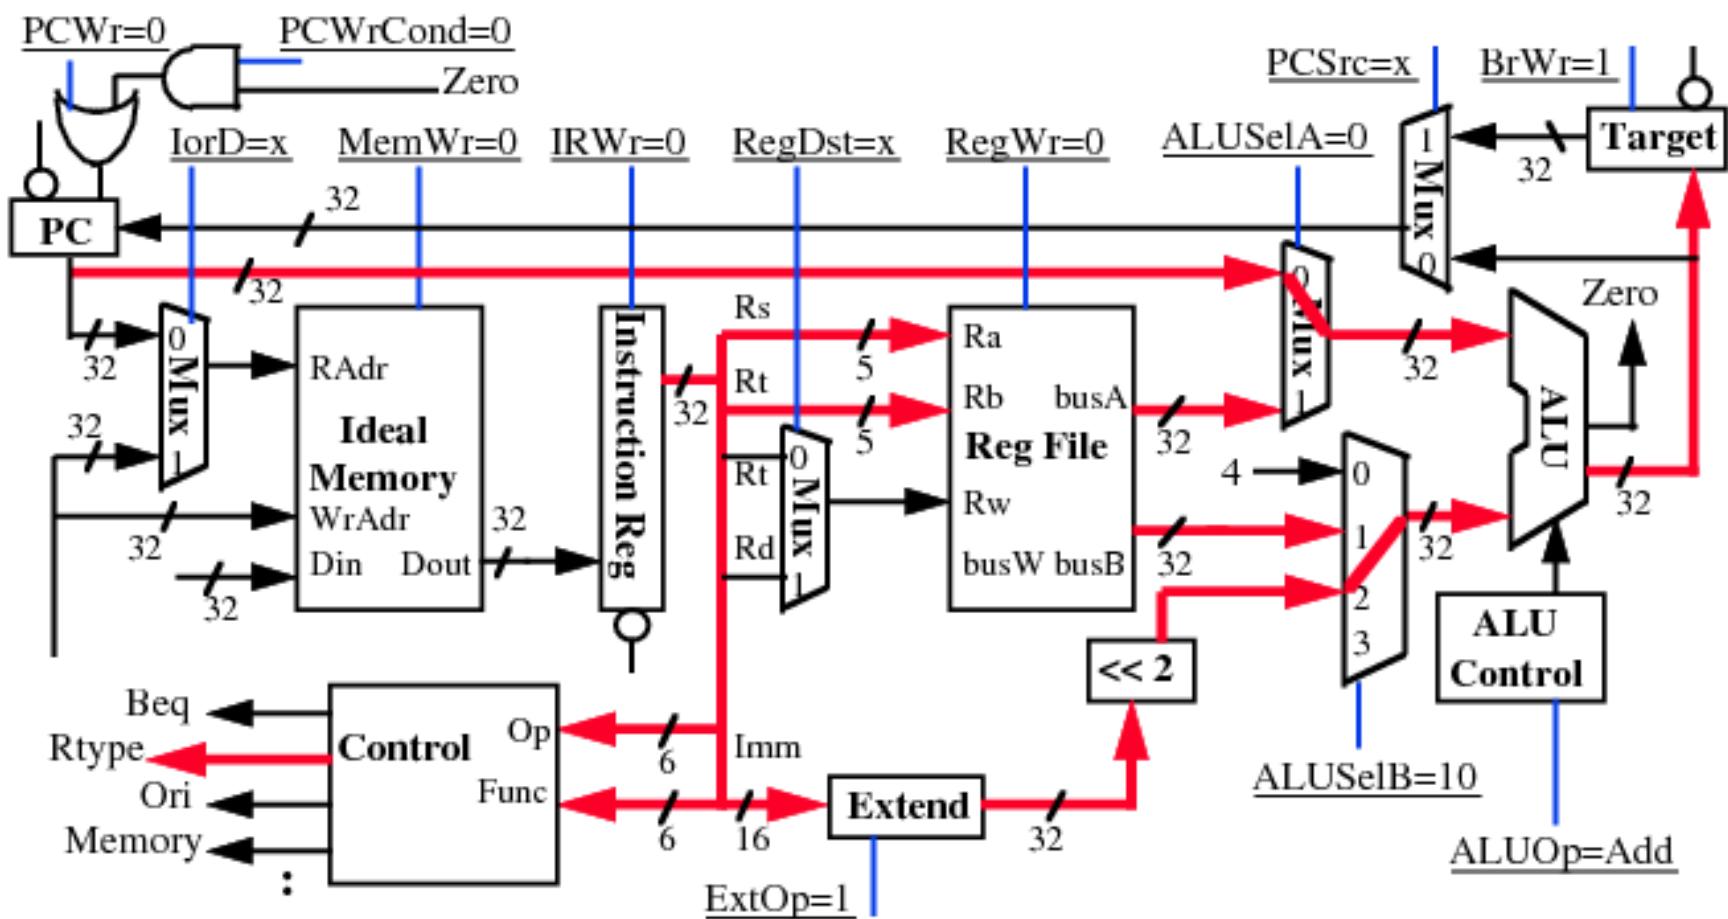
Branch Completion

- if ($\text{busA} == \text{busB}$)
 - $\text{PC} \leftarrow \text{Target}$



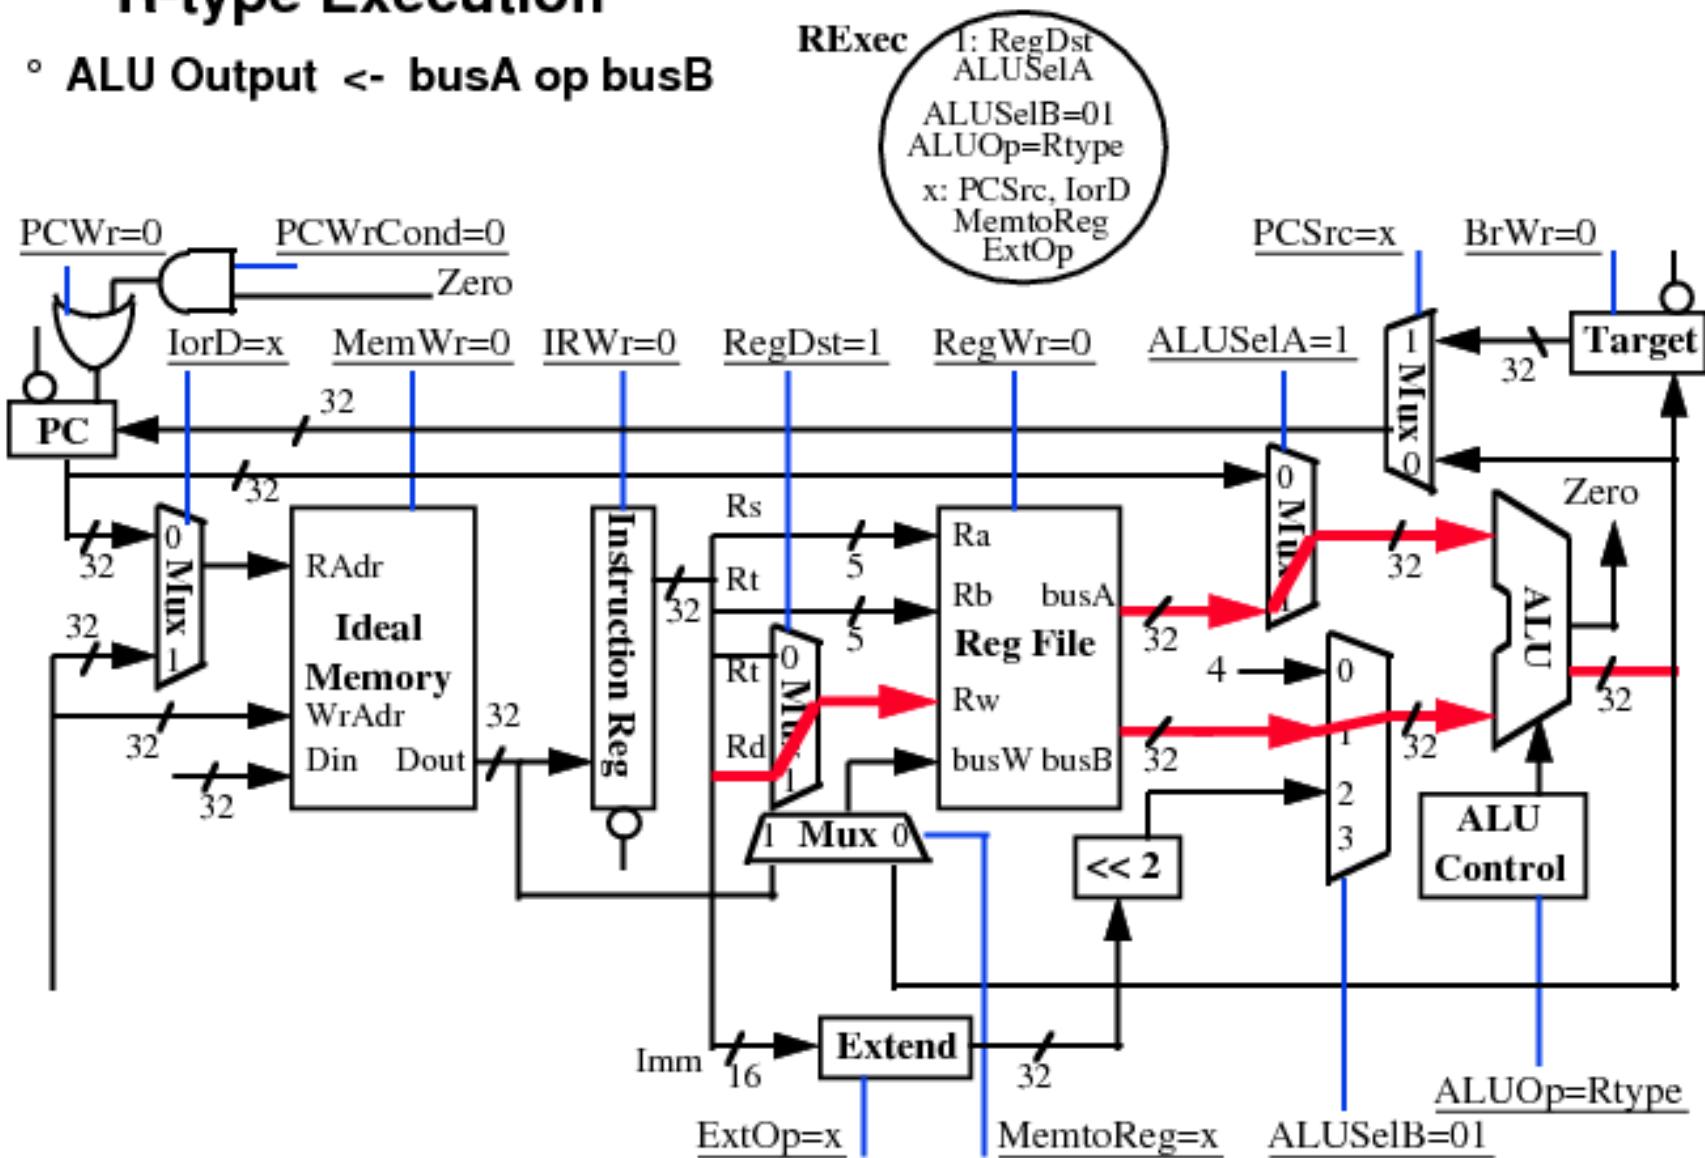
Instruction Decode: We have a R-type!

- Next Cycle: R-type Execution



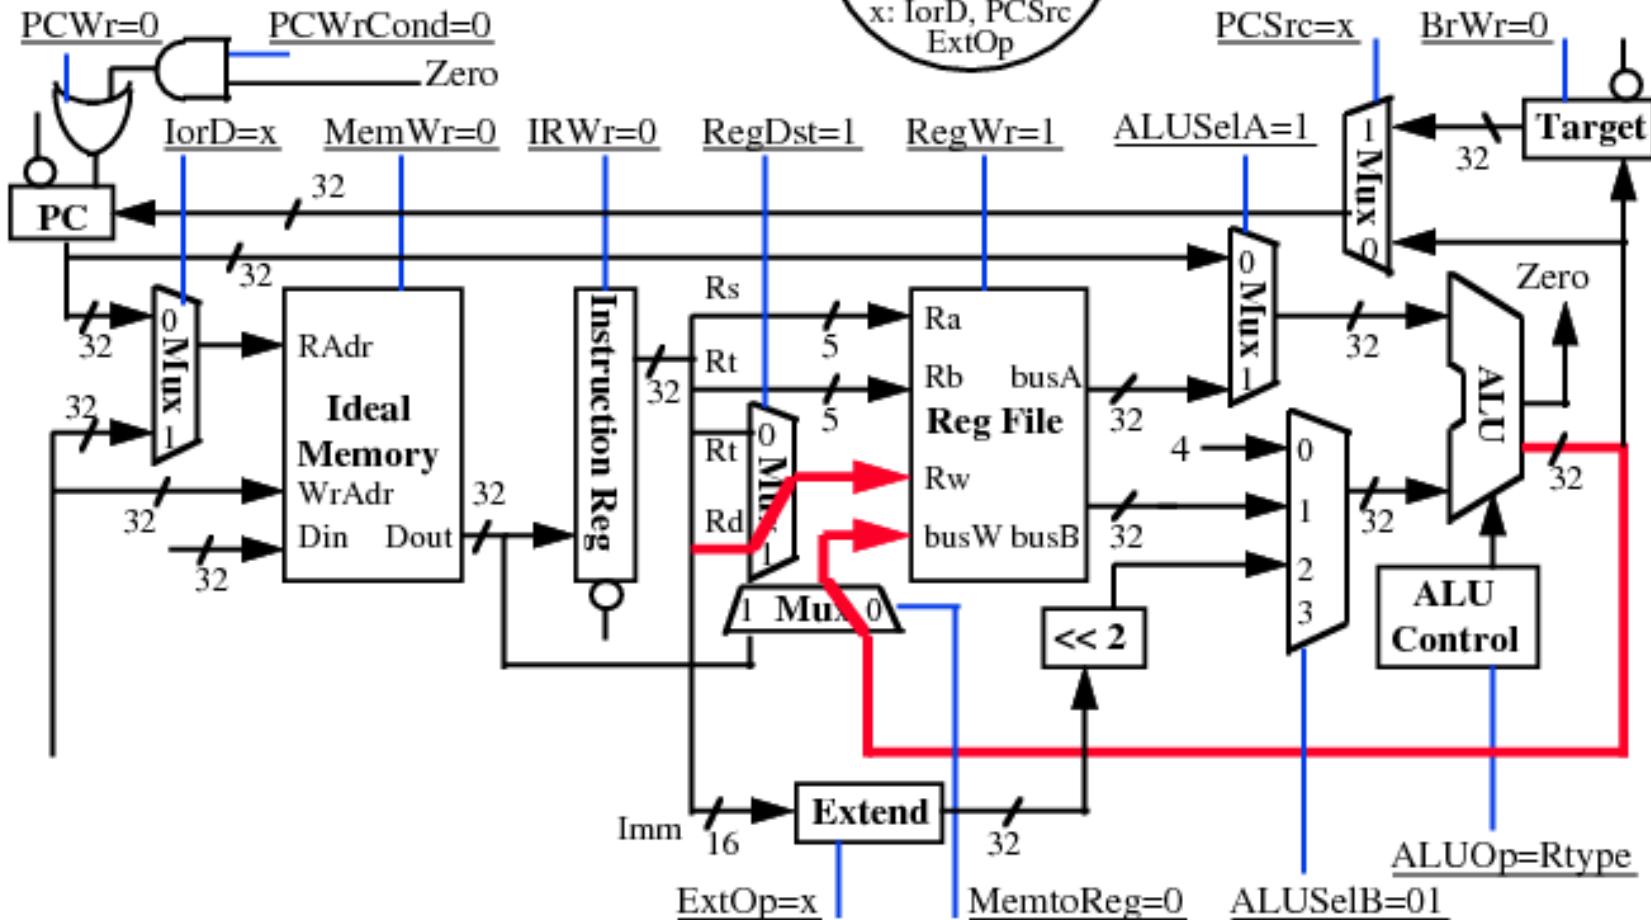
R-type Execution

- ALU Output <- busA op busB



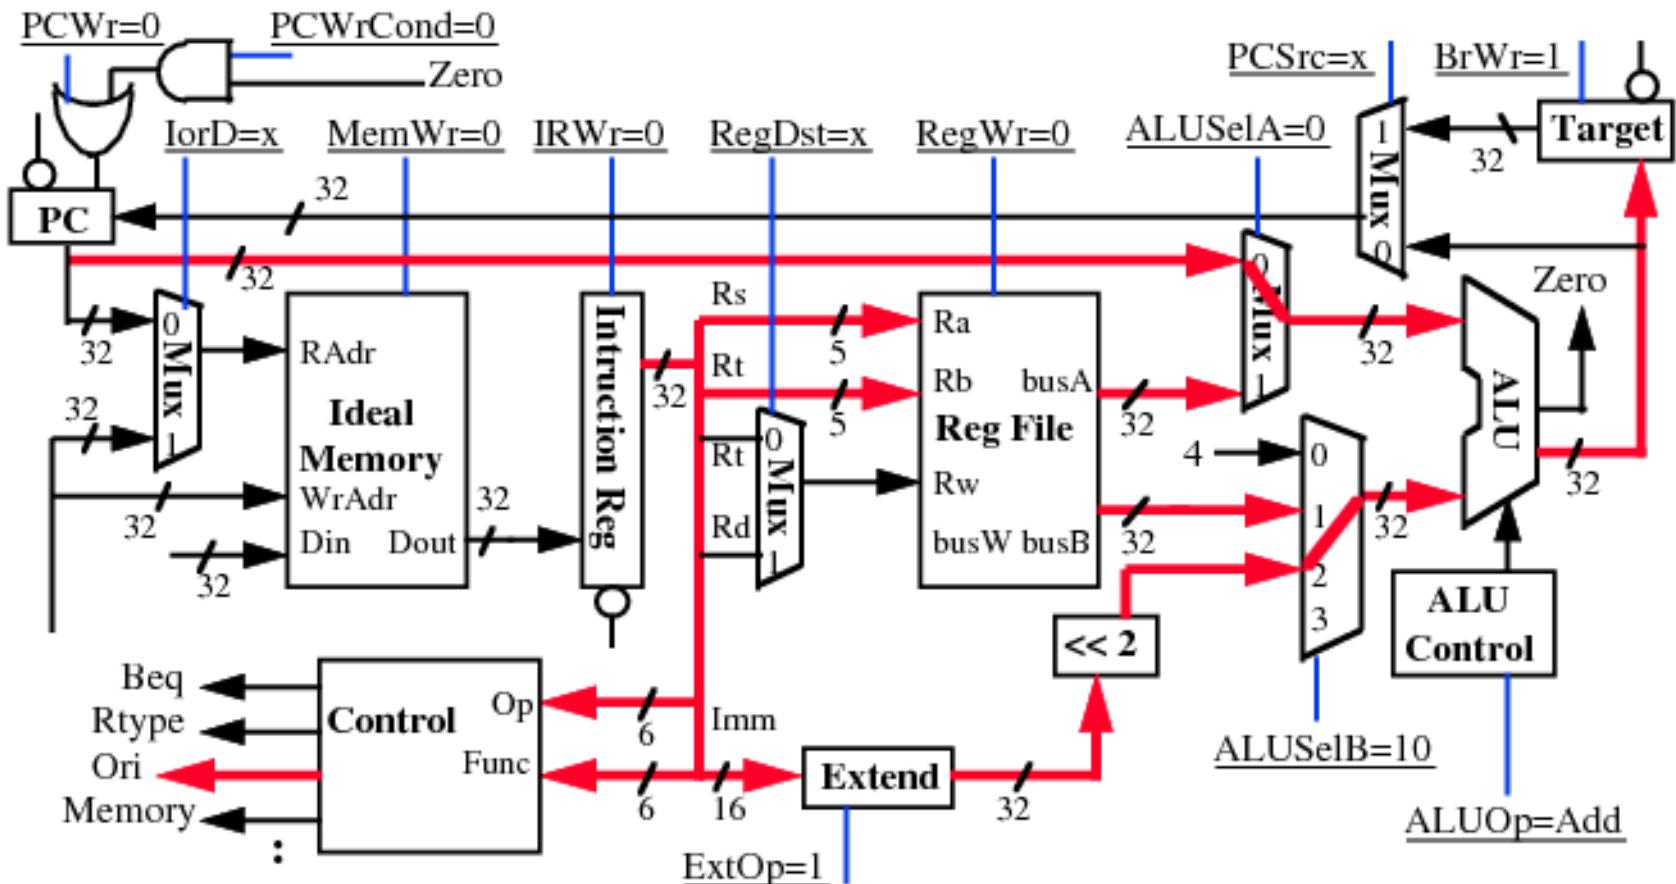
R-type Completion

- $R[rd] \leftarrow ALU\ Output$



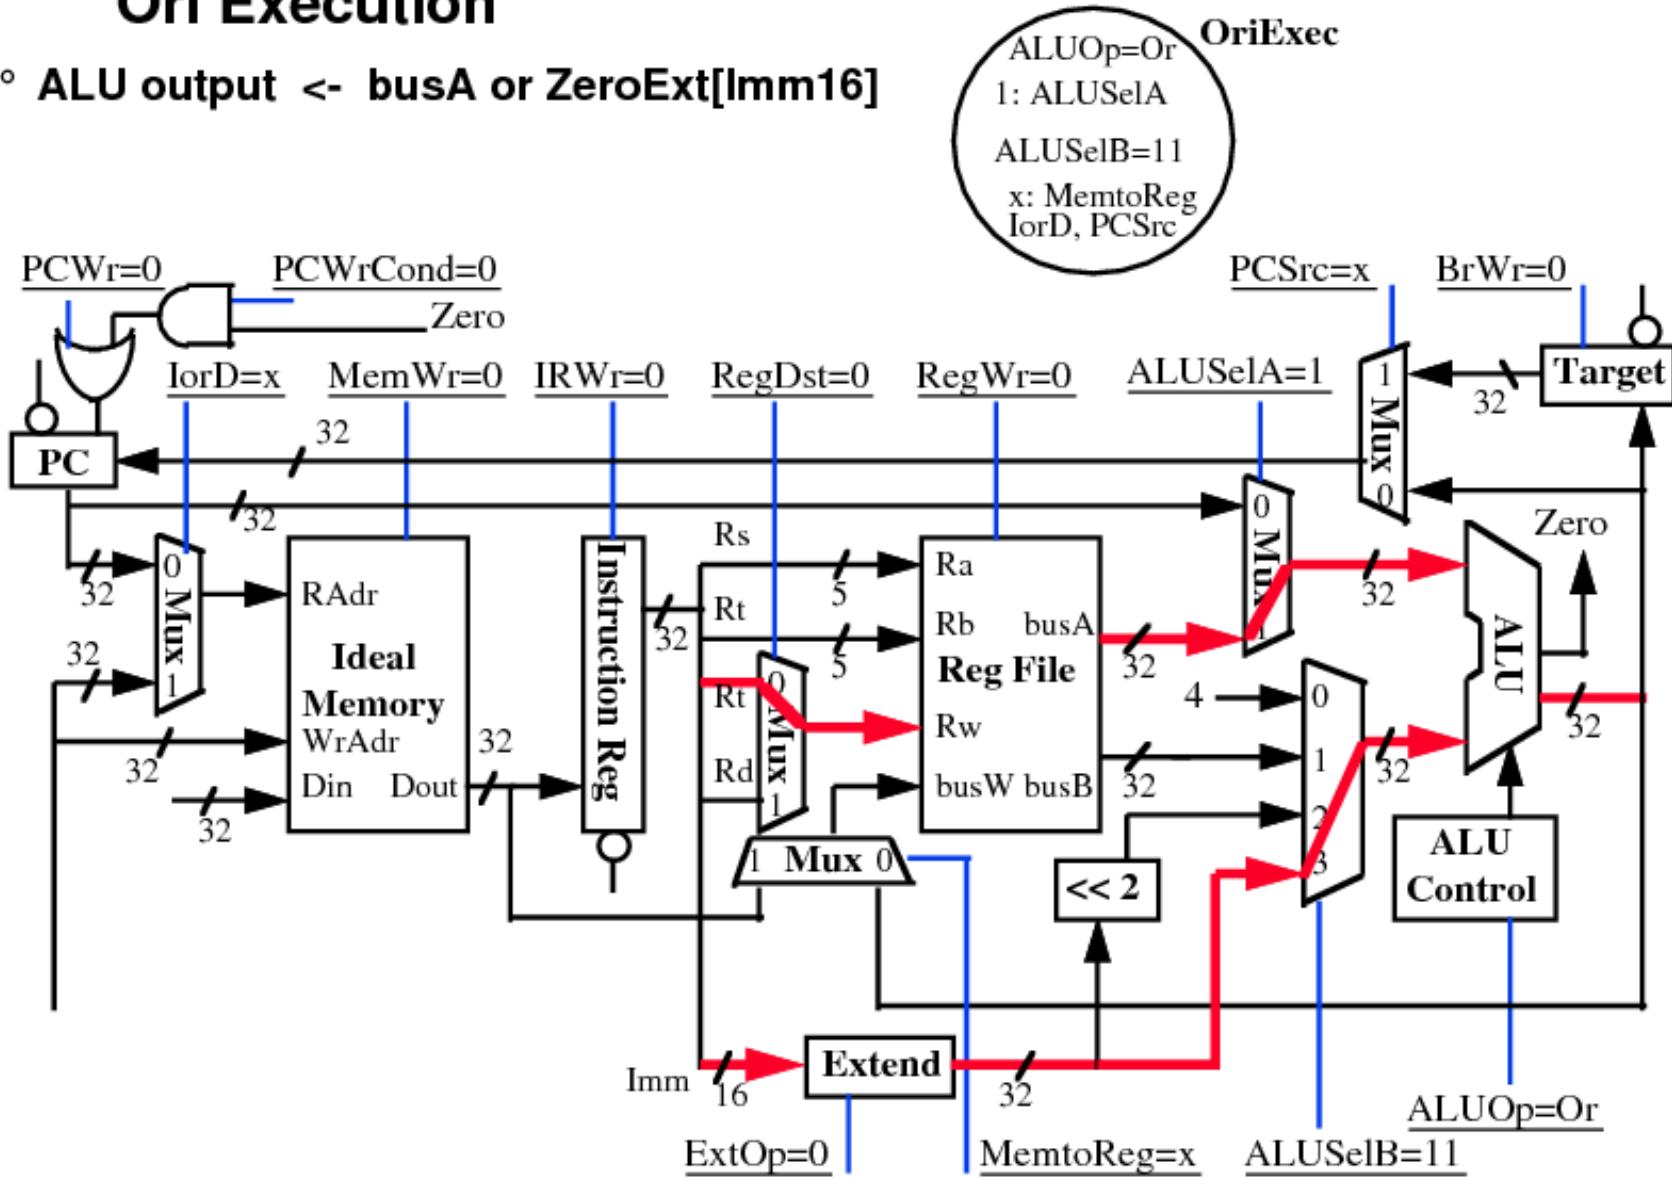
Instruction Decode: We have an Ori!

- Next Cycle: Ori Execution



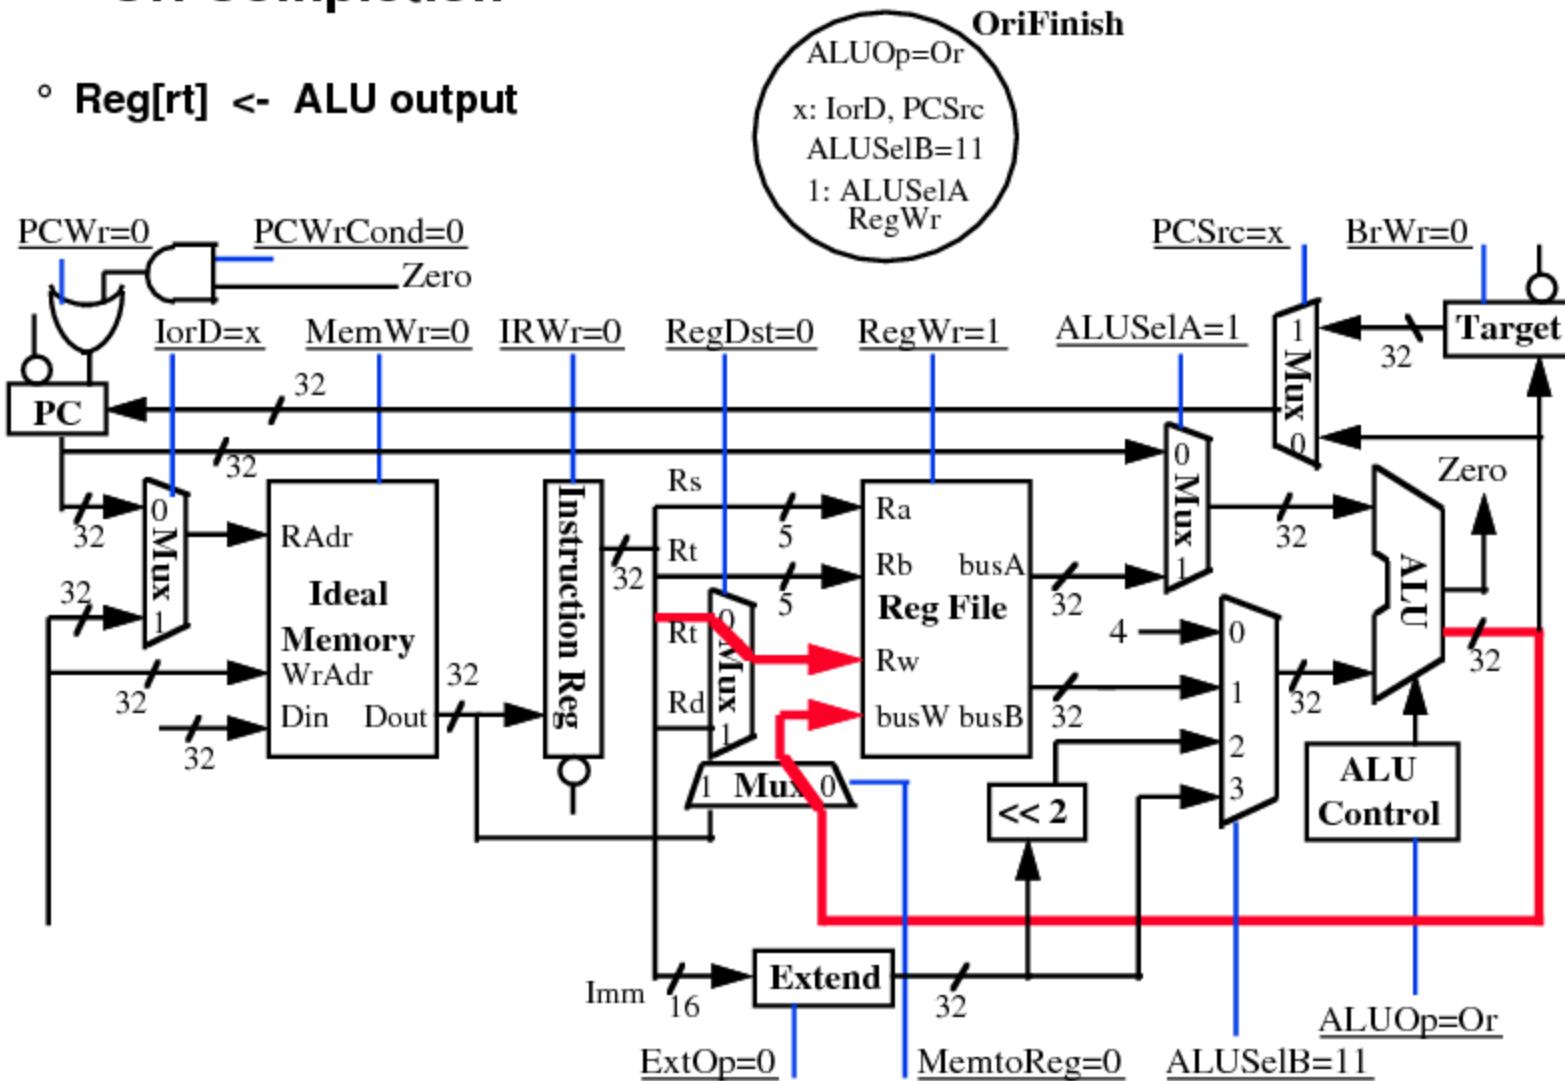
Ori Execution

- ° ALU output \leftarrow busA or ZeroExt[Imm16]



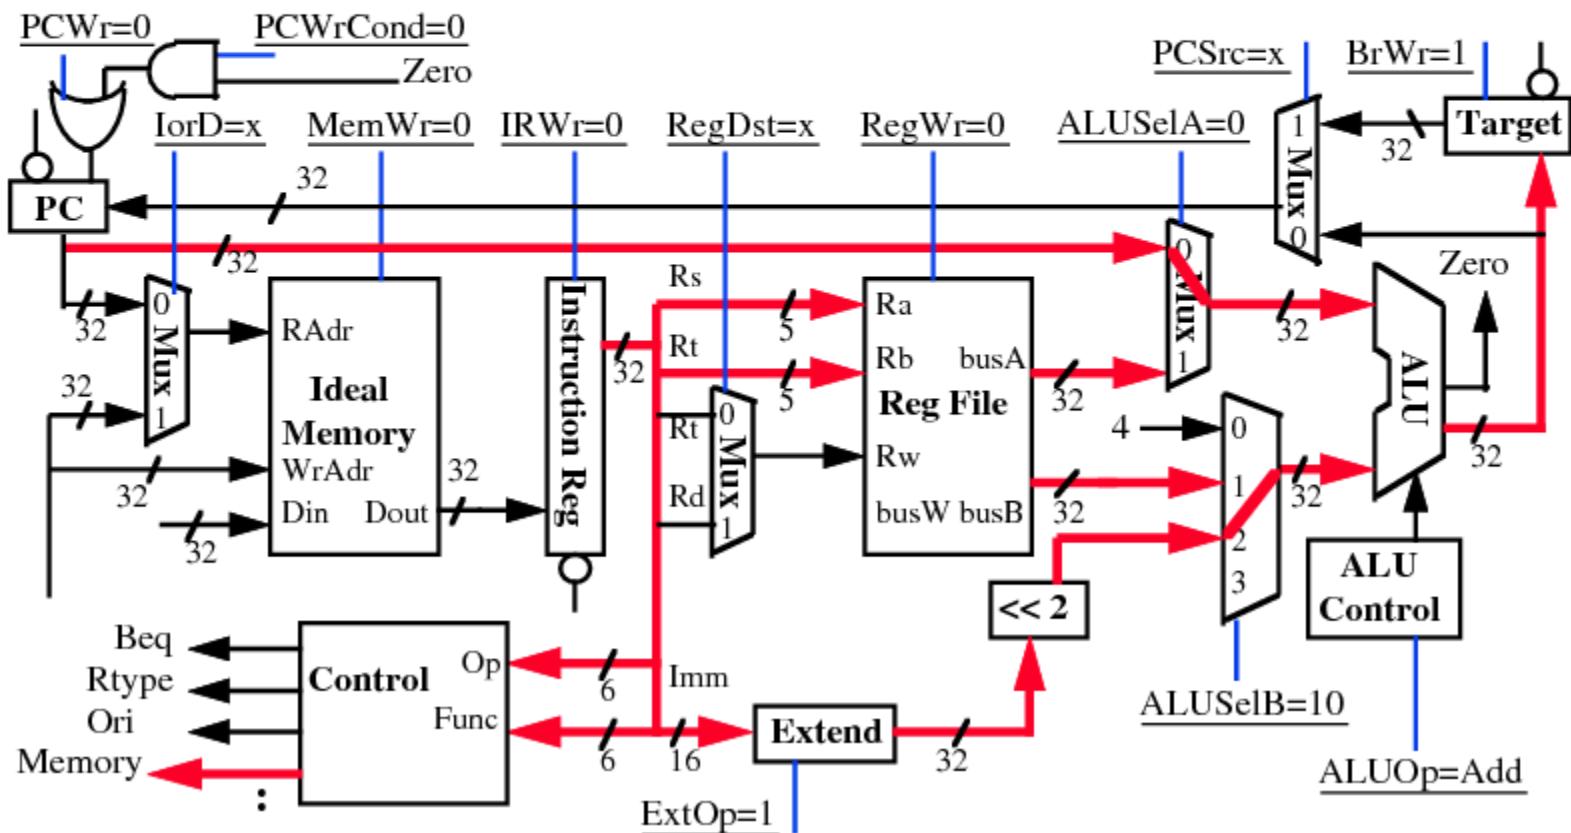
Ori Completion

- Reg[rt] <- ALU output



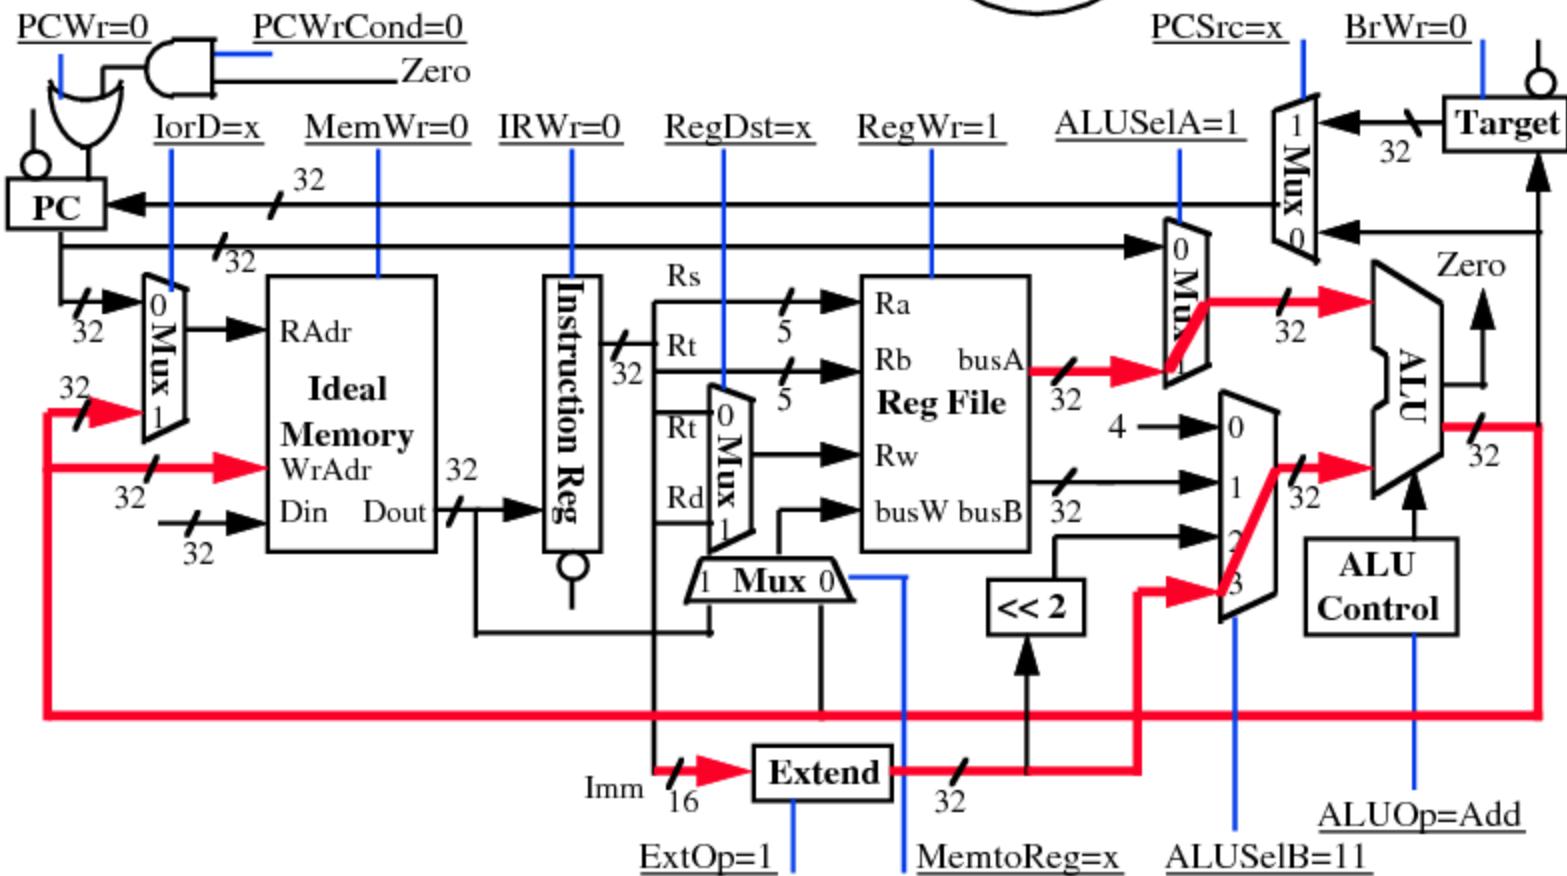
Instruction Decode: We have a Memory Access!

- Next Cycle: Memory Address Calculation



Memory Address Calculation

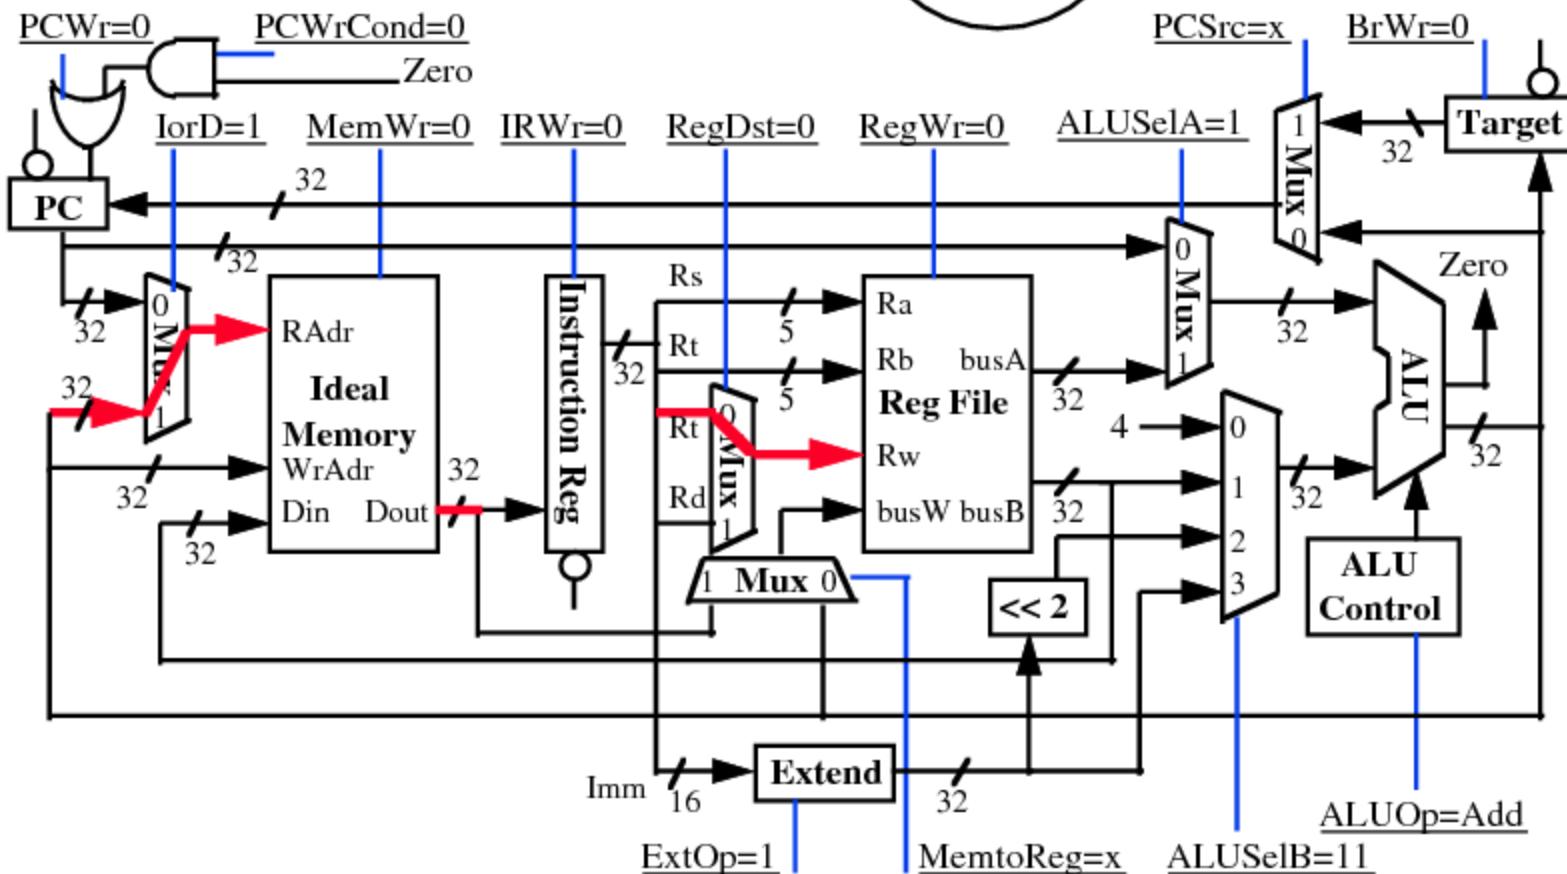
- ALU output $\leftarrow \text{busA} + \text{SignExt}[\text{Imm}16]$



Memory Access for Load

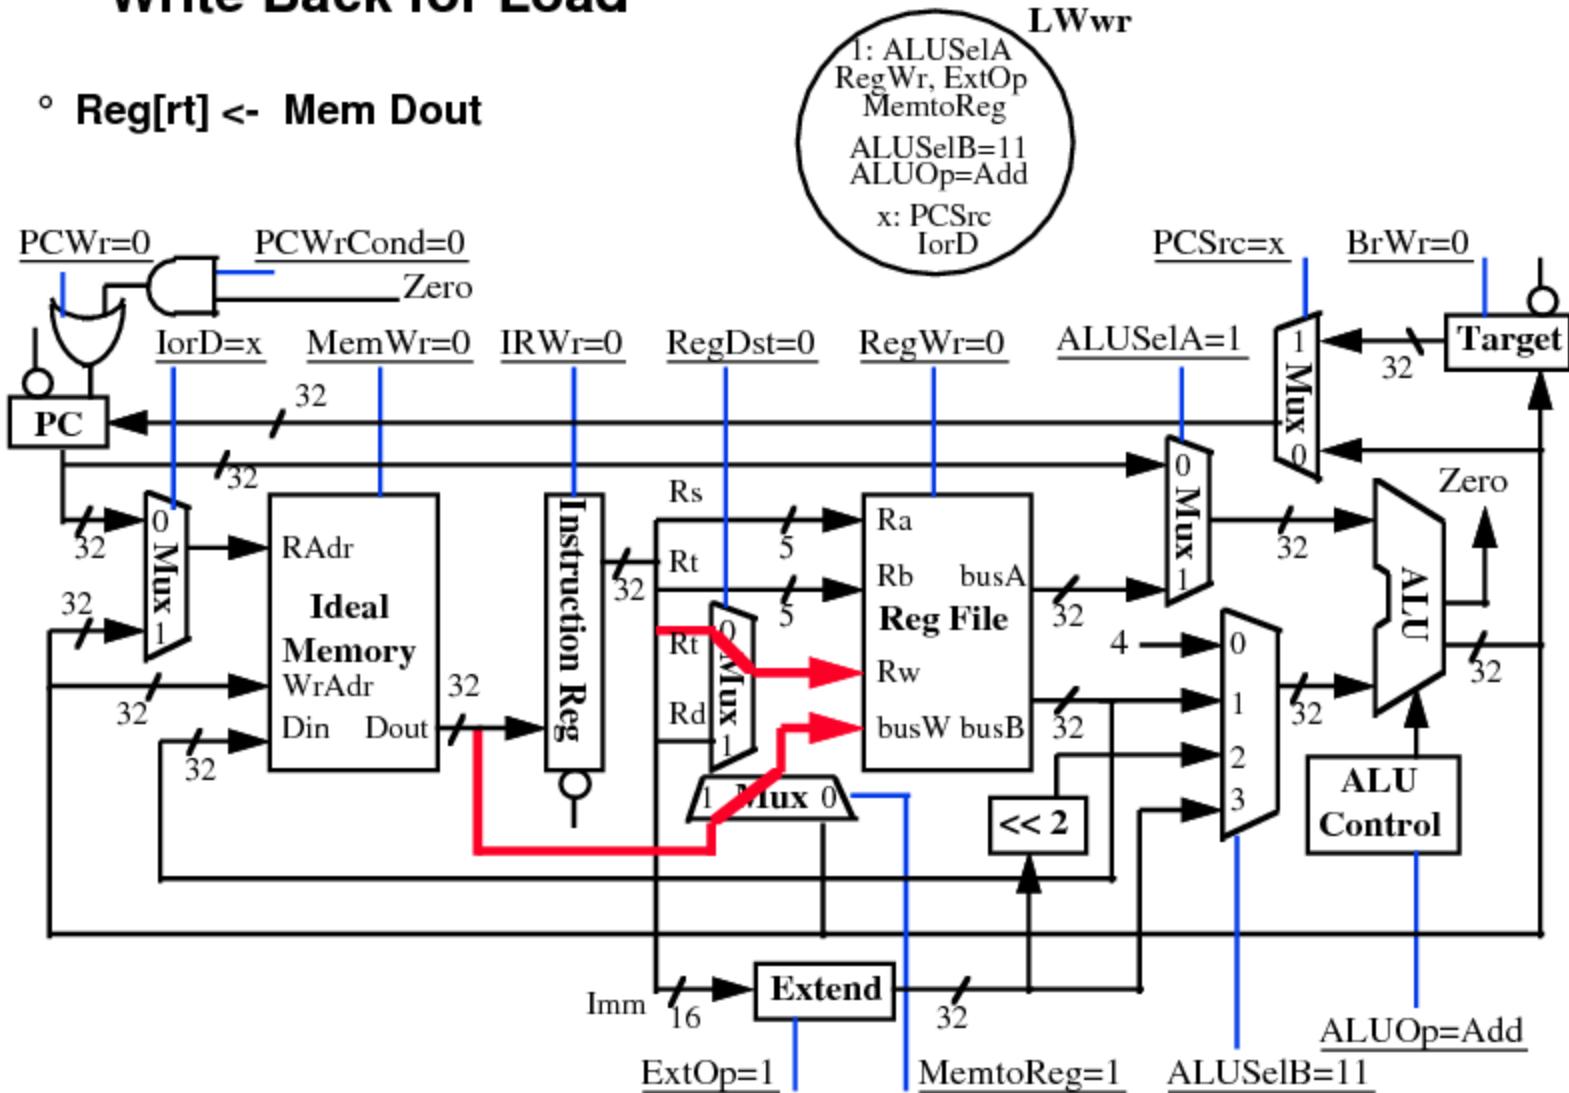
- Mem Dout <- mem[ALU output]

LWmem
 1: ExtOp
 ALUSelA, IorD
 ALUSelB=11
 ALUOp=Add
 x: MemtoReg
 PCSrc



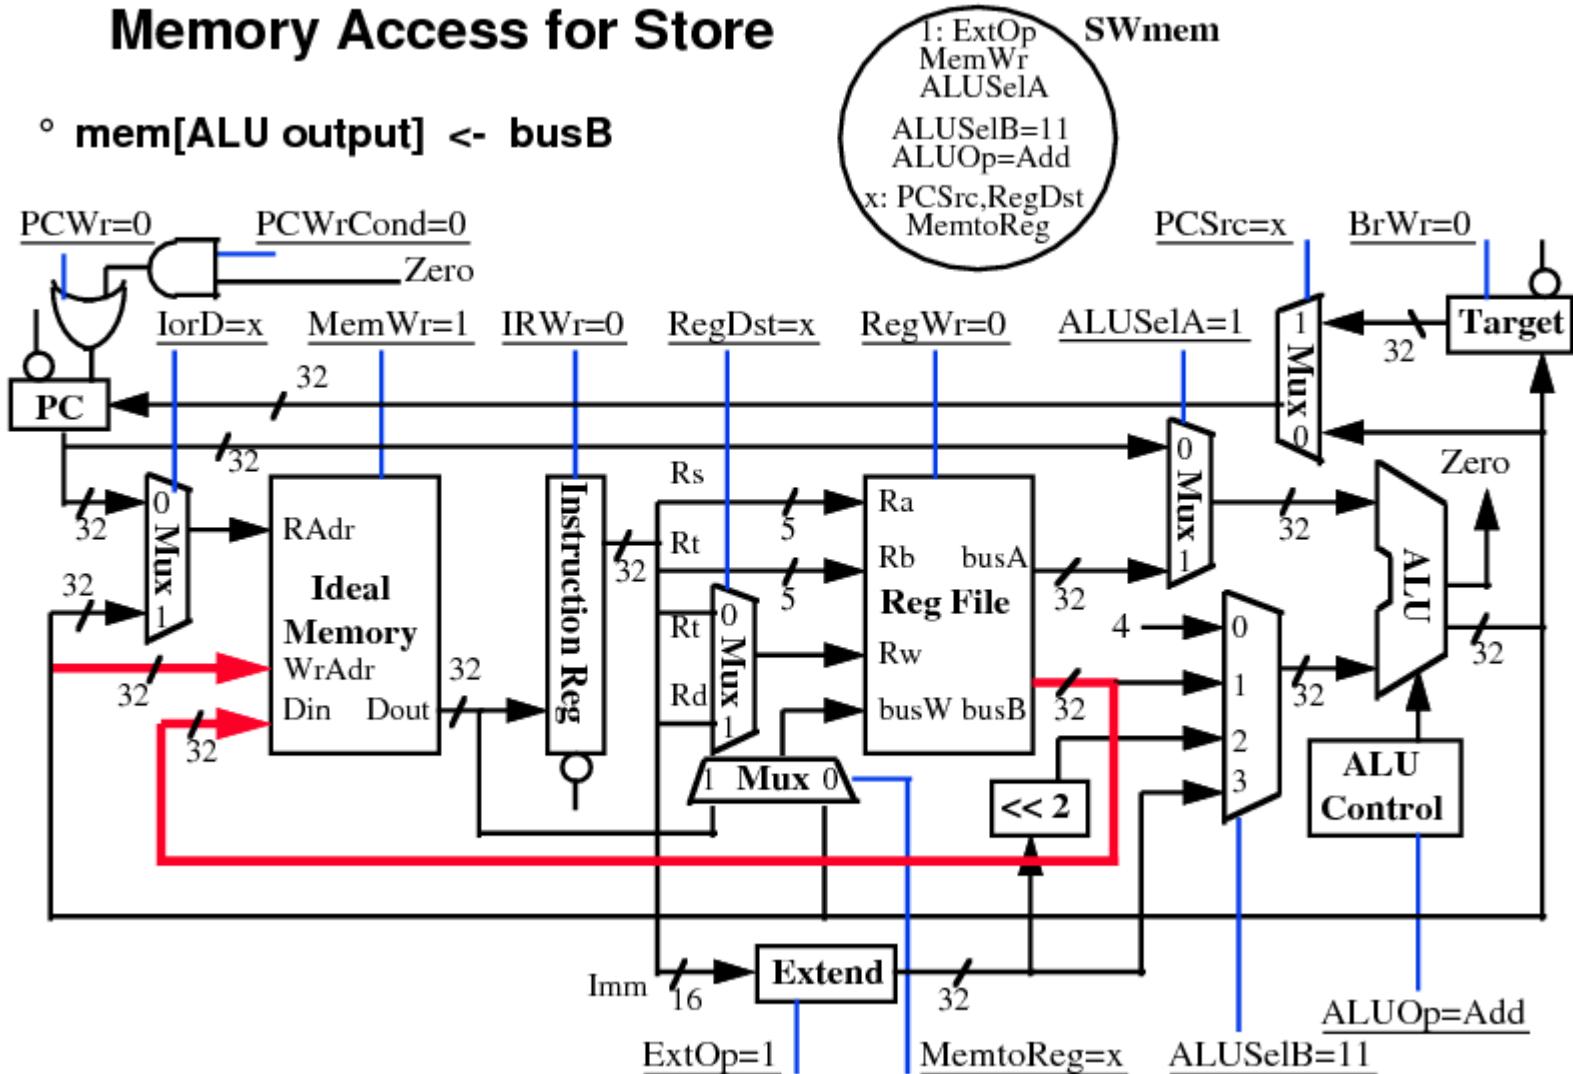
Write Back for Load

- Reg[rt] <- Mem Dout

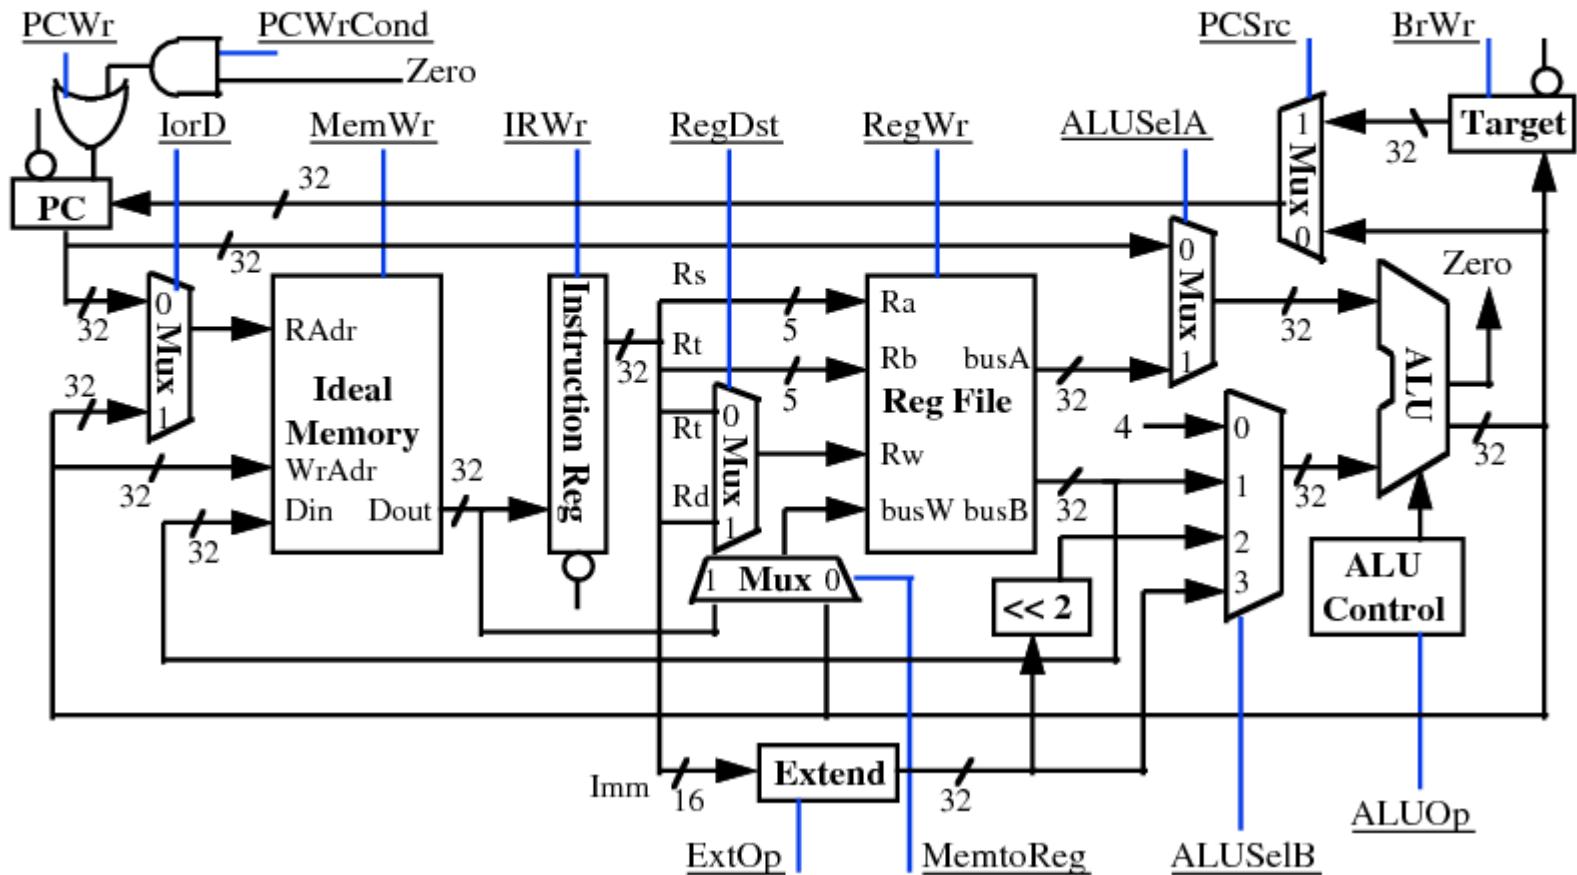


Memory Access for Store

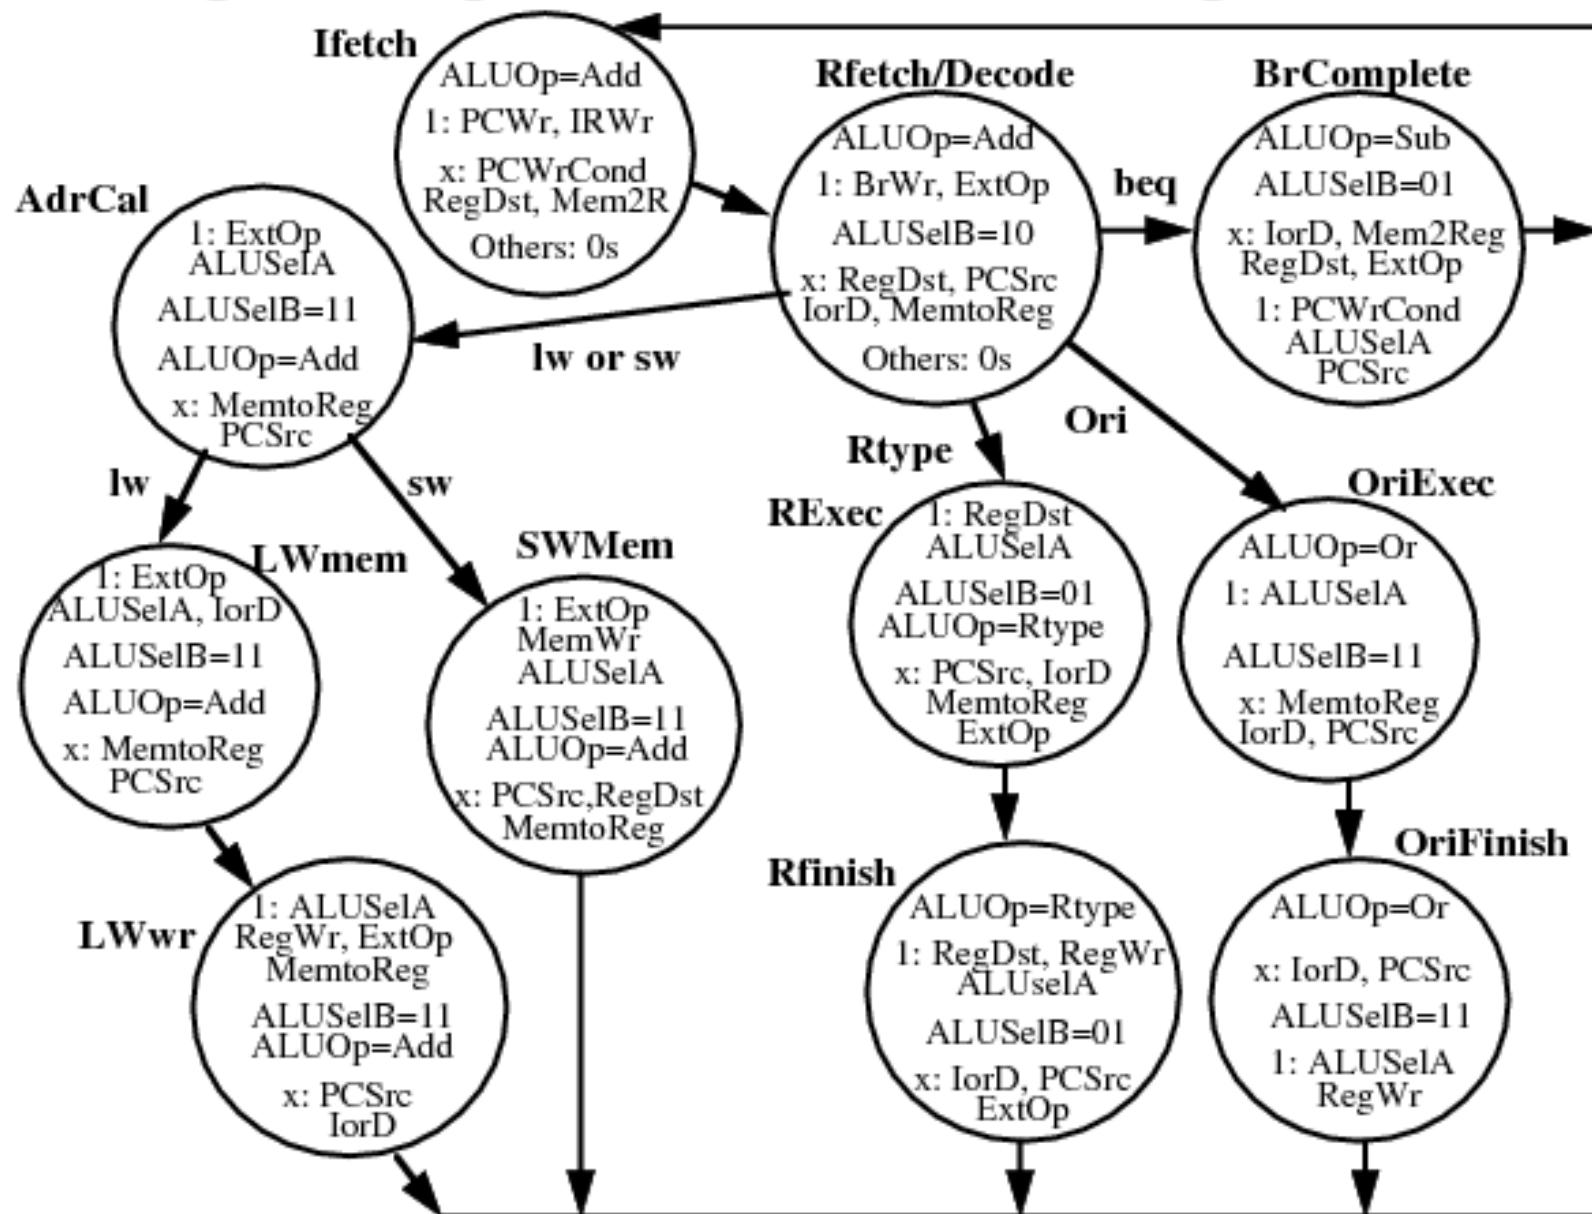
° $\text{mem}[\text{ALU output}] \leftarrow \text{busB}$



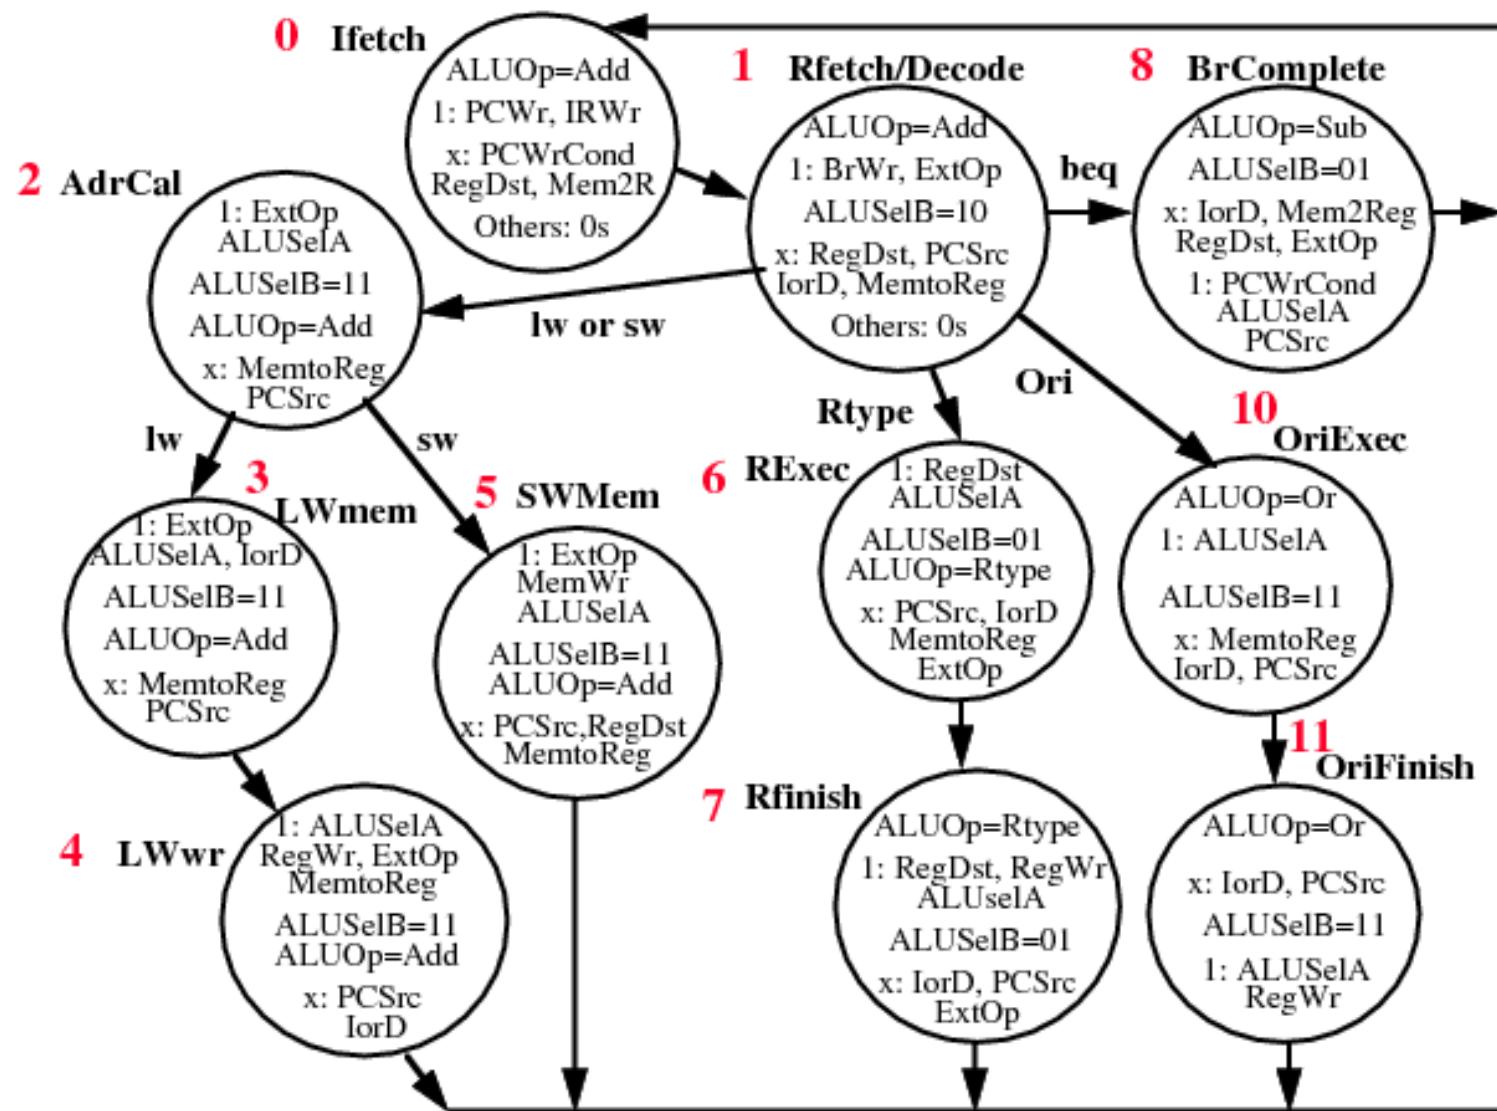
Putting it all together: Multiple Cycle Datapath



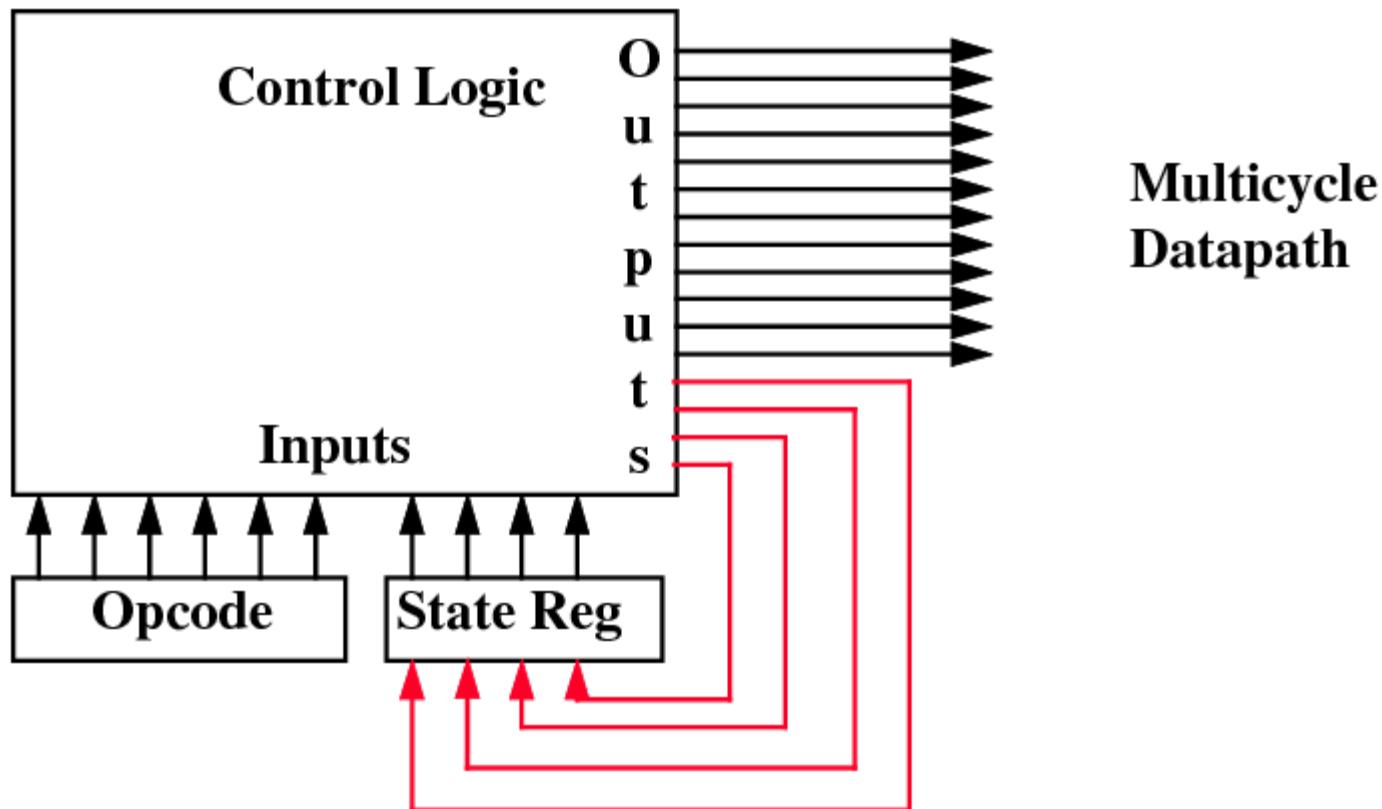
Putting it all together: Control State Diagram



Control Logic in the Form of Finite State Diagram



Sequencing Control: Explicit Next State Function

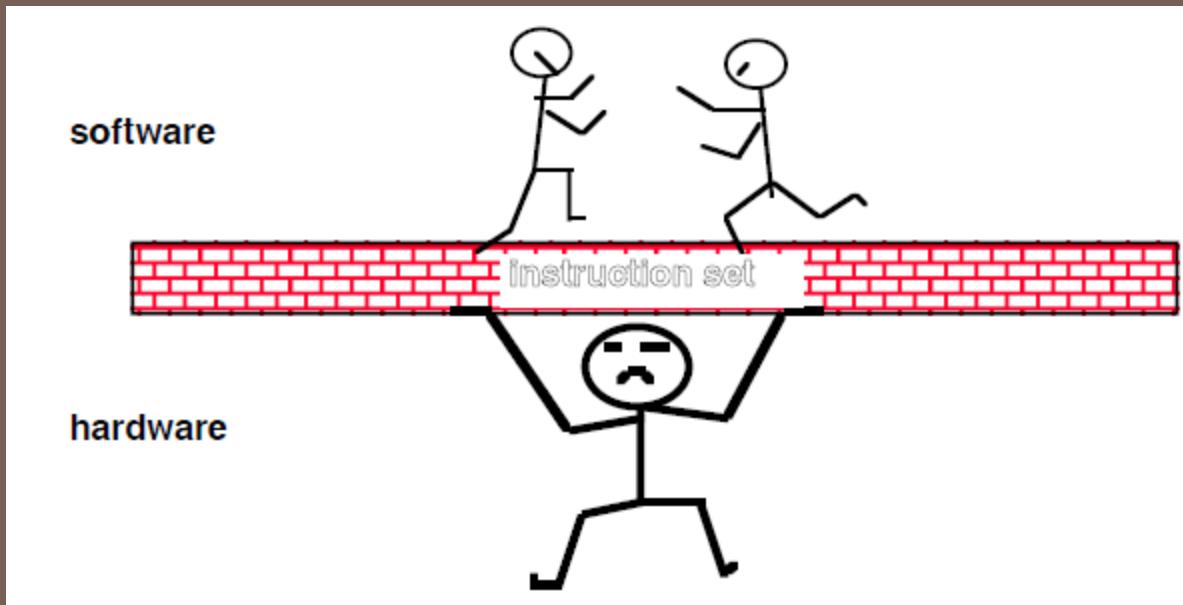


The Big Picture: Performance Perspective

- ❑ Performance of a machine was determined by
 - ❑ Instruction Count
 - ❑ Clock cycle Time
 - ❑ Clock cycles per instruction
- ❑ Processor Design (data path and control) will determine
 - ❑ Clock cycle time
 - ❑ Clock cycles per instruction
- ❑ We shall first design a Single Cycle Processor
 - ❑ Advantage: One clock cycle per instruction
 - ❑ Disadvantage: Long cycle time

Summary

- Disadvantages of the Single Cycle Processor
 - Long cycle time
 - Cycle time is too long for all instructions except load
- Multiple cycle processor
 - Divide the instruction into smaller steps
 - Execute each step (instead of the entire instruction) in one cycle



COMPUTER SYSTEMS ORGANIZATION

Acknowledgment: Almost all of these slides are based on Dave Patterson's CS152 Lecture Slides at UC, Berkeley

Single Cycle CPU Design -- Spring 2012 -- IIIT-H -- Suresh Purini

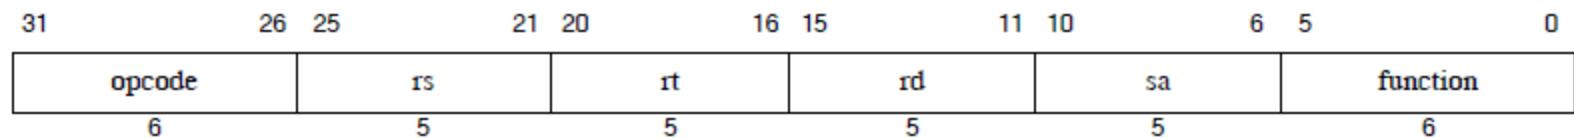
MIPS CPU Instructions

Three Types of CPU Instructions

- R-type
- I-Type
- J-Type

R-Type Instructions

□ R-type Instruction Format



R-Type Instructions: ADD

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	rs	rt	rd	0 00000	ADD 100000	

Format: ADD rd, rs, rt

MIPS32

- Format: ADD rd, rs, rt
- $R[rd] = R[rs] + R[rt]$
- 32-bit 2's Complement Addition
- Destination register will not be modified if integer overflow exceptions occurs.

R-Type Instructions: SUB

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	rs	rt	rd	0 00000	SUB 100010	

- Format: SUB rd, rs, rt
- $R[rd] = R[rs] - R[rt]$
- 32-bit signed subtraction
- Destination register will not be modified if integer overflow exceptions occurs.

R-Type Instructions: AND

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	rs	rt	rd	0 00000	AND 100100	

Format: AND rd, rs, rt

MIPS32

- Format: AND rd, rs, rt
- $R[rd] = R[rs] \& R[rt]$

R-Type Instructions: OR

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	rs	rt	rd	0 00000	OR 100101	
6	5	5	5	5	6	

Format: OR rd, rs, rt

MIPS32

- Format: OR rd, rs, rt
- $R[rd] = R[rs] \mid R[rt]$

R-Type Instructions: SLT

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	rs	rt	rd	0 00000	SLT 101010	

Format: SLT rd, rs, rt

MIPS32

- Format: SLT rd, rs, rt
- $R[rd] = R[rs] < R[rt] ? 1:0$
- Signed comparision

There are many other R-type instructions like ADDU, NOR, XOR etc.

R-Type Instructions: SLL

31	26 25	21 20	16 15	11 10	6 5	0
SPECIAL 000000	0 00000	rt	rd	sa	SLL 000000	

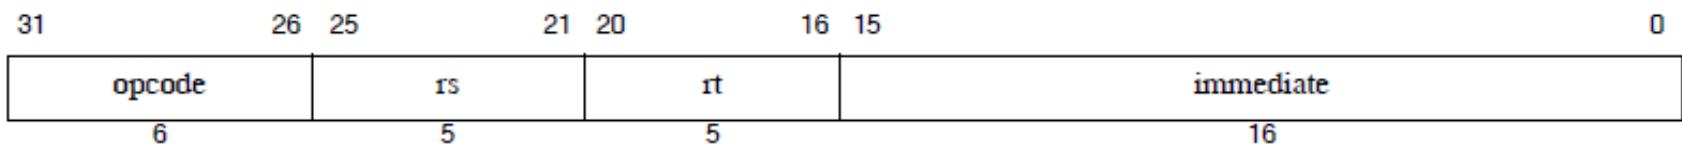
Format: SLL rd, rt, sa

MIPS32

- Format: SLL rd, rt, sa
- $R[rd] = R[rt] \ll sa$

Note: In our processor design we do not implement shift instructions.

I-type Instructions



I-type Instructions

31	26 25	21 20	16 15	0
ADDI 001000	rs	rt	immediate	

Format: ADDI rt, rs, immediate

MIPS32

- Format: ADDI rt, rs, immediate
- $R[rt] = R[rs] + \text{sign_extend(immediate)}$
- immediate is 16-bit signed immediate
- 32-bit 2'complement addition
- Destination register will not be updated if integer overflow exception occurs

I-type Instructions

31	26 25	21 20	16 15	0
ANDI 001100	rs	rt	immediate	

Format: ANDI rt, rs, immediate

MIPS32

- Format: ANDI rt, rs, immediate
- $R[rt] = R[rs] \& \text{zero_extend(immediate)}$

I-type Instructions: ORI

31	26 25	21 20	16 15	0
ORI 001101	rs	rt	immediate	

Format: ORI rt, rs, immediate

MIPS32

- Format: ORI rt, rs, immediate
- $R[rt] = R[rs] \mid \text{zero_extend(immediate)}$

I-type Instructions: SLTI

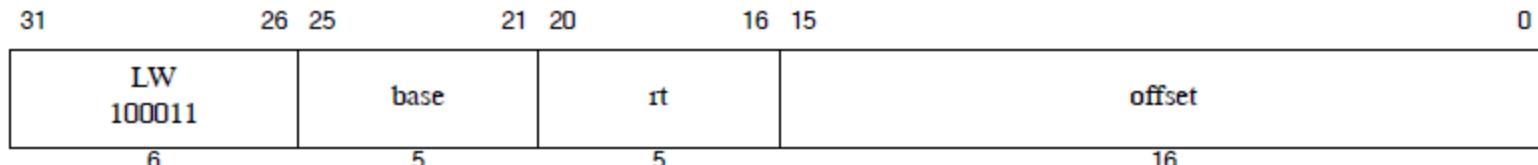
31	26 25	21 20	16 15	0
SLTI 001010	rs	rt	immediate	

Format: SLTI rt, rs, immediate

MIPS32

- Format: SLTI rt, rs, immediate
- $R[rt] = R[rs] < \text{sign_extend(immediate)} ? 1:0$

I-type Instructions: LW

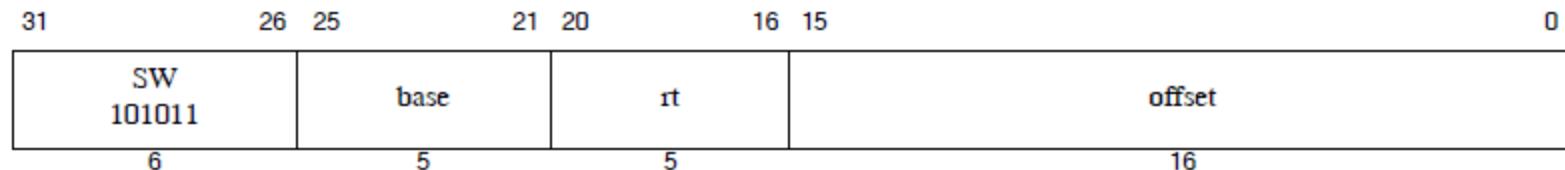


Format: LW rt, offset(base)

MIPS32

- Format: LW rt, offset(base)
- $vaddr = \text{sign_extend}(\text{offset}) + R[\text{base}]$
- $R[rt] = \text{Mem}[vaddr]$
- If vaddr is now word-aligned, an exception will be raised.

I-type Instructions: SW



Format: SW rt, offset(base)

MIPS32

- **Format:** SW rt, offset(base)
- $vaddr = \text{sign_extend}(\text{offset}) + R[\text{base}]$
- $\text{Mem}[vaddr] = R[rt]$
- If vaddr is now word-aligned, an exception will be raised.

I-type Instructions: SLTI

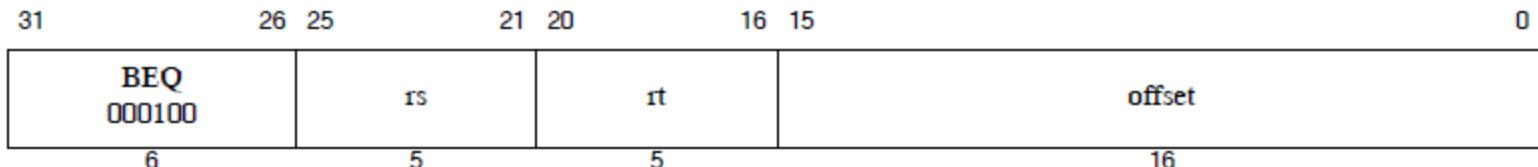
31	26 25	21 20	16 15	0
SLTI 001010	rs	rt	immediate	

Format: SLTI rt, rs, immediate

MIPS32

- Format: SLTI rt, rs, immediate
- $R[rt] = R[rs] < \text{sign_extend(immediate)} ? 1:0$

I-Type Instructions: BEQ

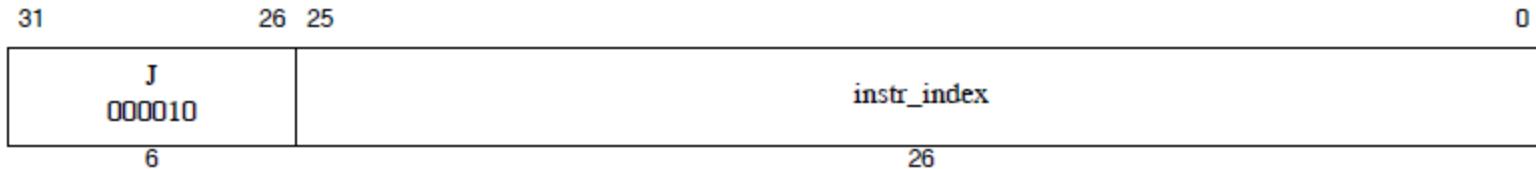


Format: BEQ rs, rt, offset

MIPS32

- Format: BEQ rs, rt, offset
- If R[rs] == R[rt] then
 - PC = addr_of_branch + 4 + sign_extend(offset << 2)

MIPS J-Type Instructions



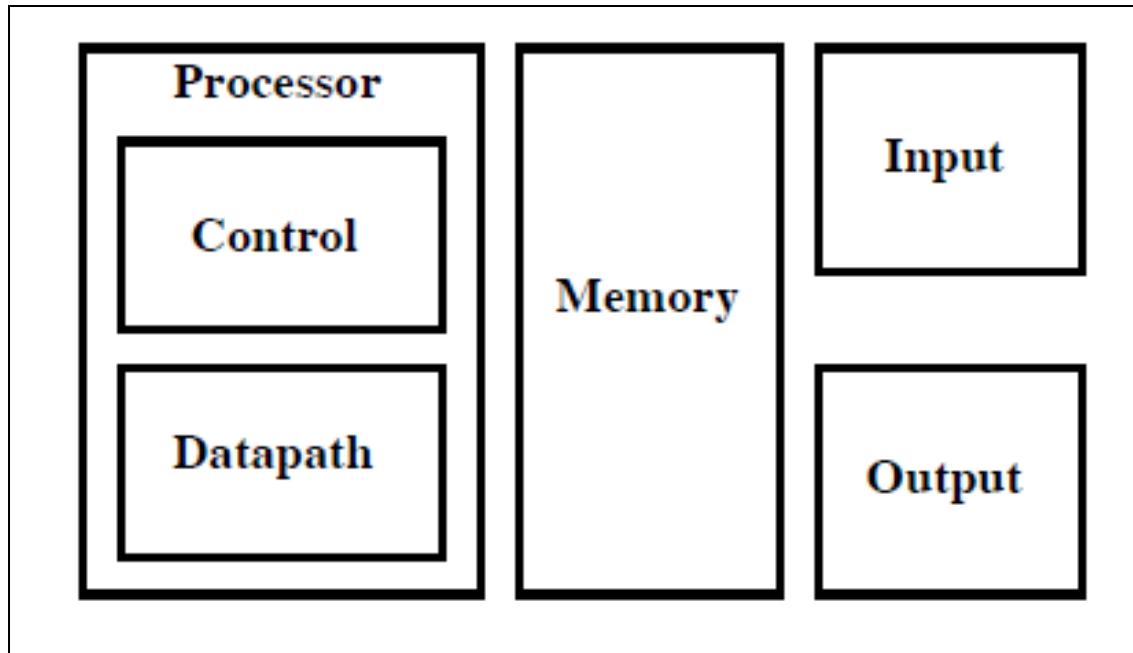
Format: J target

MIPS32

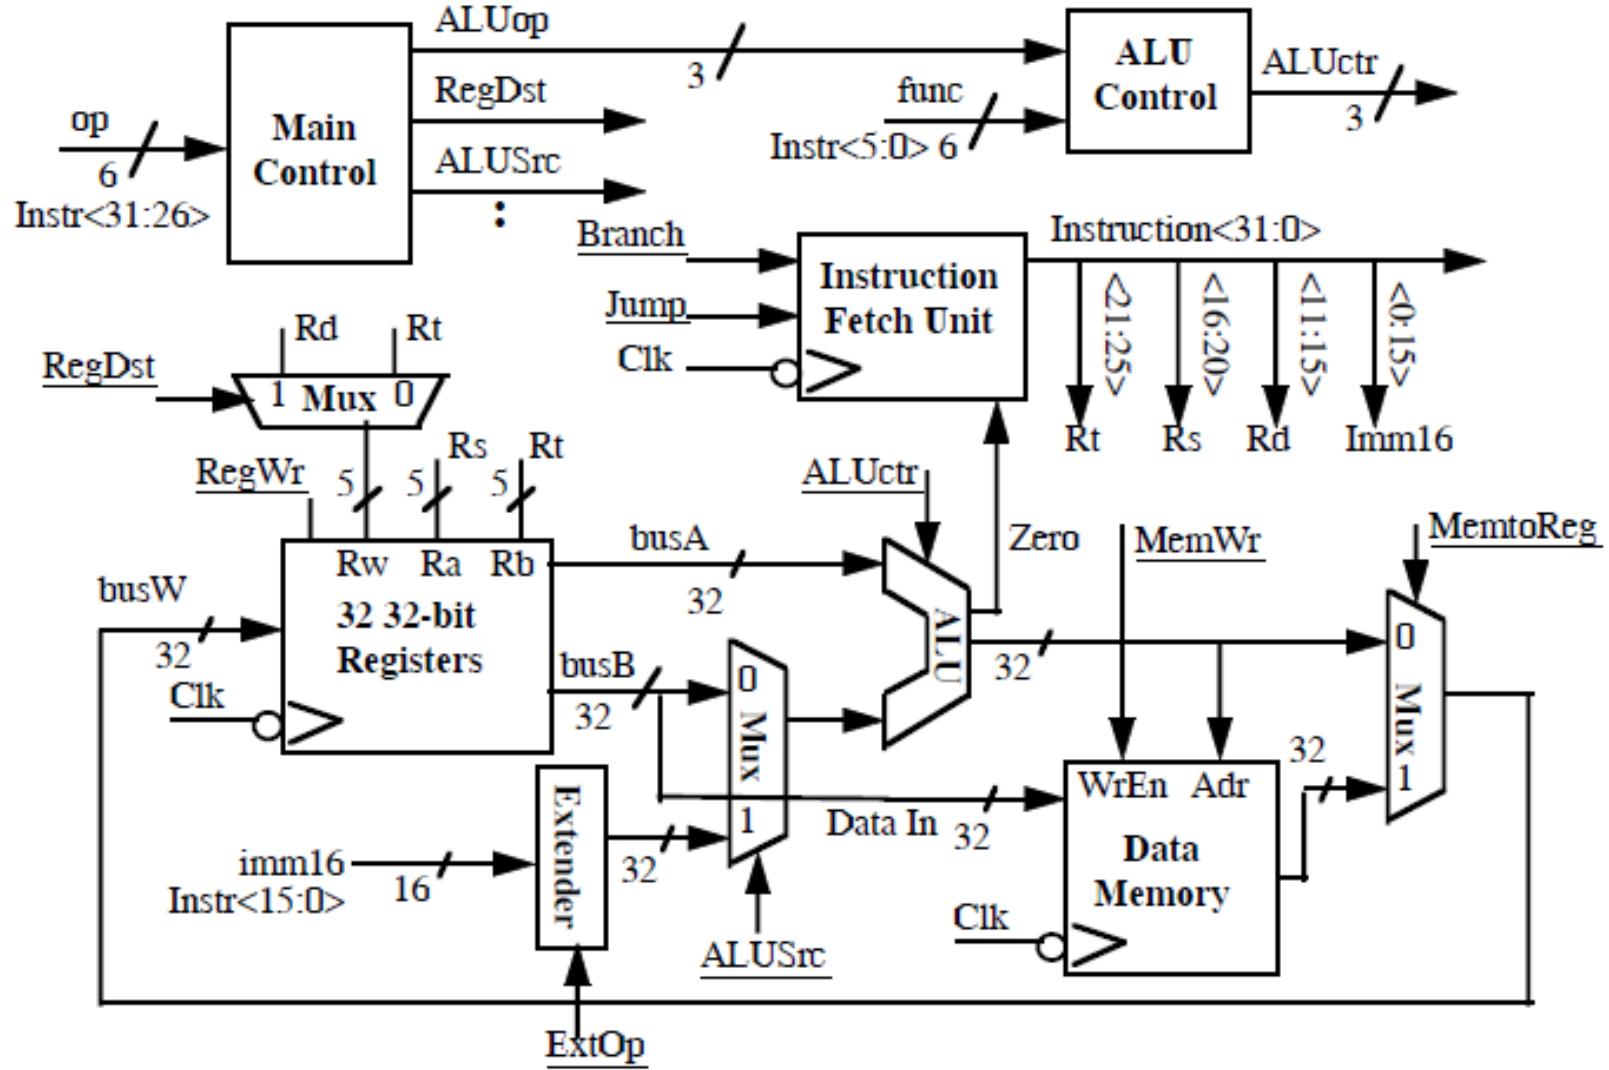
- Format: J target
- Target Address
 - Lower 28 bits: instr_index | | 00
 - Upper Four Bits: Bits 31, 30, 29, 28 of the address of the Jump Instruction.

The Big Picture: Where are We Now?

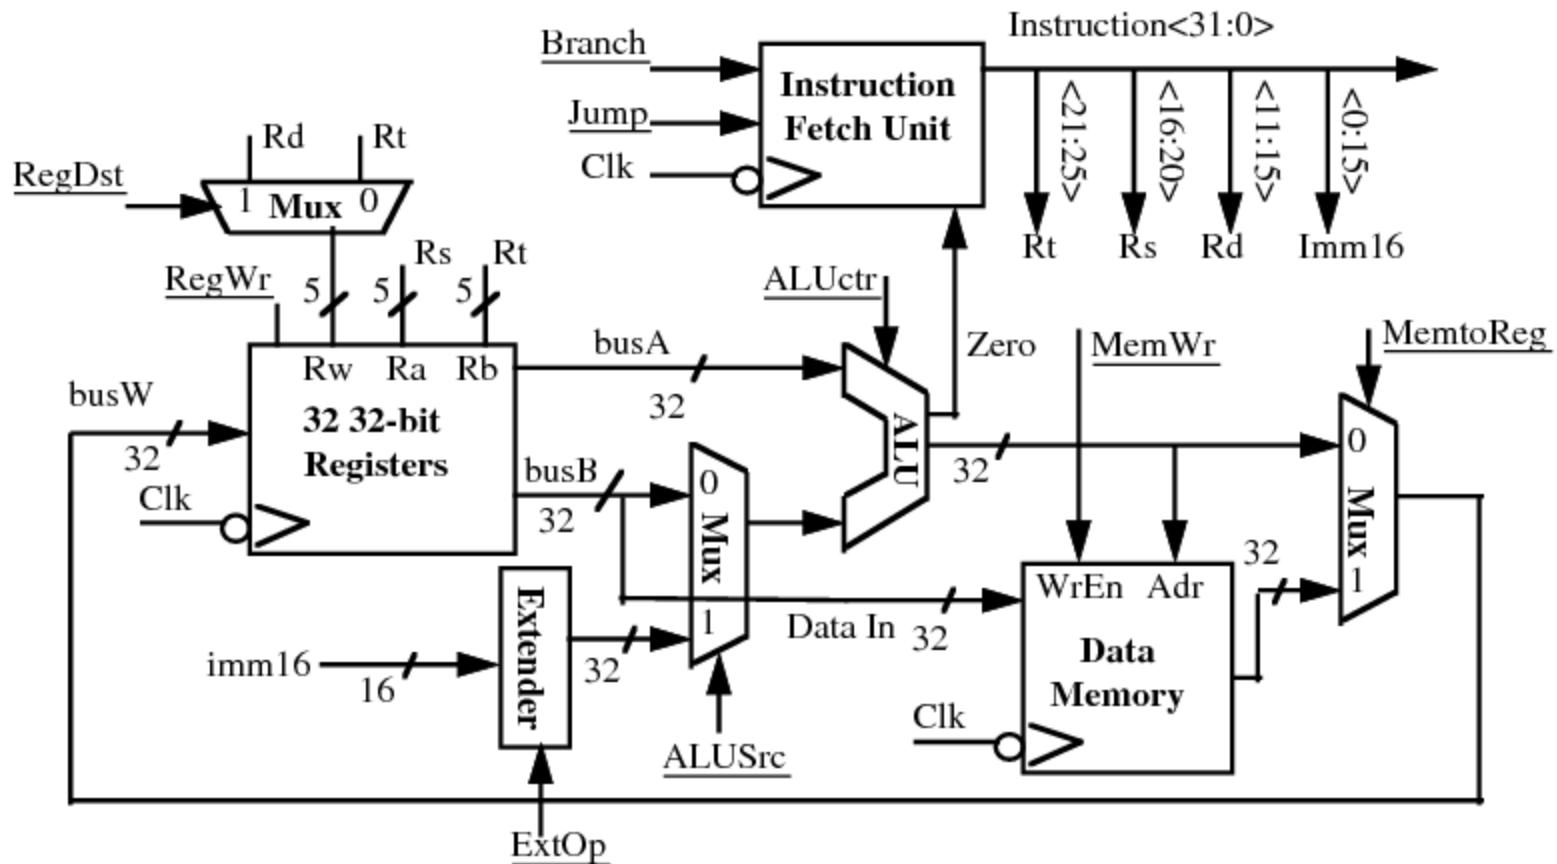
□ Five Classic Components of a Computer



MIPS Processor: Control Path + Data Path



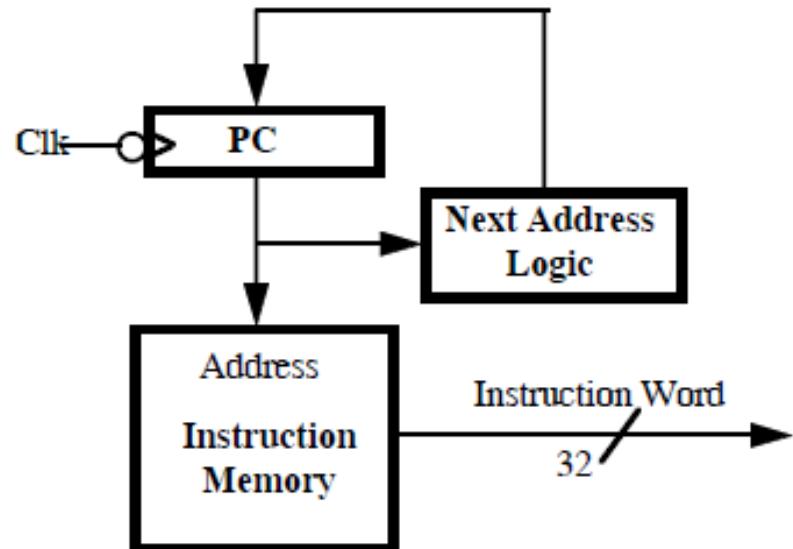
Data Path



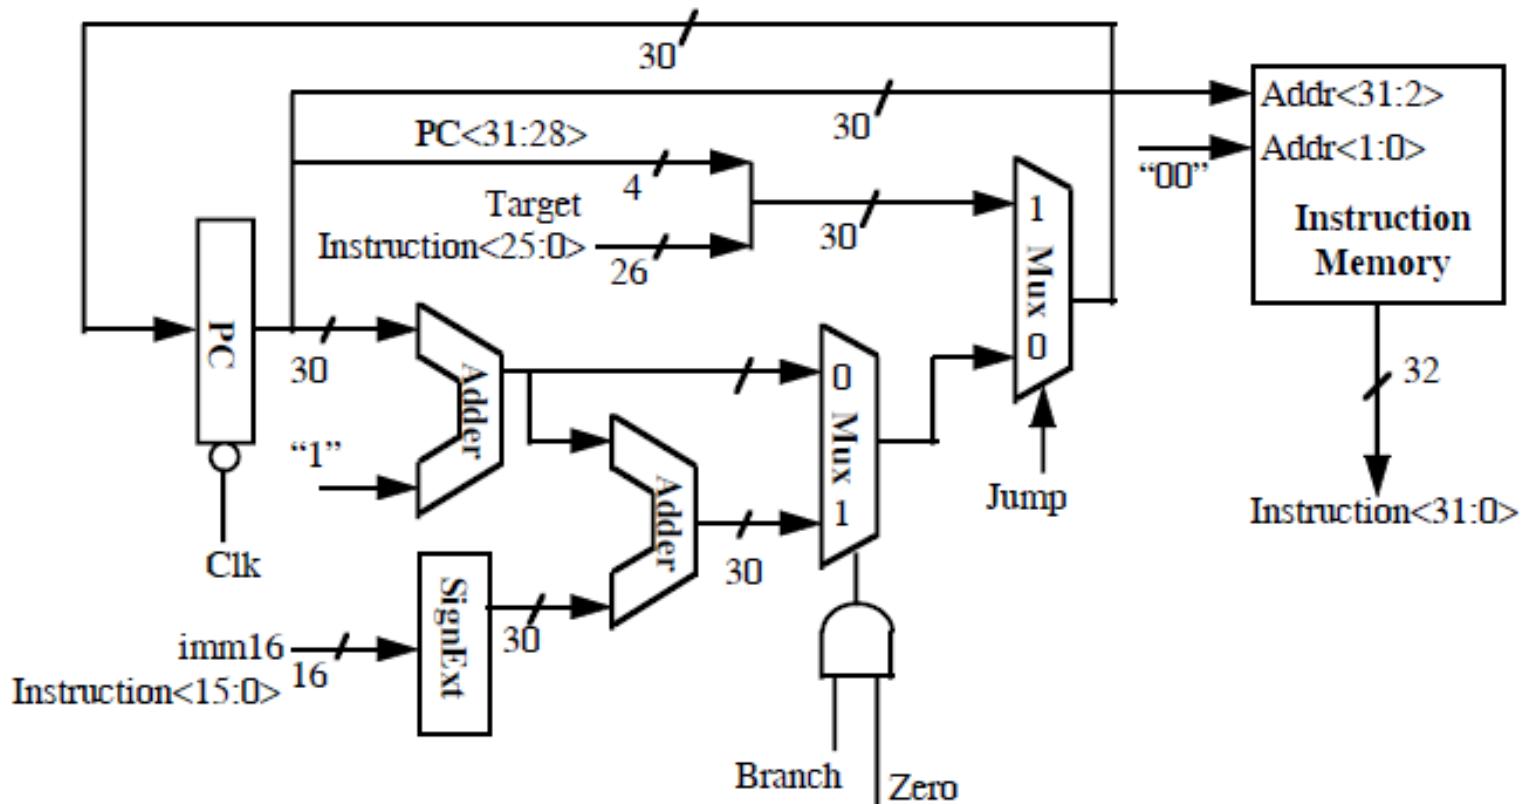
Overview of Instruction Fetch Unit

At a falling clock edge what happens:

- PC gets updated at the falling clock edge
- Fetch the Instruction from the address pointed to by PC
- Pass the PC through the next address logic
- Next value of the PC
 - Sequential Code
 - $\text{nextPC} = \text{PC} + 4$
 - Branch and Jump
 - $\text{nextPC} = \text{"something else"}$



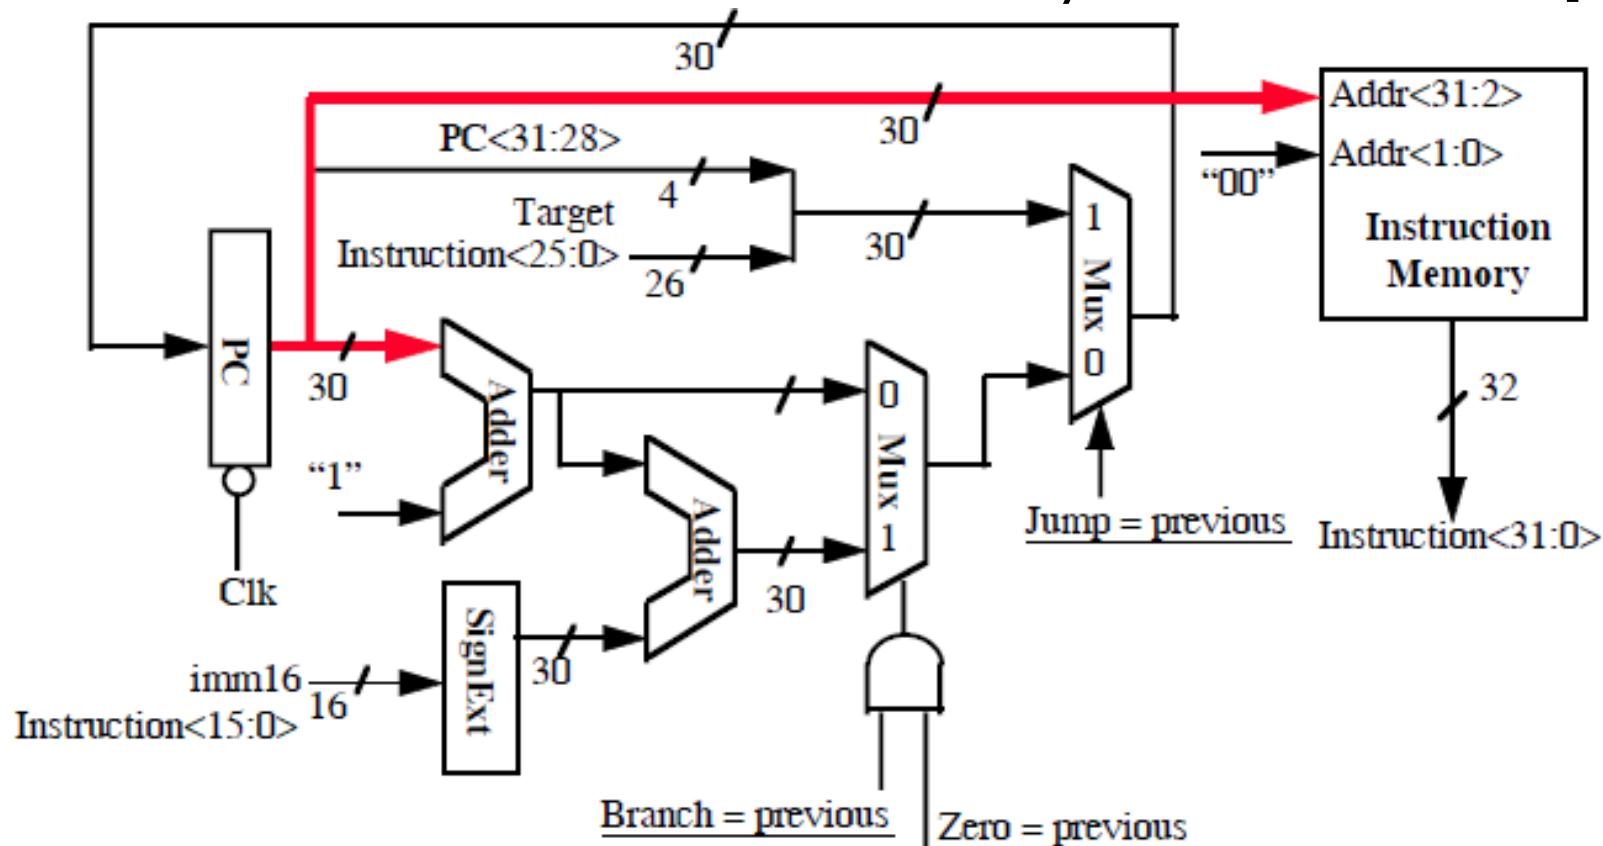
Instruction Fetch Unit



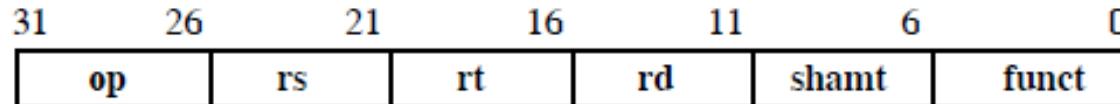
Instruction Fetch Unit at the Beginning of Add / Subtract

The following two steps are the same for all the instructions.

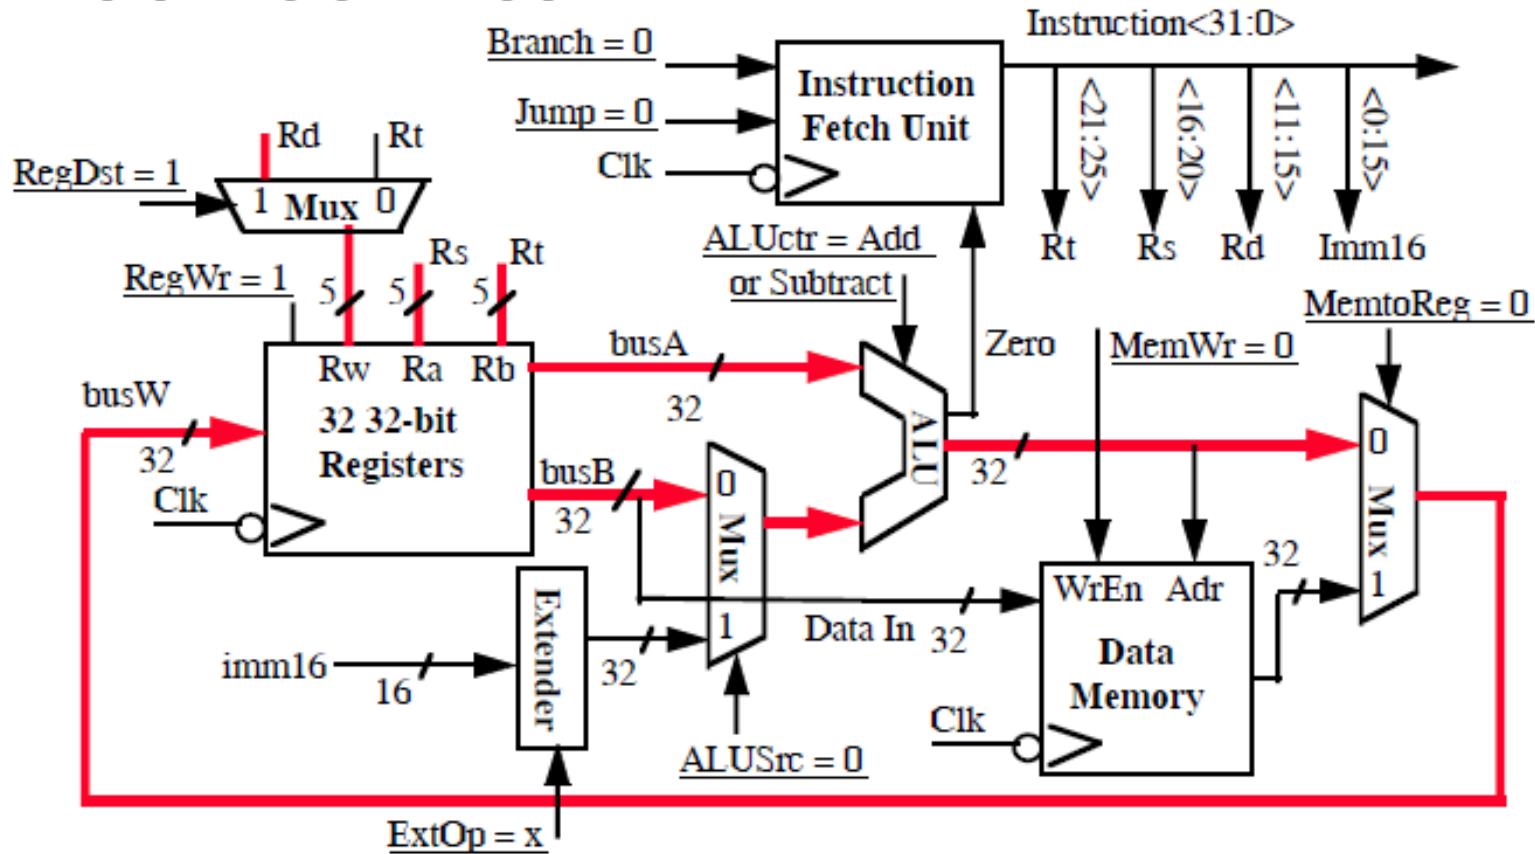
1. $PC = \text{nextPC}$
2. Fetch the instruction from Instruction memory: $\text{Instruction} = \text{mem}[PC]$



The Single Cycle Datapath during Add and Subtract

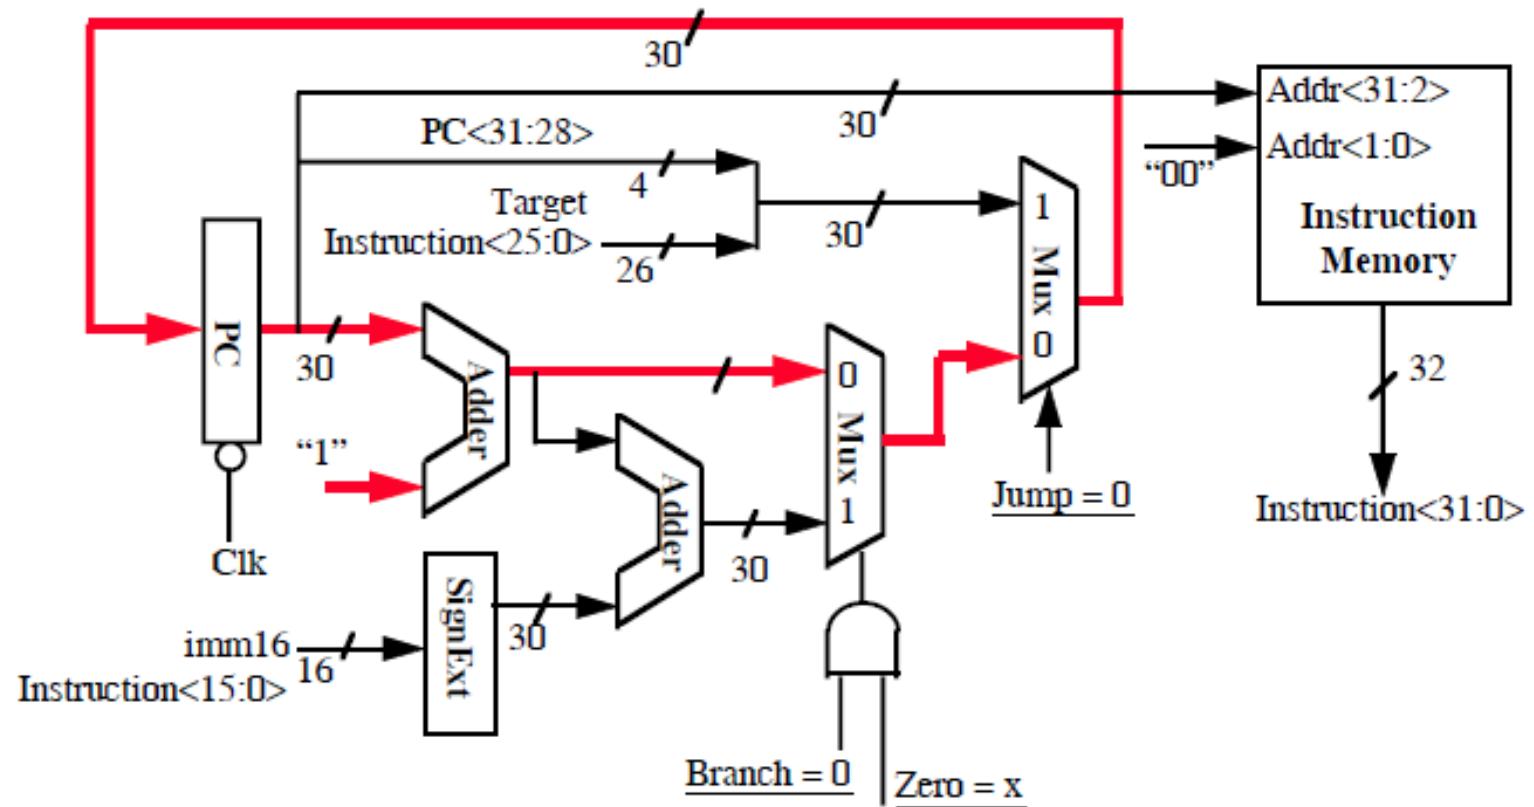


$$R[rd] \leftarrow R[rs] + / - R[rt]$$

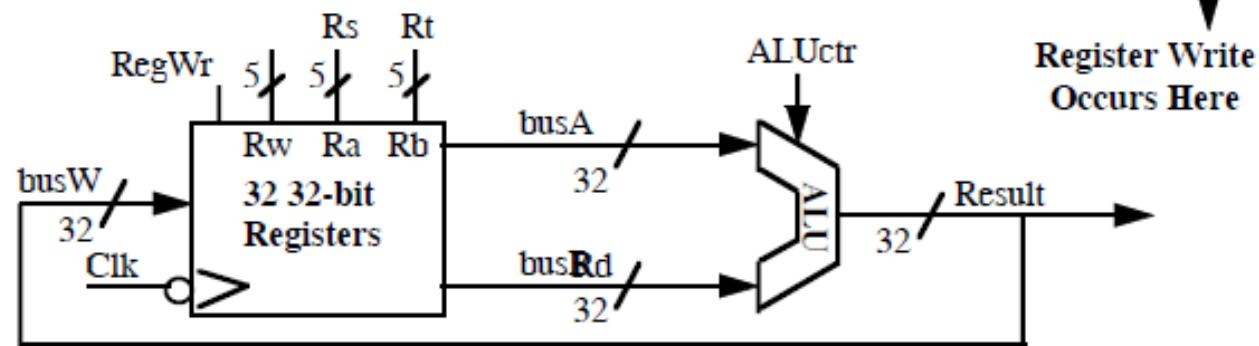
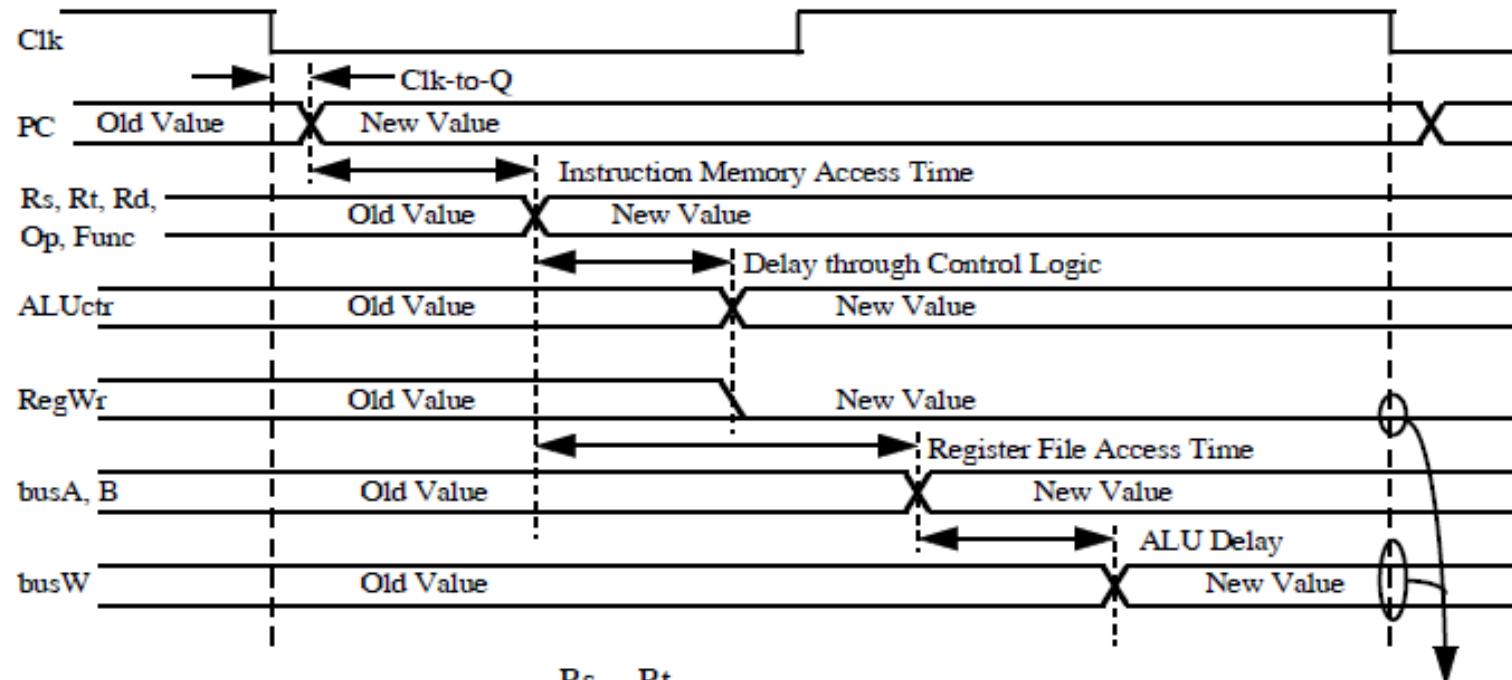


Instruction Fetch Unit at the End of Add and Subtract

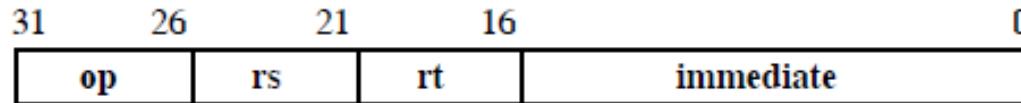
- $PC = PC + 4$
 - This is the same for all instructions except: Branch and Jump



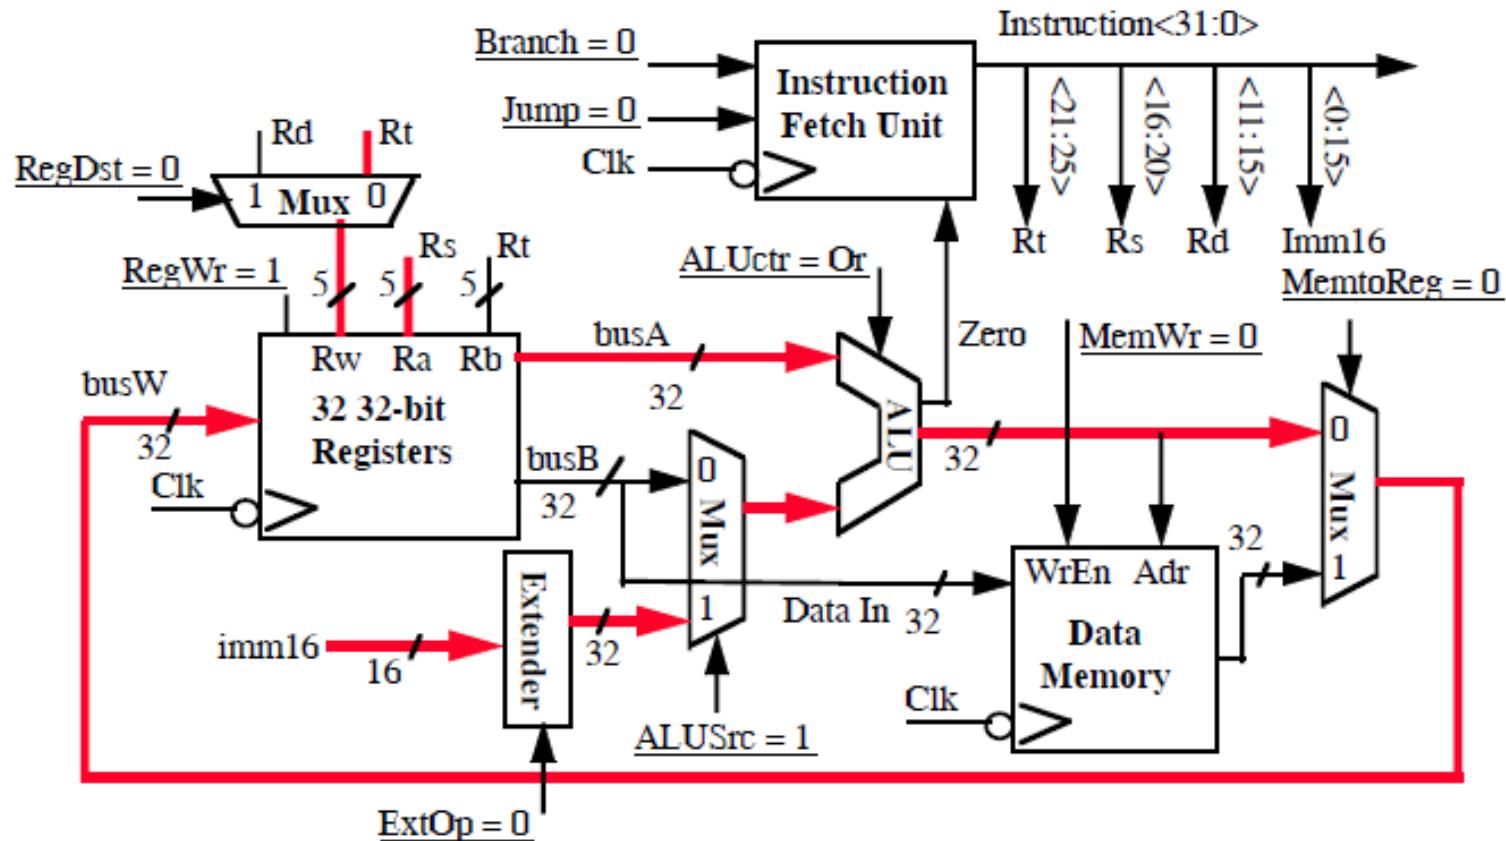
Register – Register Timing



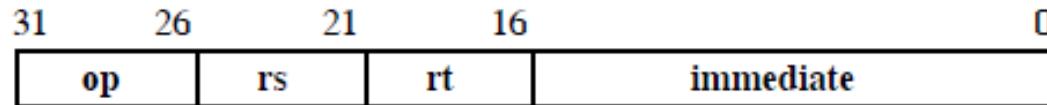
The Single Cycle Datapath during Or Immediate



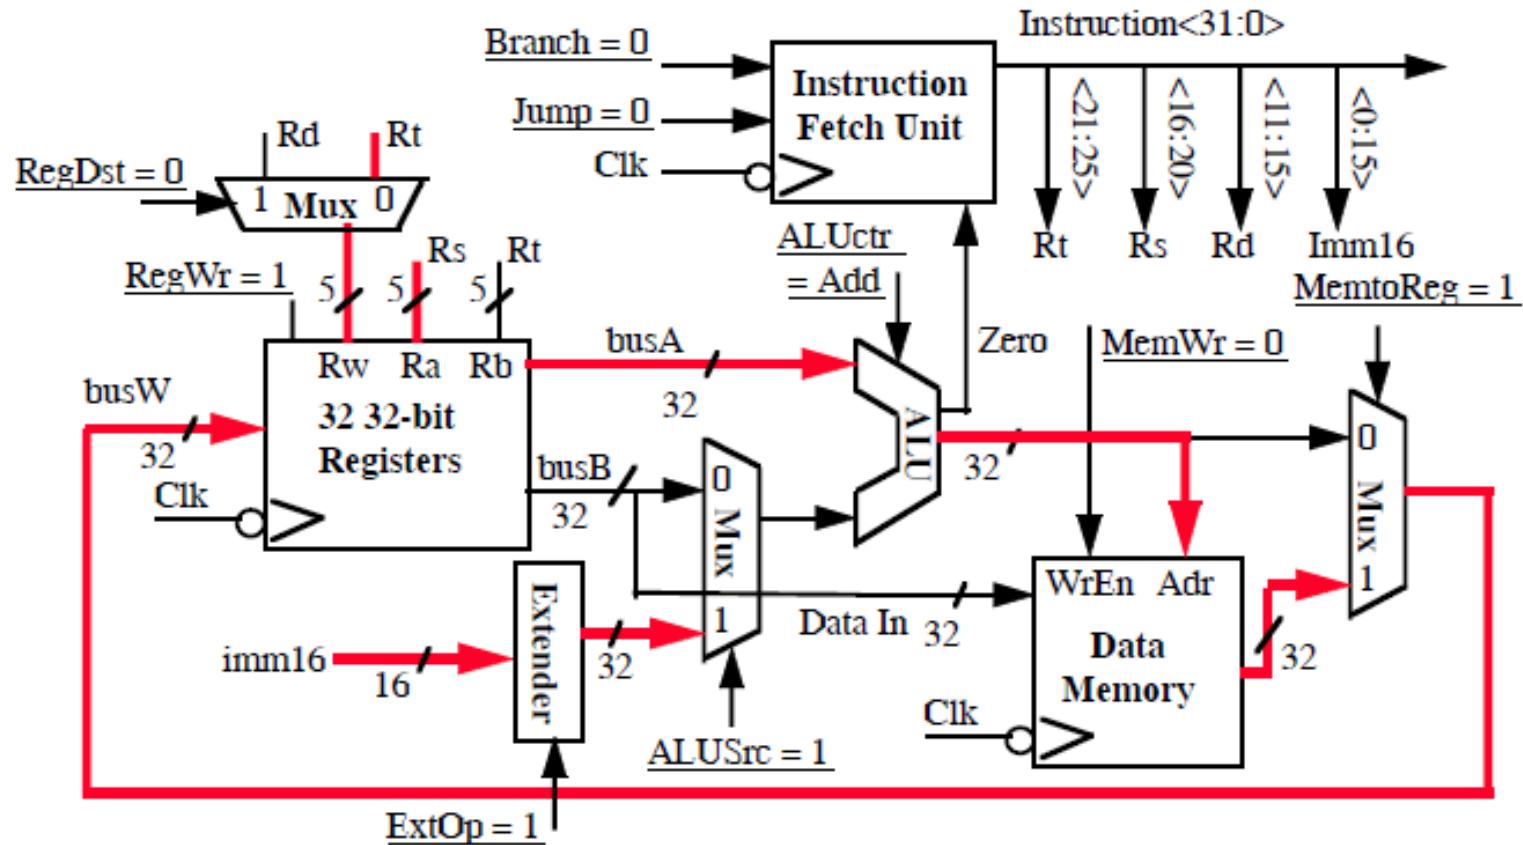
◦ $R[rt] \leftarrow R[rs] \text{ or } \text{ZeroExt}[Imm16]$



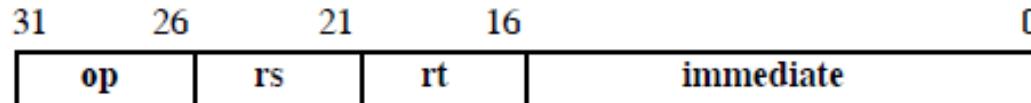
The Single Cycle Datapath during Load



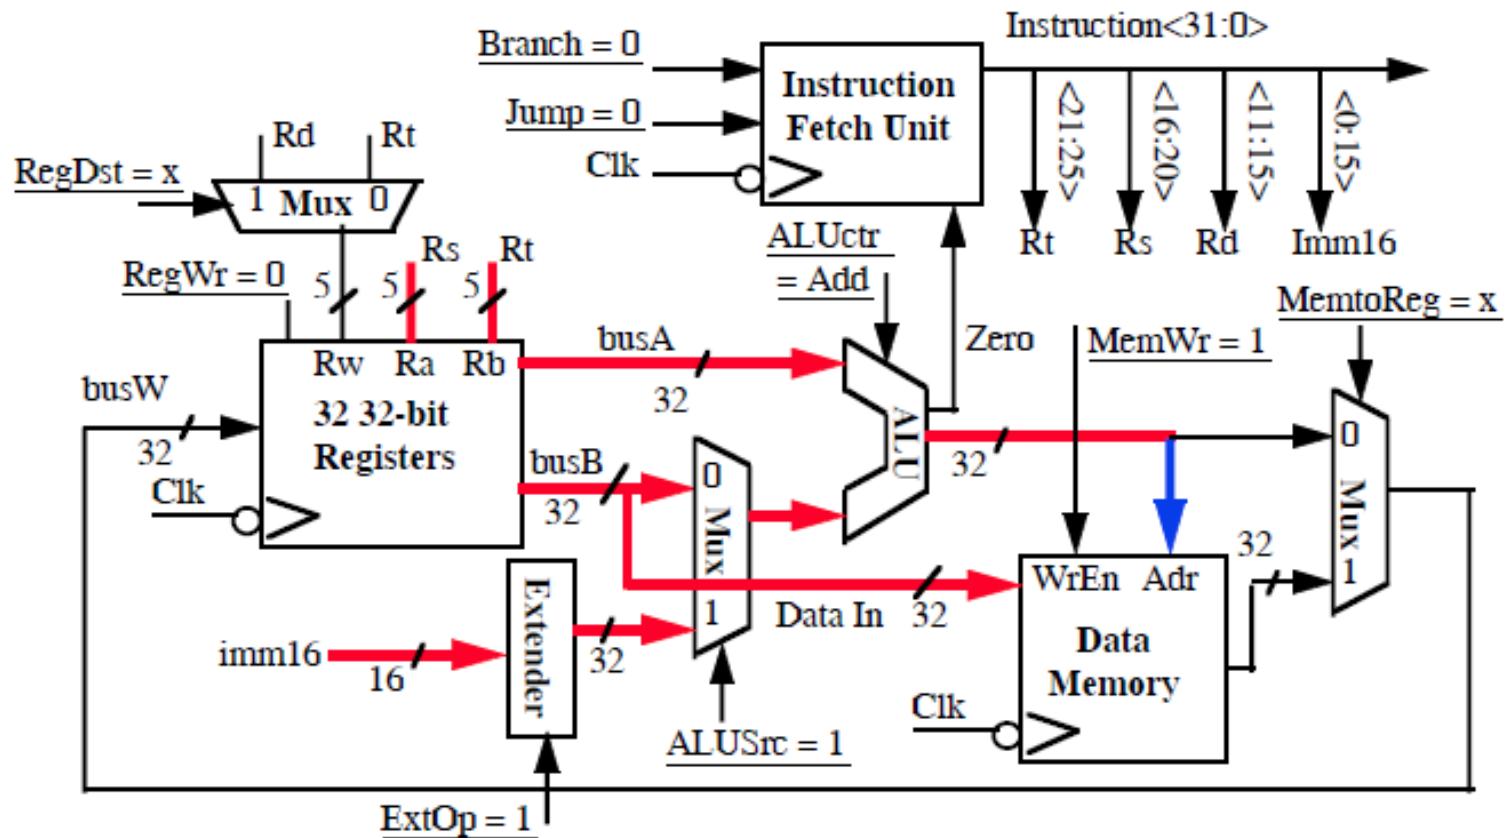
° $R[rt] \leftarrow \text{Data Memory } \{R[rs] + \text{SignExt}[imm16]\}$



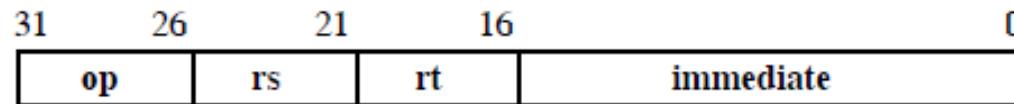
The Single Cycle Datapath during Store



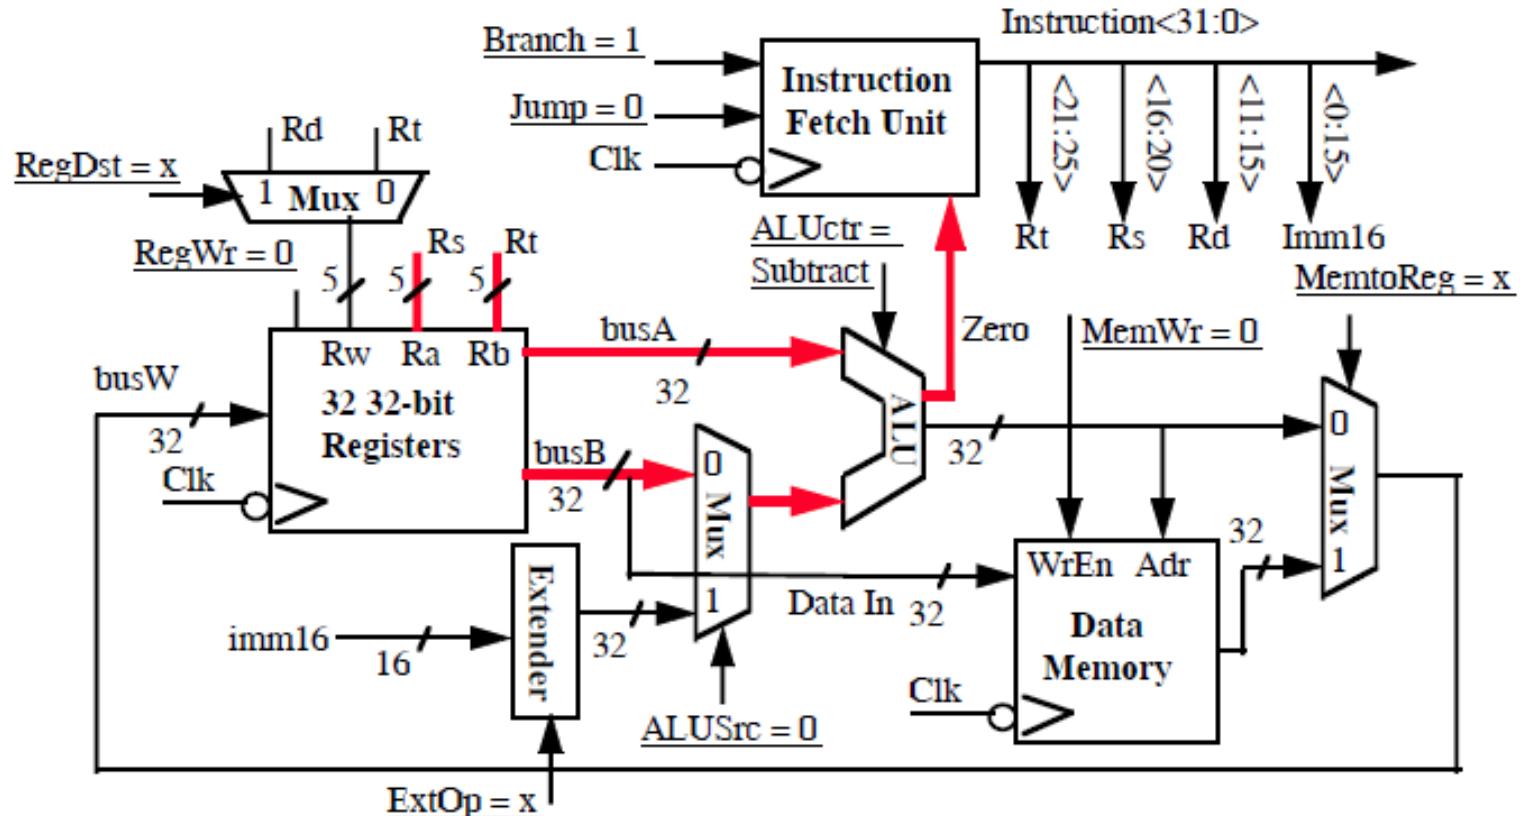
◦ Data Memory {R[rs] + SignExt[imm16]} <- R[rt]



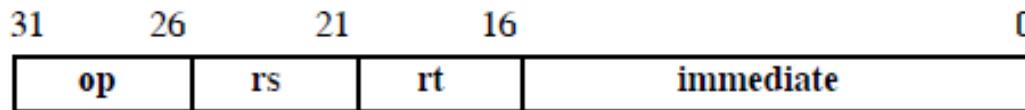
The Single Cycle Datapath during Branch



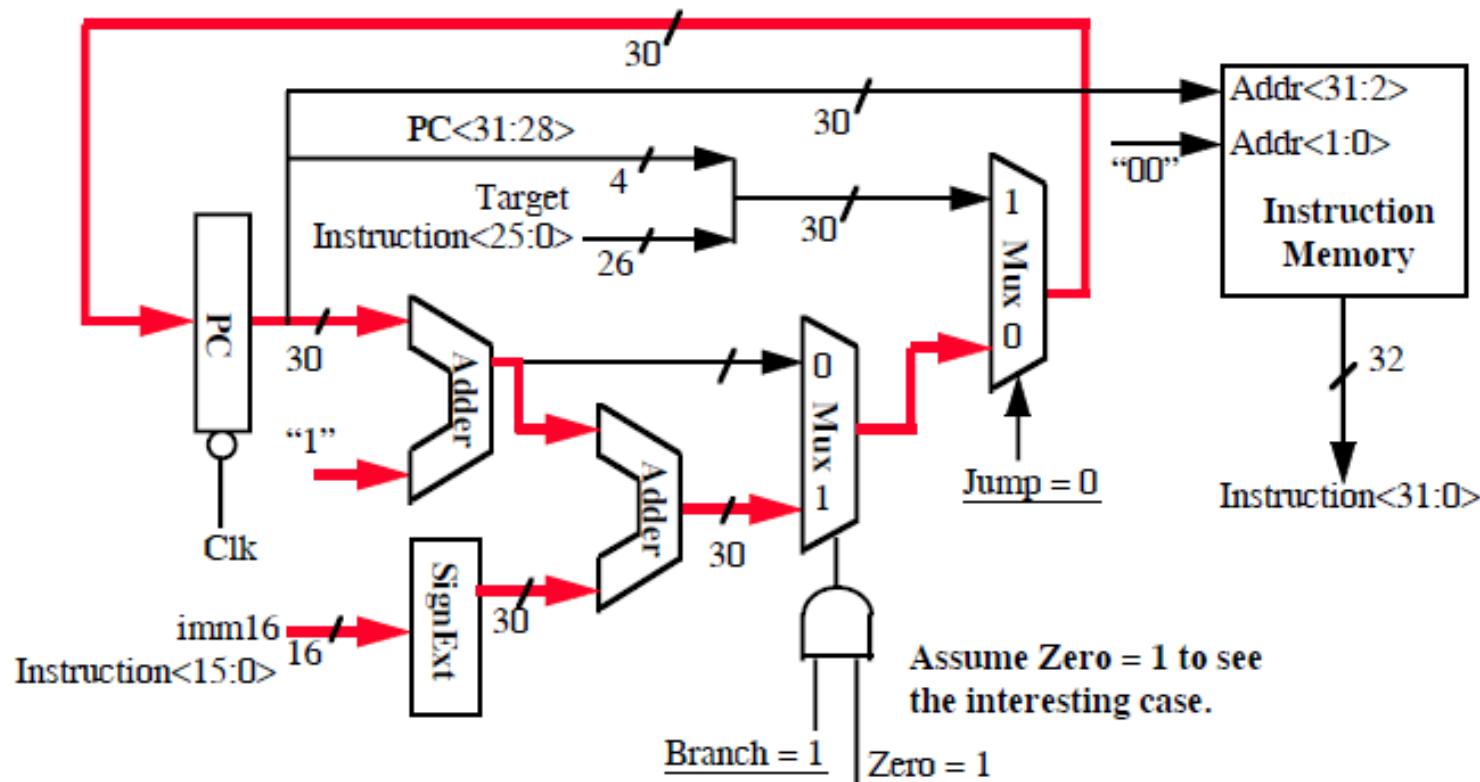
° if $(R[rs] - R[rt]) == 0$ then Zero $\leftarrow 1$; else Zero $\leftarrow 0$



Instruction Fetch Unit at the End of Branch

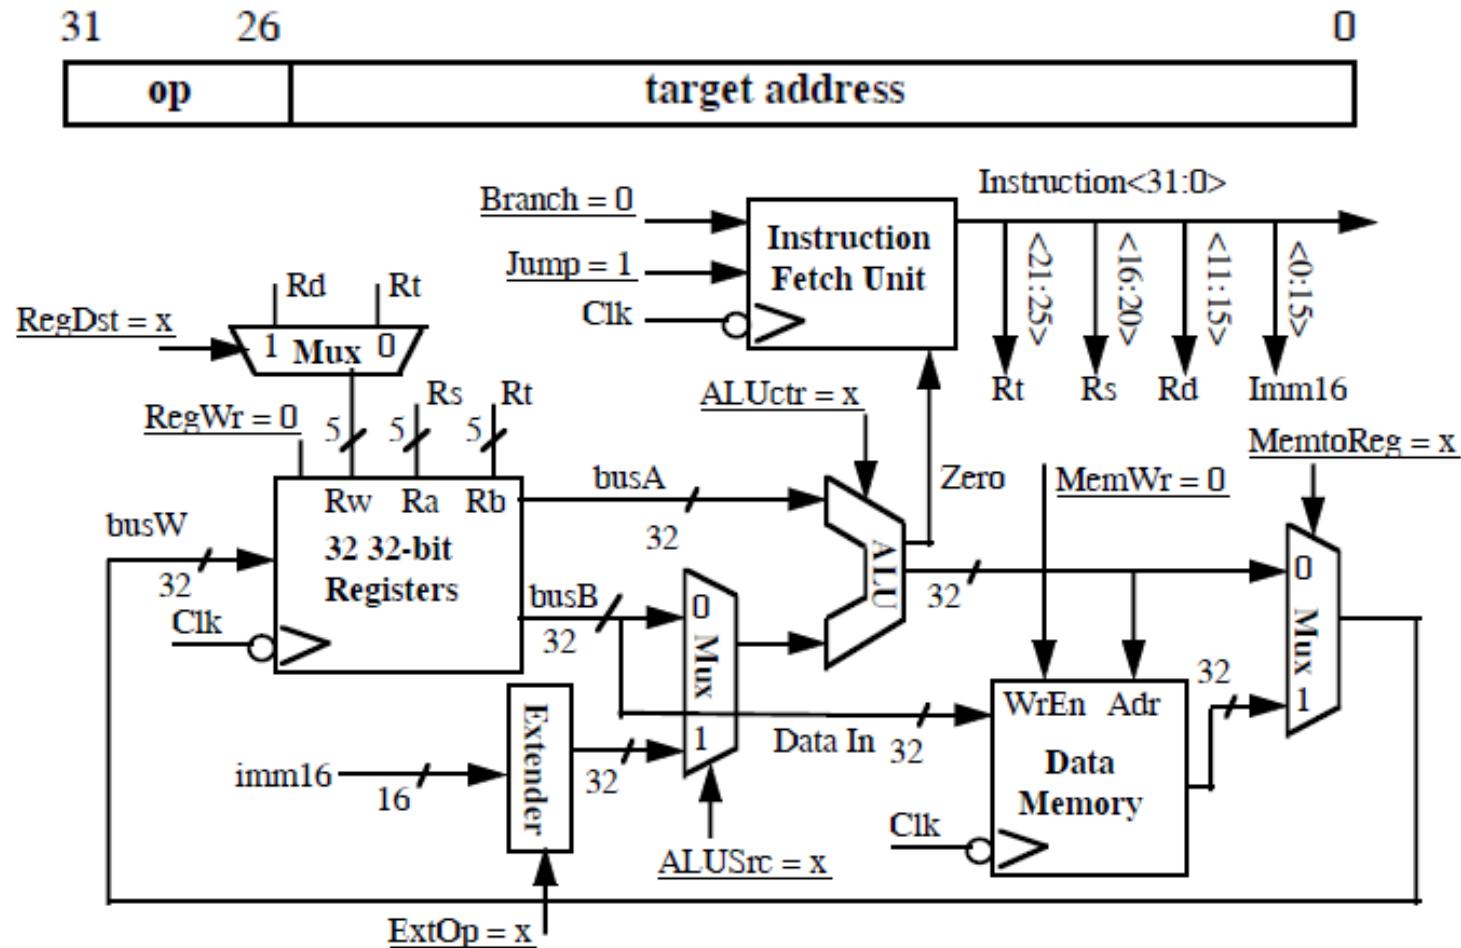


- if (Zero == 1) then $PC = PC + 4 + \text{SignExt}[imm16]*4$; else $PC = PC + 4$

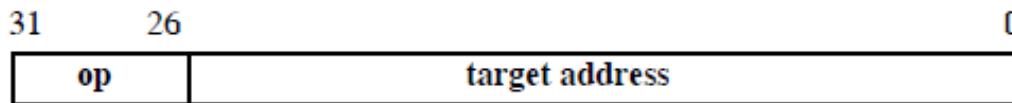


The Single Cycle Datapath during Jump

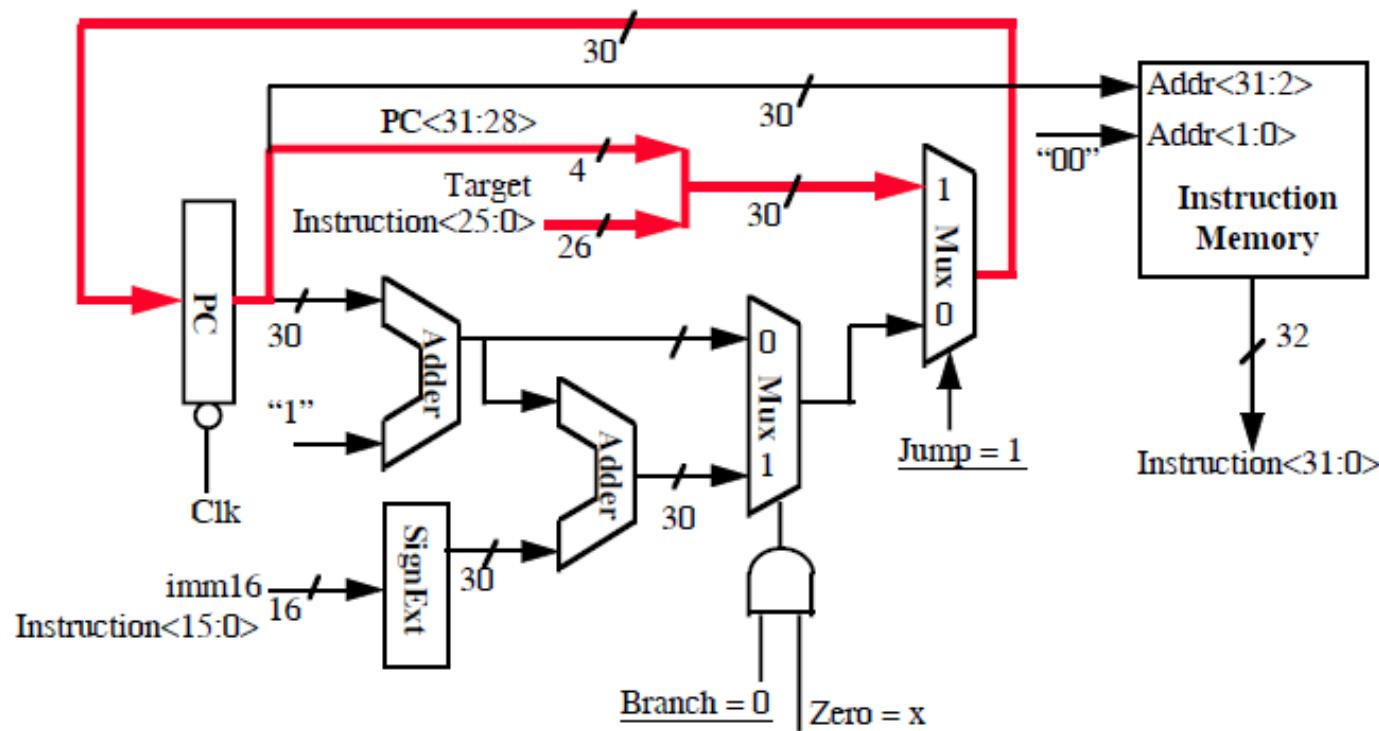
- Nothing to do! Make sure control signals are set correctly!



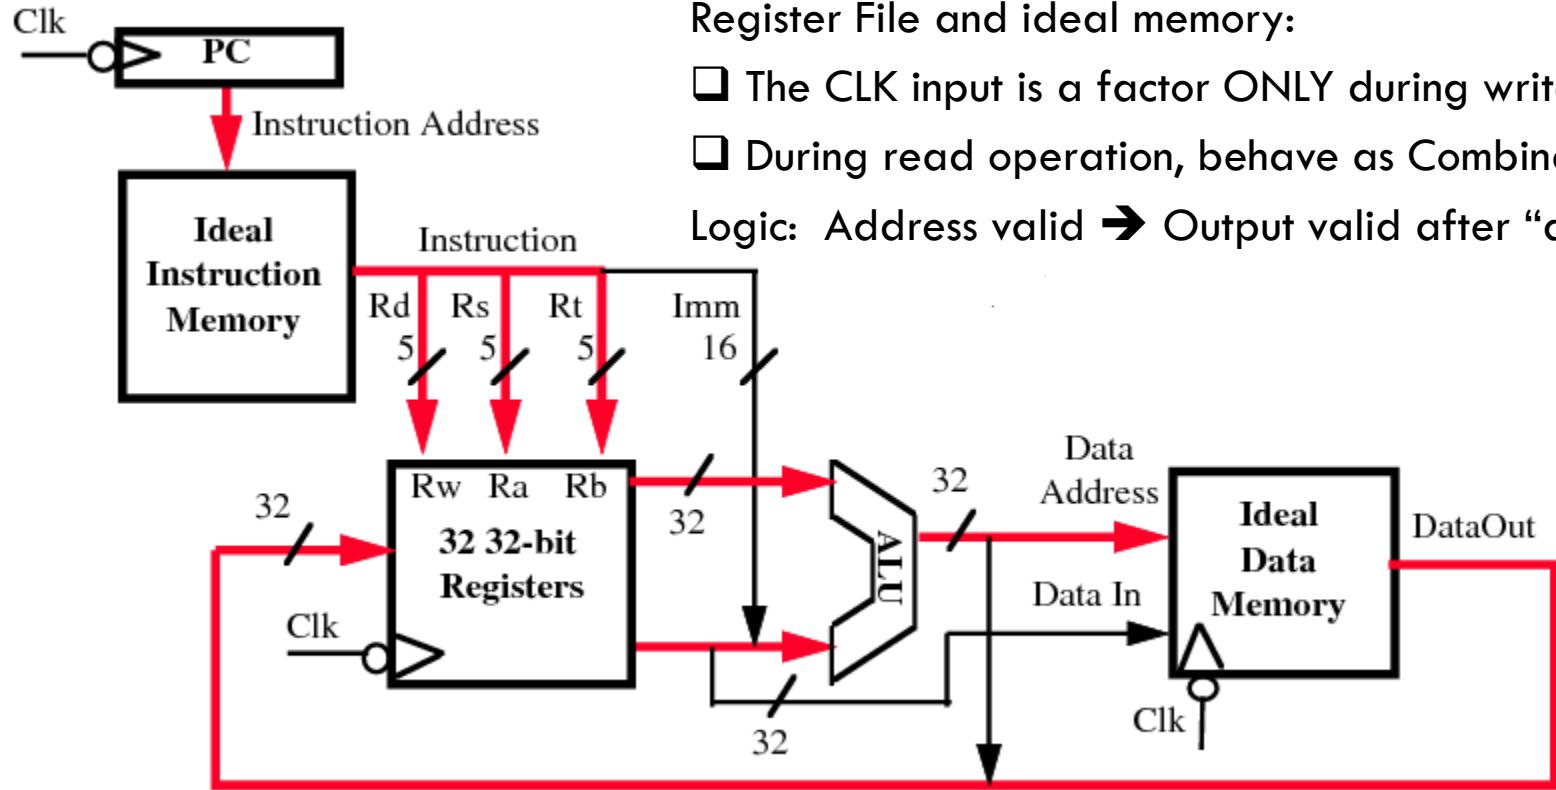
Instruction Fetch Unit at the End of Jump



◦ $PC \leftarrow PC<31:29> \text{ concat } target<25:0> \text{ concat } "00"$



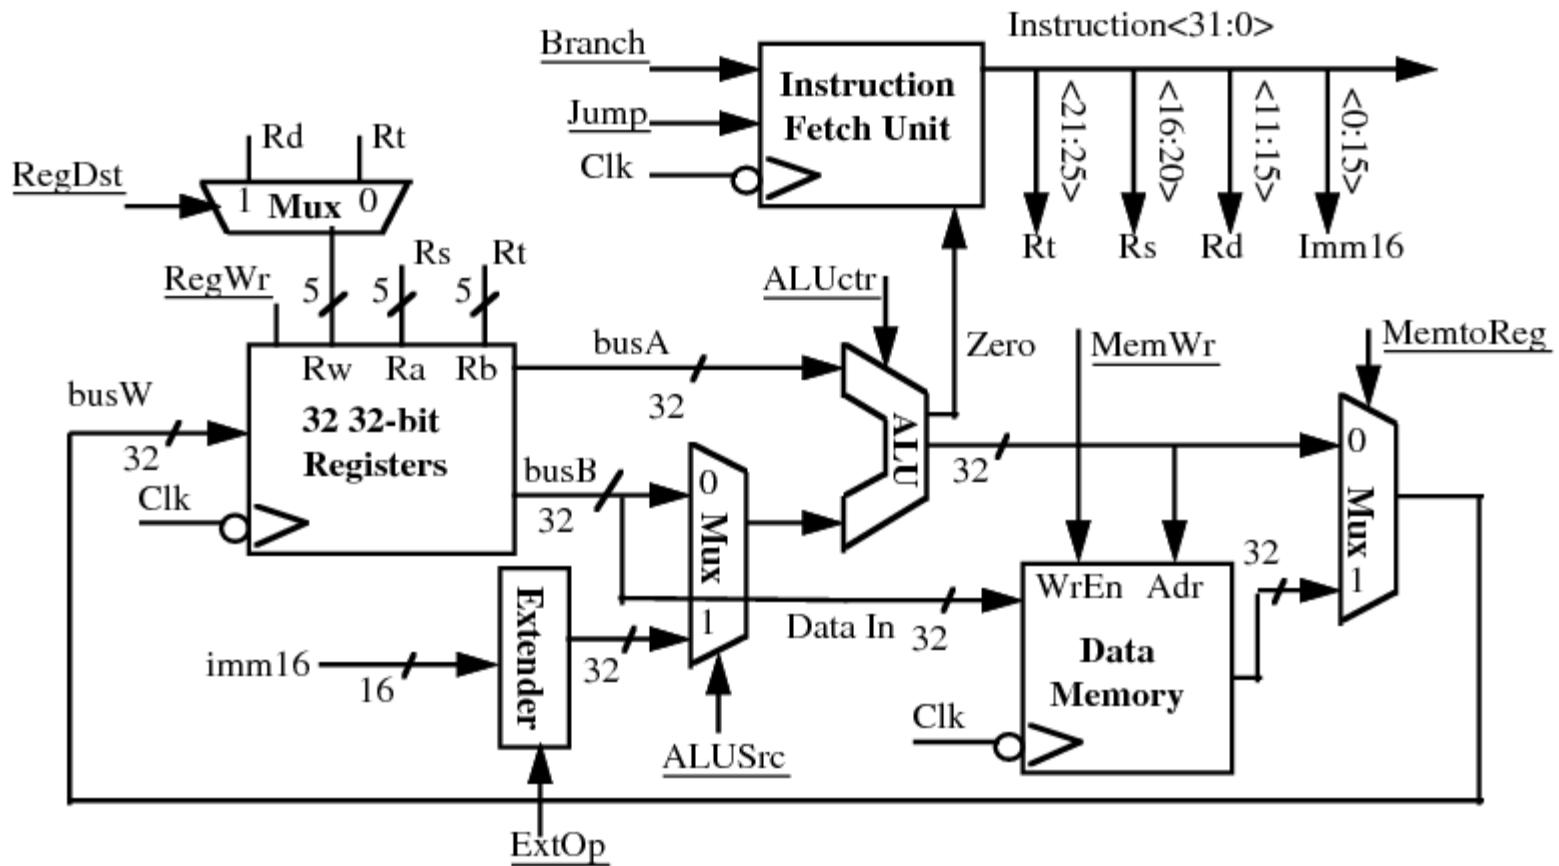
An Abstract View of the Critical Path



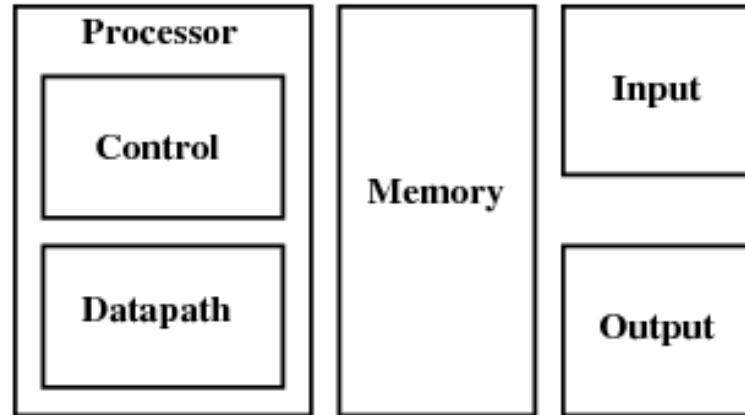
Critical Path (Load Operation) = PC's Clk-to-Q + Instruction Memory's Access Time + Register File's Access Time + ALU to Perform 32-bit Add + Data Memory Access Time + Setup Time for Register File Write

Putting it all together: A Single Cycle Datapath

We have everything except control signals (underline)



The Big Picture: Where are we Now?



- The Five Classic Components of a Computer
- Next Topic: Control Path Design

A Summary of Control Signals

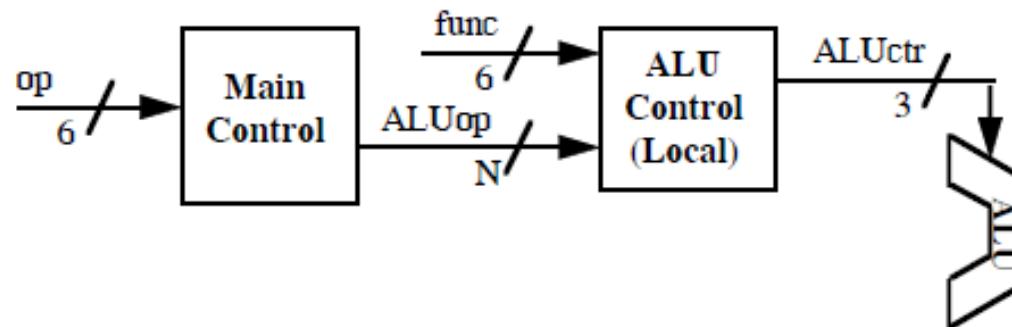
Diagram showing control signal assignments based on **func** and **op**:

	func	10 0000	10 0010	We Don't Care :-)				
	op	00 0000	00 0000	00 1101	10 0011	10 1011	00 0100	00 0010
		add	sub	ori	lw	sw	beq	jump
RegDst		1	0	0	0	x	x	x
ALUSrc		0	0	1	1	1	0	x
MemtoReg		0	0	0	1	x	x	x
RegWrite		1	1	1	1	0	0	0
MemWrite		0	0	0	0	1	0	0
Branch		0	0	0	0	0	1	0
Jump		0	0	0	0	0	0	1
ExtOp		x	x	0	1	1	x	x
ALUctr<2:0>		Add	Subtract	Or	Add	Add	Subtract	xxx

	31	26	21	16	11	6	0
R-type	op	rs	rt	rd	shamt	funct	add, sub
I-type	op	rs	rt	immediate			ori, lw, sw, beq
J-type	op	target address					jump

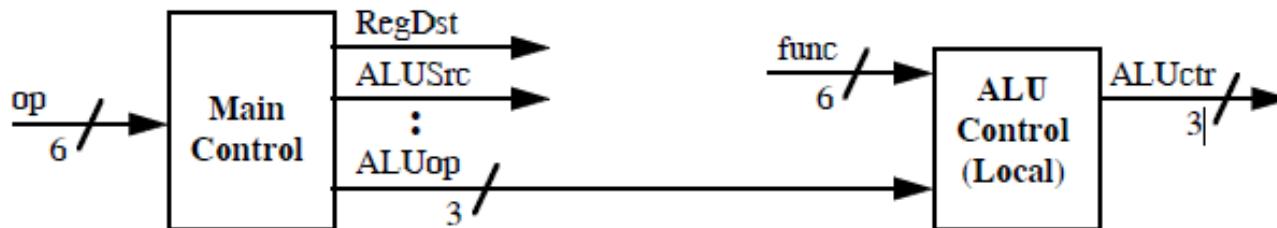
The Concept of Local Decoding

op	00 0000	00 1101	10 0011	10 1011	00 0100	00 0010
	R-type	ori	lw	sw	beq	jump
RegDst	1	0	0	x	x	x
ALUSrc	0	1	1	1	0	x
MemtoReg	0	0	1	x	x	x
RegWrite	1	1	1	0	0	0
MemWrite	0	0	0	1	0	0
Branch	0	0	0	0	1	0
Jump	0	0	0	0	0	1
ExtOp	x	0	1	1	x	x
ALUop<N:0>	"R-type"	Or	Add	Add	Subtract	xxx



Key Idea: Two levels of Control logic.

The “Truth Table” for the Main Control

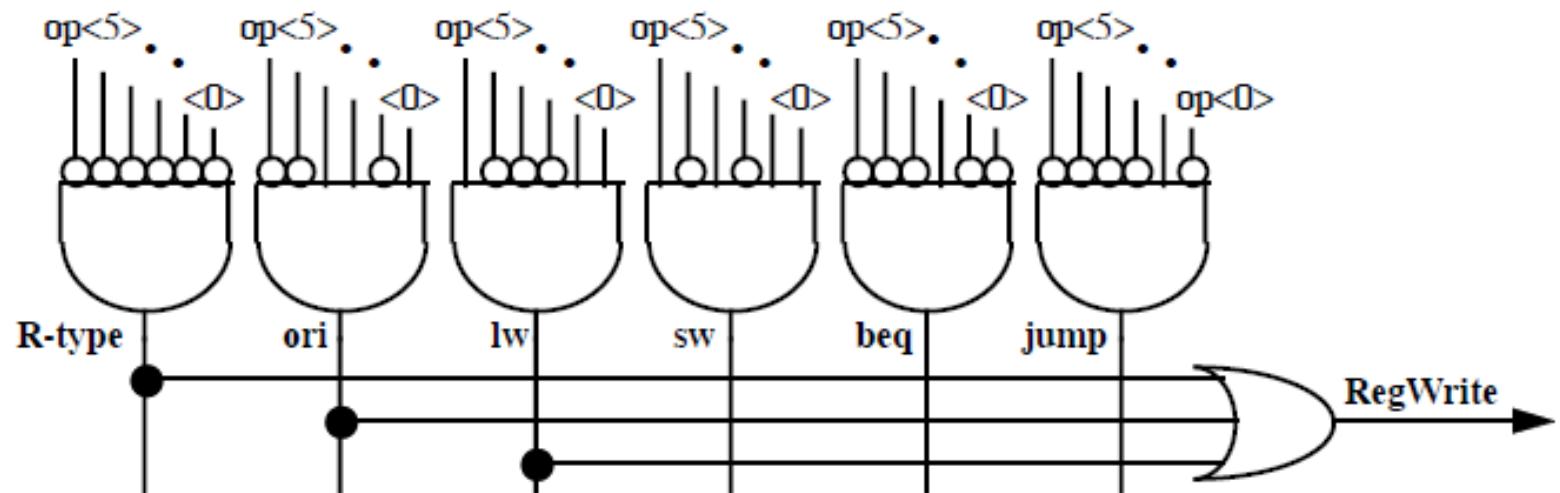


op	00 0000	00 1101	10 0011	10 1011	00 0100	00 0010
	R-type	ori	lw	sw	beq	jump
RegDst	1	0	0	x	x	x
ALUSrc	0	1	1	1	0	x
MemtoReg	0	0	1	x	x	x
RegWrite	1	1	1	0	0	0
MemWrite	0	0	0	1	0	0
Branch	0	0	0	0	1	0
Jump	0	0	0	0	0	1
ExtOp	x	0	1	1	x	x
ALUop (Symbolic)	“R-type”	Or	Add	Add	Subtract	xxx
ALUop <2>	1	0	0	0	0	x
ALUop <1>	0	1	0	0	0	x
ALUop <0>	0	0	0	0	1	x

Question: Can you write the truth table for the ALU control keeping in mind the ALU we designed in the class?

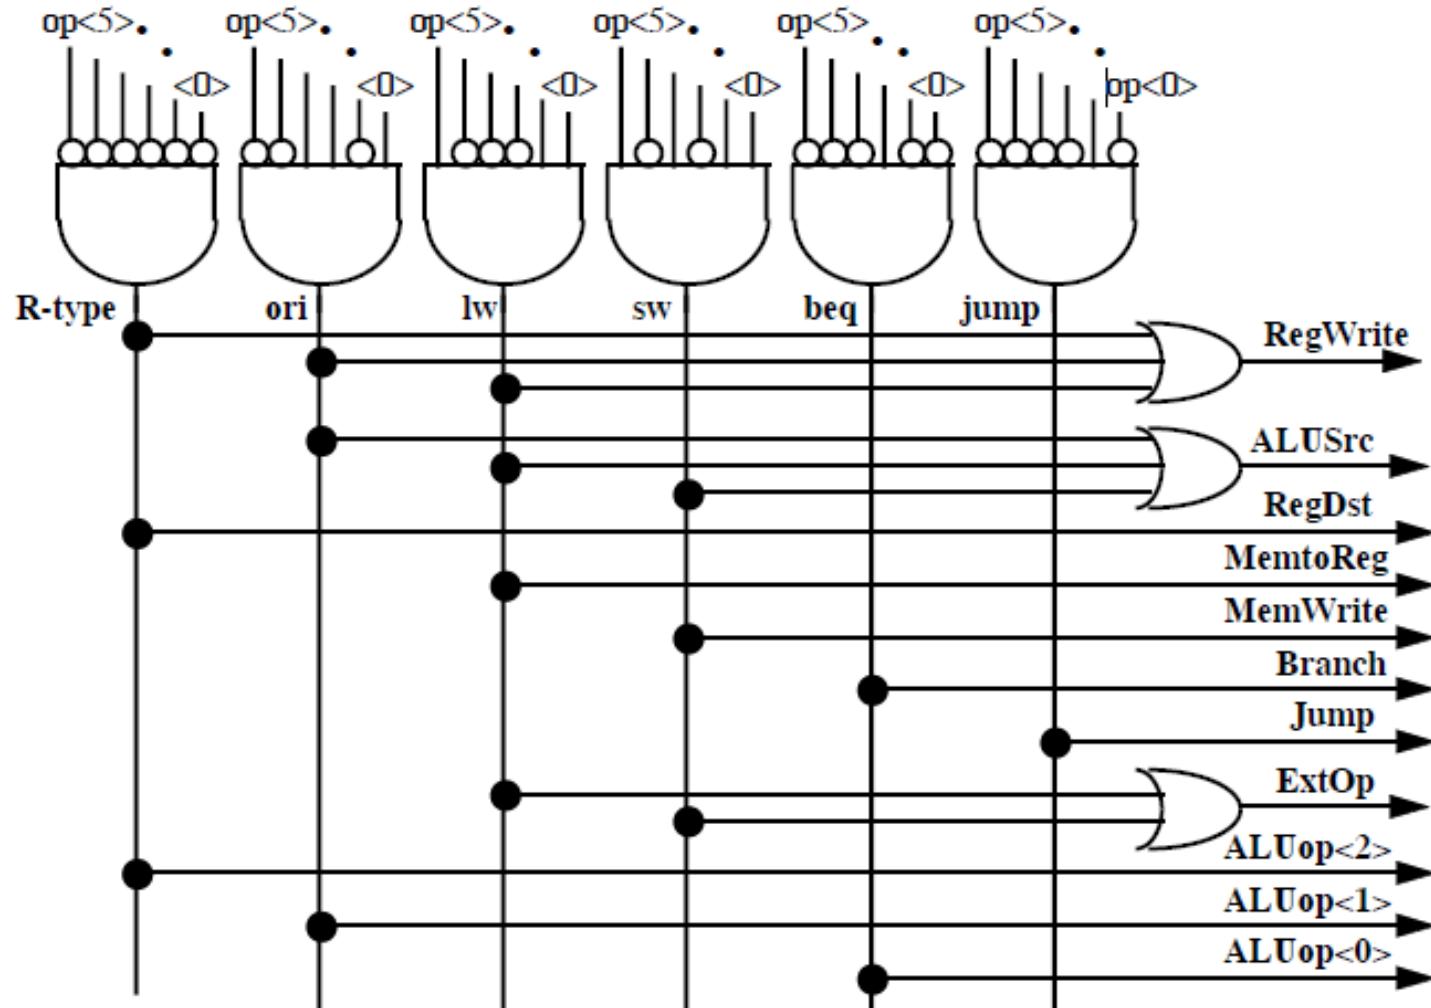
The “Truth Table” for RegWrite

op	00 0000	00 1101	10 0011	10 1011	00 0100	00 0010
R-type	R-type	ori	lw	sw	beq	jump
RegWrite	1	1	1	x	x	x

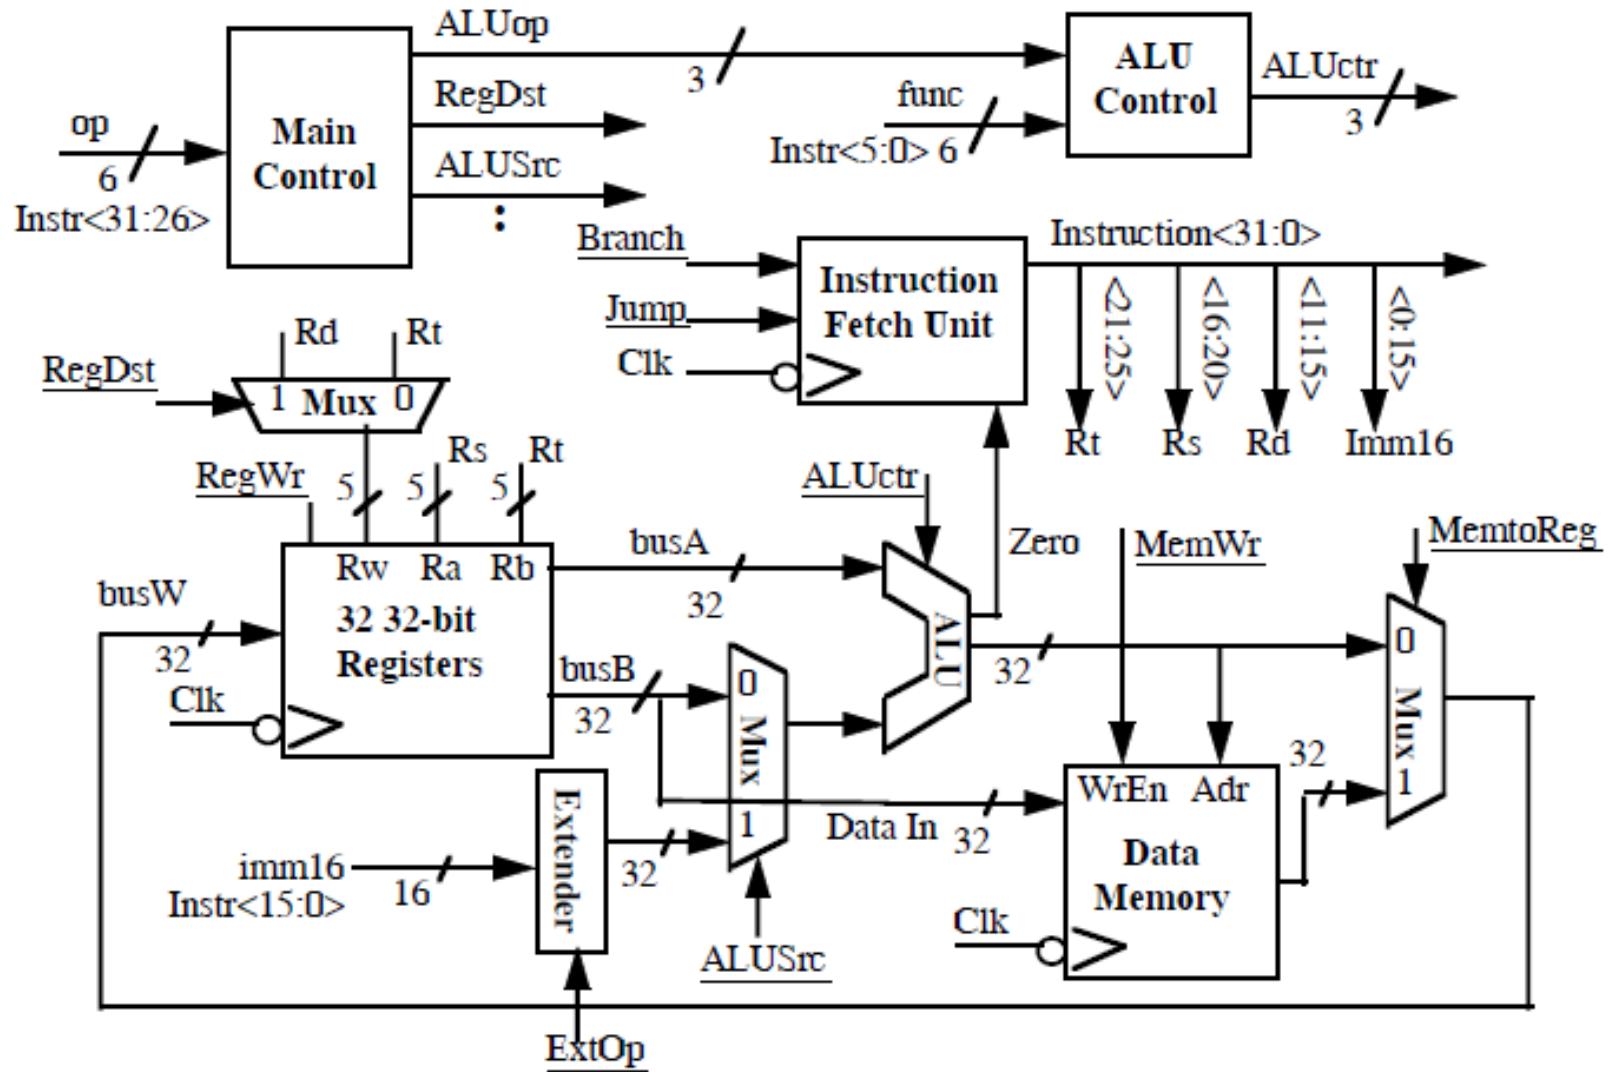


PLA Implementation of Main Control

Hmm! What is PLA?



Putting it All Together: A Single Cycle Processor

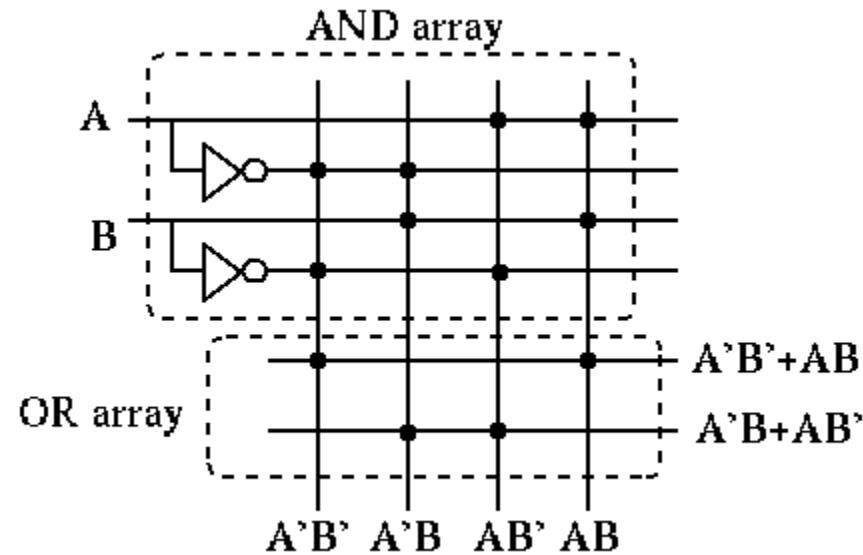
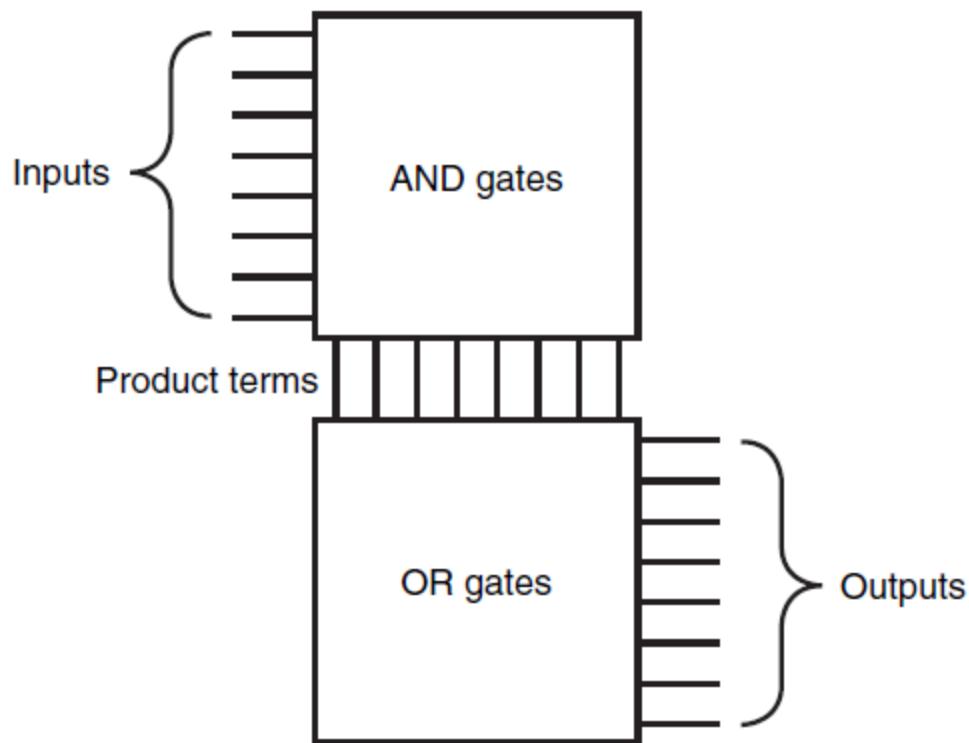


Drawback of this Single Cycle Processor

- ❑ Long cycle time:
 - ❑ Cycle time must be long enough for the load instruction:
 - PC's Clock -to-Q +
 - Instruction Memory Access Time +
 - Register File Access Time +
 - ALU Delay (address calculation) +
 - Data Memory Access Time +
 - Register File Setup Time
 - ❑ Cycle time is much longer than needed for all other instructions
- We are assuming
Clock Skew is zero

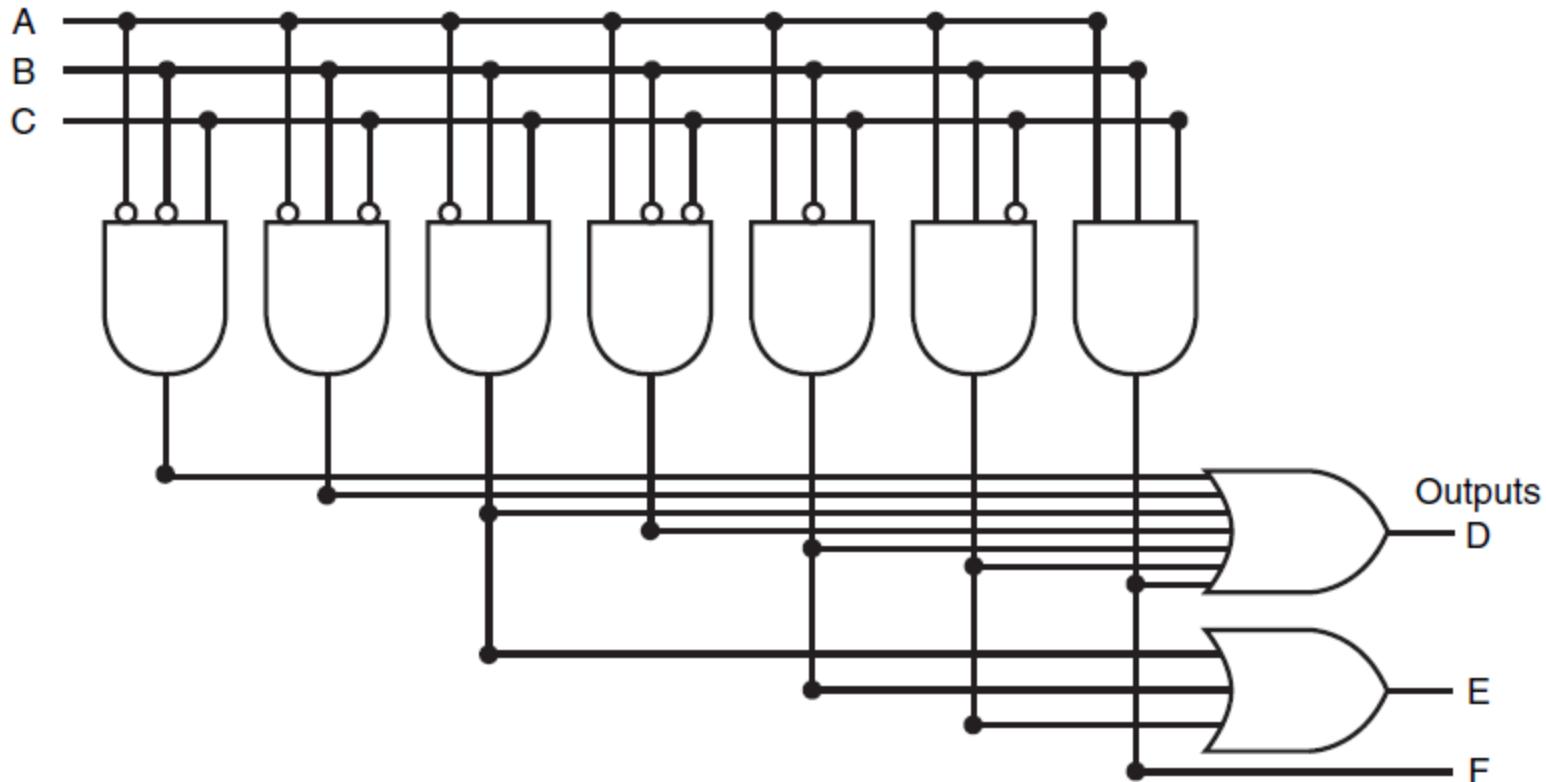
Programmable Logic Arrays

- PLAs can be used to realize combinational circuits



PLAs

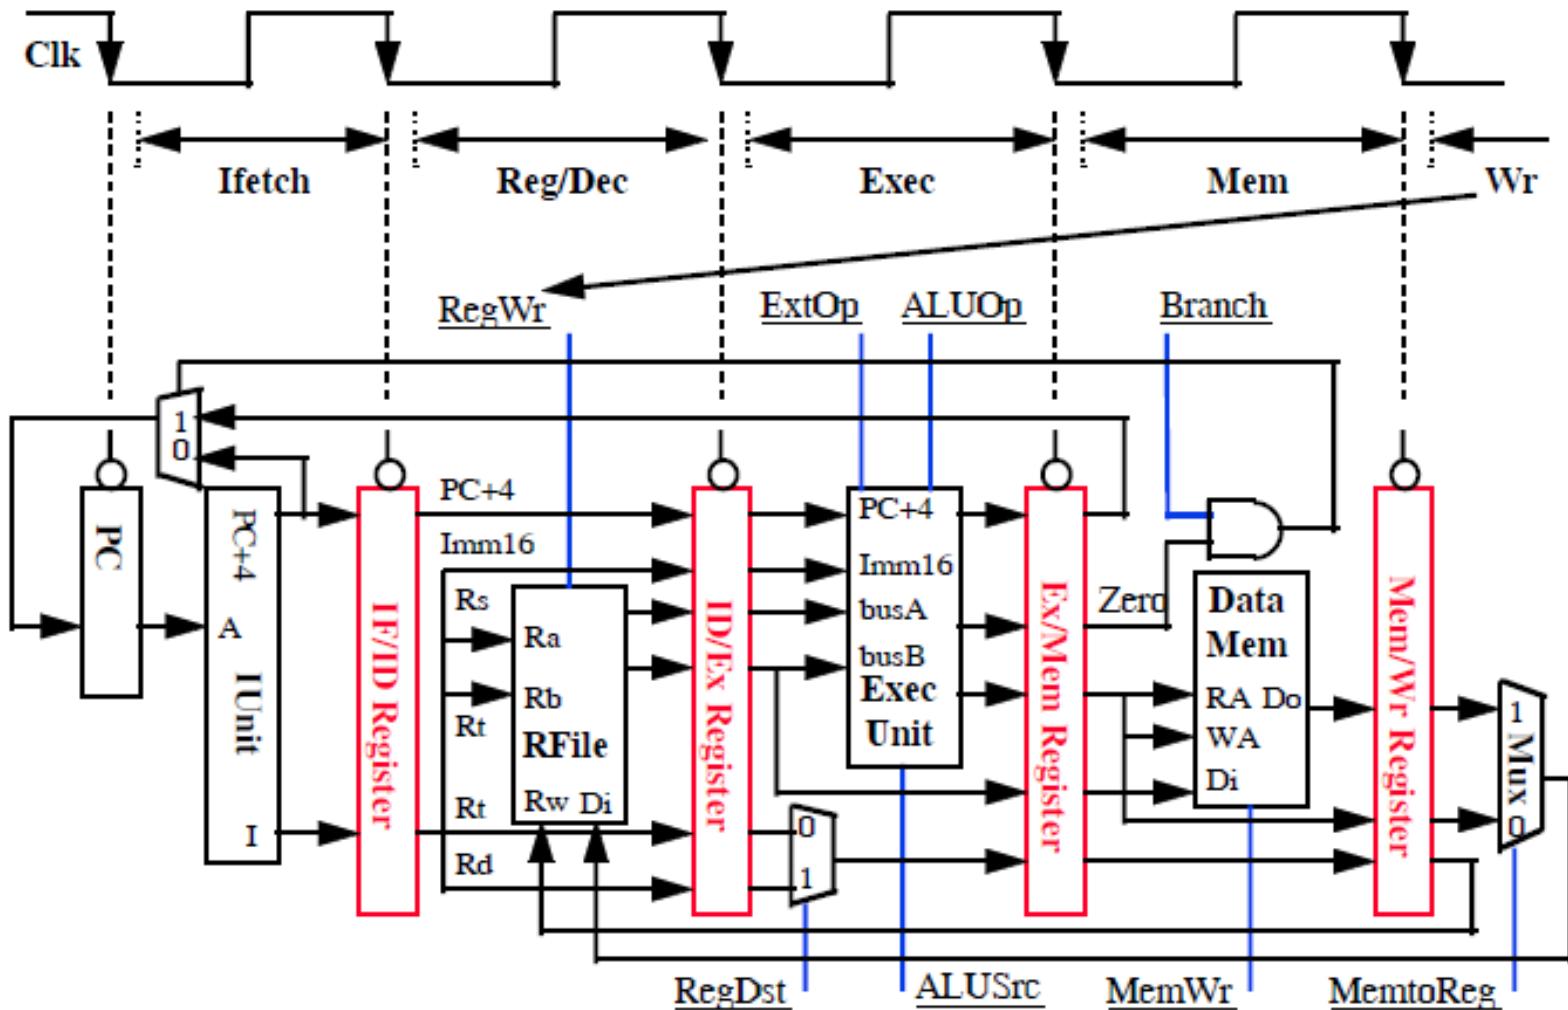
Inputs



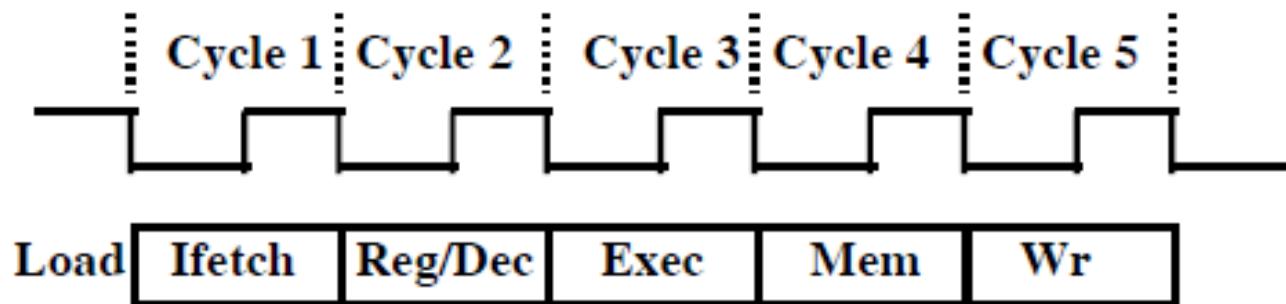
Acknowledgment: Almost all of these slides are based on Dave Patterson's CS152 Lecture Slides at UC, Berkeley.

COMPUTER SYSTEMS ORGANIZATION

A Pipelined Datapath



The Five Stages of Load

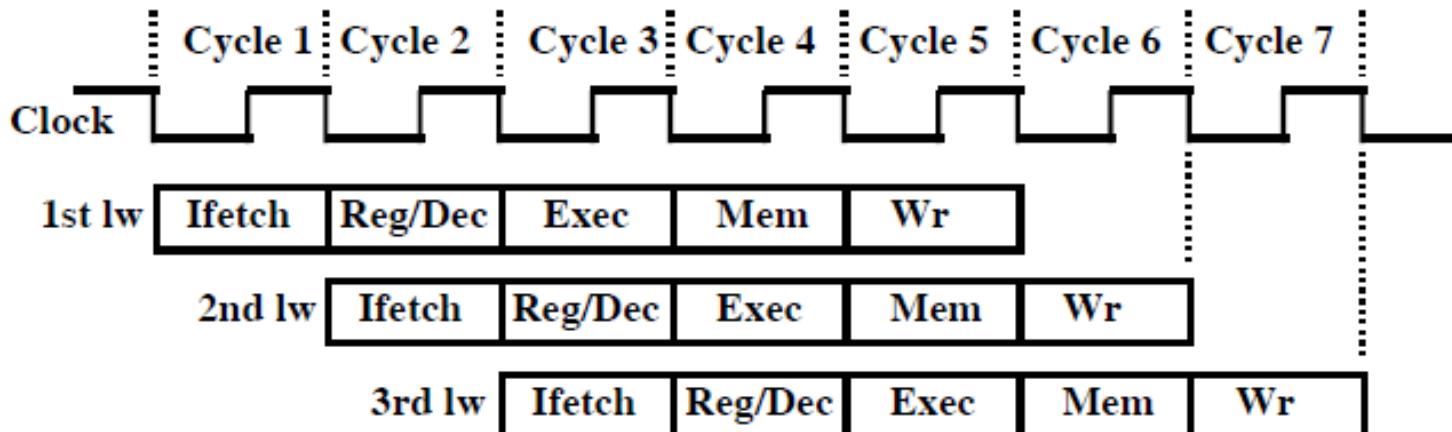


- **Ifetch: Instruction Fetch**
 - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: Calculate the memory address**
- **Mem: Read the data from the Data Memory**
- **Wr: Write the data back to the register file**

Key Ideas Behind Pipelining

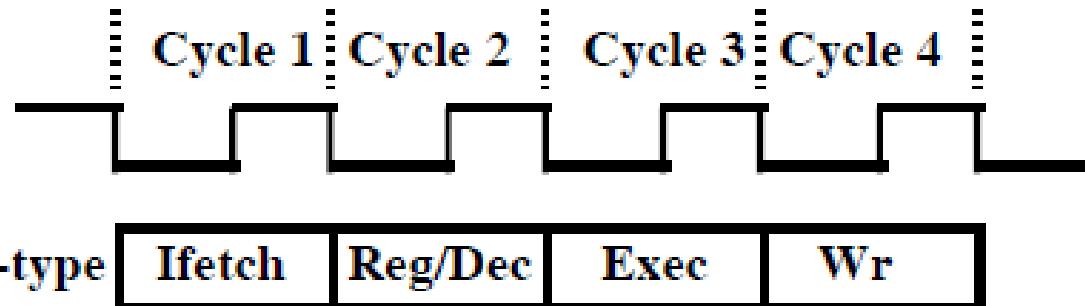
- ❑ The load instruction has 5 stages:
 - ❑ Five independent functional units to work on each stage
 - ❑ Each functional unit is used only once
- ❑ The 2nd load can start as soon as the 1st finishes its **Ifetch** stage
- ❑ Each load still takes five cycles to complete
- ❑ The throughput, however, is much higher

Pipelining the Load Instruction



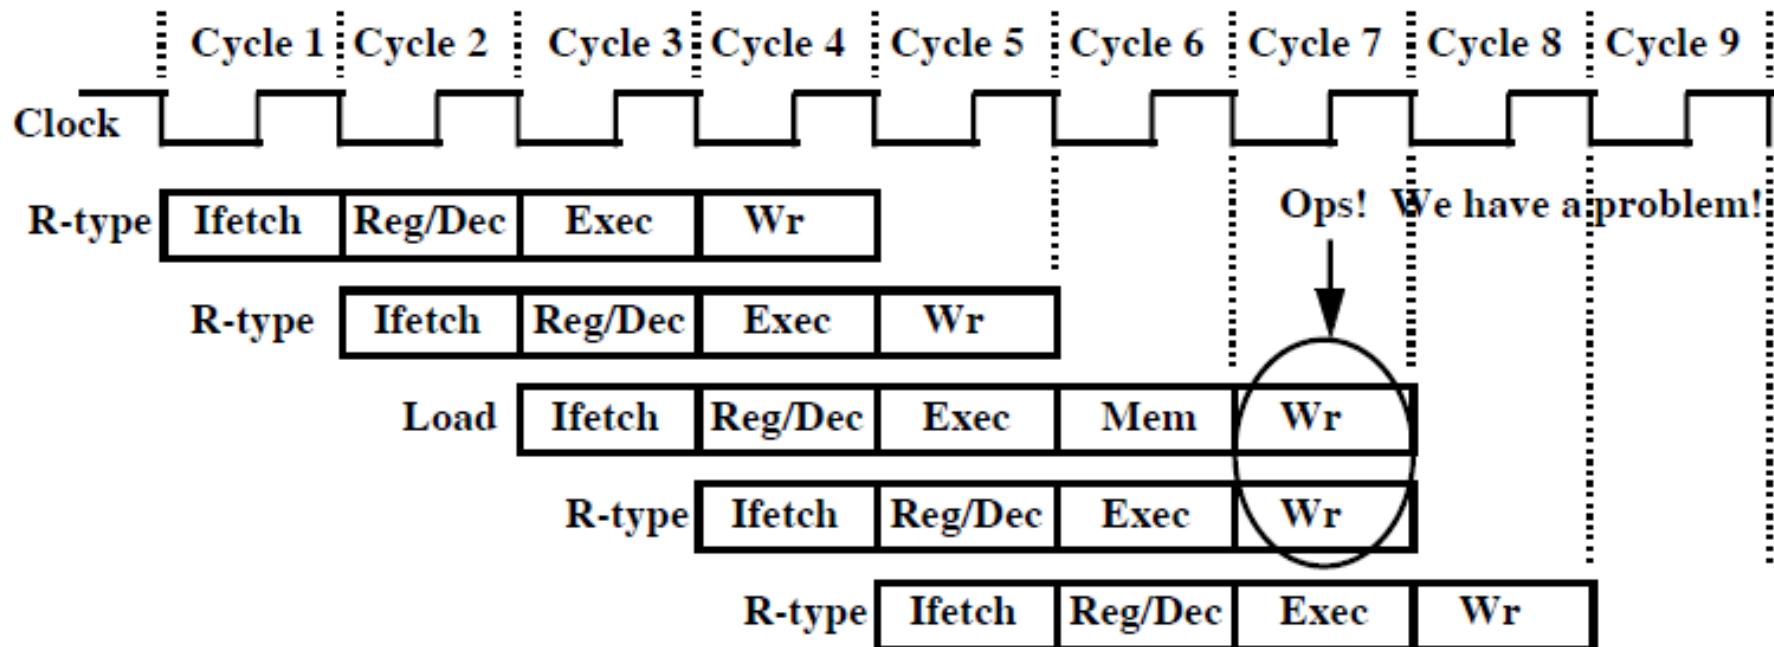
- The five independent functional units in the pipeline datapath are:
 - Instruction Memory for the Ifetch stage
 - Register File’s Read ports (bus A and busB) for the Reg/Dec stage
 - ALU for the Exec stage
 - Data Memory for the Mem stage
 - Register File’s Write port (bus W) for the Wr stage
- One instruction enters the pipeline every cycle
 - One instruction comes out of the pipeline (complete) every cycle
 - The “Effective” Cycles per Instruction (CPI) is 1

The Four Stages of R-type



- **Ifetch: Instruction Fetch**
 - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: ALU operates on the two register operands**
- **Wr: Write the ALU output back to the register file**

Pipelining the R-type and Load Instruction

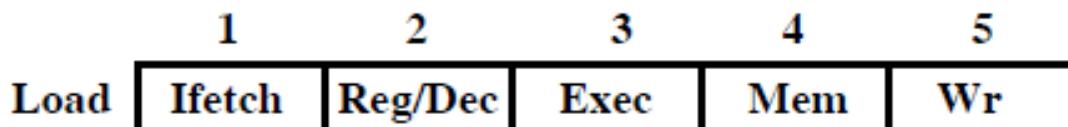


- ° We have a problem:
 - Two instructions try to write to the register file at the same time!

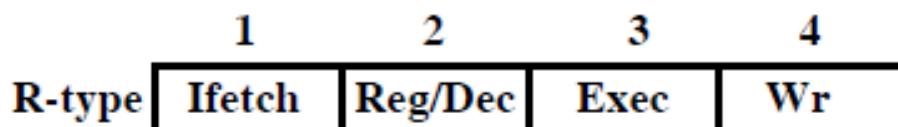
Structural Hazard: Two instructions require access to the same functional unit.

Important Observation

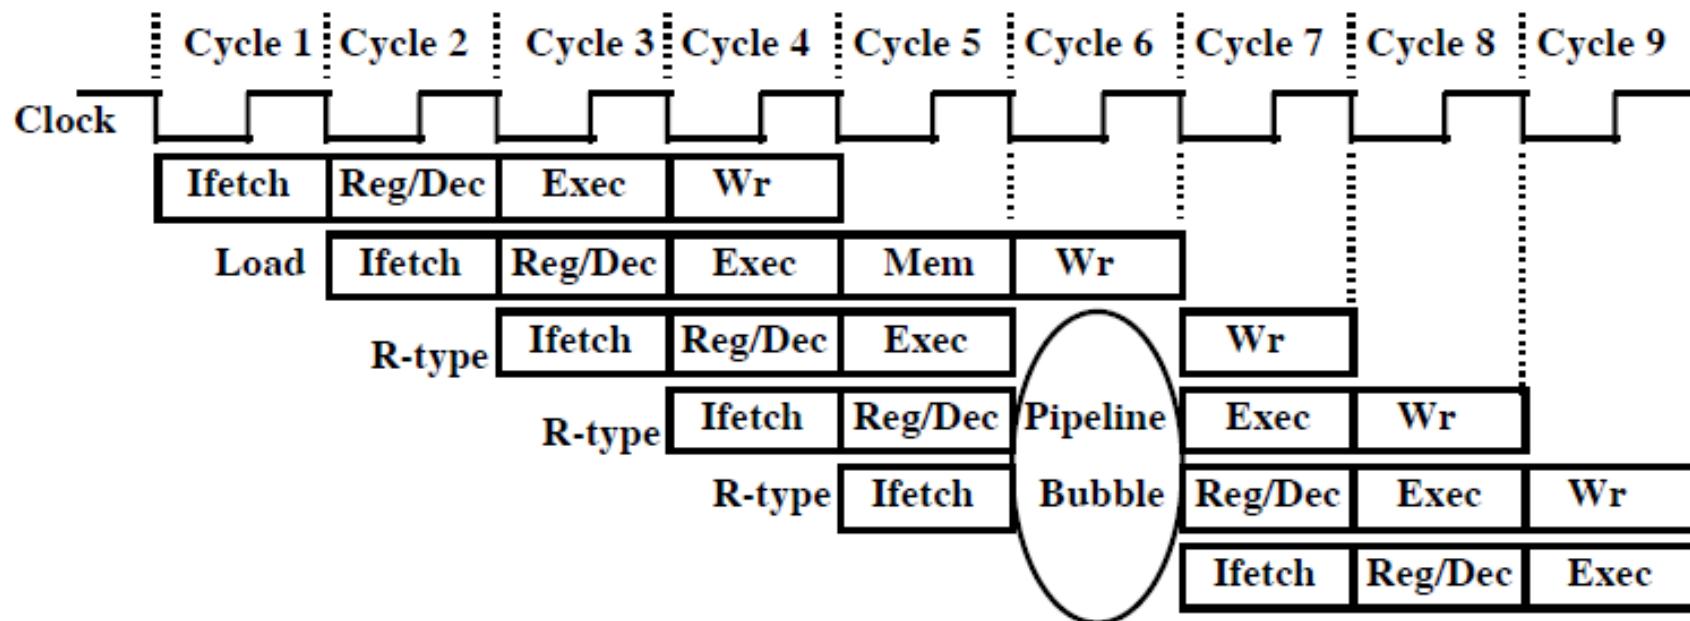
- ° Each functional unit can only be used once per instruction
- ° Each functional unit must be used at the same stage for all instructions:
 - Load uses Register File's Write Port during its 5th stage



- R-type uses Register File's Write Port during its 4th stage



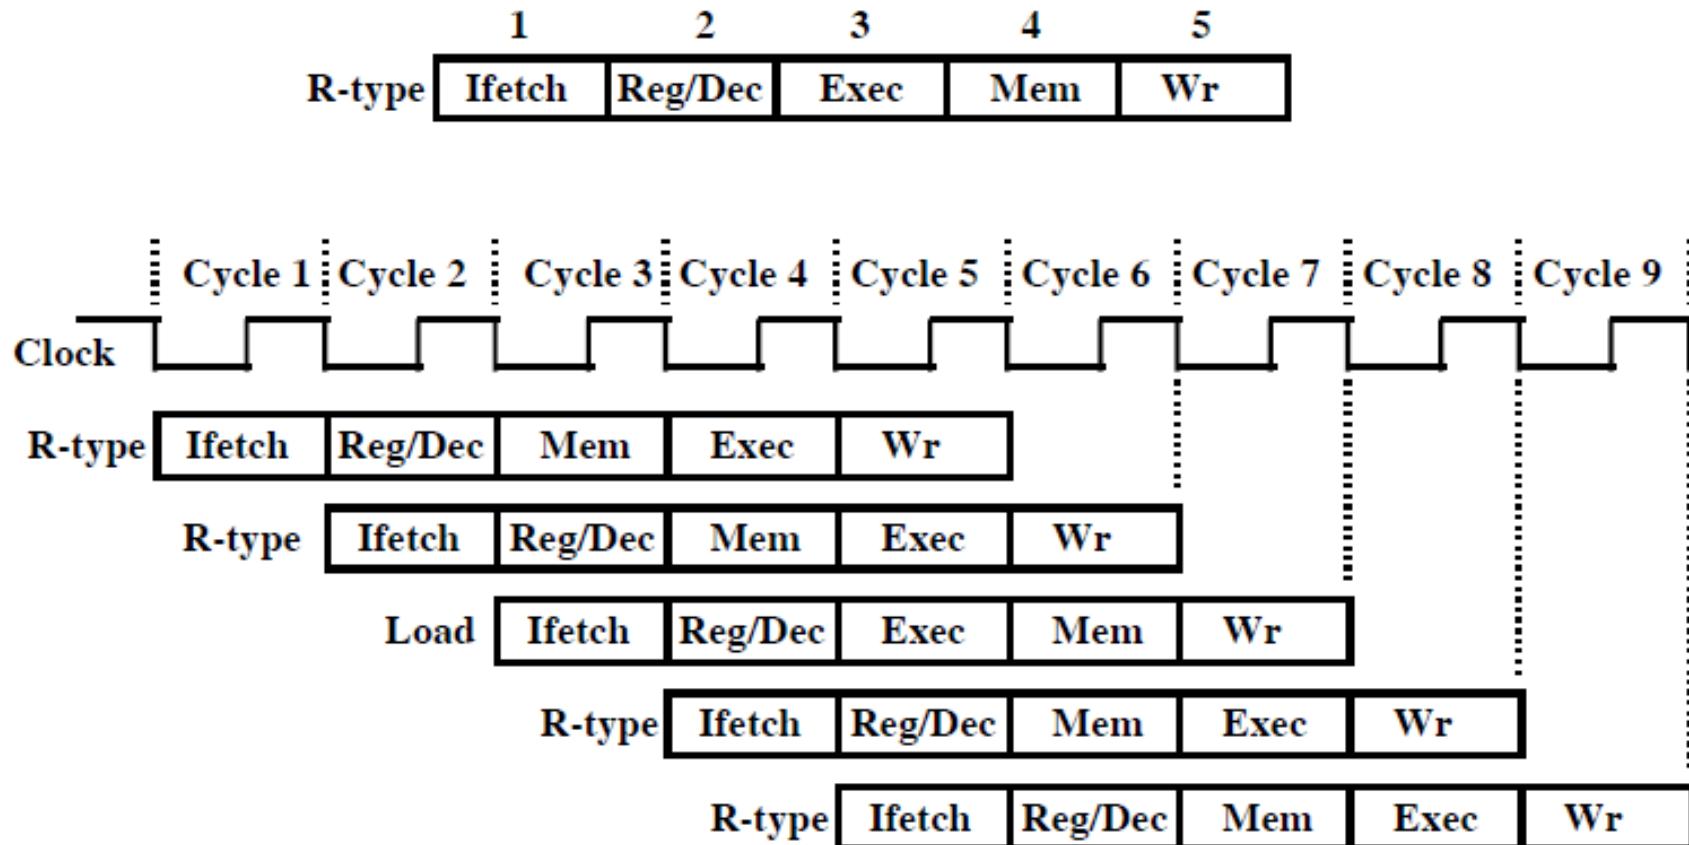
Solution 1: Insert “Bubble” into the Pipeline



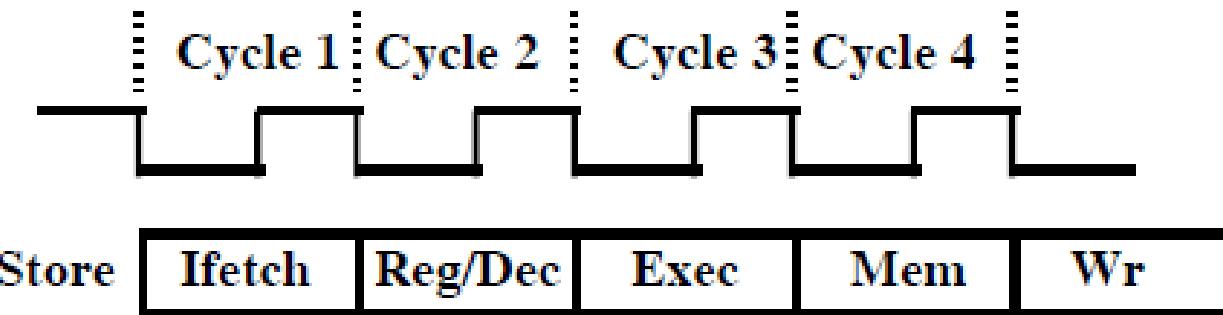
- Insert a “bubble” into the pipeline to prevent 2 writes at the same cycle
 - The control logic can be complex
- No instruction is completed during Cycle 5:
 - The “Effective” CPI for load is 2

Solution 2: Delay R-type's Write by One Cycle

- Delay R-type's register write by one cycle:
 - Now R-type instructions also use Reg File's write port at Stage 5
 - Mem stage is a NOOP stage: nothing is being done

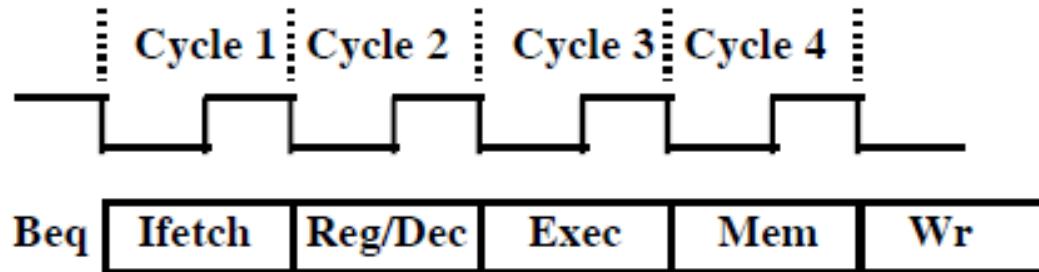


The Four Stages of Store



- **Ifetch: Instruction Fetch**
 - **Fetch the instruction from the Instruction Memory**
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: Calculate the memory address**
- **Mem: Write the data into the Data Memory**

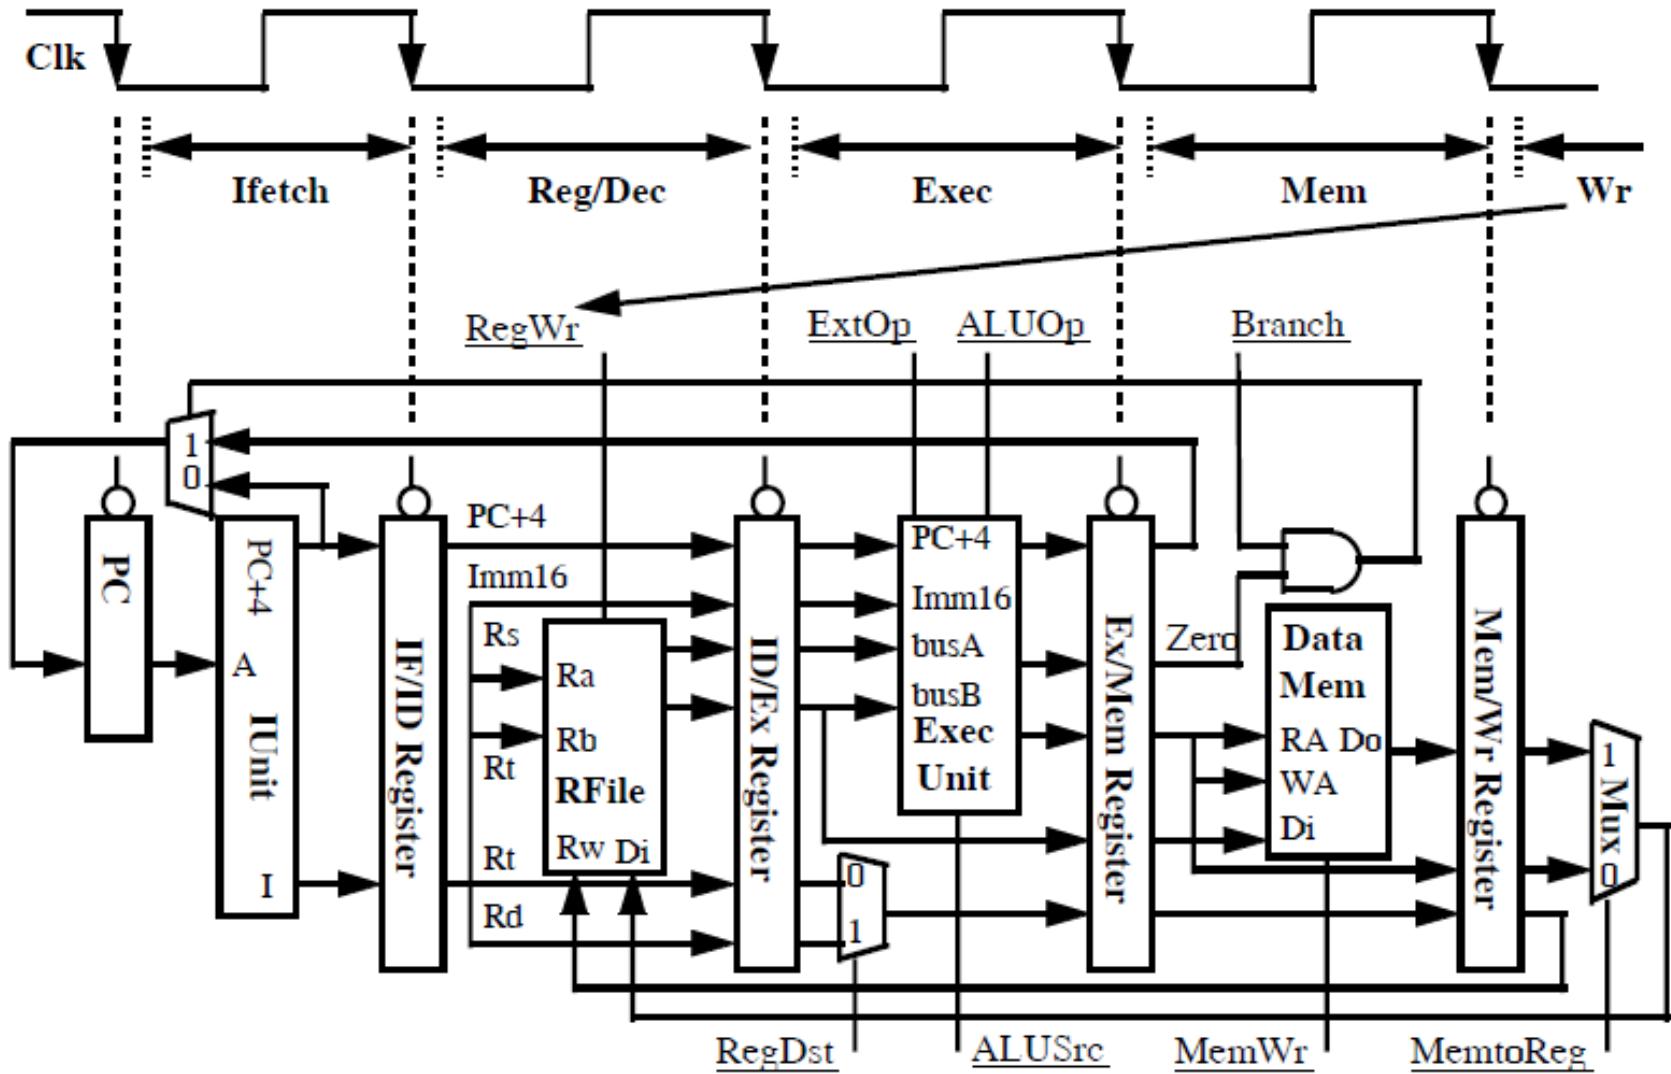
The Four Stages of Beq



- **Ifetch: Instruction Fetch**
 - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: ALU compares the two register operands**
 - Adder calculates the branch target address
- **Mem: If the registers we compared in the Exec stage are the same,**
 - Write the branch target address into the PC

Hey, this happens in
second cycle in our
Multi-Cycle CPU Design

A Pipelined Datapath

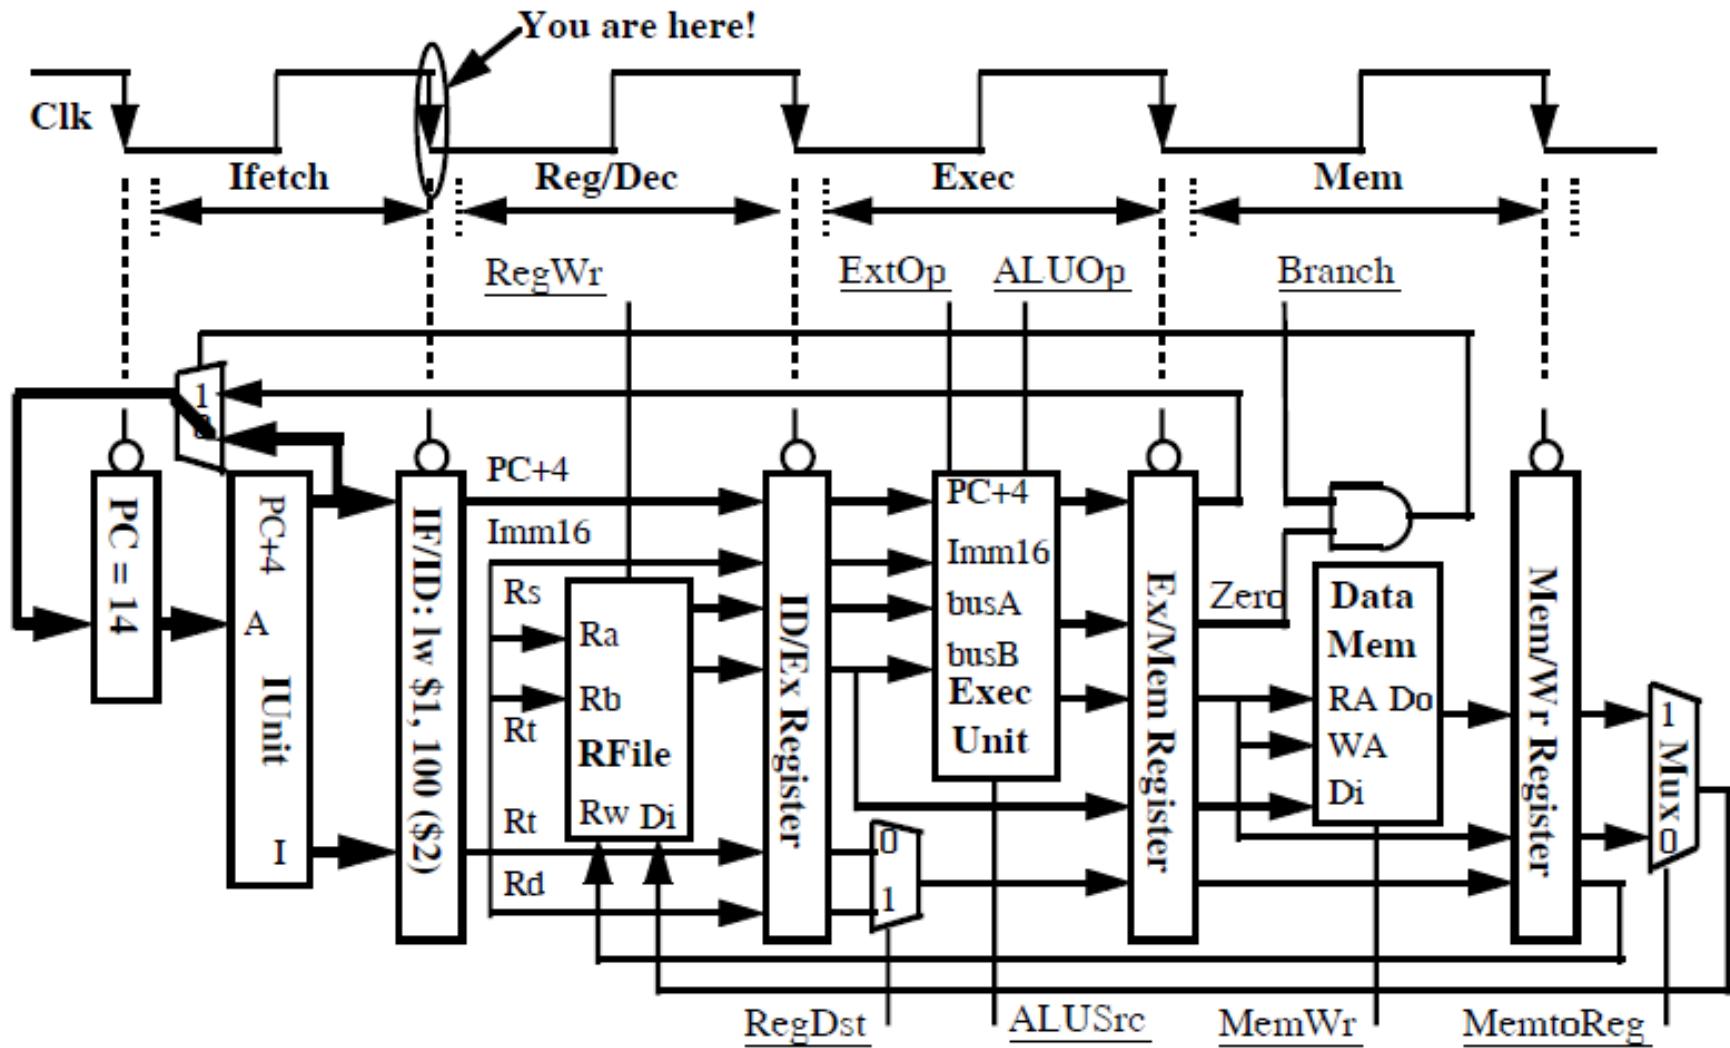


Fields of IF/ID Register:

1. 32-bits to store instruction
2. 32-bits to store PC+4

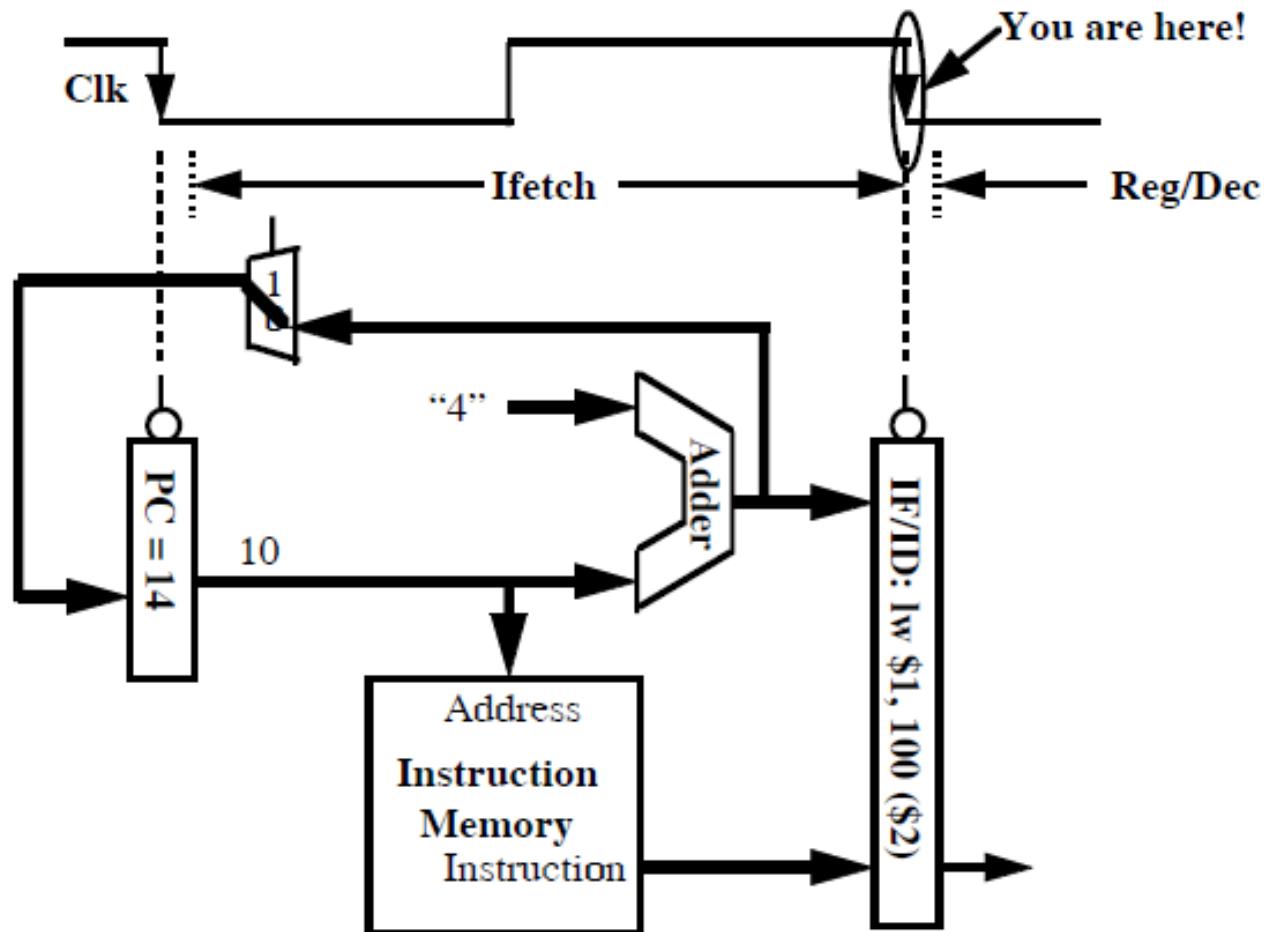
The Instruction Fetch Stage

- Location 10: lw \$1, 0x100(\$2) $\$1 \leftarrow \text{Mem}[(\$2) + 0x100]$



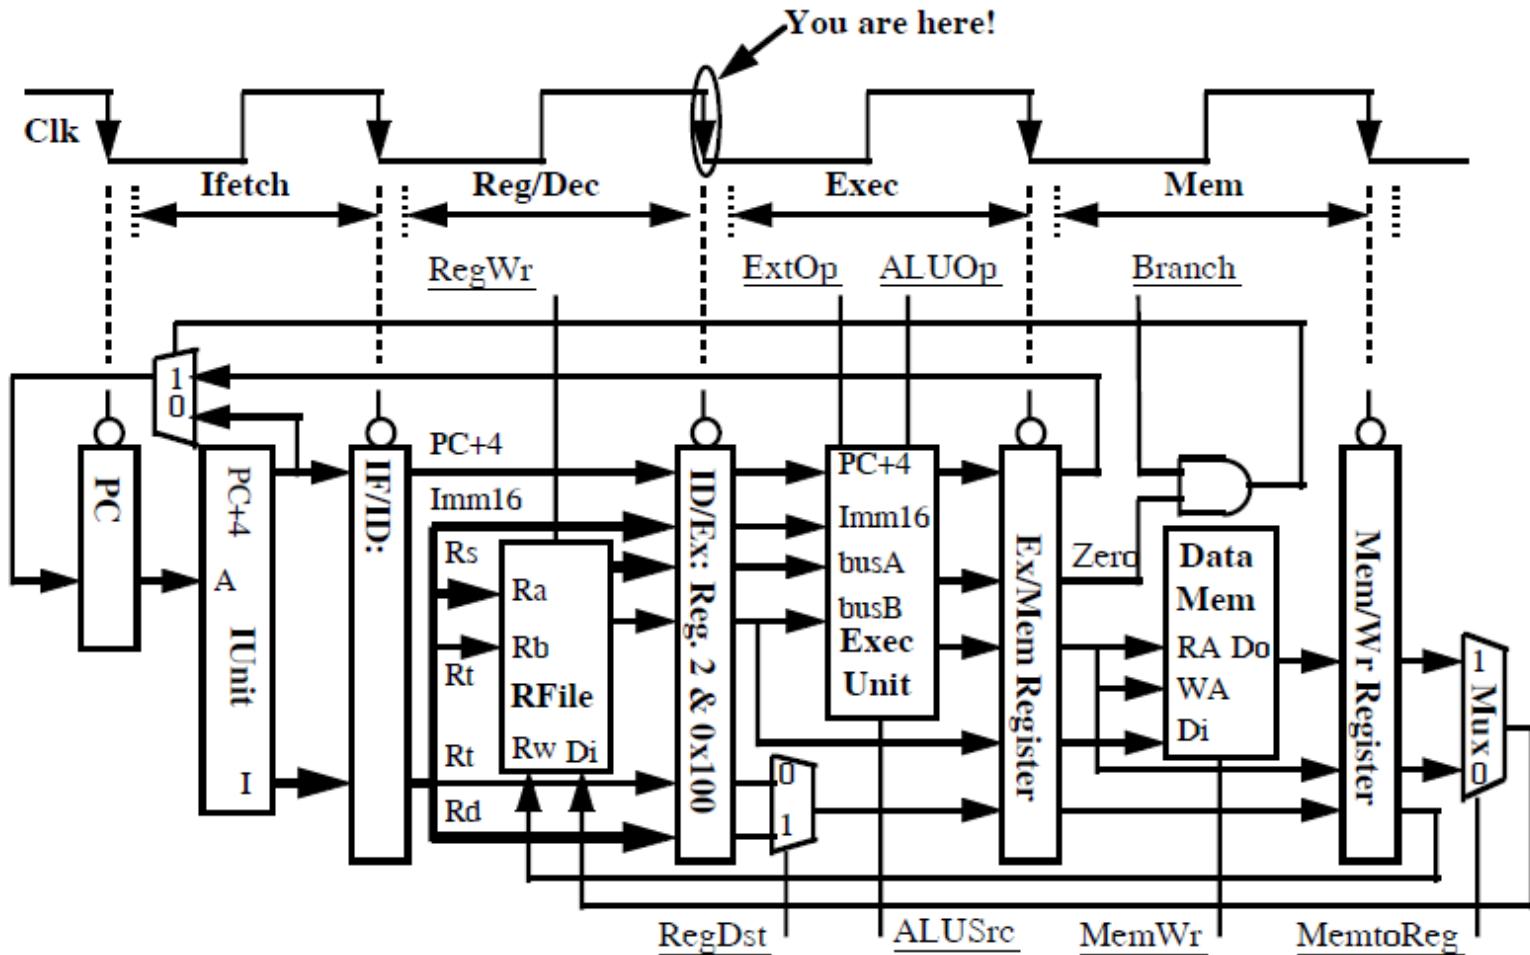
A Detail View of the Instruction Unit

- Location 10: lw \$1, 0x100(\$2)



The Decode / Register Fetch Stage

- Location 10: lw \$1, 0x100(\$2) $\$1 \leftarrow \text{Mem}[(\$2) + 0x100]$



Fields of ID/Ex Register:

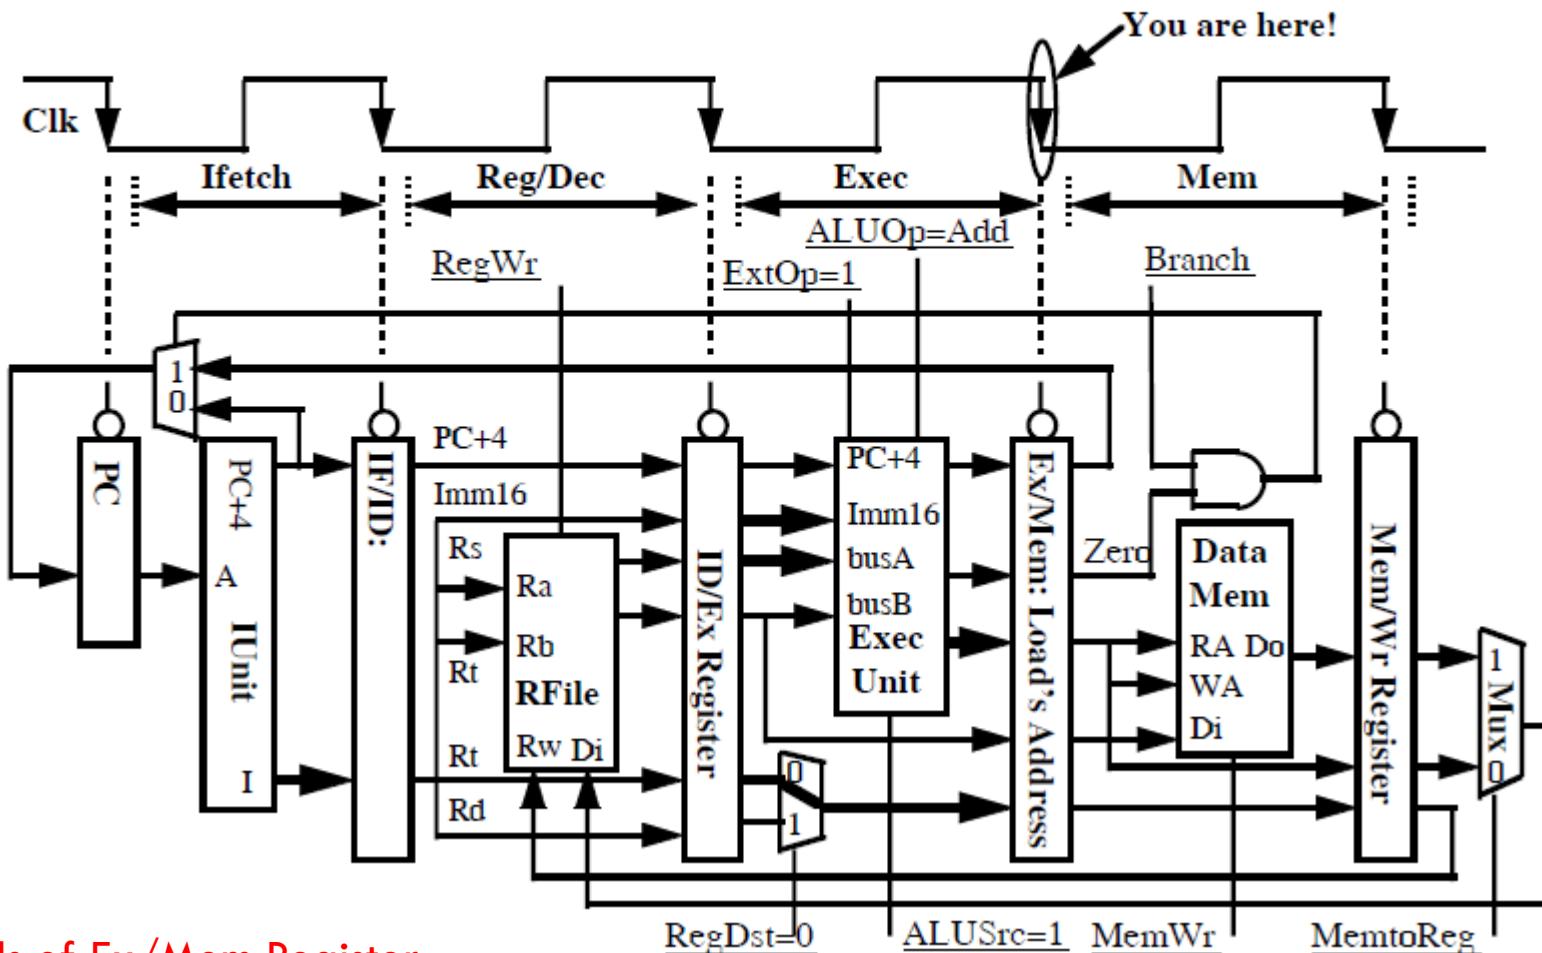
- 32-bits for PC+4
- 16-bits for Imm16
- 32-bits for M[Rs]

Fields of ID/Ex Register (continued):

- 32-bits for M[Rt]
- 5-bits for Rt
- 5-bits for Rd

Load's Address Calculation Stage

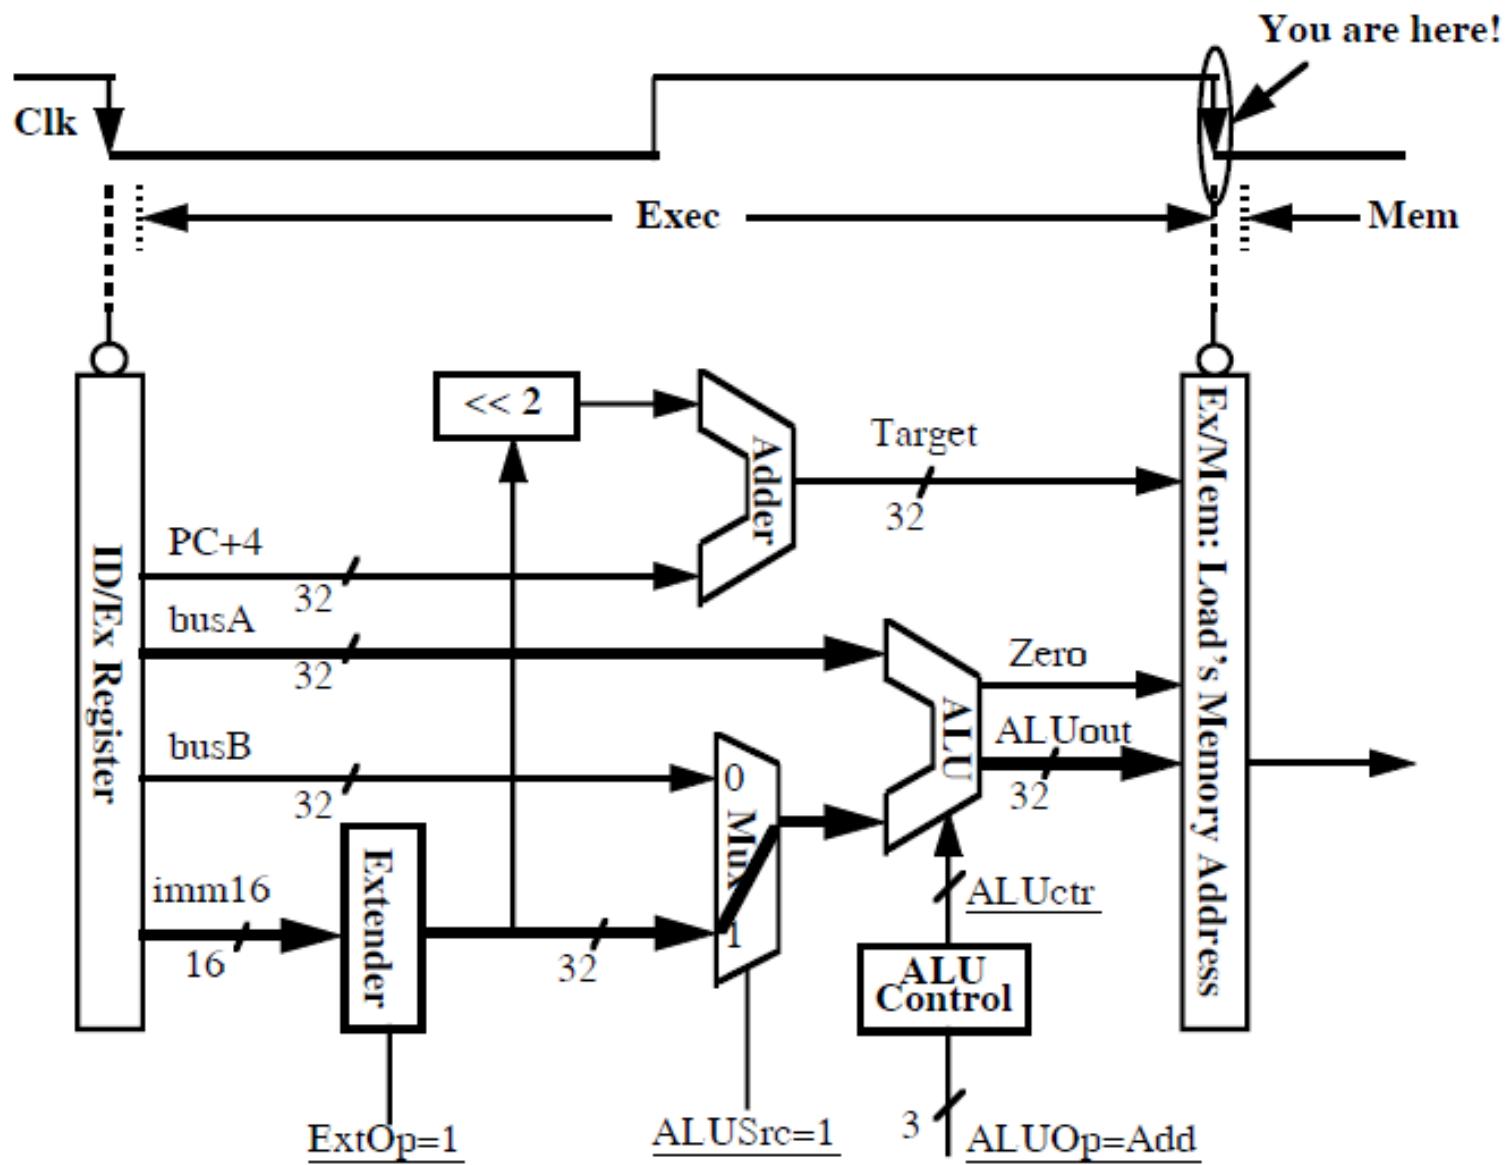
- Location 10: lw \$1, 0x100(\$2) $\$1 \leftarrow \text{Mem}[(\$2) + 0x100]$



Fields of Ex/Mem Register:

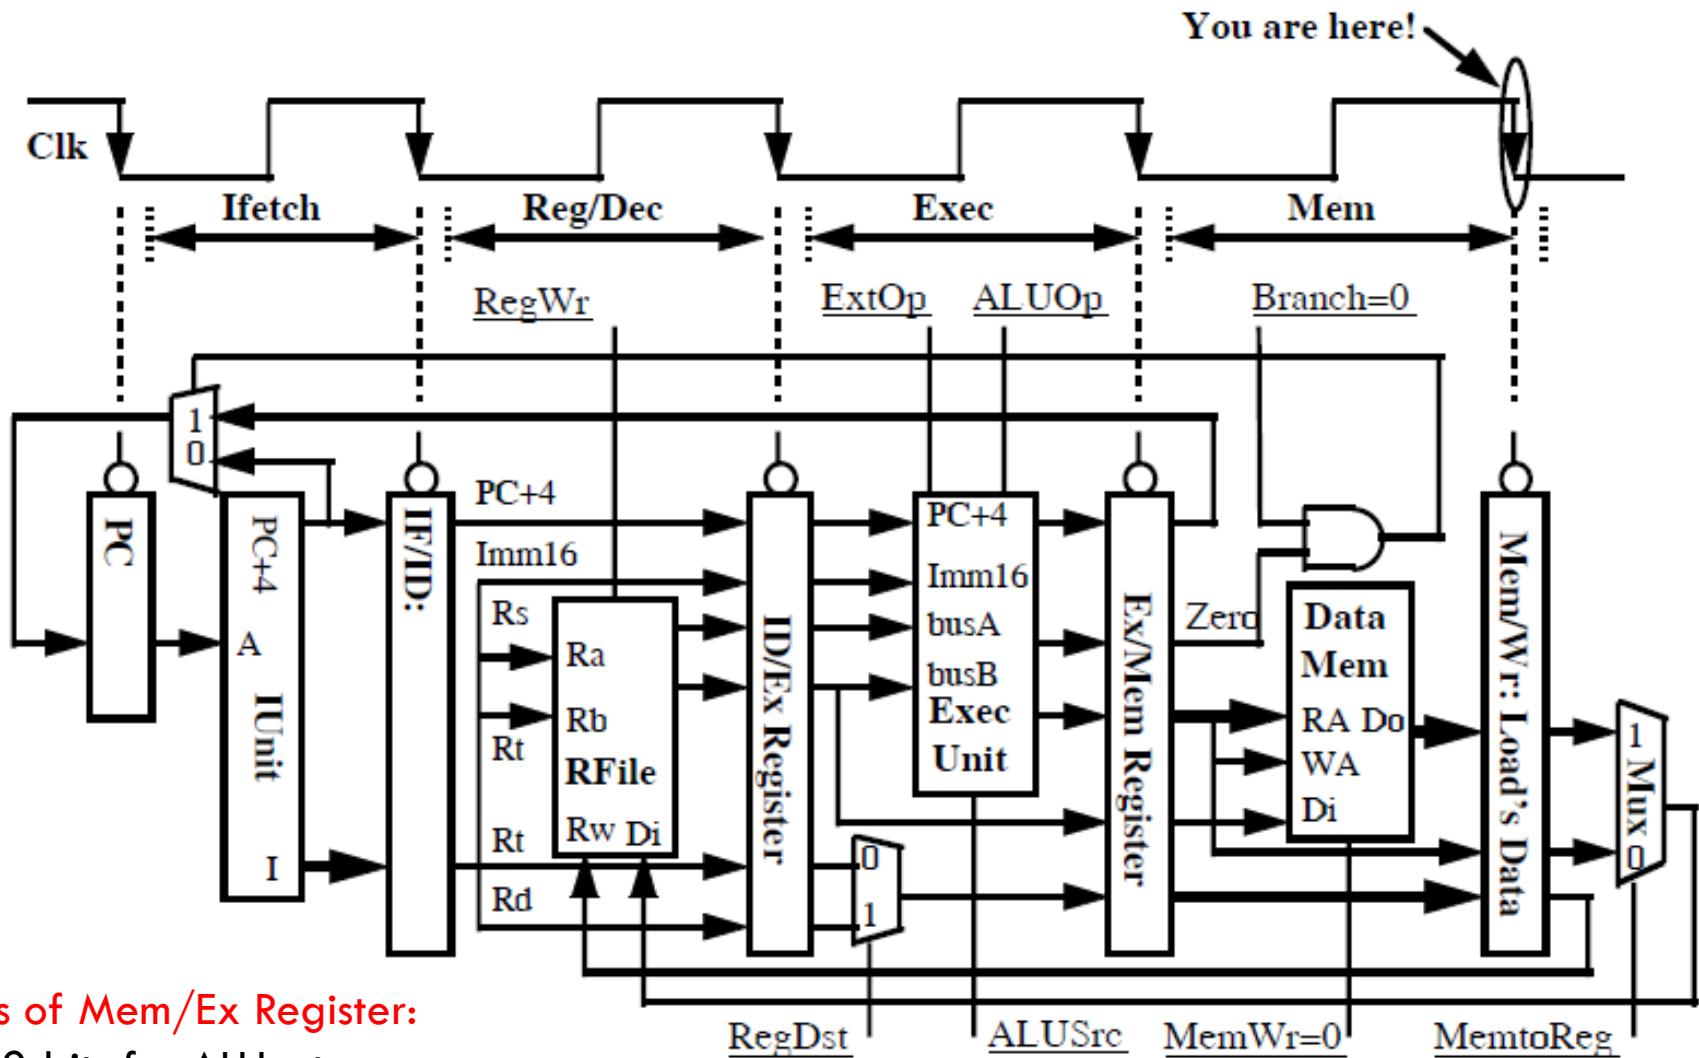
1. 32-bits for Branch Target Address
2. 32-bits for ALUout
3. 1-bit for Zero Flag
4. 32-bits for Mem[Rt]
5. 5-bits for RegDest: Rt or Rd

A Detail View of the Execution Unit



Load's Memory Access Stage

- Location 10: lw \$1, 0x100(\$2) $\$1 \leftarrow \text{Mem}[(\$2) + 0x100]$



Fields of Mem/Ex Register:

- 32-bits for ALUout
- 32-bits for Data Memory
- 5-bits for RegDest: Rt or Rd

RegDst

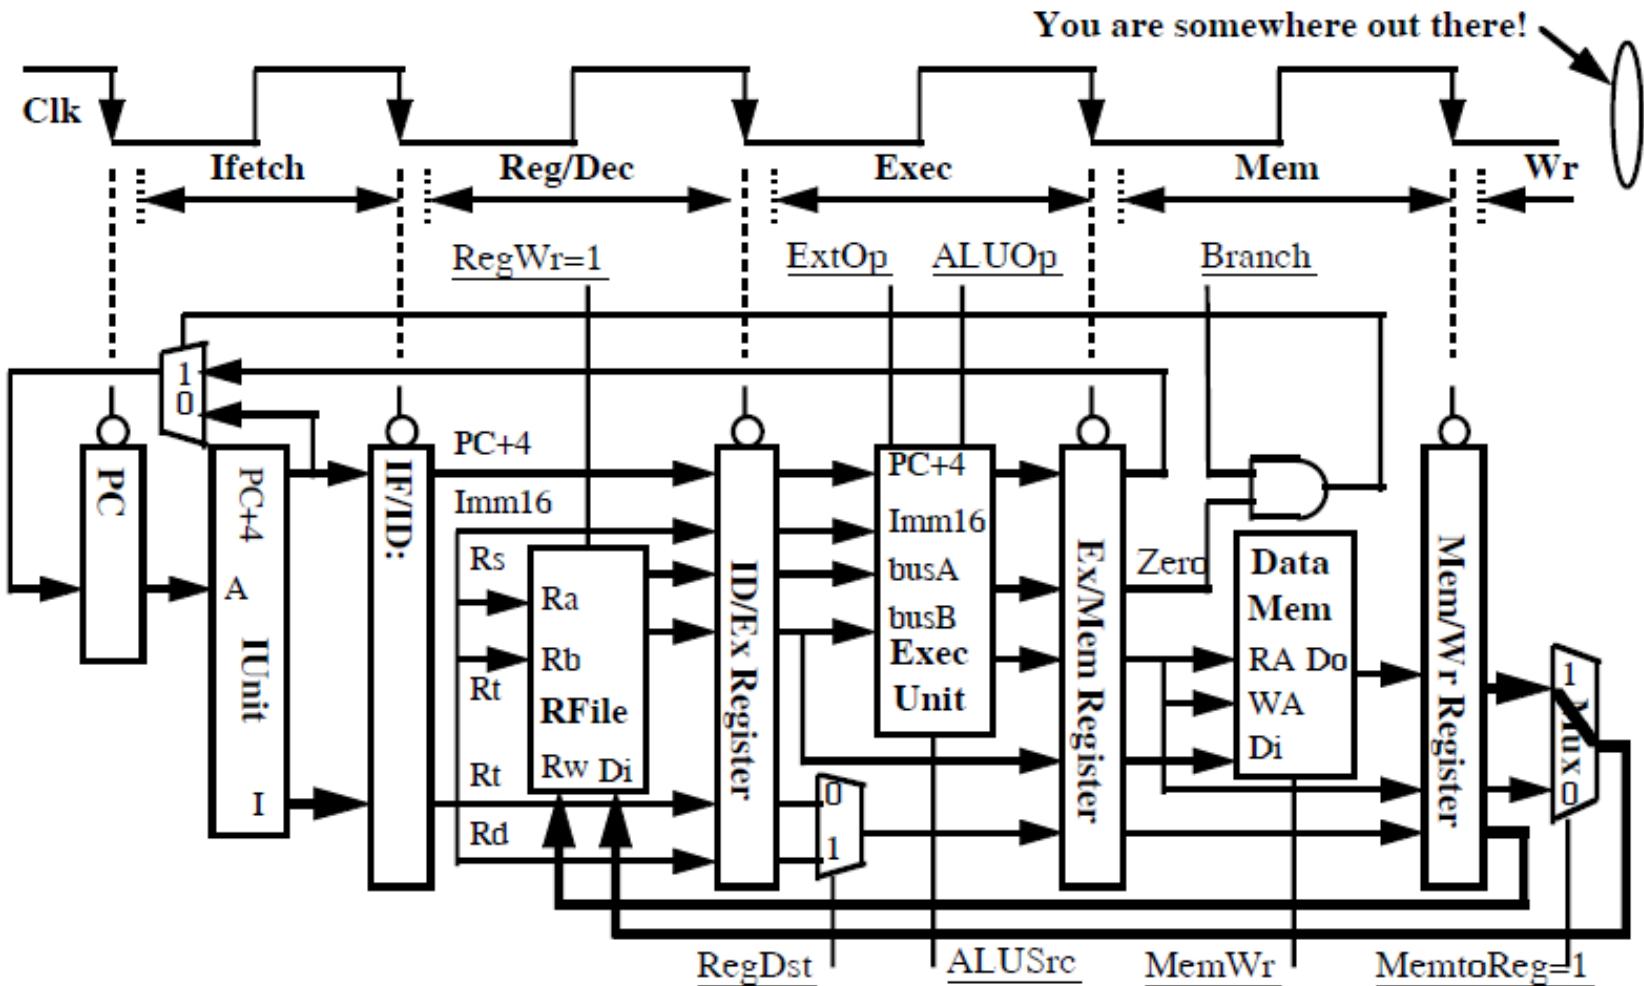
ALUSrc

MemWr=0

MemtoReg

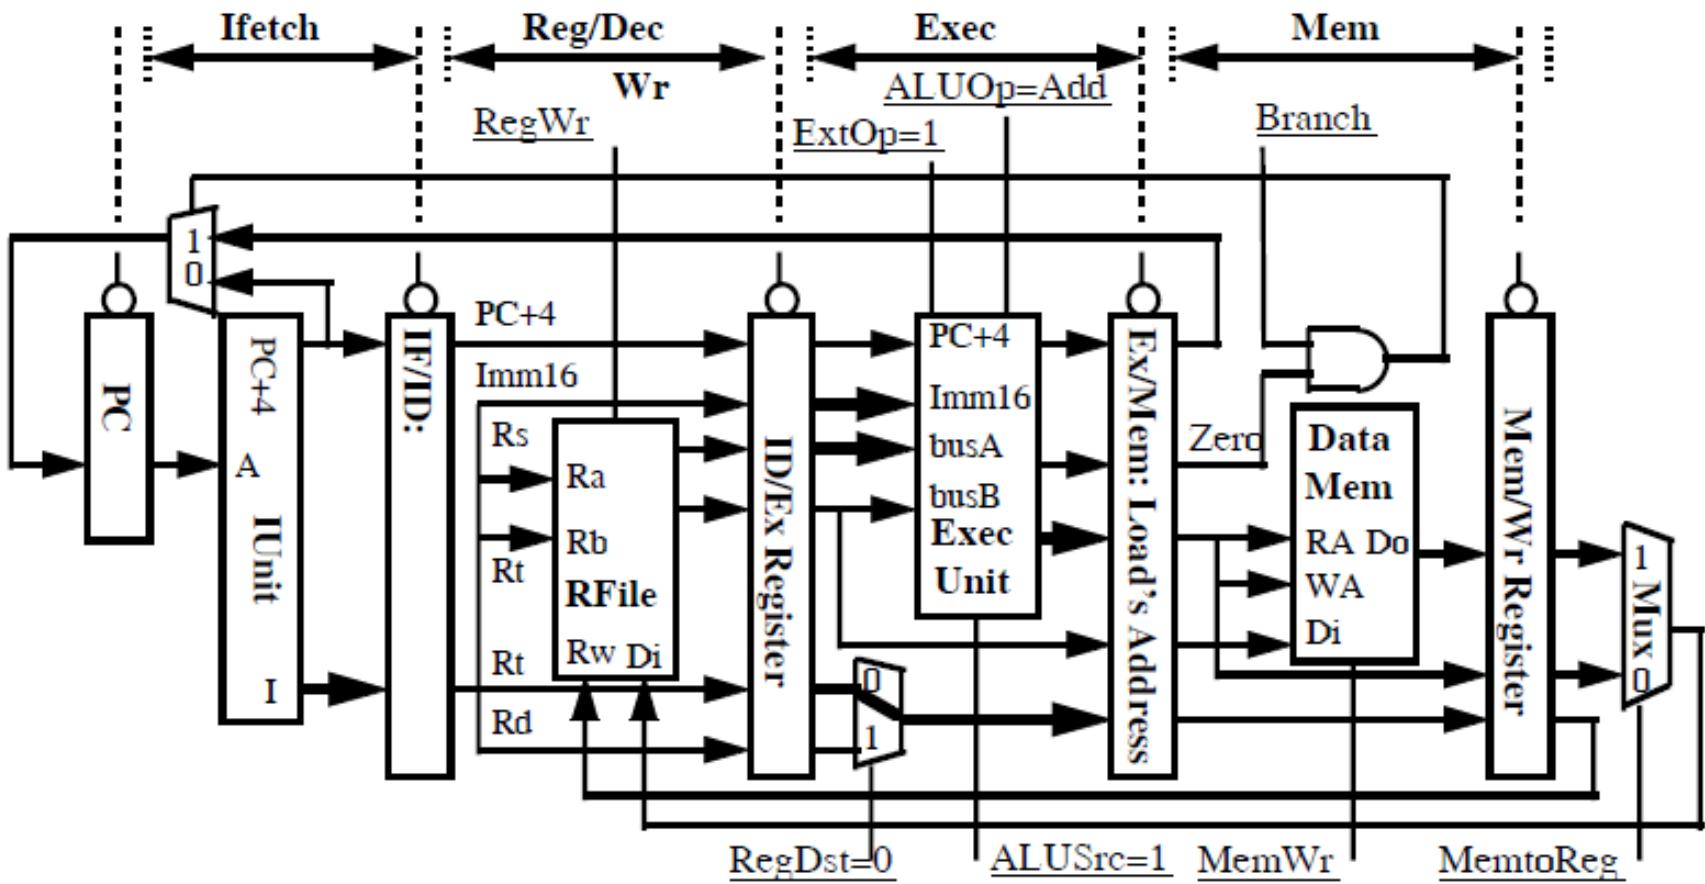
Load's Write Back Stage

- Location 10: lw \$1, 0x100(\$2) $S1 \leftarrow \text{Mem}[(S2) + 0x100]$



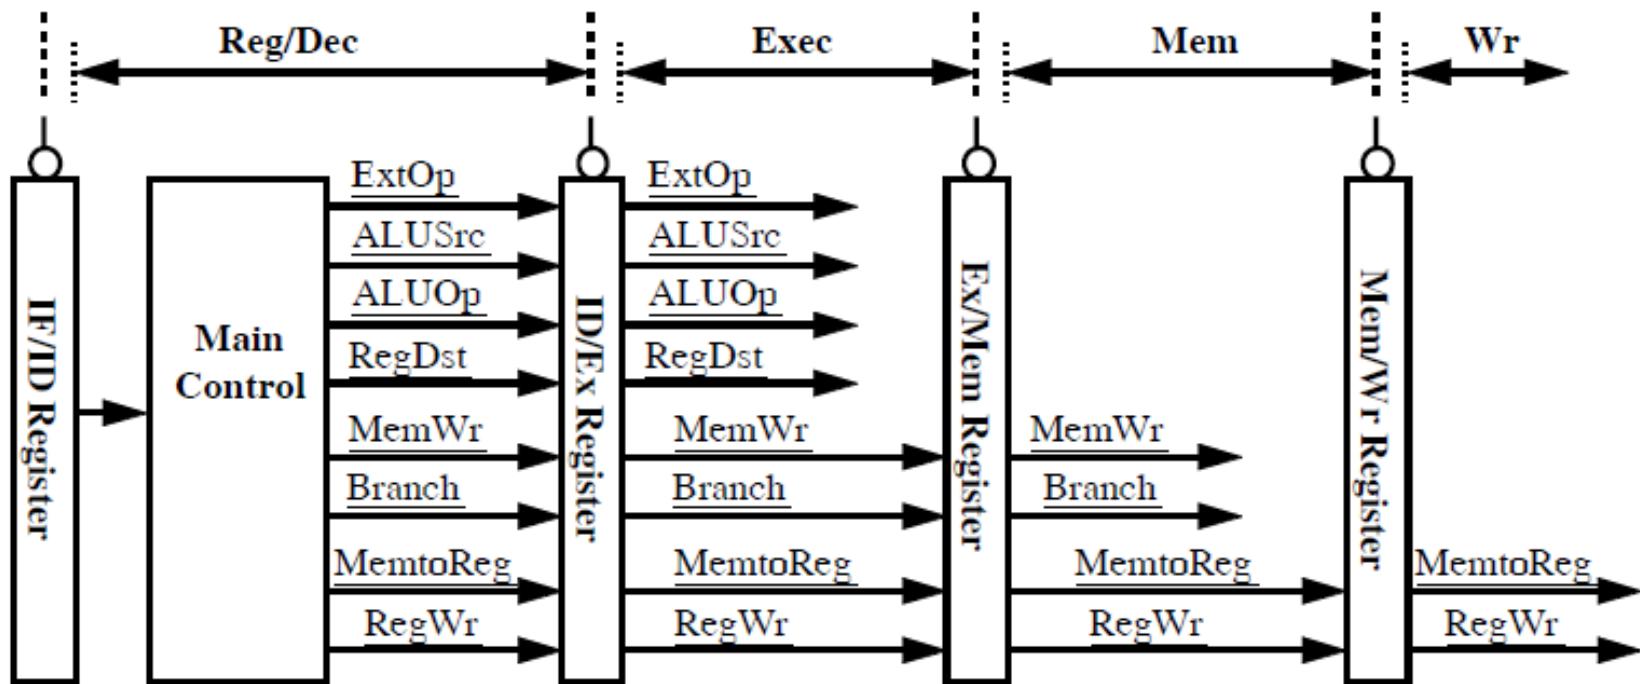
How About Control Signals?

- Key Observation: Control Signals at Stage N = Func (Instr. at Stage N)
 - N = Exec, Mem, or Wr
- Example: Controls Signals at Exec Stage = Func(Load's Exec)



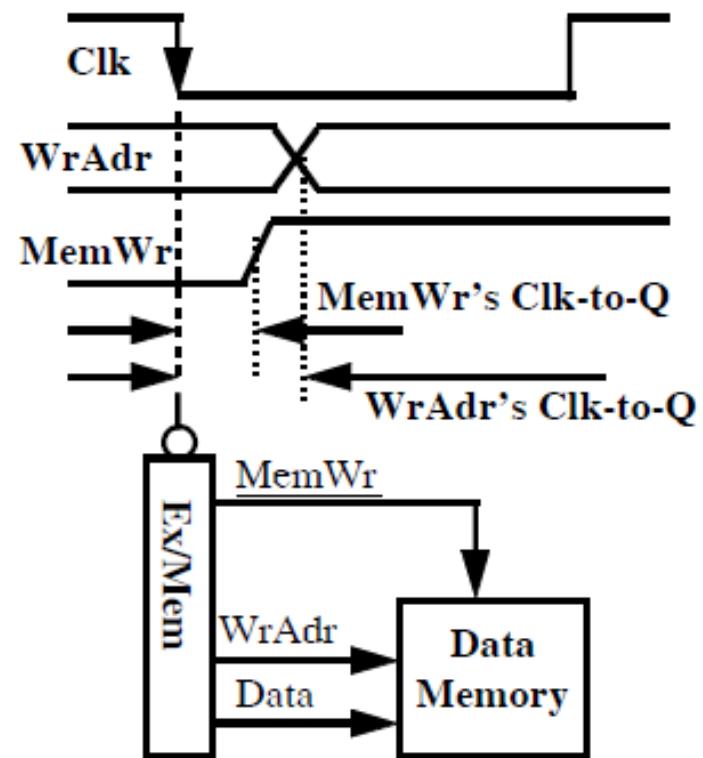
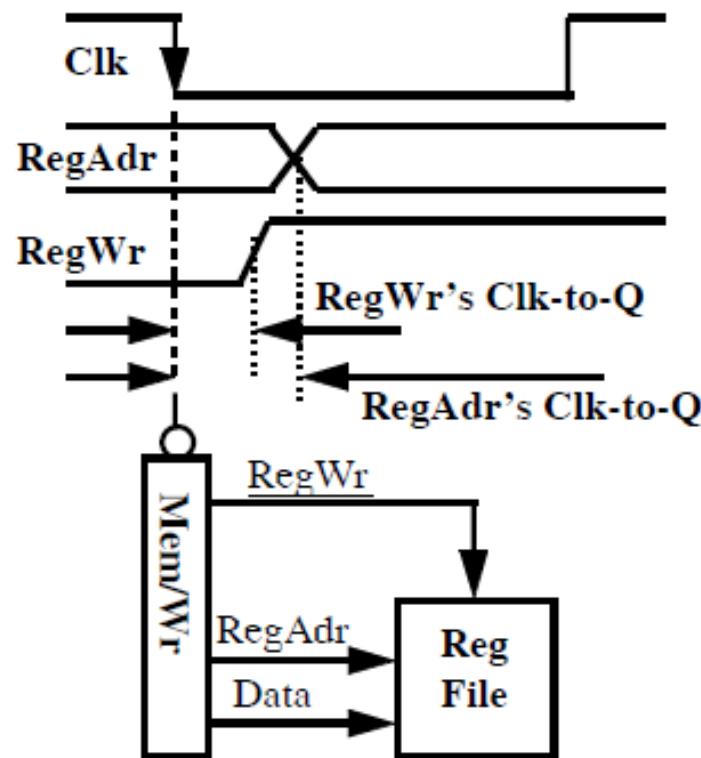
Pipeline Control

- The Main Control generates the control signals during Reg/Dec
 - Control signals for Exec (ExtOp, ALUSrc, ...) are used 1 cycle later
 - Control signals for Mem (MemWr Branch) are used 2 cycles later
 - Control signals for Wr (MemtoReg MemWr) are used 3 cycles later



We need to add more fields for holding control bits to the intermediate register structures we defined in the previous fields.

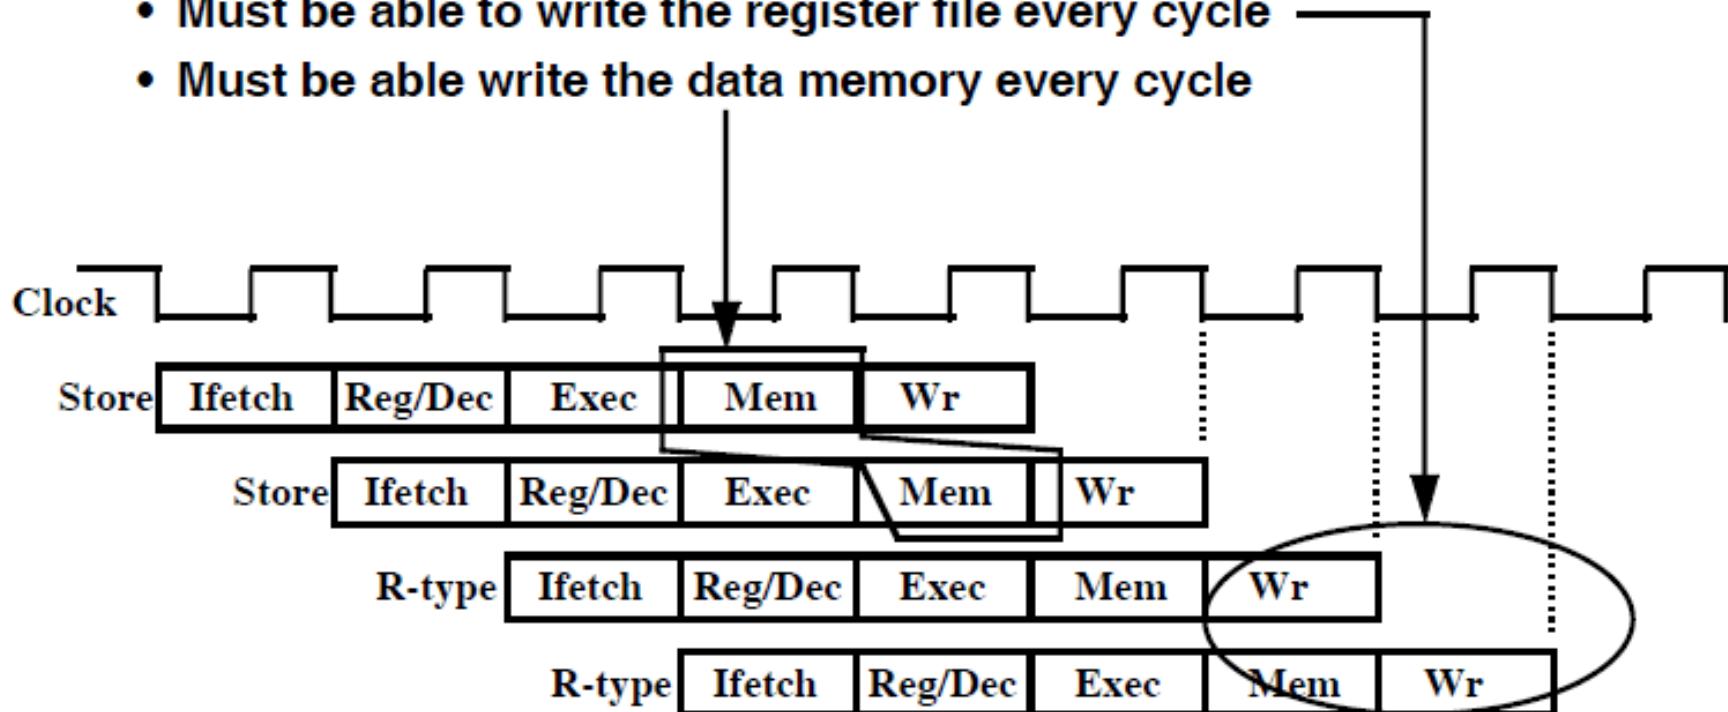
Beginning of the Wr's Stage: A Real World Problem



- At the beginning of the Wr stage, we have a problem if:
 - RegAddr's (Rd or Rt) Clk-to-Q > RegWr's Clk-to-Q
- Similarly, at the beginning of the Mem stage, we have a problem if:
 - WrAddr's Clk-to-Q > MemWr's Clk-to-Q
- We have a race condition between Address and Write Enable!

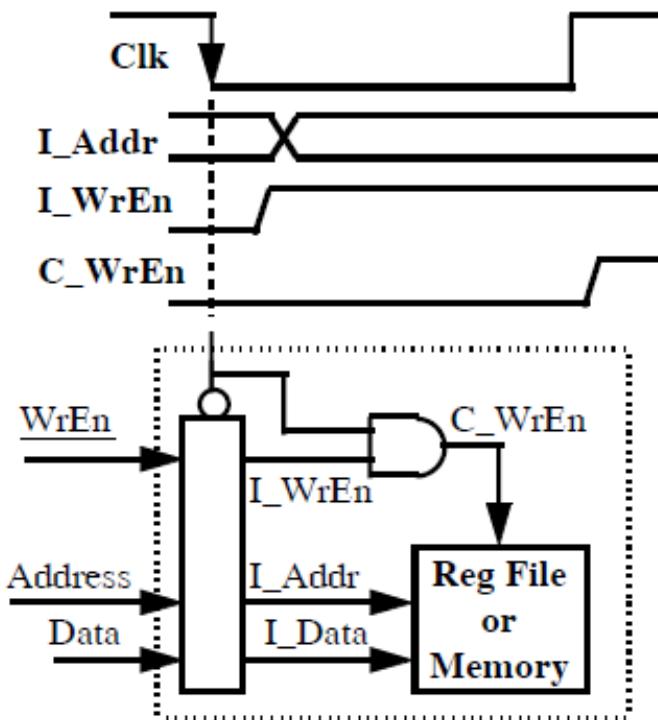
The Pipeline Problem

- Multiple Cycle design prevents race condition between Addr and WrEn:
 - Make sure Address is stable by the end of Cycle N
 - Asserts WrEn during Cycle N + 1
- This approach can NOT be used in the pipeline design because:
 - Must be able to write the register file every cycle
 - Must be able write the data memory every cycle



Synchronize Register File & Synchronize Memory

- Solution: And the Write Enable signal with the Clock
 - This is the ONLY place where gating the clock is used
 - MUST consult circuit expert to ensure no timing violation:
 - Example: Clock High Time > Write Access Delay

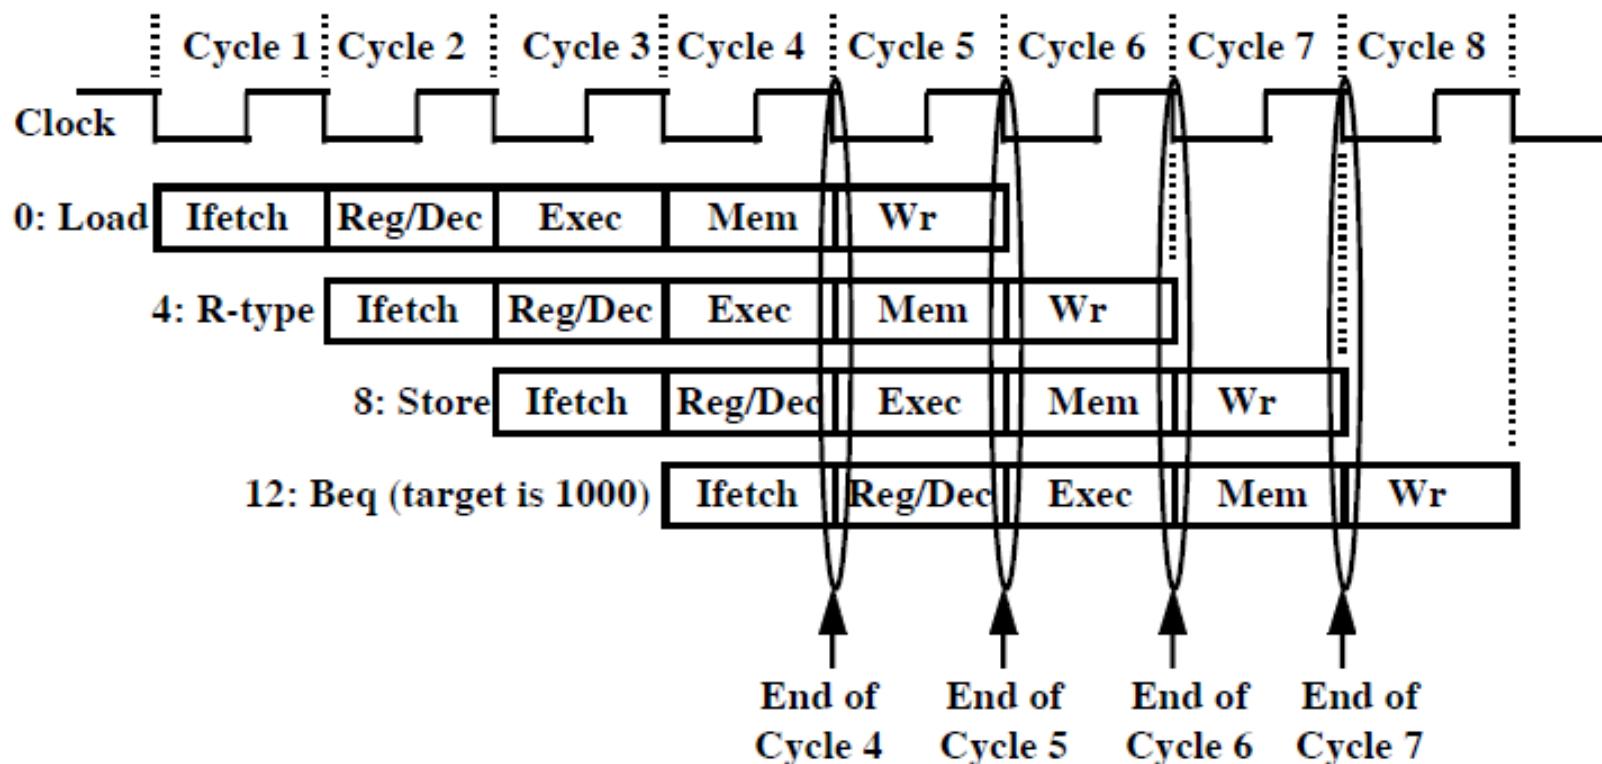


Synchronize Memory and Register File

Address, Data, and WrEn must be stable at least 1 set-up time before the Clk edge

Write occurs at the cycle following the clock edge that captures the signals

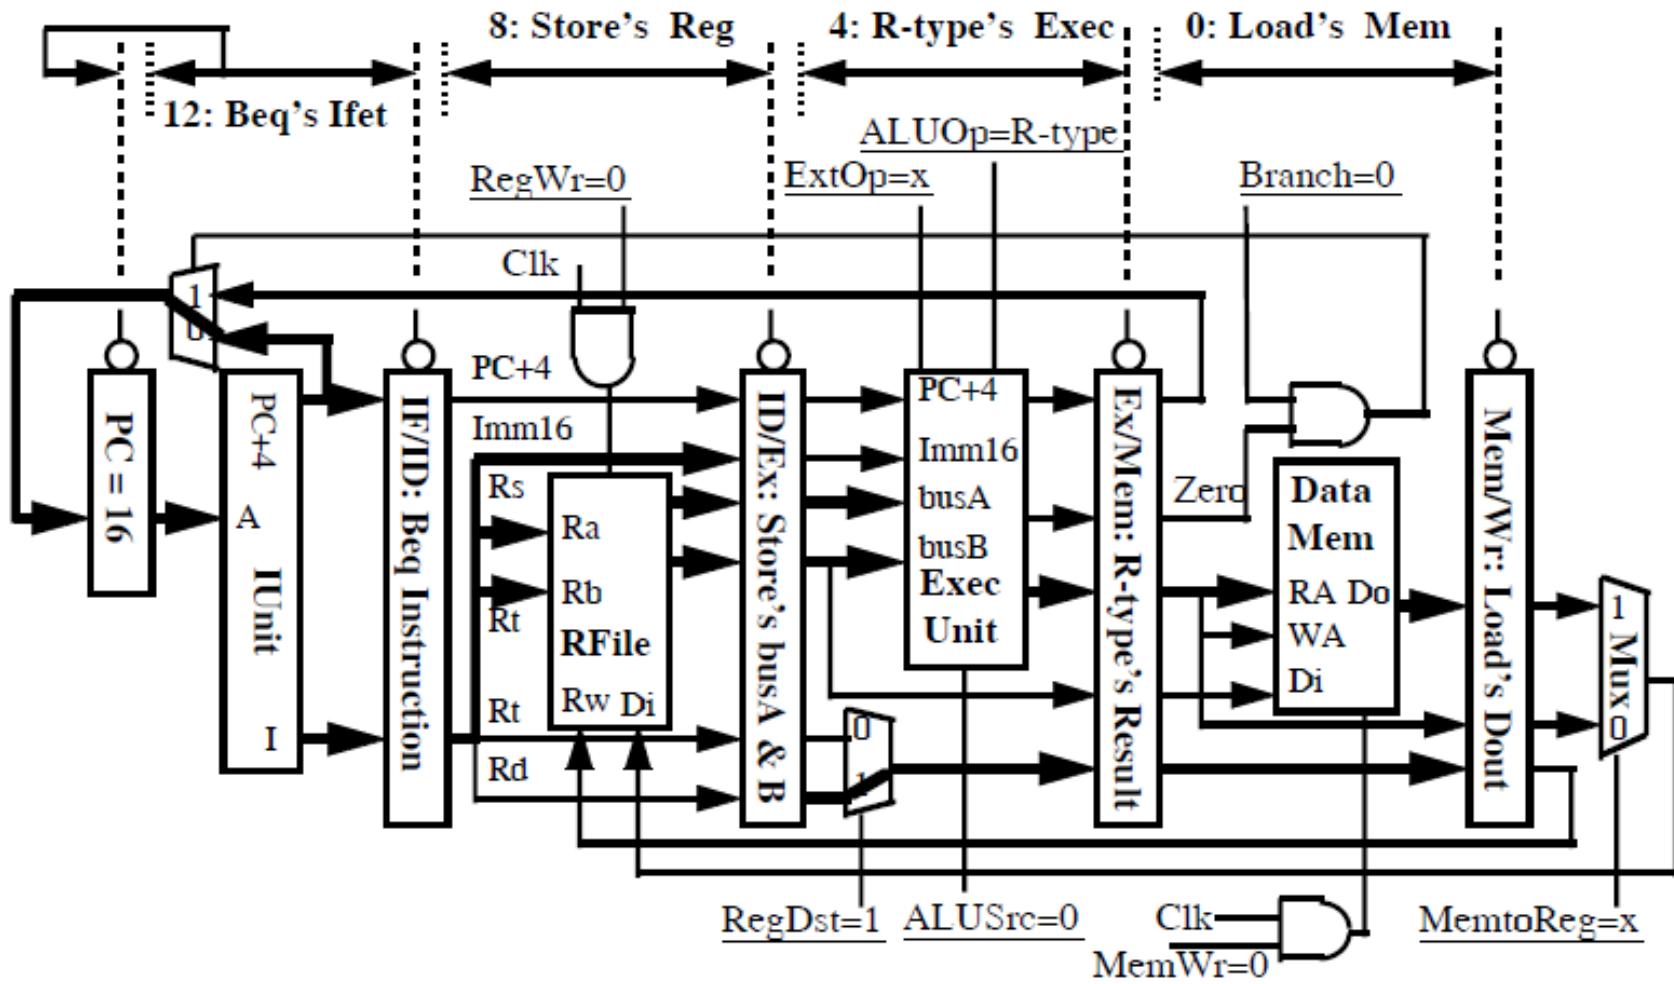
A More Extensive Pipelining Example



- End of Cycle 4: Load’s Mem, R-type’s Exec, Store’s Reg, Beq’s Ifetch
- End of Cycle 5: Load’s Wr, R-type’s Mem, Store’s Exec, Beq’s Reg
- End of Cycle 6: R-type’s Wr, Store’s Mem, Beq’s Exec
- End of Cycle 7: Store’s Wr, Beq’s Mem

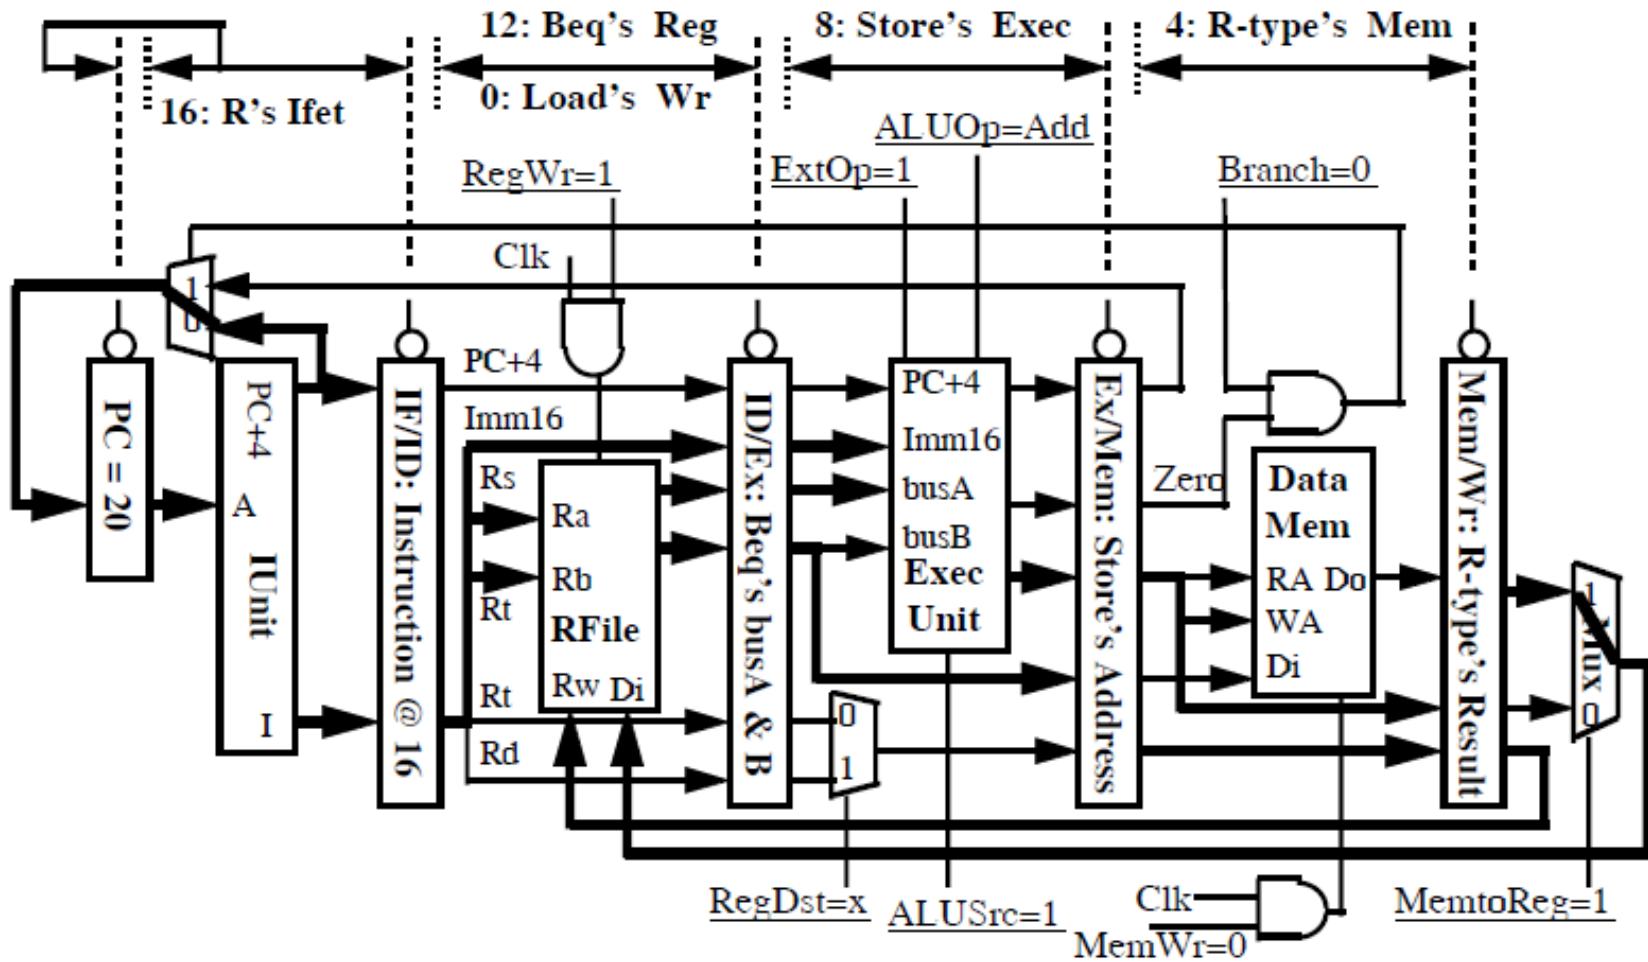
Pipelining Example: End of Cycle 4

- 0: Load's Mem 4: R-type's Exec 8: Store's Reg 12: Beq's Ifetch



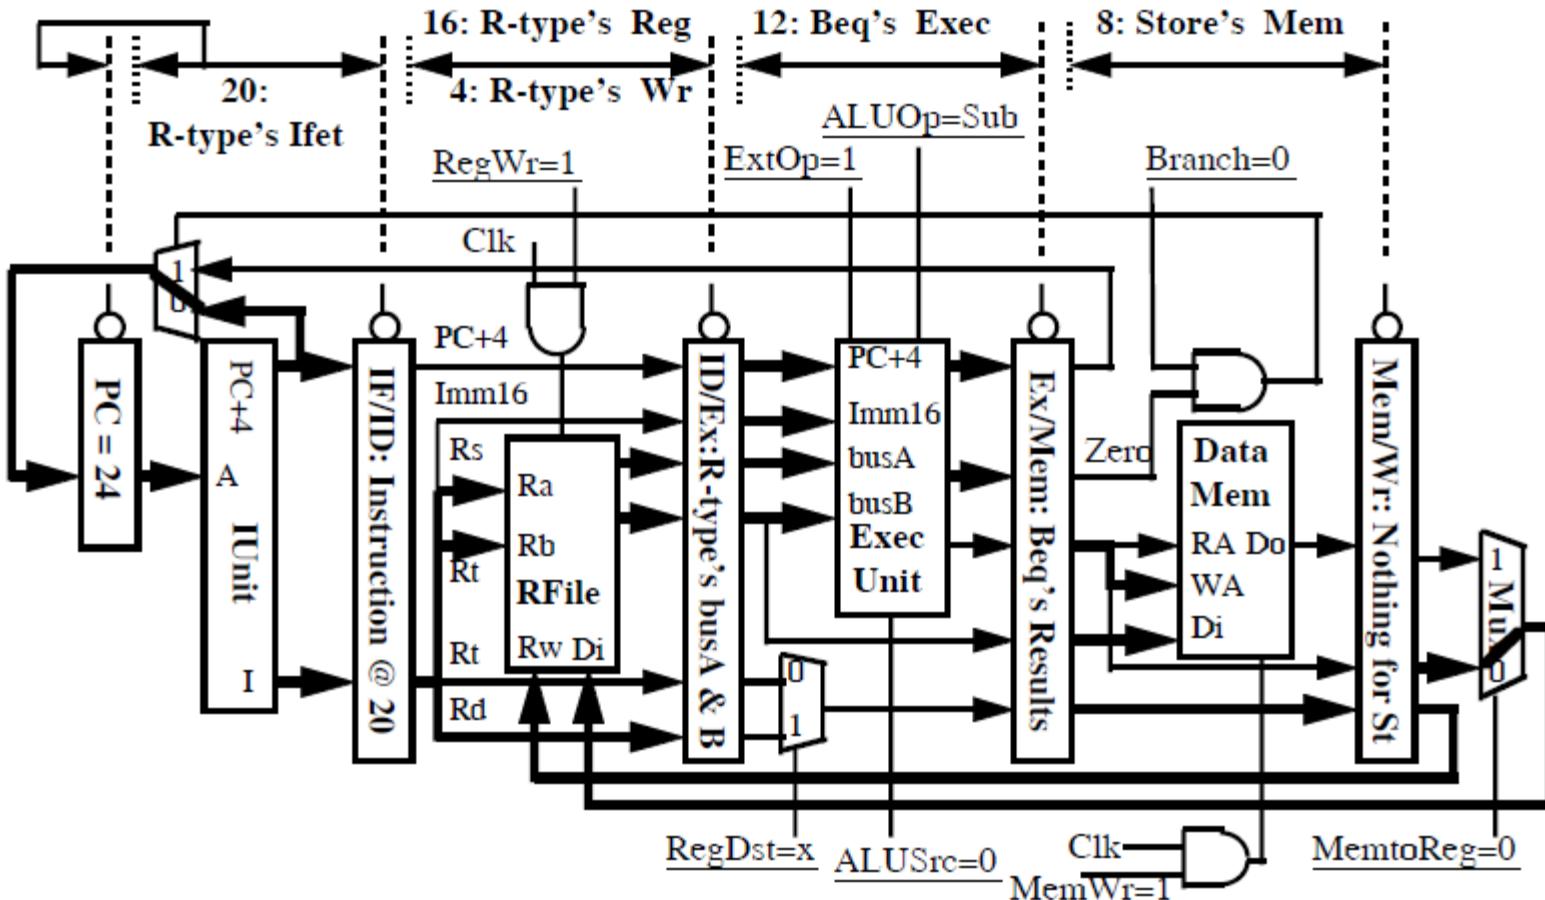
Pipelining Example: End of Cycle 5

- 0: Lw's Wr 4: R's Mem 8: Store's Exec 12: Beq's Reg 16: R's Ifetch



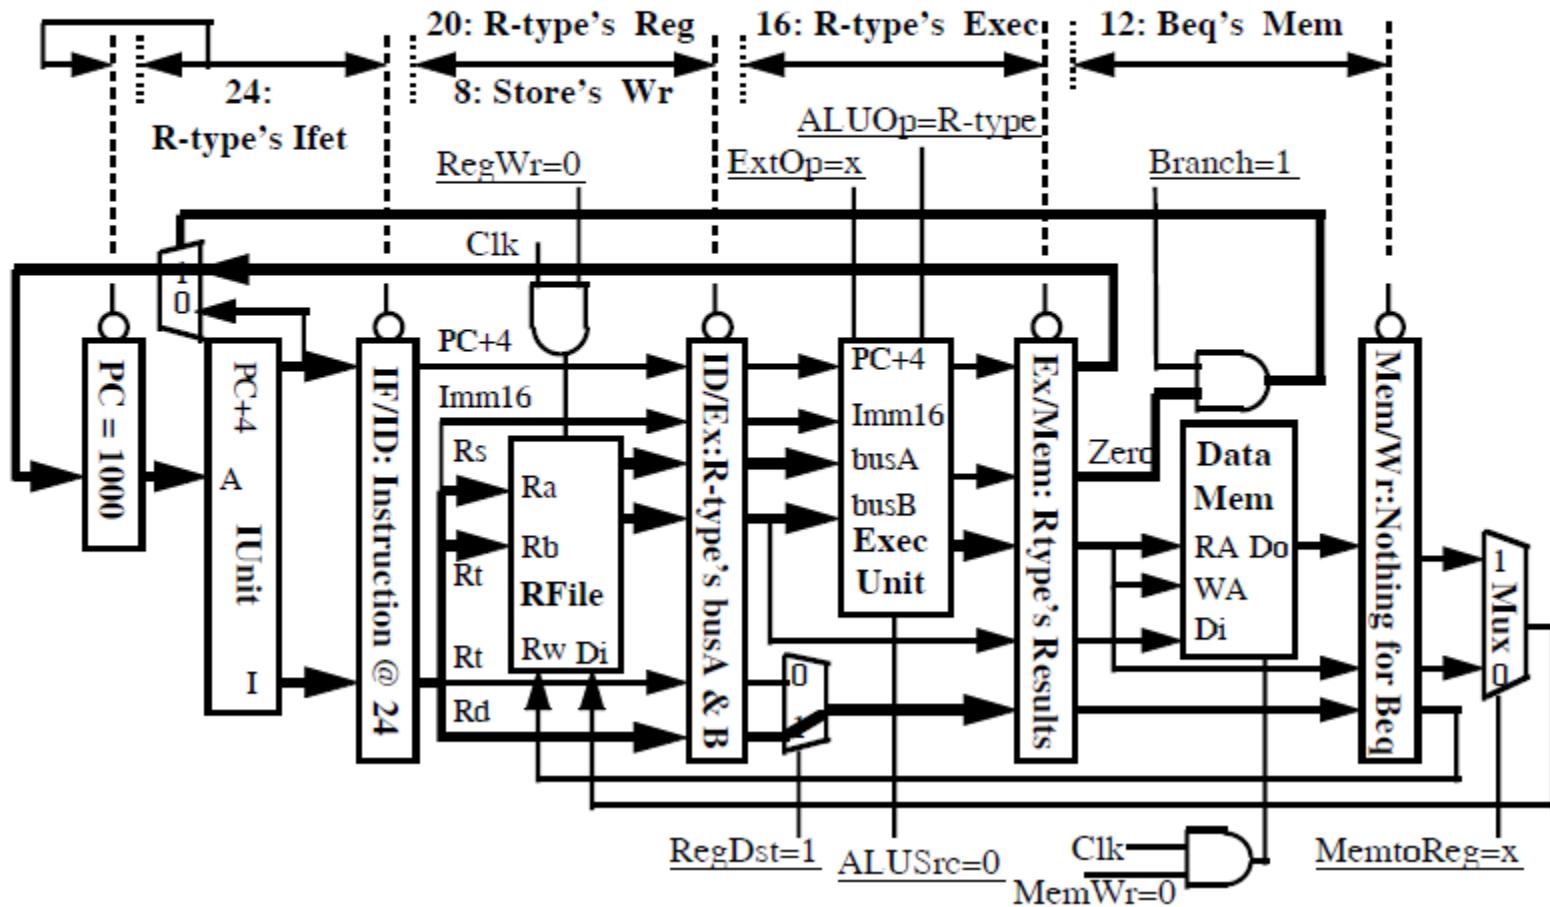
Pipelining Example: End of Cycle 6

- 4: R's Wr 8: Store's Mem 12: Beq's Exec 16: R's Reg 20: R's Ifet

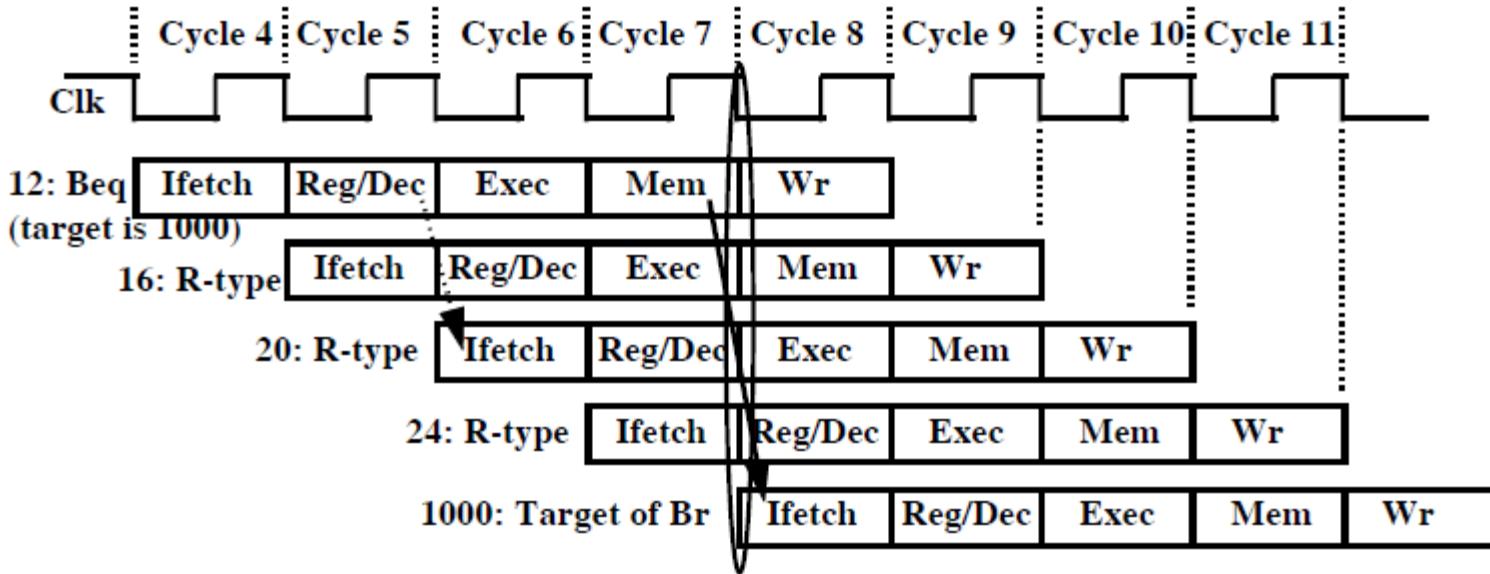


Pipelining Example: End of Cycle 7

- 8: Store's Wr 12: Beq's Mem 16: R's Exec 20: R's Reg 24: R's Ifet

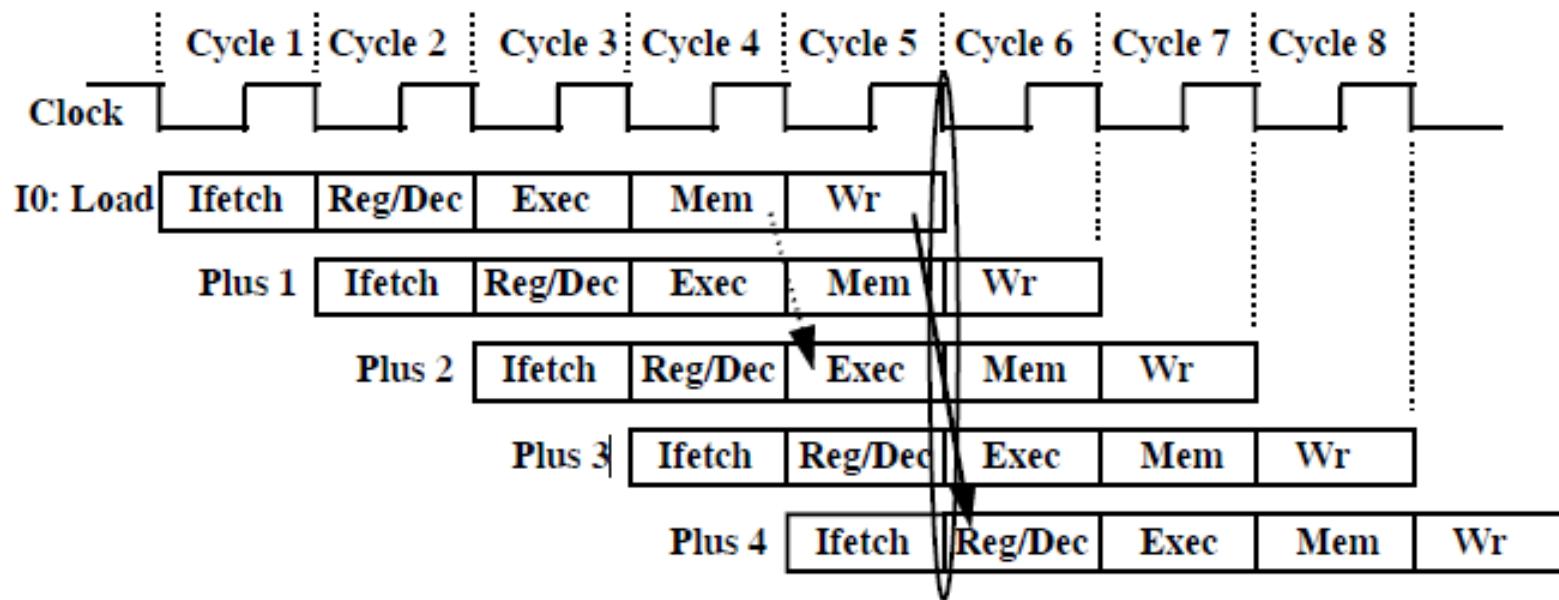


The Delay Branch Phenomenon



- ° Although Beq is fetched during Cycle 4:
 - Target address is NOT written into the PC until the end of Cycle 7
 - Branch's target is NOT fetched until Cycle 8
 - 3-instruction delay before the branch take effect
- ° This is referred to as Branch Hazard:
 - Clever design techniques can reduce the delay to ONE instruction

The Delay Load Phenomenon

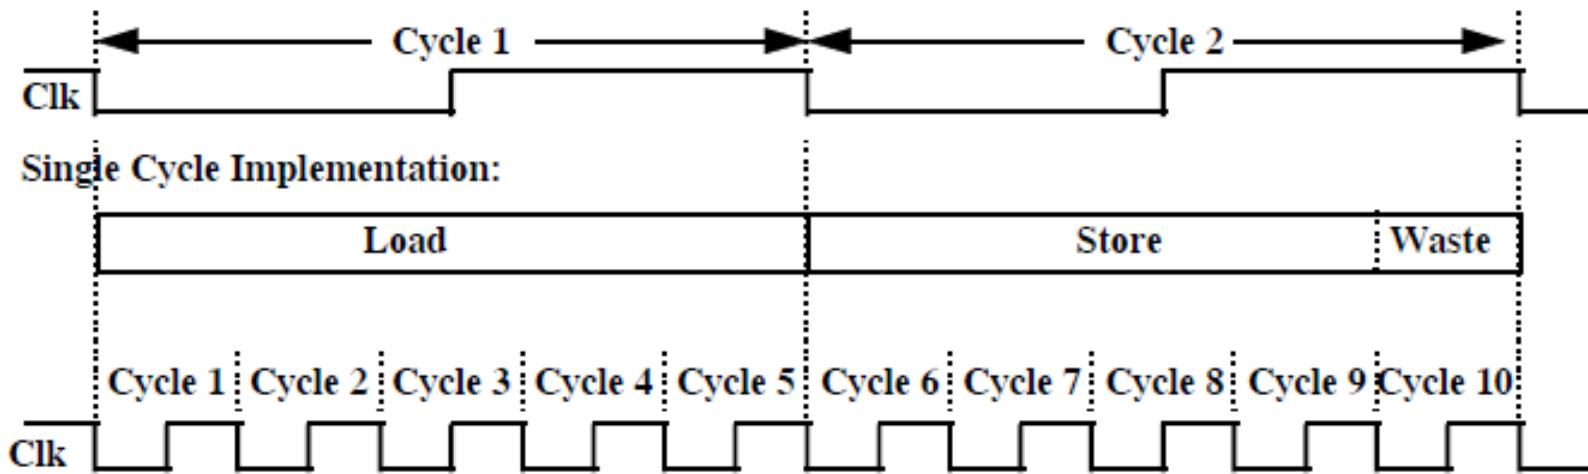


- ° Although Load is fetched during Cycle 1:
 - The data is NOT written into the Reg File until the end of Cycle 5
 - We cannot read this value from the Reg File until Cycle 6
 - 3-instruction delay before the load take effect
- ° This is referred to as Data Hazard:
 - Clever design techniques can reduce the delay to ONE instruction

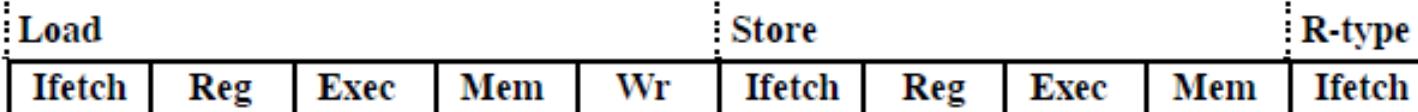
Summary

- Disadvantages of the Single Cycle Processor
 - Long cycle time
 - Cycle time is too long for all instructions except the Load
- Multiple Clock Cycle Processor:
 - Divide the instructions into smaller steps
 - Execute each step (instead of the entire instruction) in one cycle
- Pipeline Processor:
 - Natural enhancement of the multiple clock cycle processor
 - Each functional unit can only be used once per instruction
 - If a instruction is going to use a functional unit:
 - it must use it at the same stage as all other instructions
 - Pipeline Control:
 - Each stage's control signal depends ONLY on the instruction that is currently in that stage

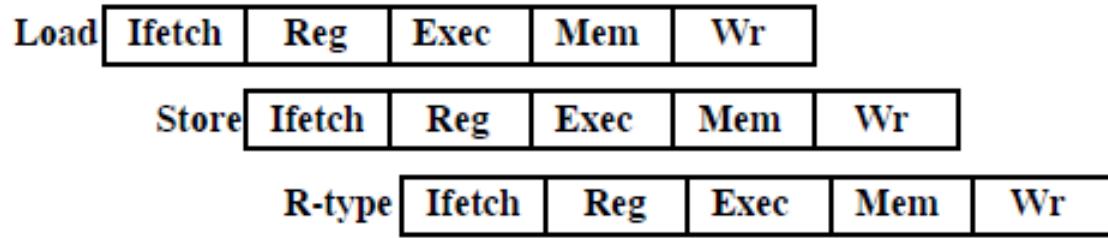
Single Cycle, Multiple Cycle, vs. Pipeline



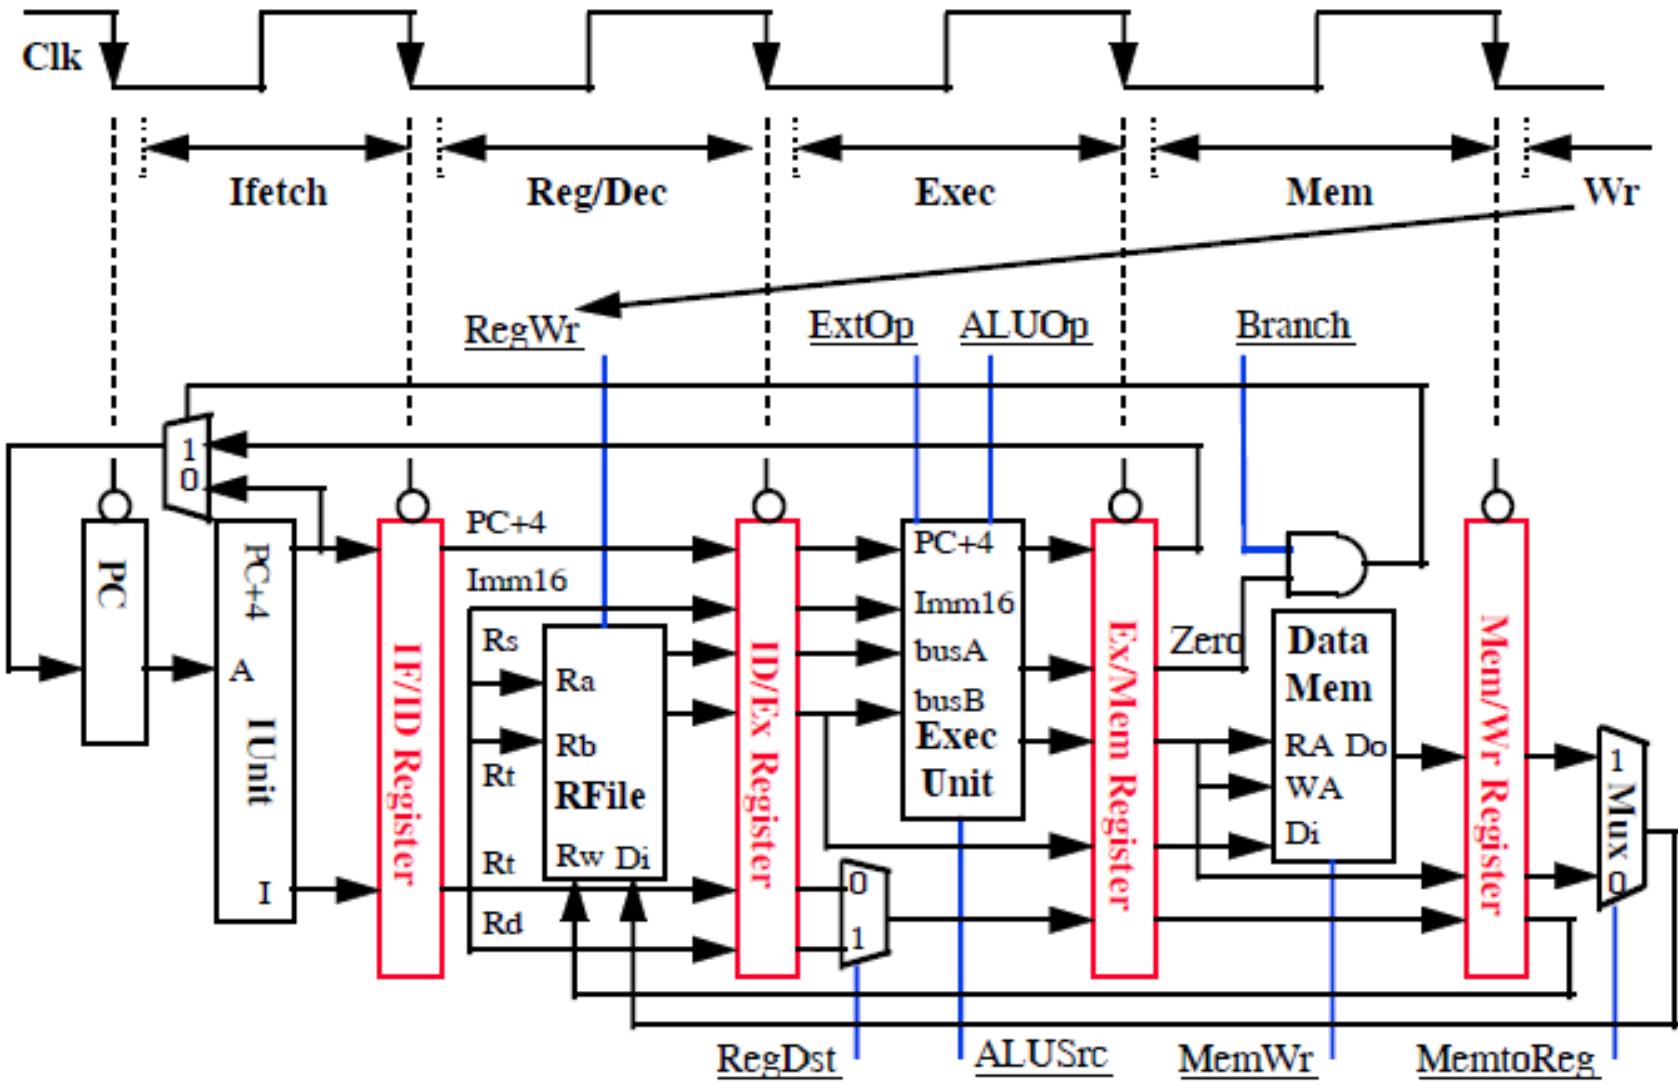
Multiple Cycle Implementation:



Pipeline Implementation:

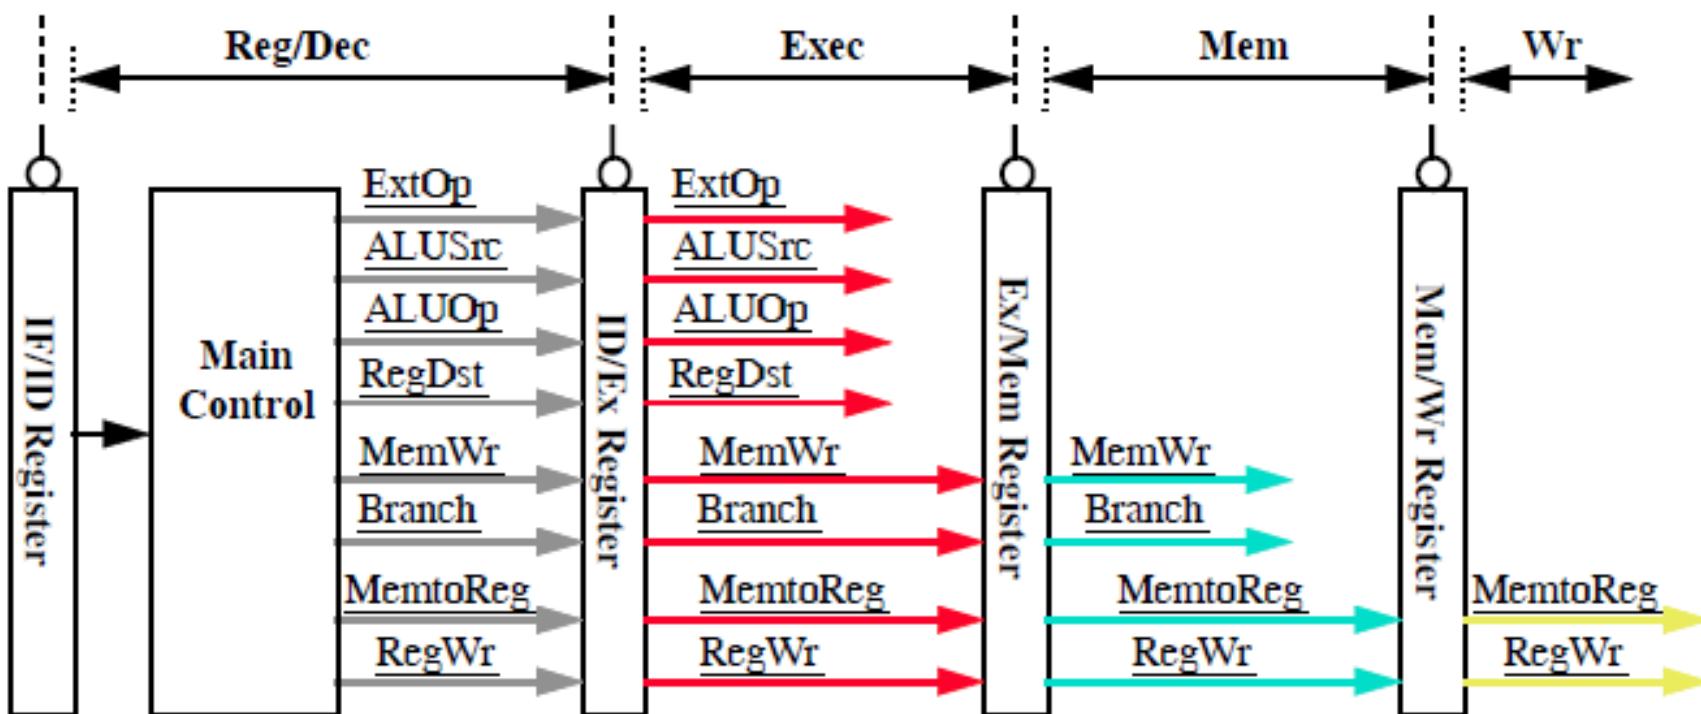


A Pipelined Datapath



Pipeline Control “Data Stationary Control”

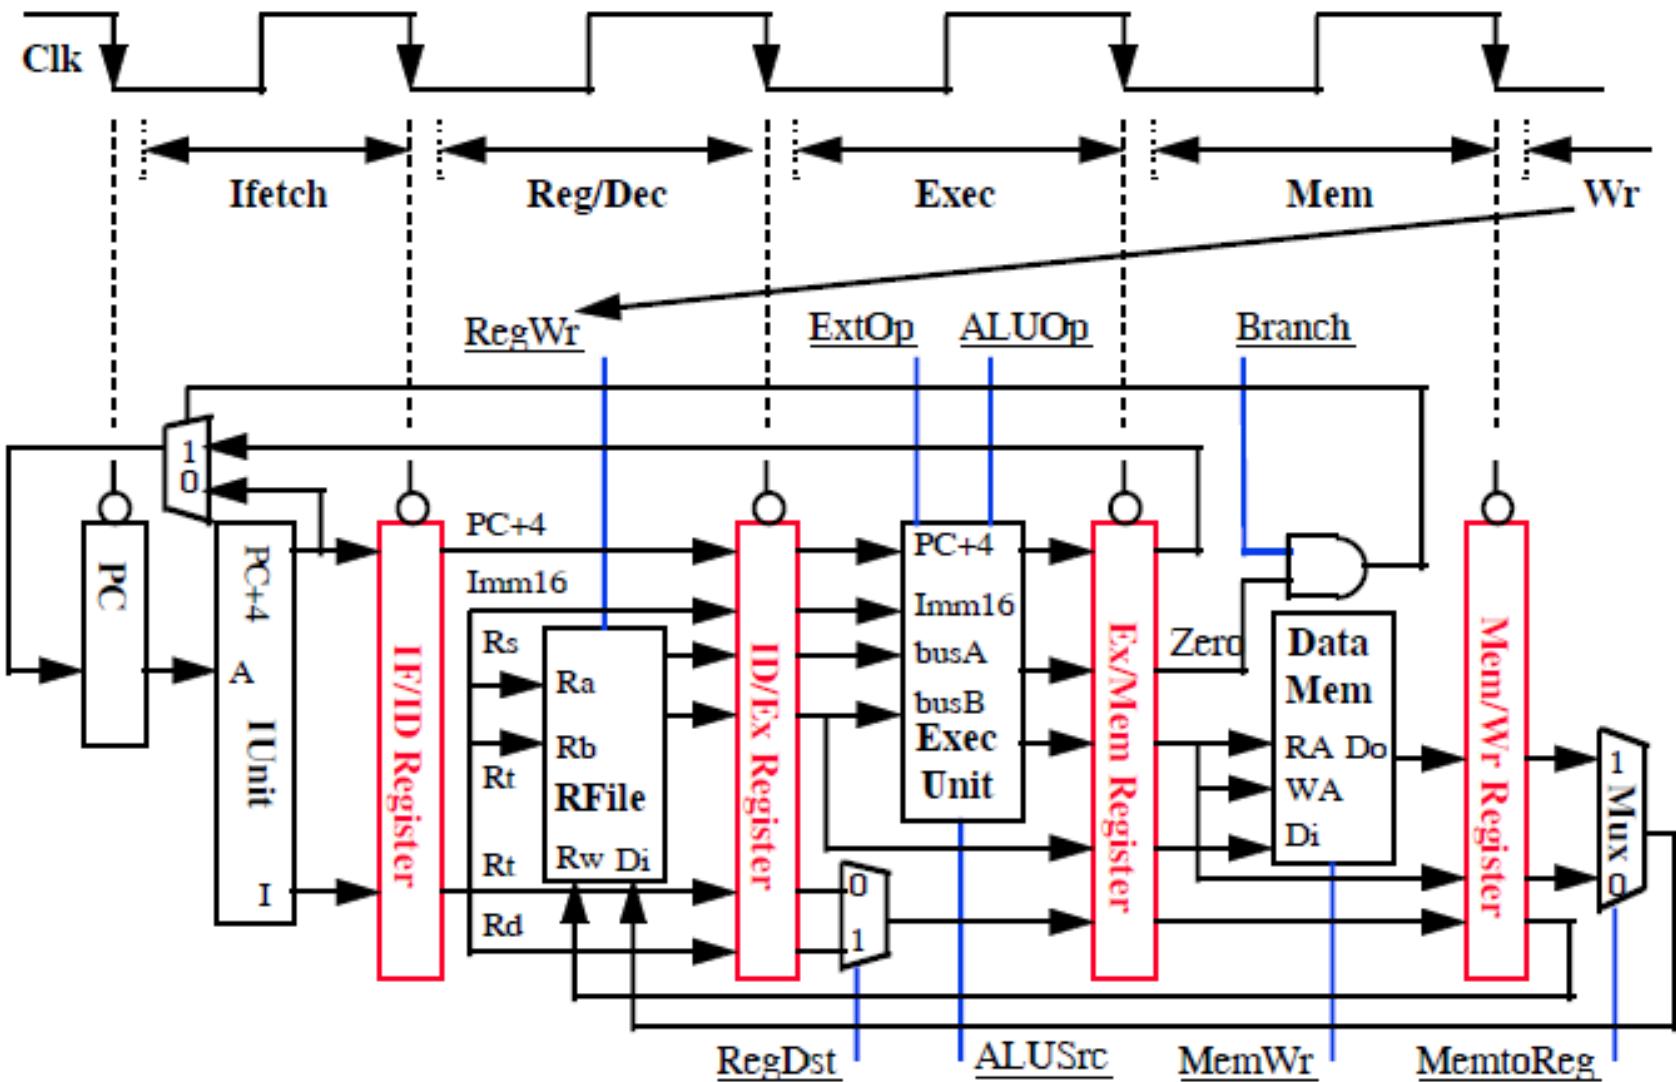
- The Main Control generates the control signals during Reg/Dec
 - Control signals for Exec (ExtOp, ALUSrc, ...) are used 1 cycle later
 - Control signals for Mem (MemWr Branch) are used 2 cycles later
 - Control signals for Wr (MemtoReg MemWr) are used 3 cycles later



Acknowledgment: Almost all of these slides are based on Dave Patterson's CS152 Lecture Slides at UC, Berkeley.

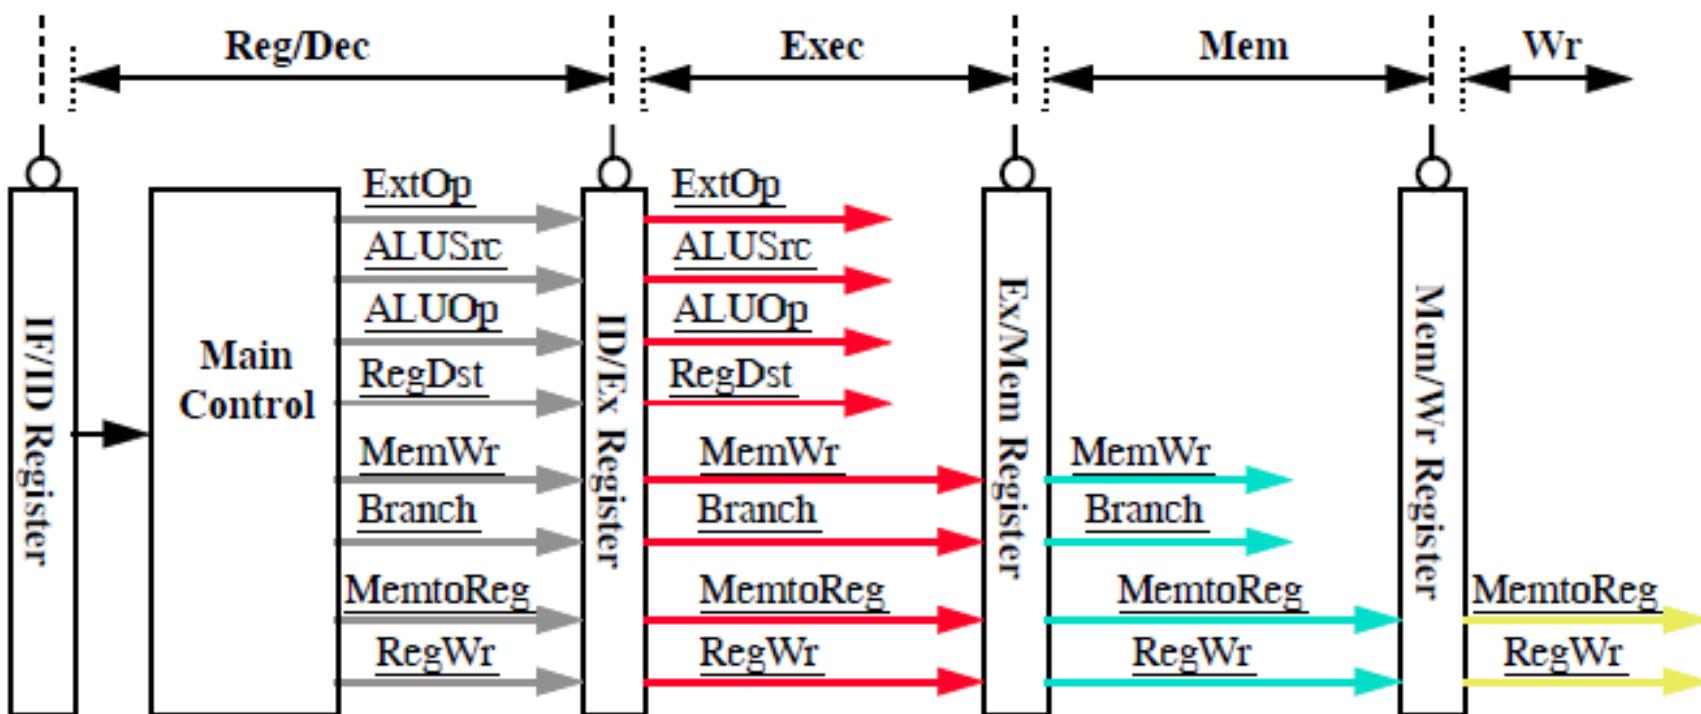
COMPUTER SYSTEMS ORGANIZATION

A Pipelined Datapath



Pipeline Control “Data Stationary Control”

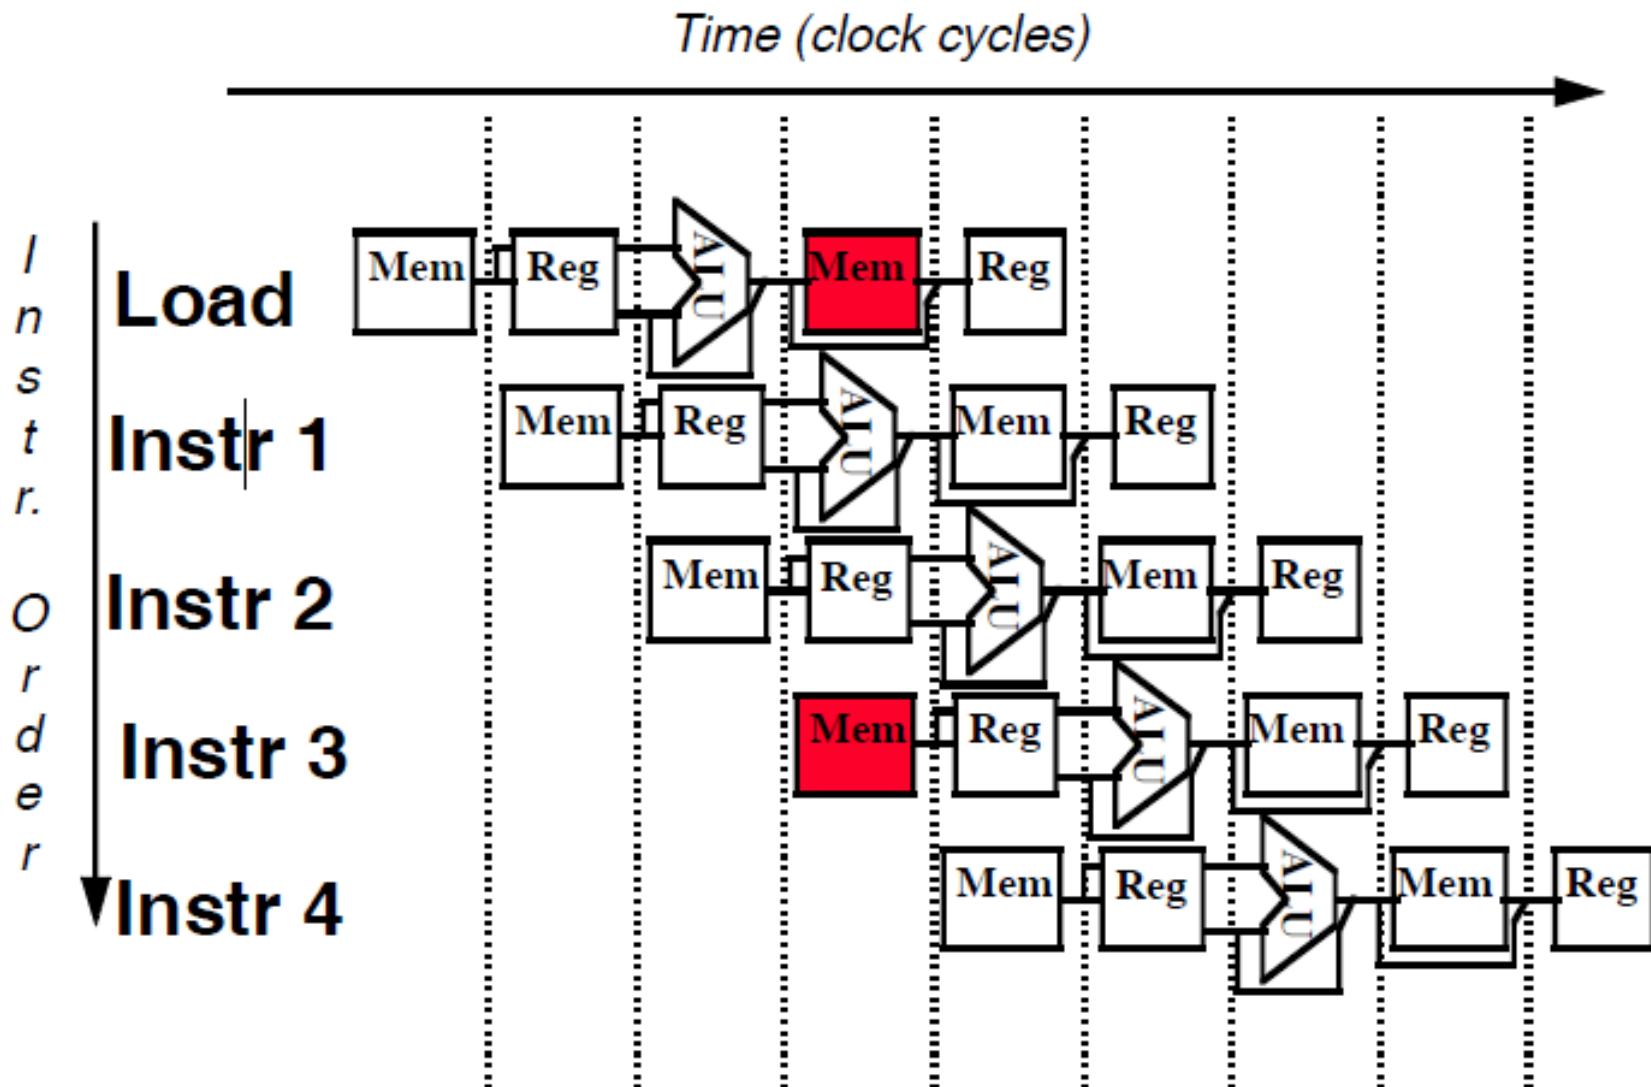
- The Main Control generates the control signals during Reg/Dec
 - Control signals for Exec (ExtOp, ALUSrc, ...) are used 1 cycle later
 - Control signals for Mem (MemWr Branch) are used 2 cycles later
 - Control signals for Wr (MemtoReg MemWr) are used 3 cycles later



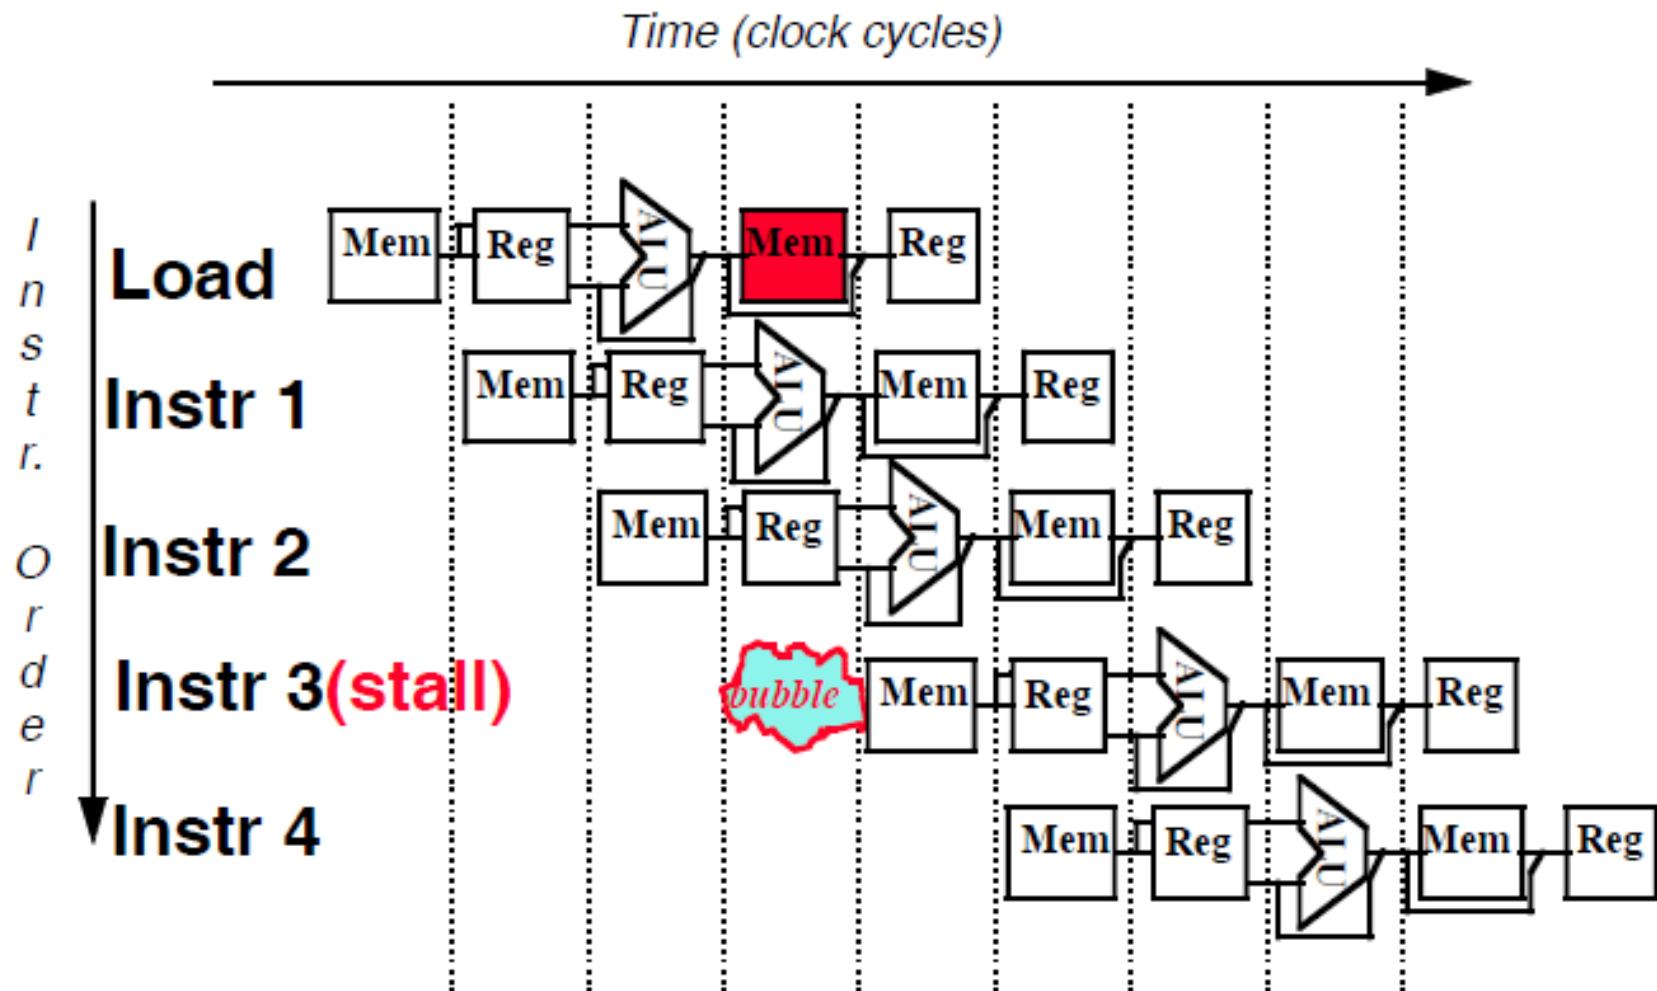
It's not that easy for computers

- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
 - **structural hazards**: HW cannot support this combination of instructions
 - **data hazards**: instruction depends on result of prior instruction still in the pipeline
 - **control hazards**: pipelining of branches & other instructions that change the PC
- Common solution is to **stall** the pipeline until the hazard is resolved, inserting one or more “**bubbles**” in the pipeline

Single Memory is a Structural Hazard

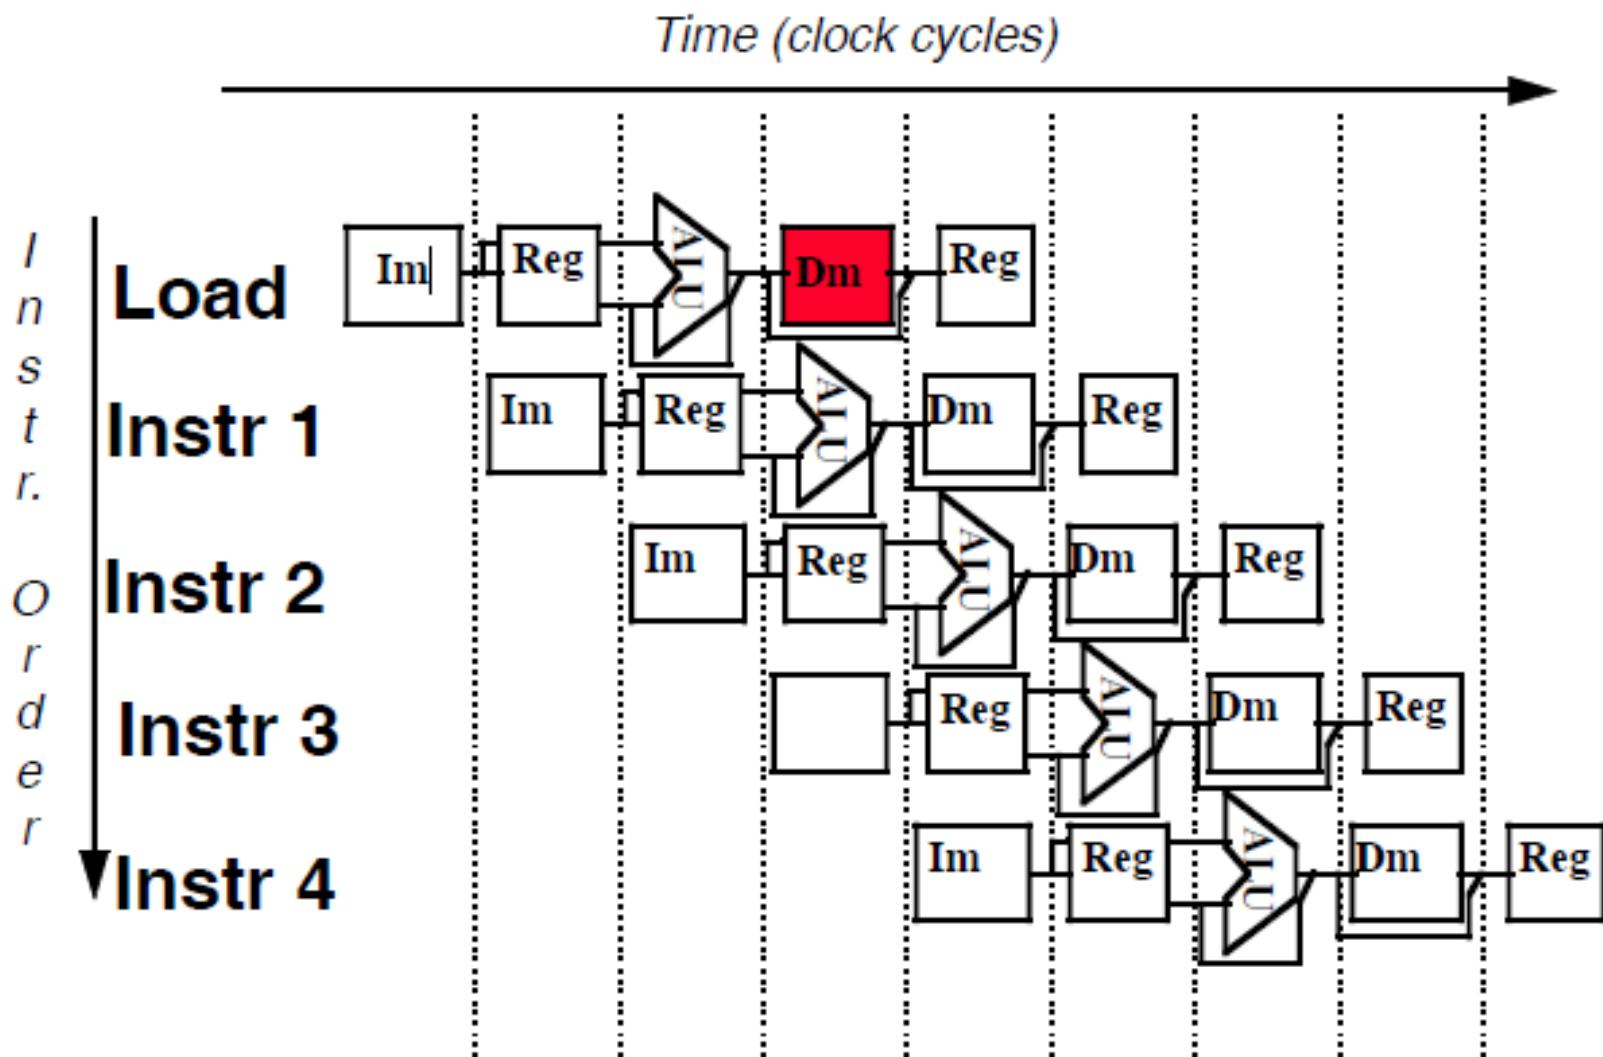


Option 1: Stall to resolve Memory Structural Hazard



Option 2: Duplicate to Resolve Structural Hazard

- Separate Instruction Cache (Im) & Data Cache (Dm)



Data Hazard on r1

add r1,r2,r3

sub r4, r1,r3

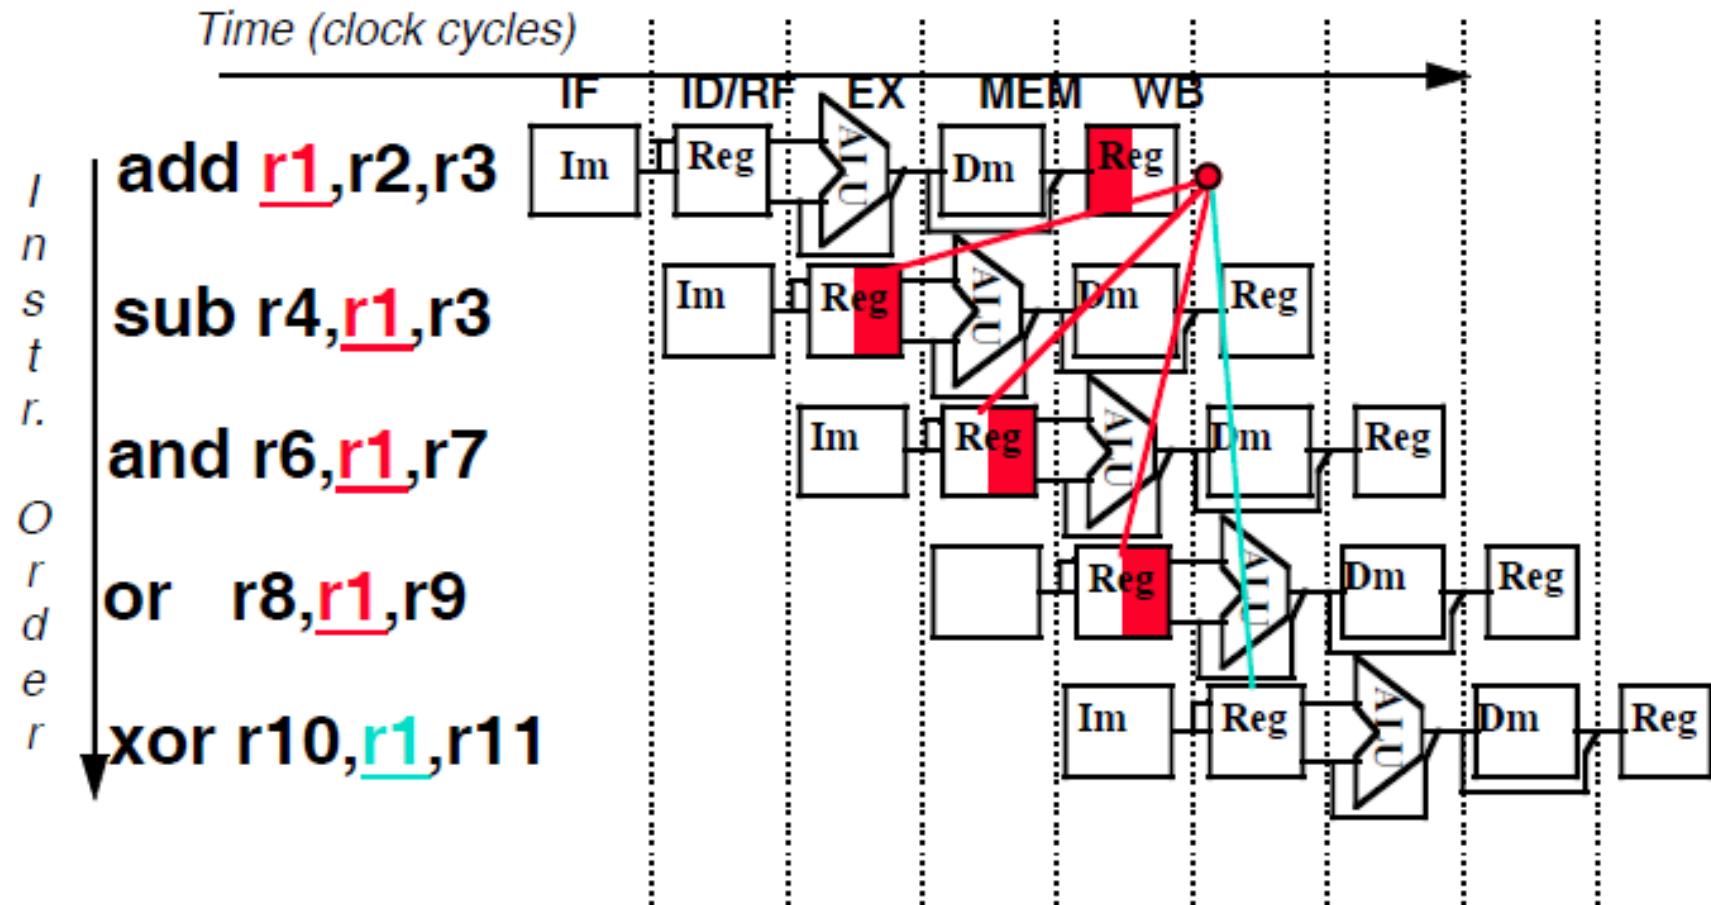
and r6, r1,r7

or r8, r1,r9

xor r10, r1,r11

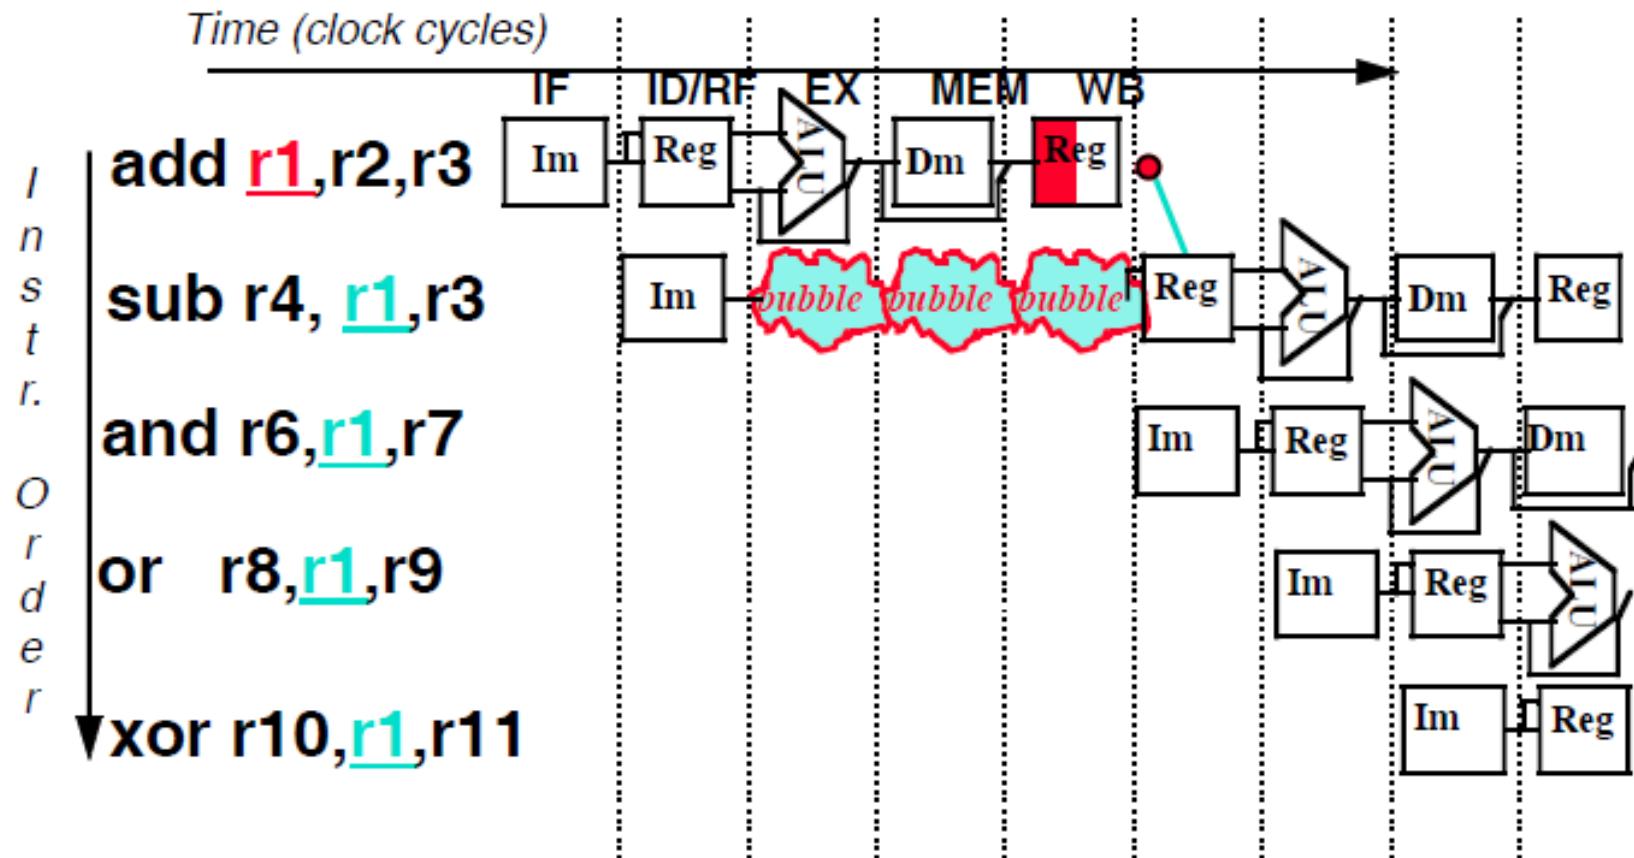
Data Hazard on r1:

- Dependencies backwards in time are hazards



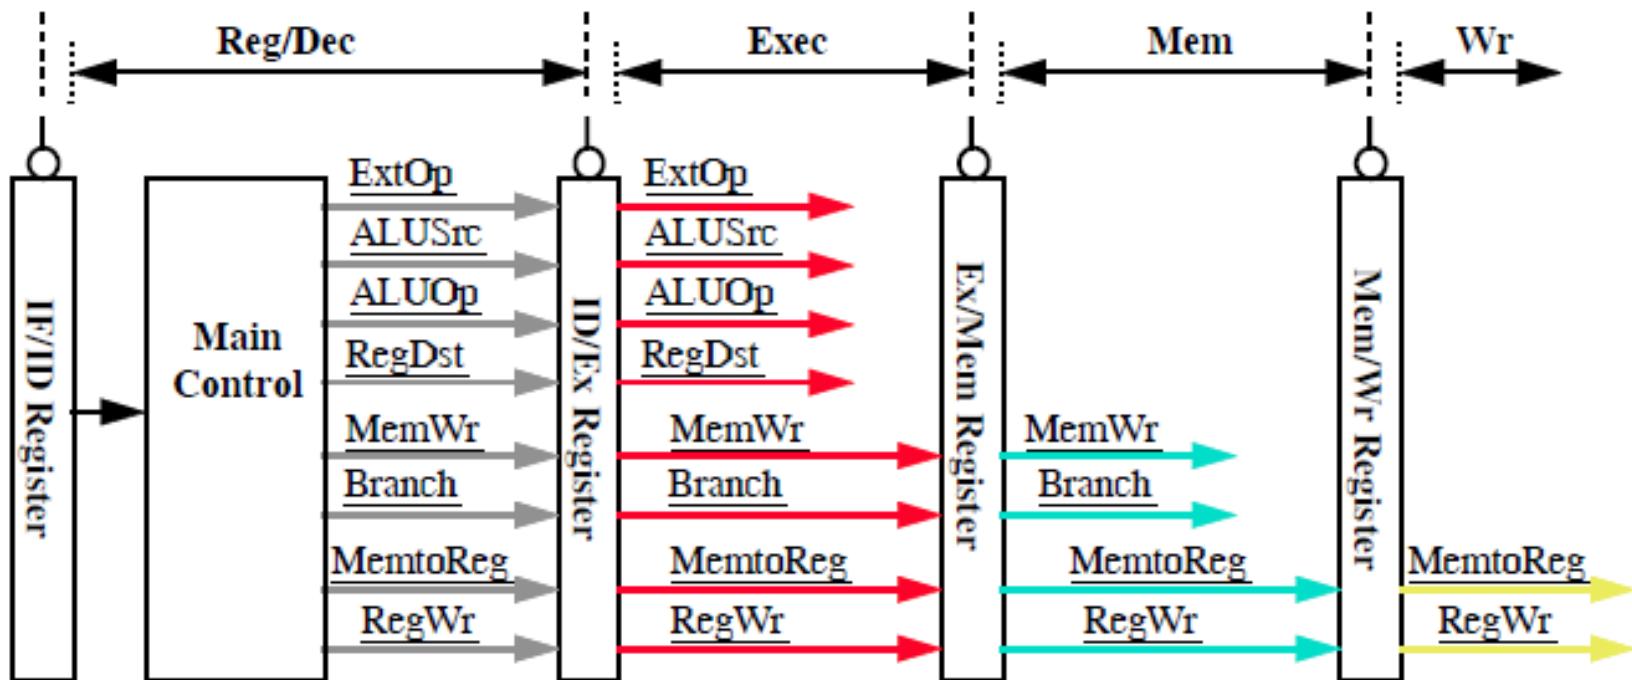
Option1: HW Stalls to Resolve Data Hazard

- Dependencies backwards in time are hazards



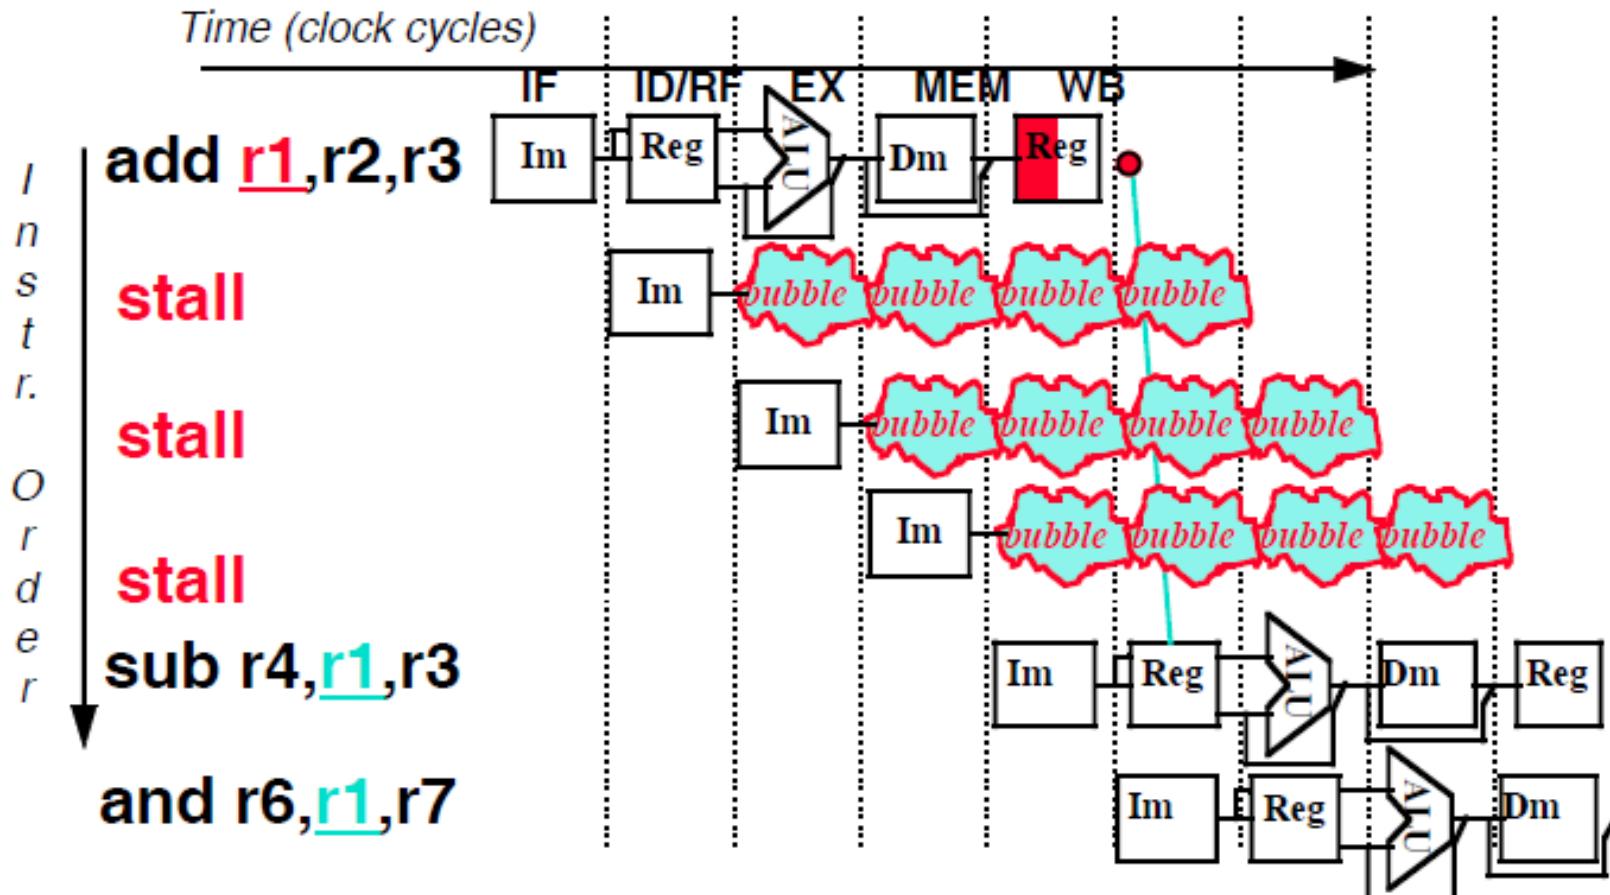
But recall use of “Data Stationary Control”

- The Main Control generates the control signals during Reg/Dec
 - Control signals for Exec (ExtOp, ALUSrc, ...) are used 1 cycle later
 - Control signals for Mem (MemWr Branch) are used 2 cycles later
 - Control signals for Wr (MemtoReg MemWr) are used 3 cycles later



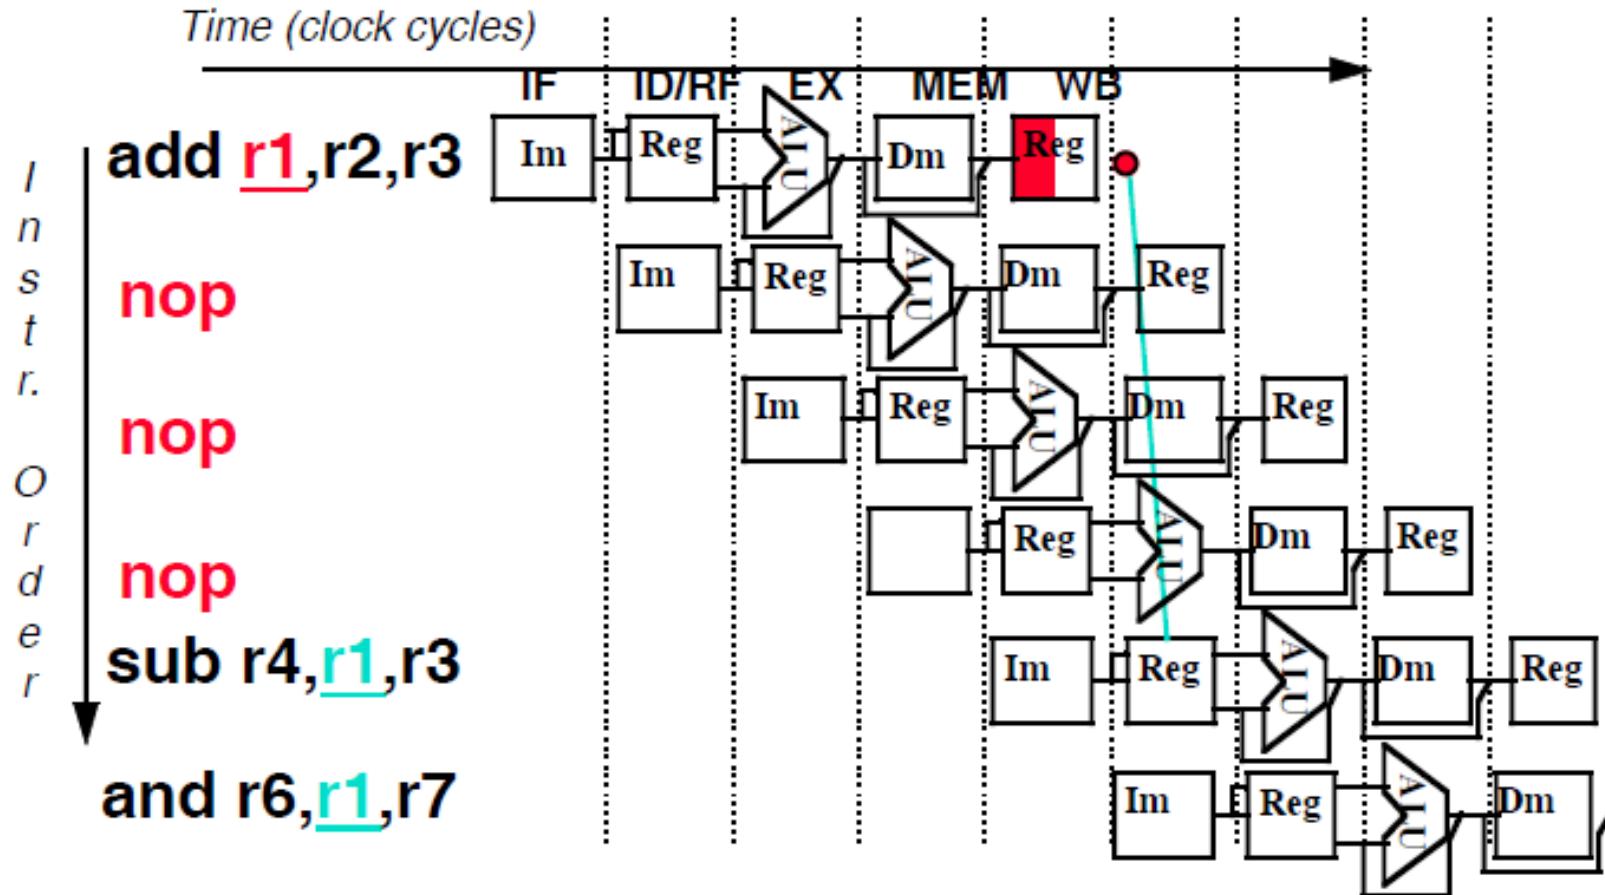
Option 1: How HW really stalls pipeline

- HW doesn't change PC => keeps fetching same instruction & sets control signals to benign values (0)



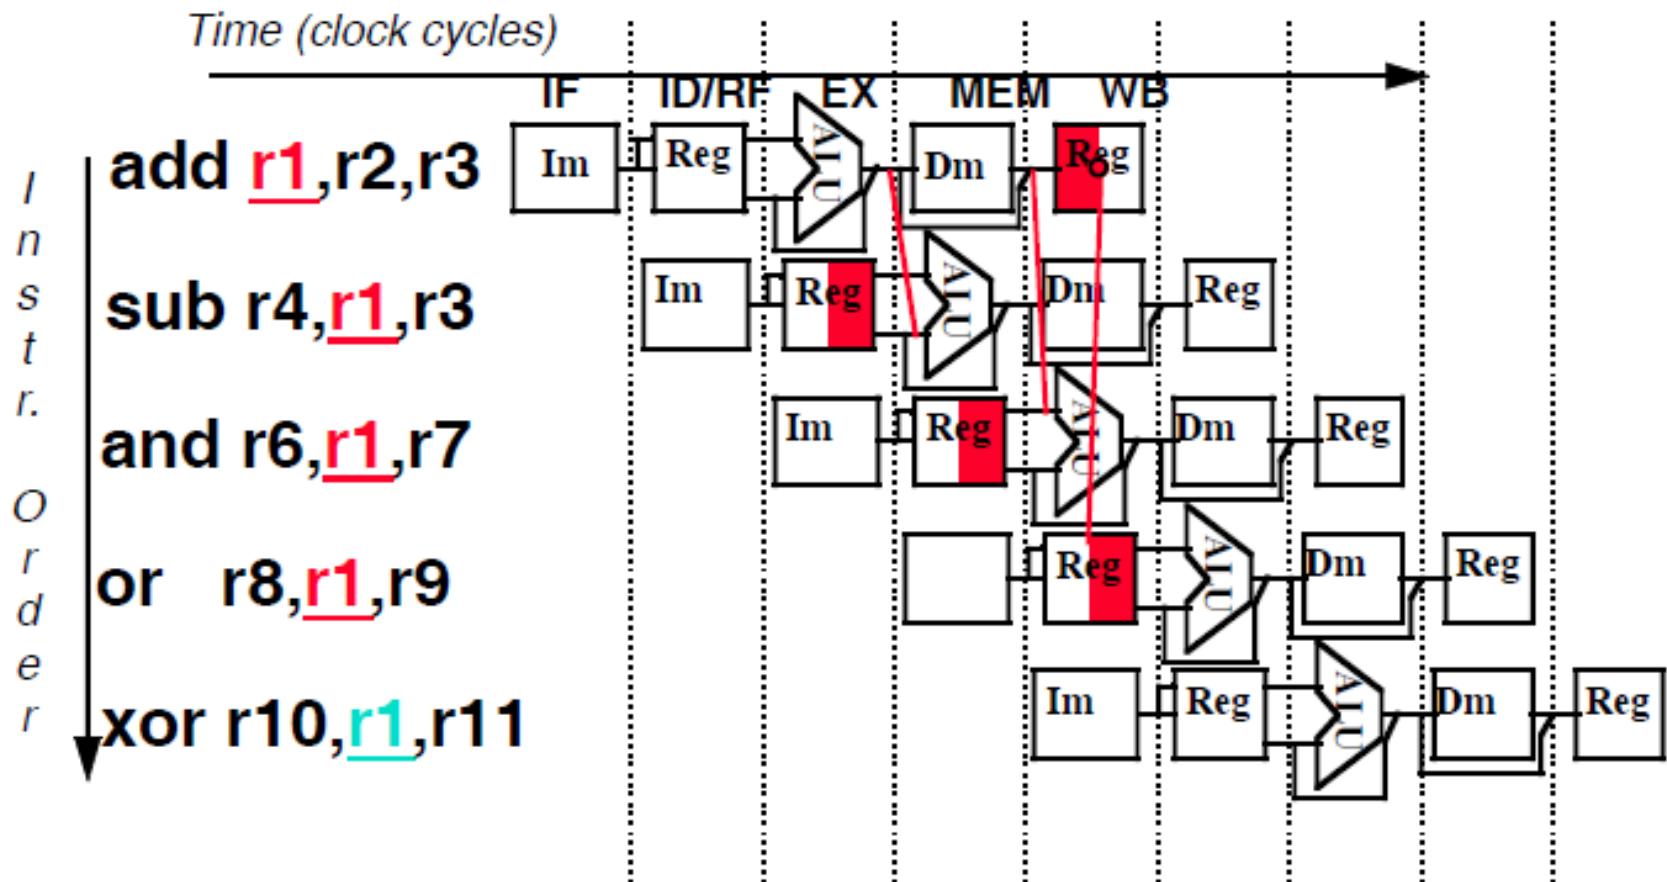
Option 2: SW inserts independent instructions

- Worst case inserts NOP instructions



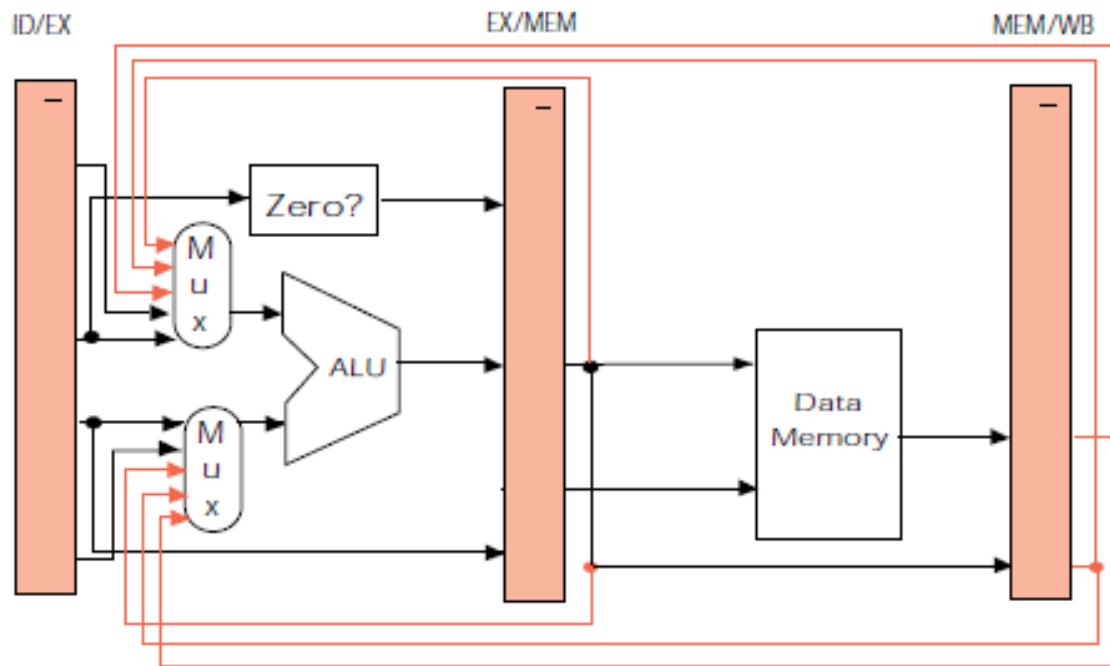
Option 3 Insight: Data is available!

- Pipeline registers already contain needed data

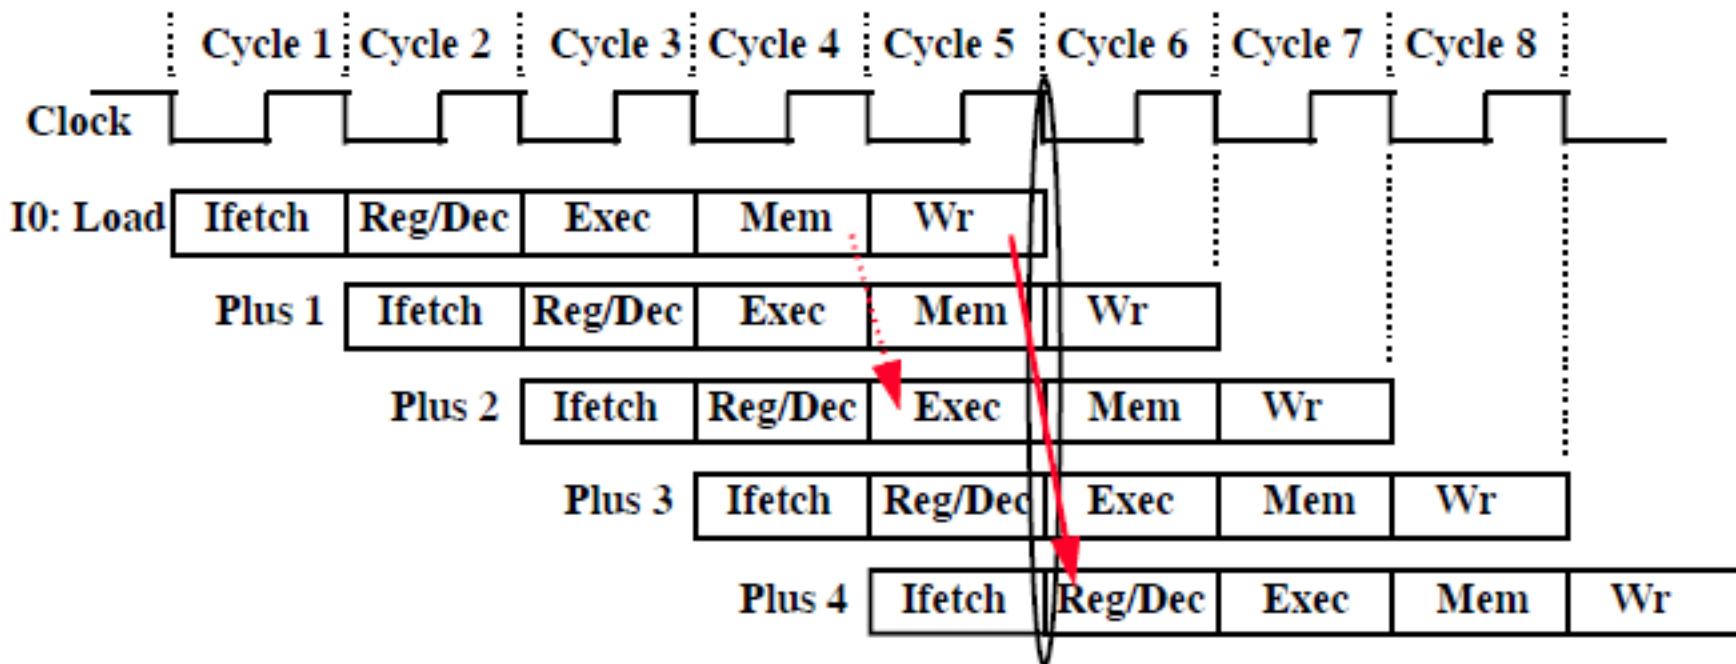


HW Change for “Forwarding” (Bypassing):

- Increase multiplexors to add paths from pipeline registers
- Assumes register read during write gets new value
(otherwise more results to be forwarded)

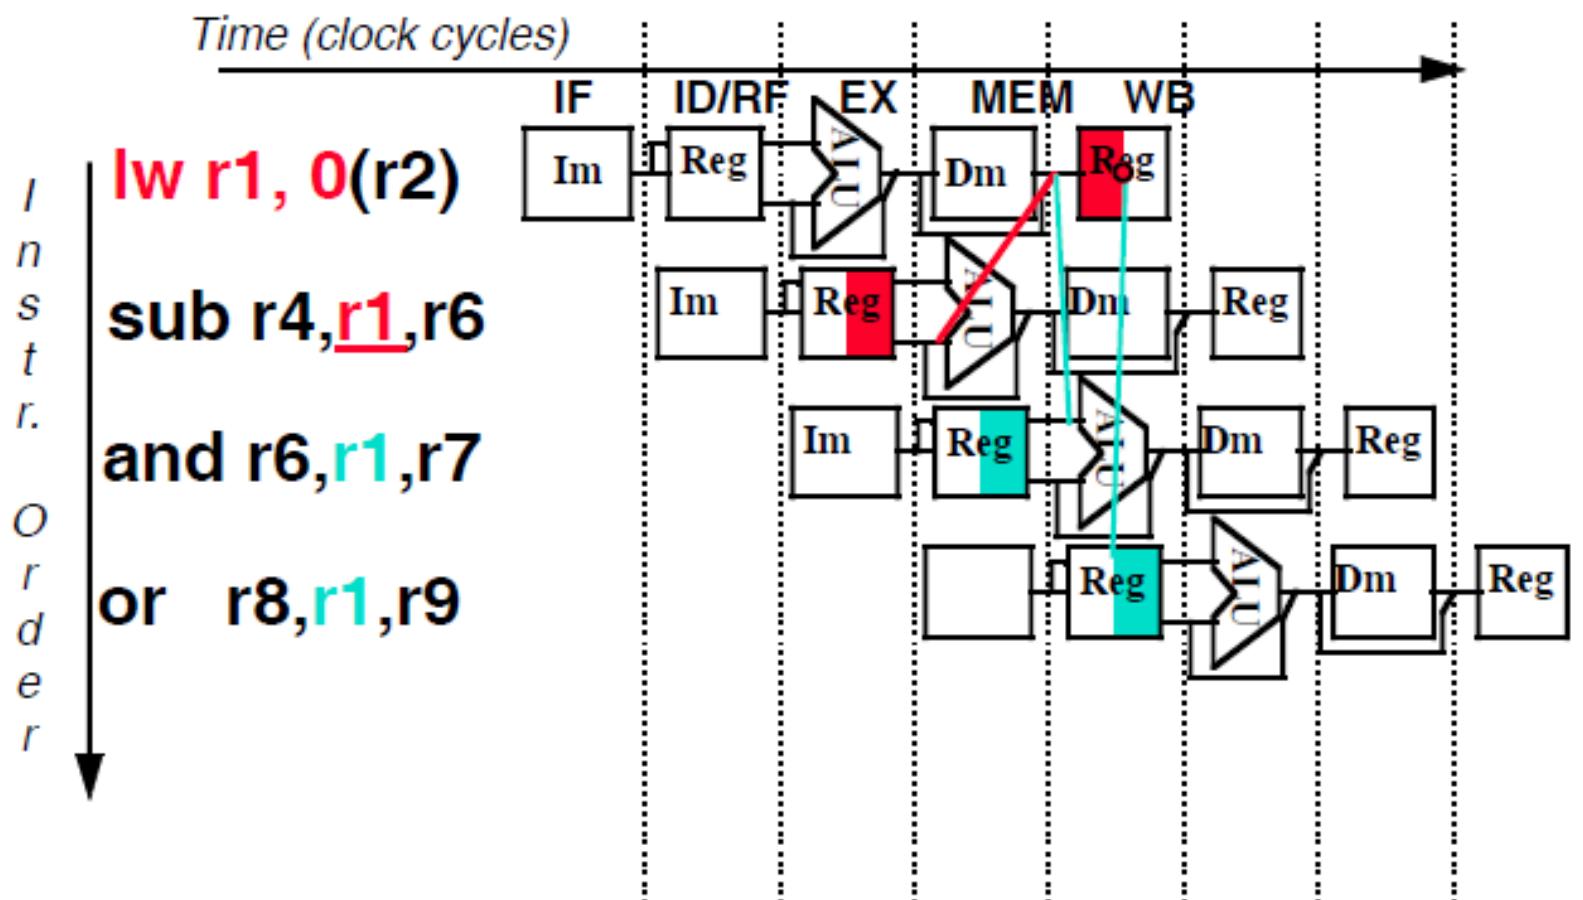


From Last Lecture: The Delay Load Phenomenon



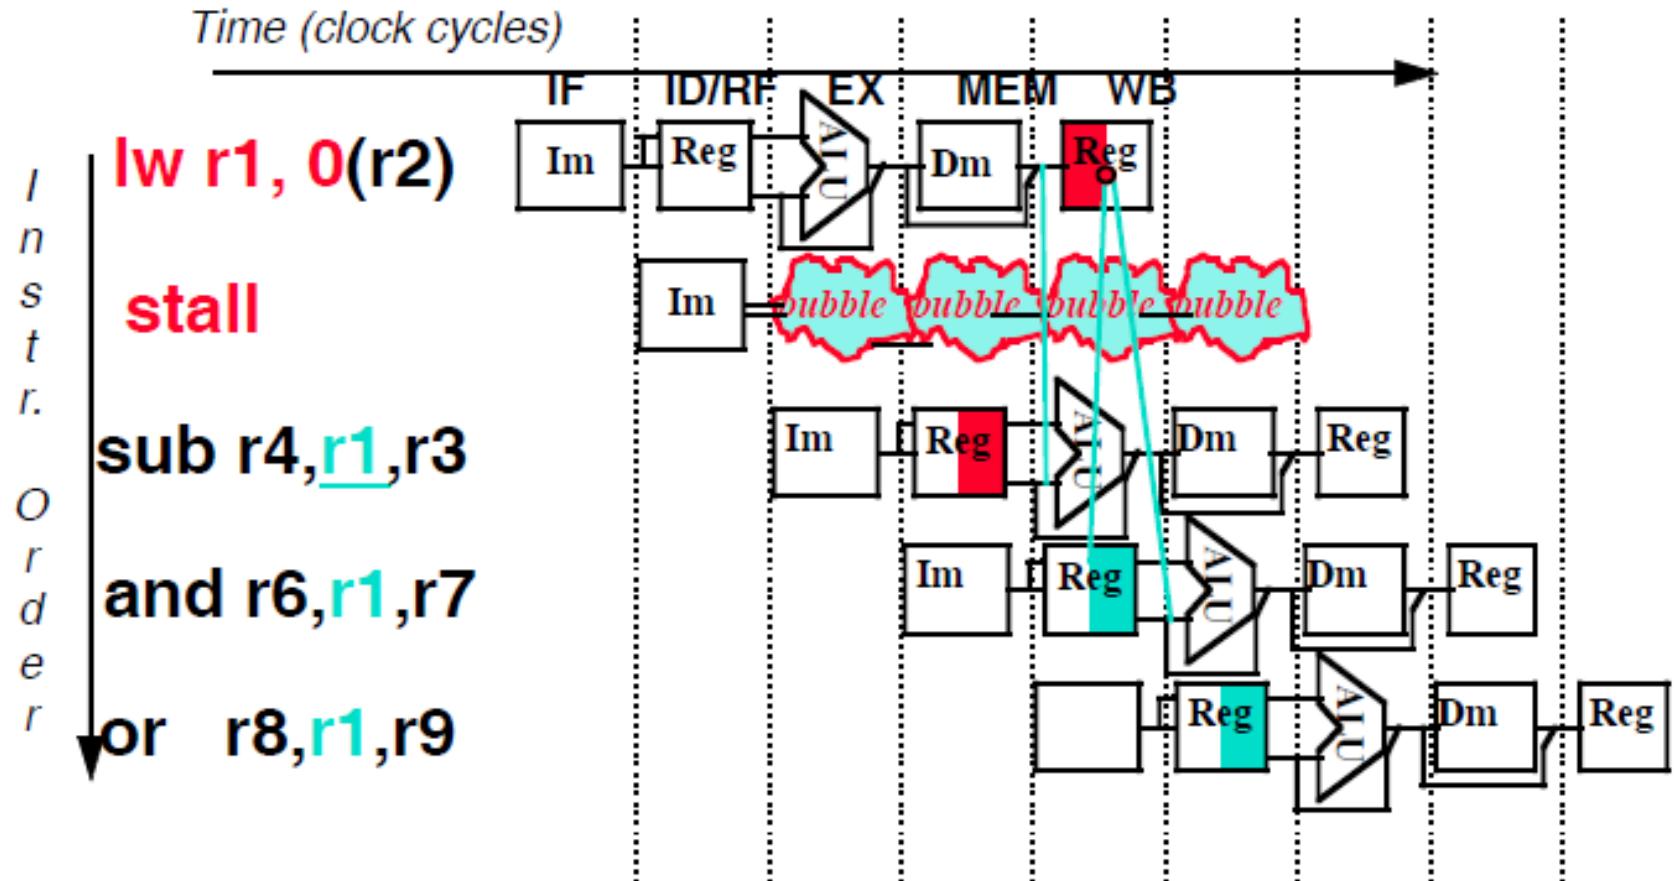
- Although Load is fetched during Cycle 1:
 - The data is NOT written into the Reg File until the end of Cycle 5
 - We cannot read this value from the Reg File until Cycle 6
 - 3-instruction delay before the load take effect

Forwarding reduces Data Hazard to 1 cycle:



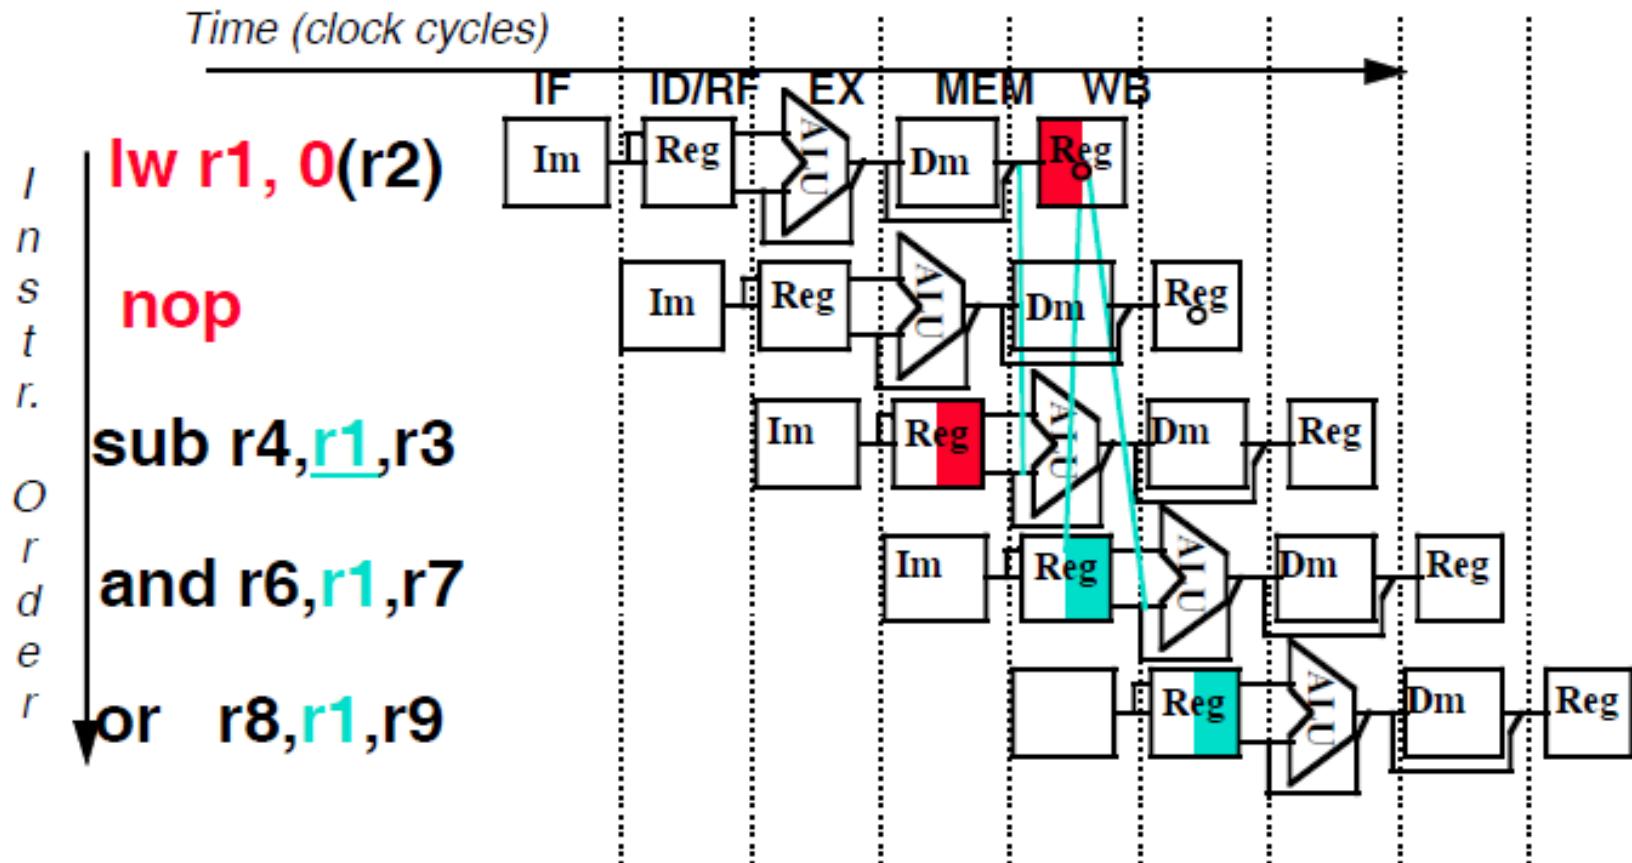
Option1: HW Stalls to Resolve Data Hazard

- “Interlock”: checks for hazard & stalls



Option 2: SW inserts independent instructions

- Worst case inserts NOP instructions
- MIPS I solution: No HW checking



Software Scheduling to Avoid Load Hazards

Try producing fast code for

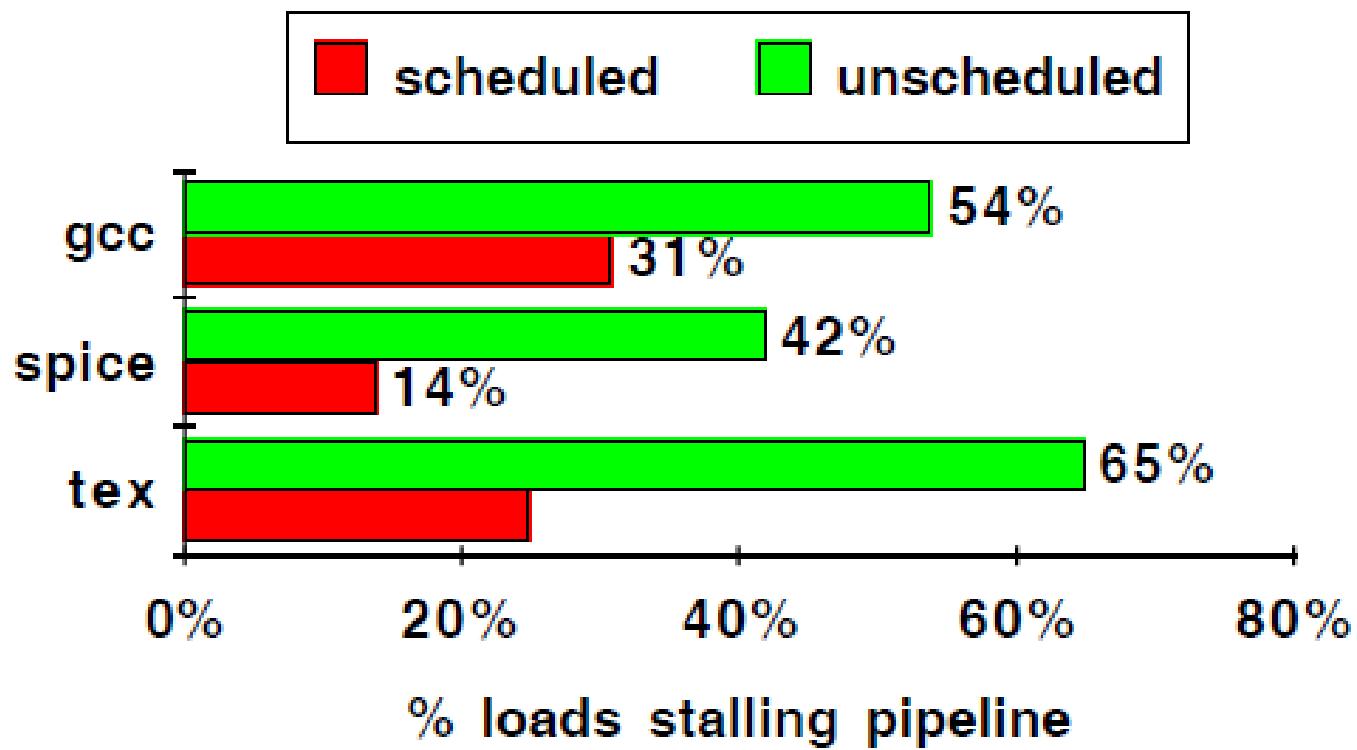
```
a = b + c;  
d = e - f;
```

assuming a, b, c, d ,e, and f
in memory.

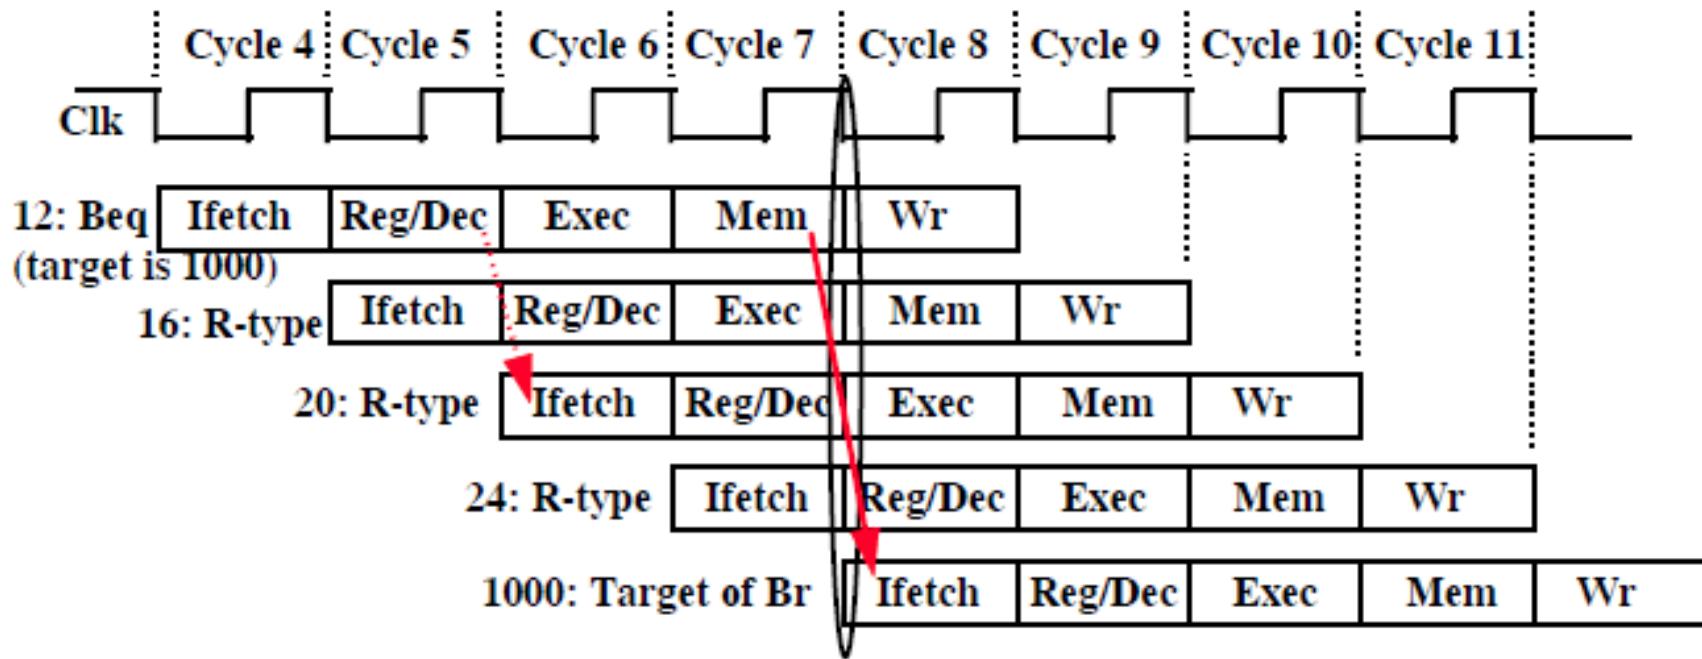
Slow code:

```
LW    Rb,b  
LW    Rc,c  
ADD   Ra,Rb,Rc  
SW    a,Ra  
LW    Re,e  
LW    Rf,f  
SUB   Rd,Re,Rf  
SW    d,Rd
```

Compiler Avoiding Load Stalls:

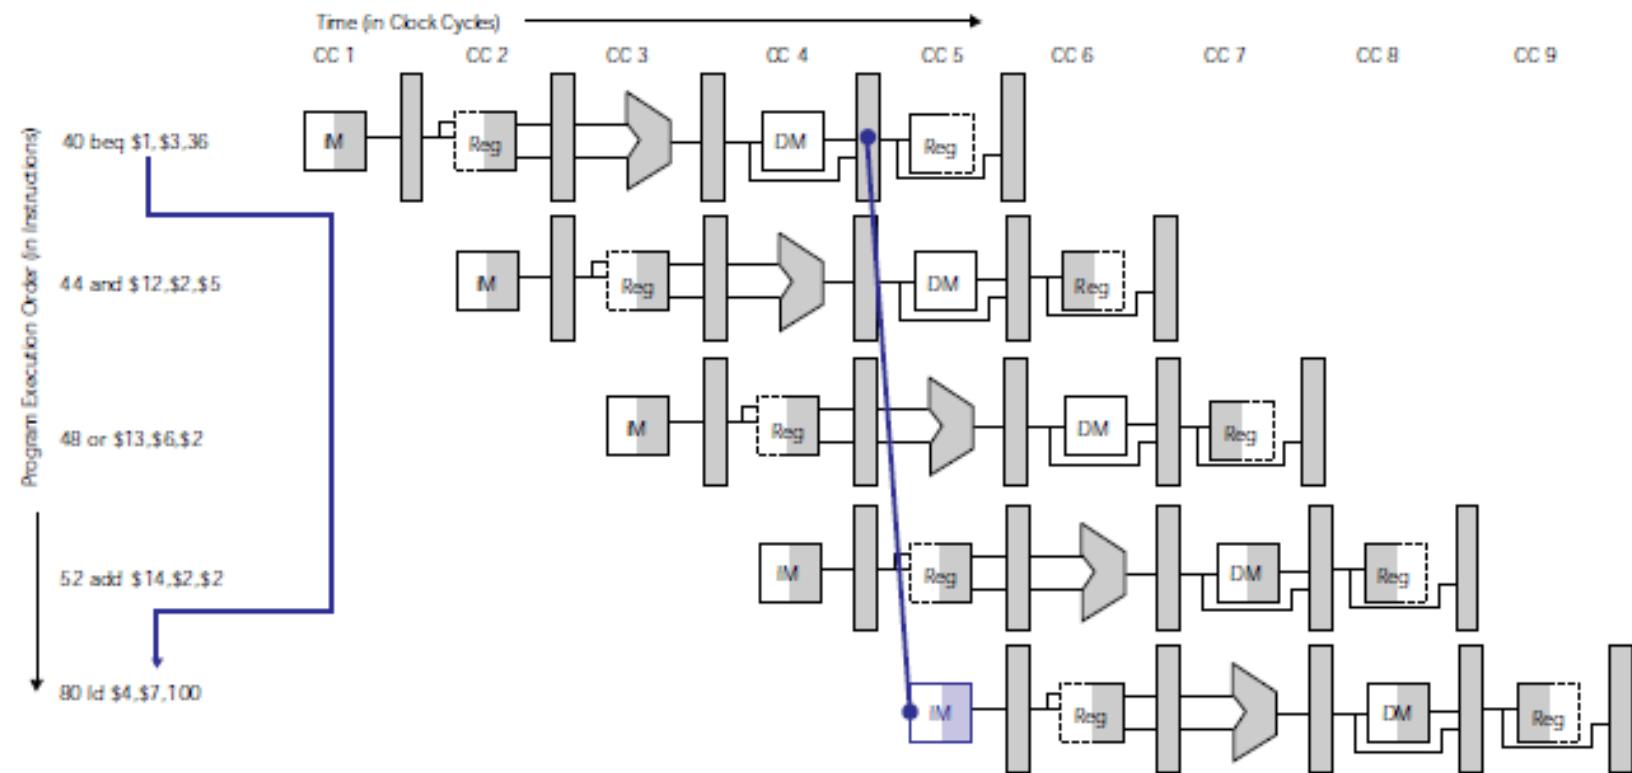


From Last Lecture: The Delay Branch Phenomenon



- ° Although Beq is fetched during Cycle 4:
 - Target address is NOT written into the PC until the end of Cycle 7
 - Branch's target is NOT fetched until Cycle 8
 - 3-instruction delay before the branch take effect

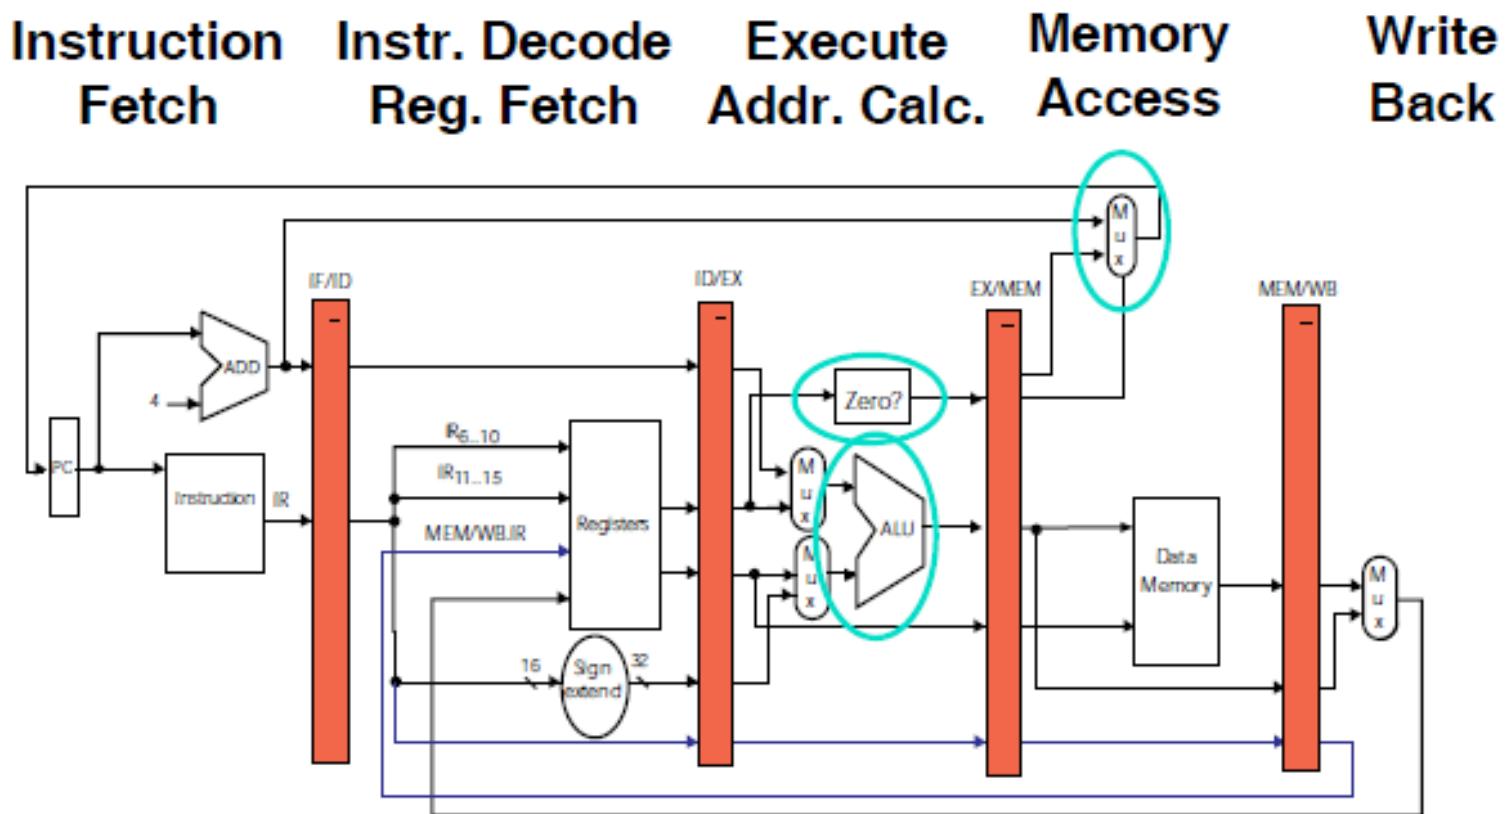
Control Hazard on Branches: 3 stage stall



Branch Stall Impact

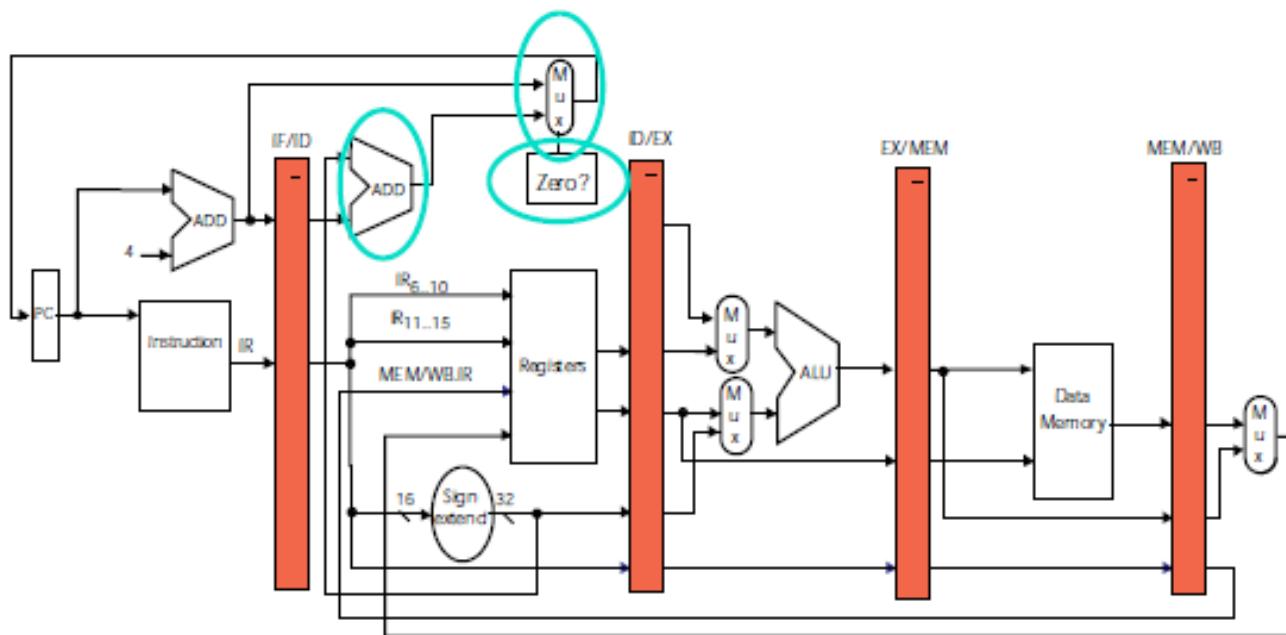
- ° If CPI = 1, 30% branch, Stall 3 cycles => new CPI = 1.9!
- ° 2 part solution:
 - Determine branch taken or not sooner, AND
 - Compute taken branch address earlier
- ° MIPS branch tests = 0 or $\neq 0$
- ° Solution Option 1:
 - Move Zero test to ID/RF stage
 - Adder to calculate new PC in ID/RF stage
 - 1 clock cycle penalty for branch vs. 3

Option 1: move HW forward to reduce branch delay



Branch Delay now 1 clock cycle

Instruction Fetch	Instr. Decode Reg. Fetch	Execute Addr. Calc.	Memory Access	Write Back
-------------------	--------------------------	---------------------	---------------	------------



Option 2: Define Branch as Delayed

- Worst case, SW inserts NOP into branch delay
- Where get instructions to fill branch delay slot?
 - Before branch instruction
 - From the target address: only valuable when branch
 - From fall through: only valuable when don't branch
- Compiler effectiveness for single branch delay slot:
 - Fills about 60% of branch delay slots
 - About 80% of instructions executed in branch delay slots useful in computation
 - about 50% ($60\% \times 80\%$) of slots usefully filled

When is pipelining hard?

- **Interrupts**: 5 instructions executing in 5 stage pipeline
 - How to stop the pipeline?
 - Restart?
 - Who caused the interrupt?

Stage Problem interrupts occurring

IF	Page fault on instruction fetch; misaligned memory access; memory-protection violation
ID	Undefined or illegal opcode
EX	Arithmetic interrupt
MEM	Page fault on data fetch; misaligned memory access; memory-protection violation

- Load with data page fault, Add with instruction page fault?
- Solution 1: interrupt vector/instruction[], check last stage
- Solution 2: interrupt ASAP, restart everything incomplete

When is pipelining hard?

- Complex Addressing Modes and Instructions
- Address modes: Autoincrement causes register change during instruction execution
 - Interrupts?
 - Now worry about write hazards since write no longer last stage
 - Write After Read (WAR): Write occurs before independent read
 - Write After Write (WAW): Writes occur in wrong order, leaving wrong result in registers
 - (Previous data hazard called RAW, for Read After Write)
- Memory-memory Move instructions
 - Multiple page faults
 - make progress?

When is pipelining hard?

- Floating Point: long execution time
- Also, may pipeline FP execution unit so that can initiate new instructions without waiting full latency

<i>FP Instruction</i>	<i>Latency</i>	<i>Initiation Rate</i>	(MIPS R4000)
Add, Subtract	4	3	
Multiply	8	4	
Divide	36	35	
Square root	112	111	
Negate	2	1	
Absolute value	2	1	
FP compare	3	2	

- Divide, Square Root take \approx 10X to \approx 30X longer than Add
 - Exceptions?
 - Adds WAR and WAW hazards since pipelines are no longer same length

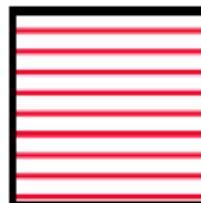
Hazard Detection

Suppose instruction i is about to be issued and a predecessor instruction j is in the instruction pipeline.

$\text{Rregs}(i)$ = Registers read by instruction i

$\text{Wregs}(i)$ = Registers written by instruction i

- A RAW hazard exists on register ρ if $\exists \rho, \rho \in \text{Rregs}(i) \cap \text{Wregs}(j)$
 - Keep a record of pending writes (for inst's in the pipe) and compare with operand regs of current instruction.
 - When instruction issues, reserve its result register.
 - When on operation completes, remove its write reservation.



- A WAW hazard exists on register ρ if $\exists \rho, \rho \in \text{Wregs}(i) \cap \text{Wregs}(j)$
- A WAR hazard exists on register ρ if $\exists \rho, \rho \in \text{Wregs}(i) \cap \text{Rregs}(j)$

First Generation RISC Pipelines

- All instructions follow same pipeline order (“static schedule”).
- Register write in last stage
 - Avoid WAW hazards
- All register reads performed in first stage after issue.
 - Avoid WAR hazards
- Memory access in stage 4
 - Avoid all memory hazards
- Control hazards resolved by delayed branch (with fast path)
- RAW hazards resolved by bypass, except on load results which are resolved by fiat (delayed load).

Substantial pipelining with very little cost or complexity.

Machine organization is (slightly) exposed!

Relies very heavily on "hit assumption" of memory accesses in cache

Review: Summary of Pipelining Basics

- Speed Up \leq Pipeline Depth; if ideal CPI is 1, then:

$$\text{Speedup} = \frac{\text{Pipeline depth}}{1 + \text{Pipeline stall cycles per instruction}} \times \frac{\text{Clock cycle unpipelined}}{\text{Clock cycle pipelined}}$$

- Hazards limit performance on computers:

- structural: need more HW resources
- data: need forwarding, compiler scheduling
- control: early evaluation & PC, delayed branch, prediction

- Increasing length of pipe increases impact of hazards since pipelining helps instruction bandwidth, not latency
- Compilers key to reducing cost of data and control hazards
 - load delay slots
 - branch delay slots
- Exceptions, Instruction Set, FP makes pipelining harder
- Longer pipelines => Branch prediction, more instruction parallelism?